

Introduction to Stream Processing



Guido Schmutz
Frankfurt - 21.2.2019



@gschmutz



guidoschmutz.wordpress.com

BASEL ▪ BERN ▪ BRUGG ▪ DÜSSELDORF ▪ FRANKFURT A.M. ▪ FREIBURG I.BR. ▪ GENF
HAMBURG ▪ KOPENHAGEN ▪ LAUSANNE ▪ MÜNCHEN ▪ STUTTGART ▪ WIEN ▪ ZÜRICH

@gschmutz

trivadis
makes IT easier.

■ Agenda

1. Motivation for Stream Processing?
2. Capabilities for Stream Processing
3. Implementing Stream Processing Solutions
4. Demo
5. Summary

■ Guido Schmutz

Working at Trivadis for more than 22 years

Oracle Groundbreaker Ambassador & Oracle ACE Director



Consultant, Trainer Software Architect for Java, Oracle, SOA and Big Data / Fast Data

Head of Trivadis Architecture Board

Technology Manager @ Trivadis

More than 30 years of software development experience

Contact: guido.schmutz@trivadis.com

Blog: <http://guidoschmutz.wordpress.com>

Slideshare: <http://www.slideshare.net/gschmutz>

Twitter: [@gschmutz](https://twitter.com/gschmutz)

gschmutz 22:08 on April 18, 2017
Topic: flink (12), kafka (19), kafka-connect (4), kafka-streams (17), spark-streaming (12), storm (10), streamsets (4)

This is the 62nd edition of my blog series blog series around Stream Processing and Analytics!

Every week I'm also updating the following two lists with the presentations/videos of the current week:

- [Presentations from Slideshare](#)
- [Videos from YouTube](#)

As usual, find below the new 145th edition

News and Blog [Tools](#)

General

- [Multi-Master Replication For Geo-Distributed Data: It's more than you think](#) by Ellen Friedman
- [Understanding Indicators of Attack \(IOAs\): The Power of Event Stream Processing in CrowdStrike Falcon](#) by Dan Brown
- [Stream processing and messaging systems for the IoT age](#) by Ben Lorica

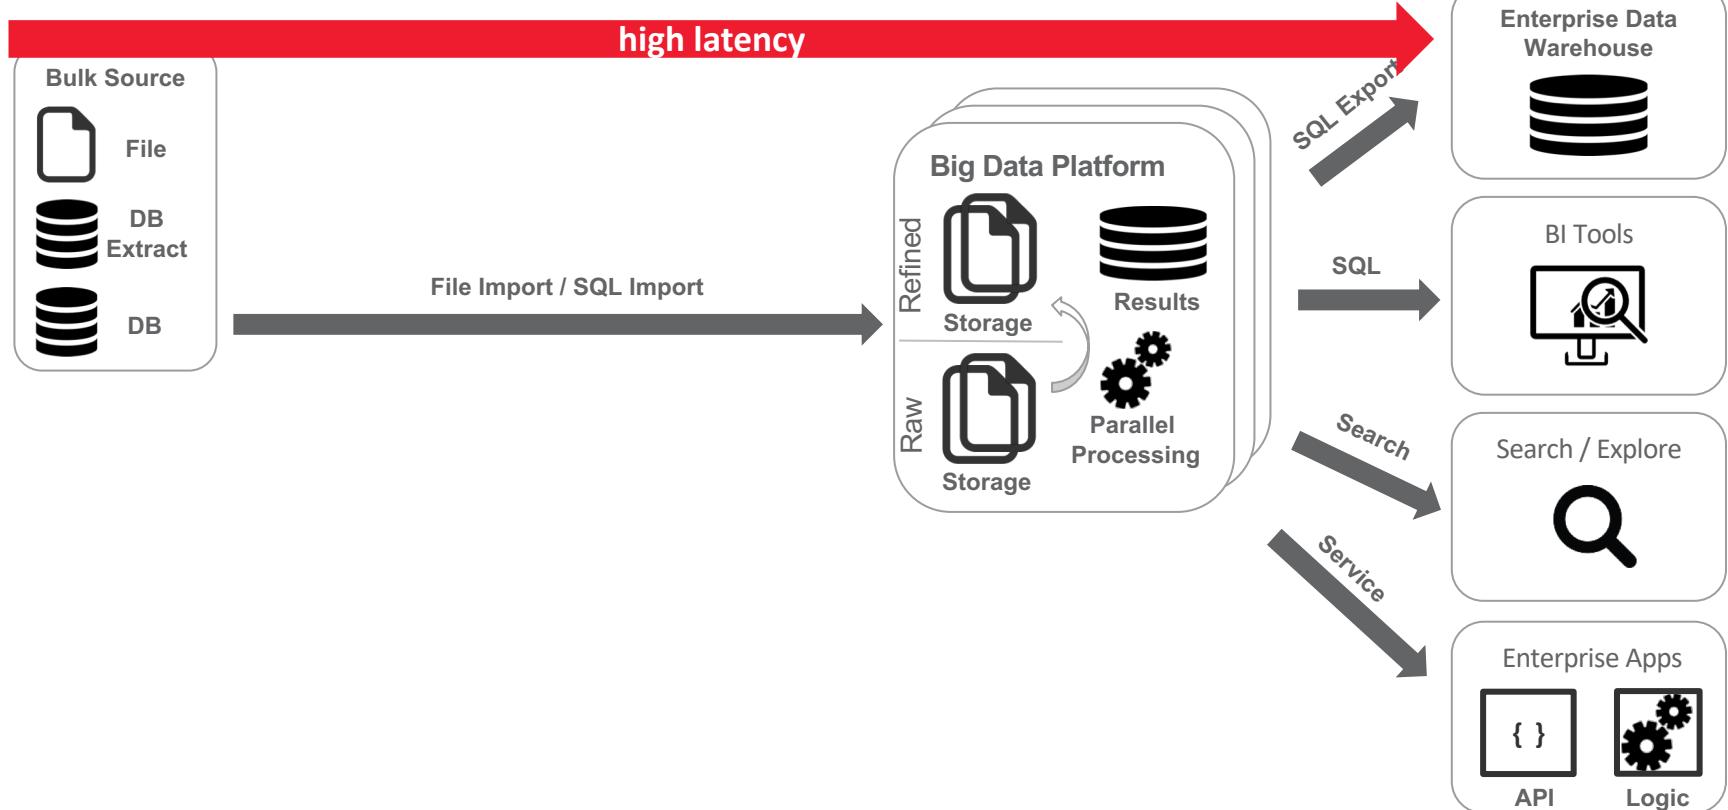
Apache Kafka / Kafka Streams / Confluent Platform

- [Creating a Data Pipeline with Kafka Connect API - From Architecture to Operations](#) by Alexandra Wang
- [Streaming Spring Boot Application Logs to ELK Stack—Part 1](#) by kasyadymuthu
- [Streaming Spring Boot Application Logs to Apache Kafka – ELK/Kafka Stack—Part 2](#) by kasyadymuthu

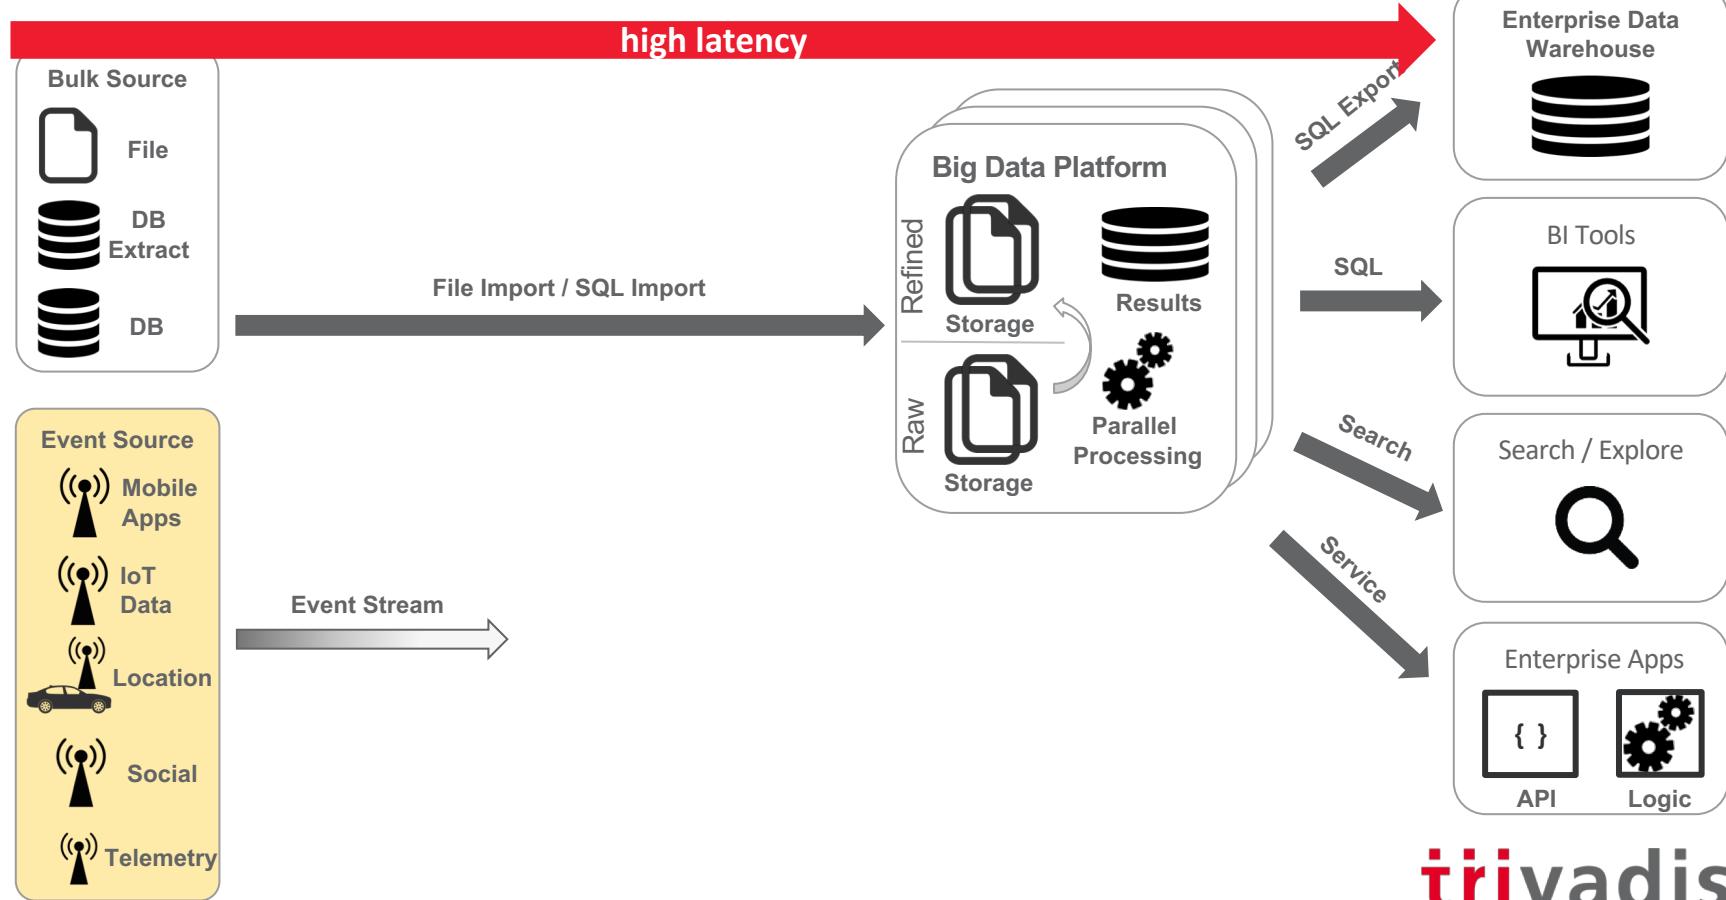


Motivation for Stream Processing?

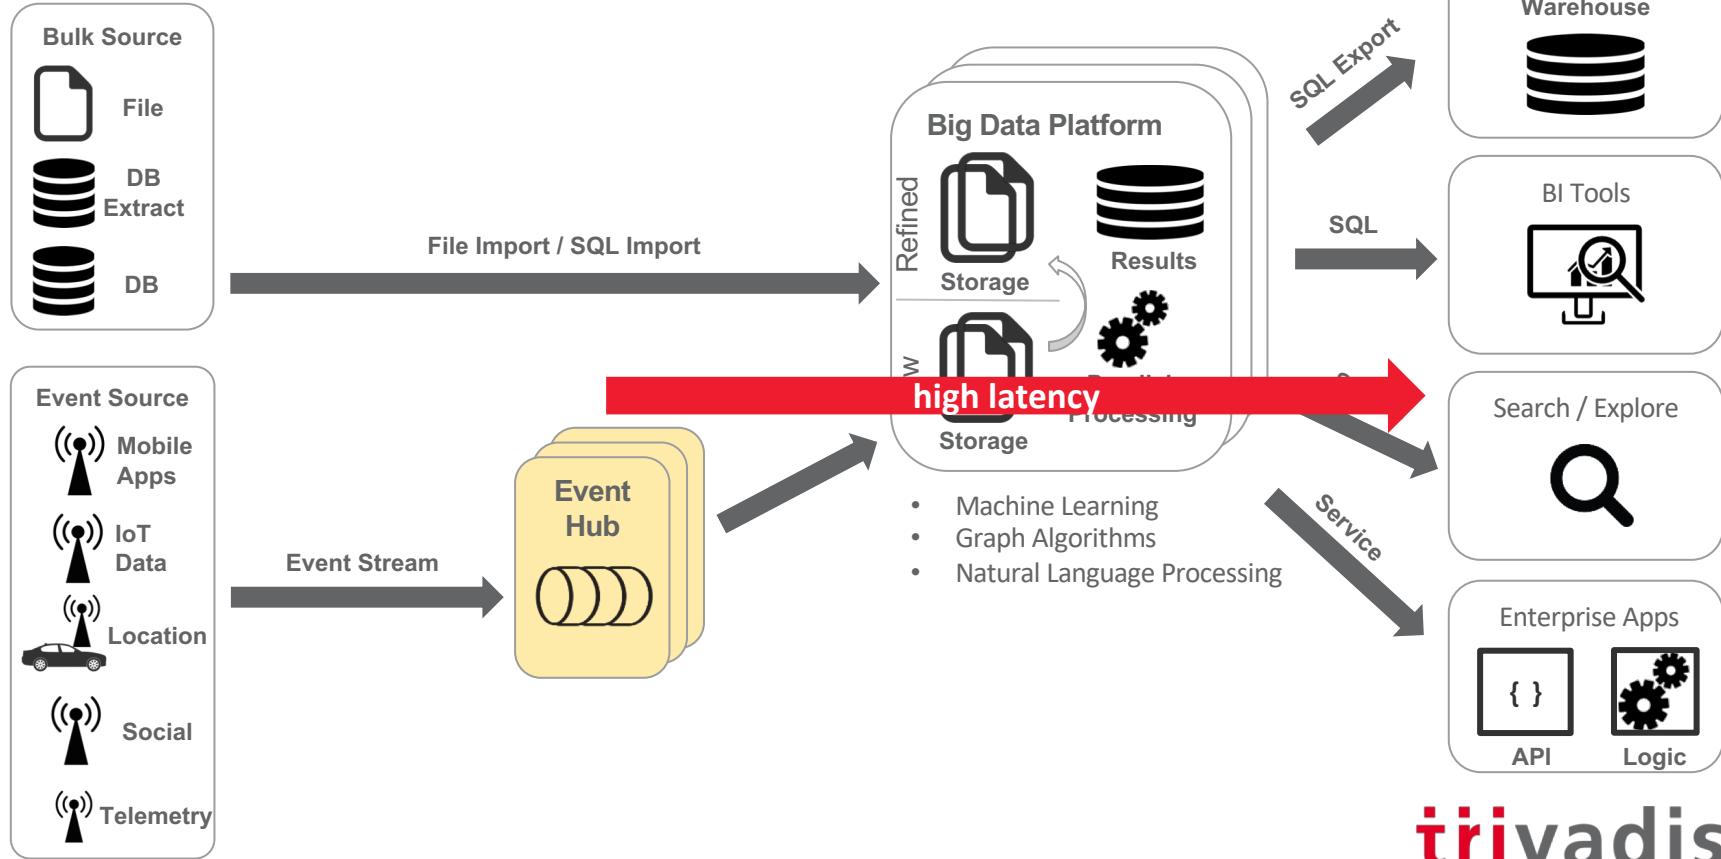
■ Big Data solves Volume and Variety – not Velocity



■ Big Data solves Volume and Variety – not Velocity

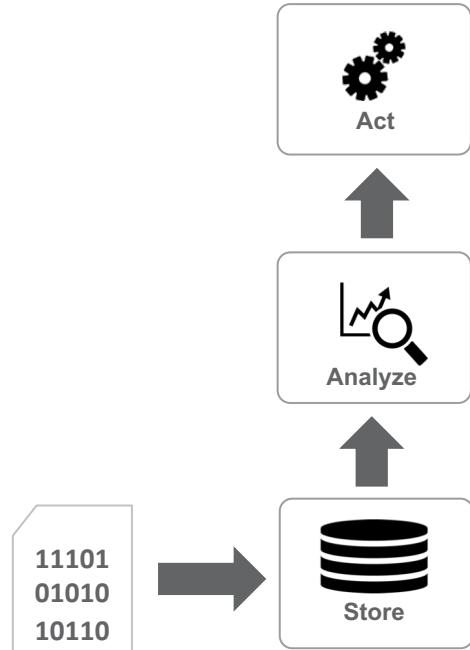


■ Big Data solves Volume and Variety – not Velocity

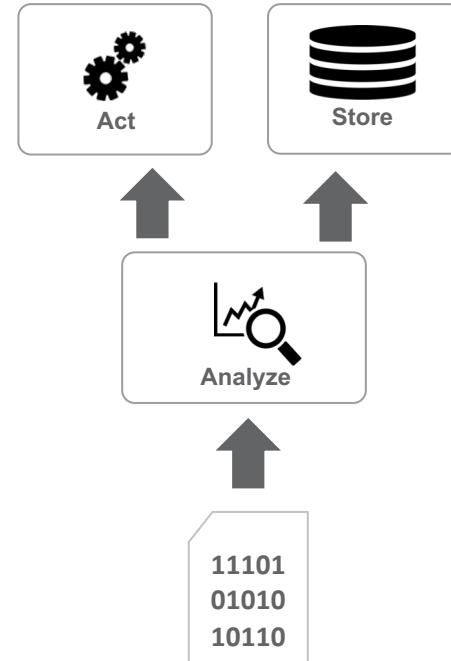


■ "Data at Rest" vs. "Data in Motion"

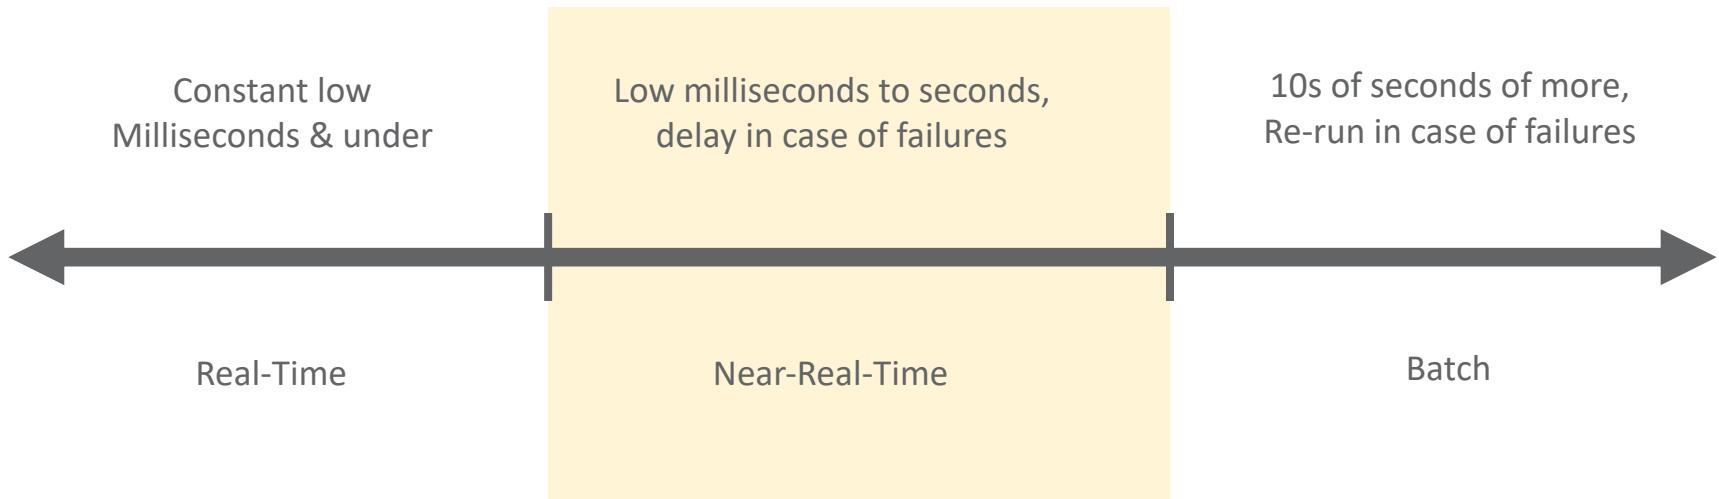
Data at Rest



Data in Motion

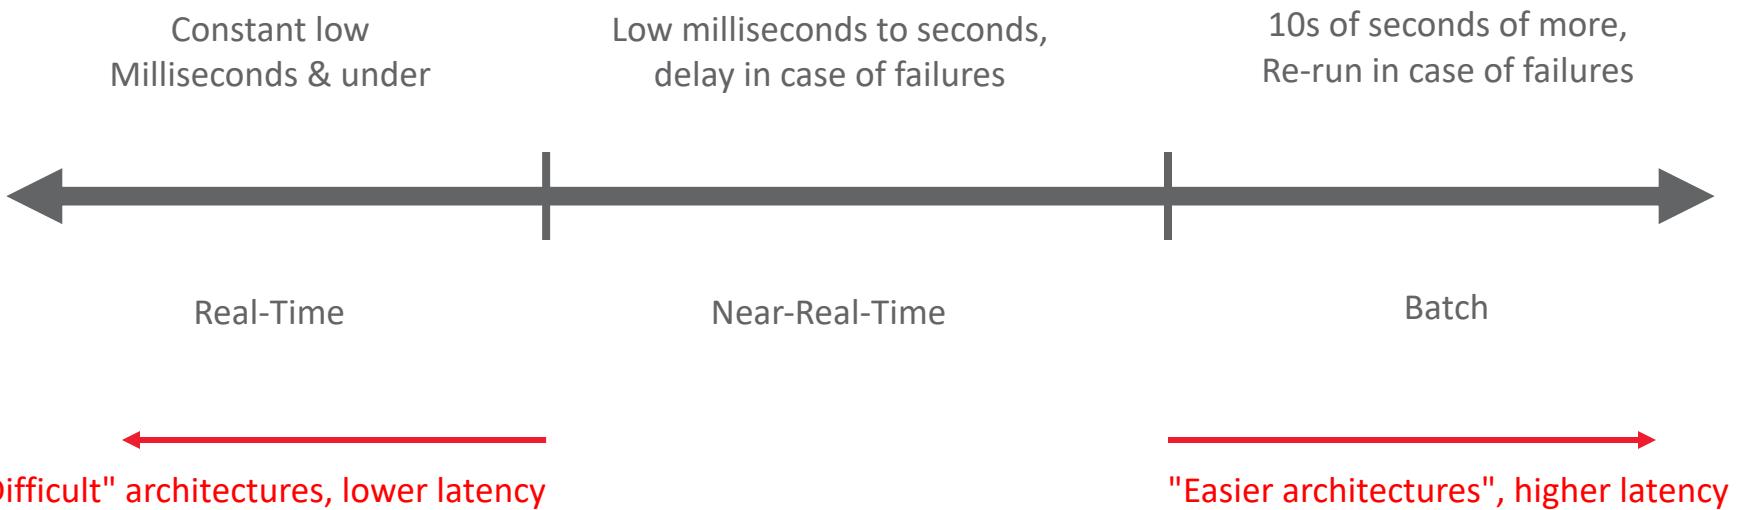


■ When to Stream / When not?

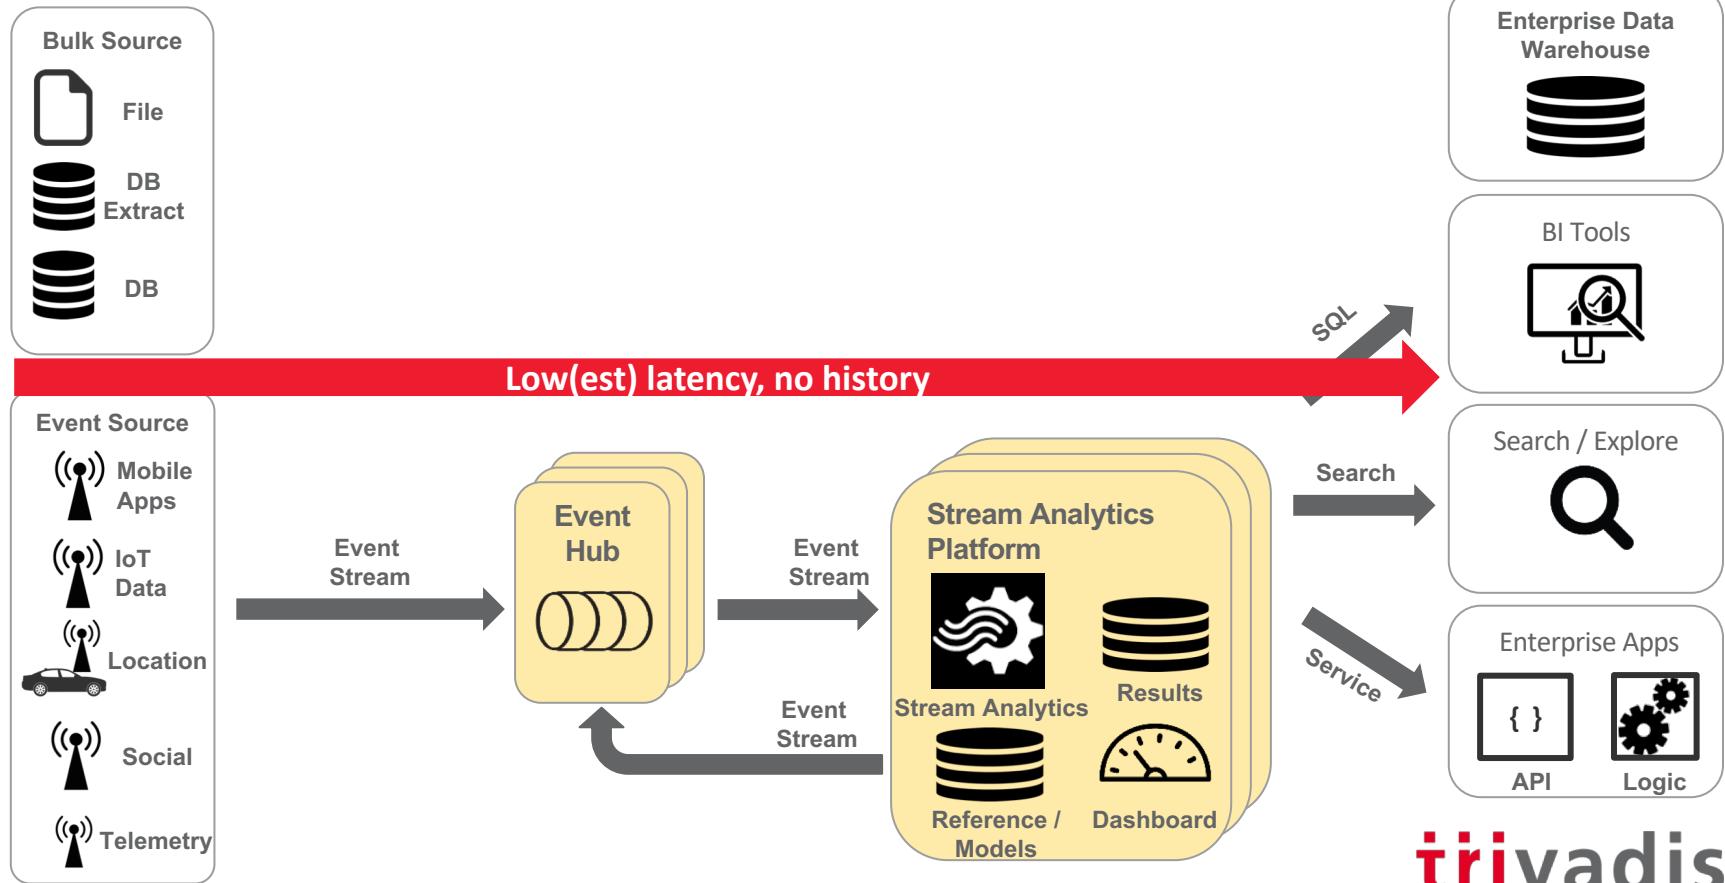


Source: adapted from Cloudera

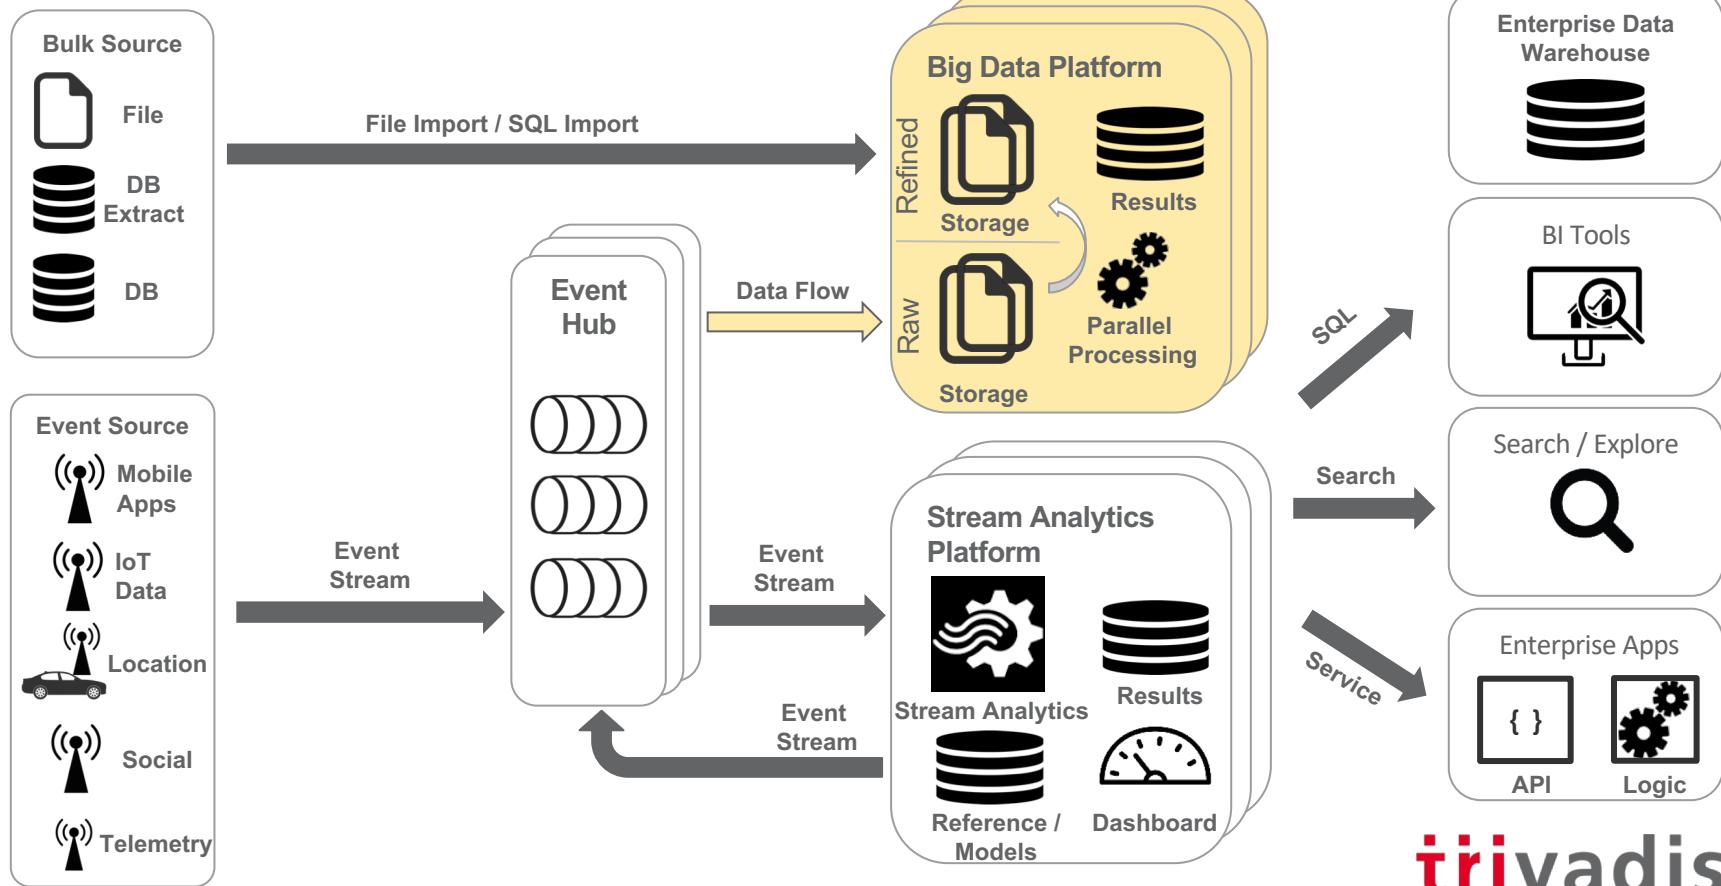
"No free lunch"



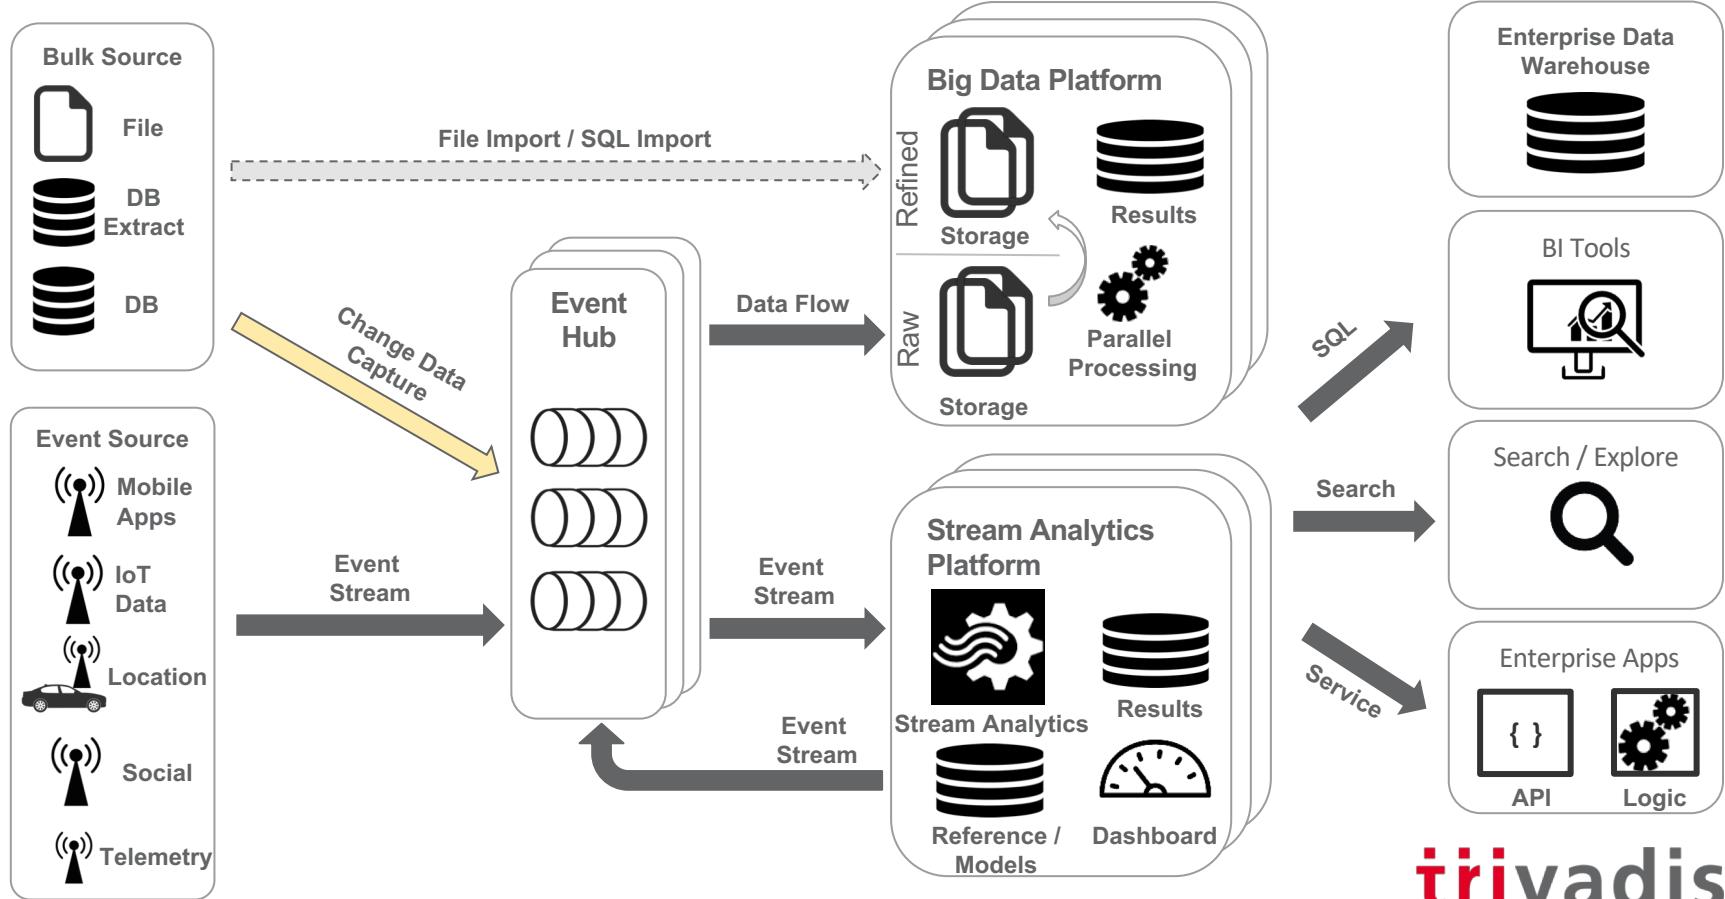
■ Stream Processing Architecture solves Velocity



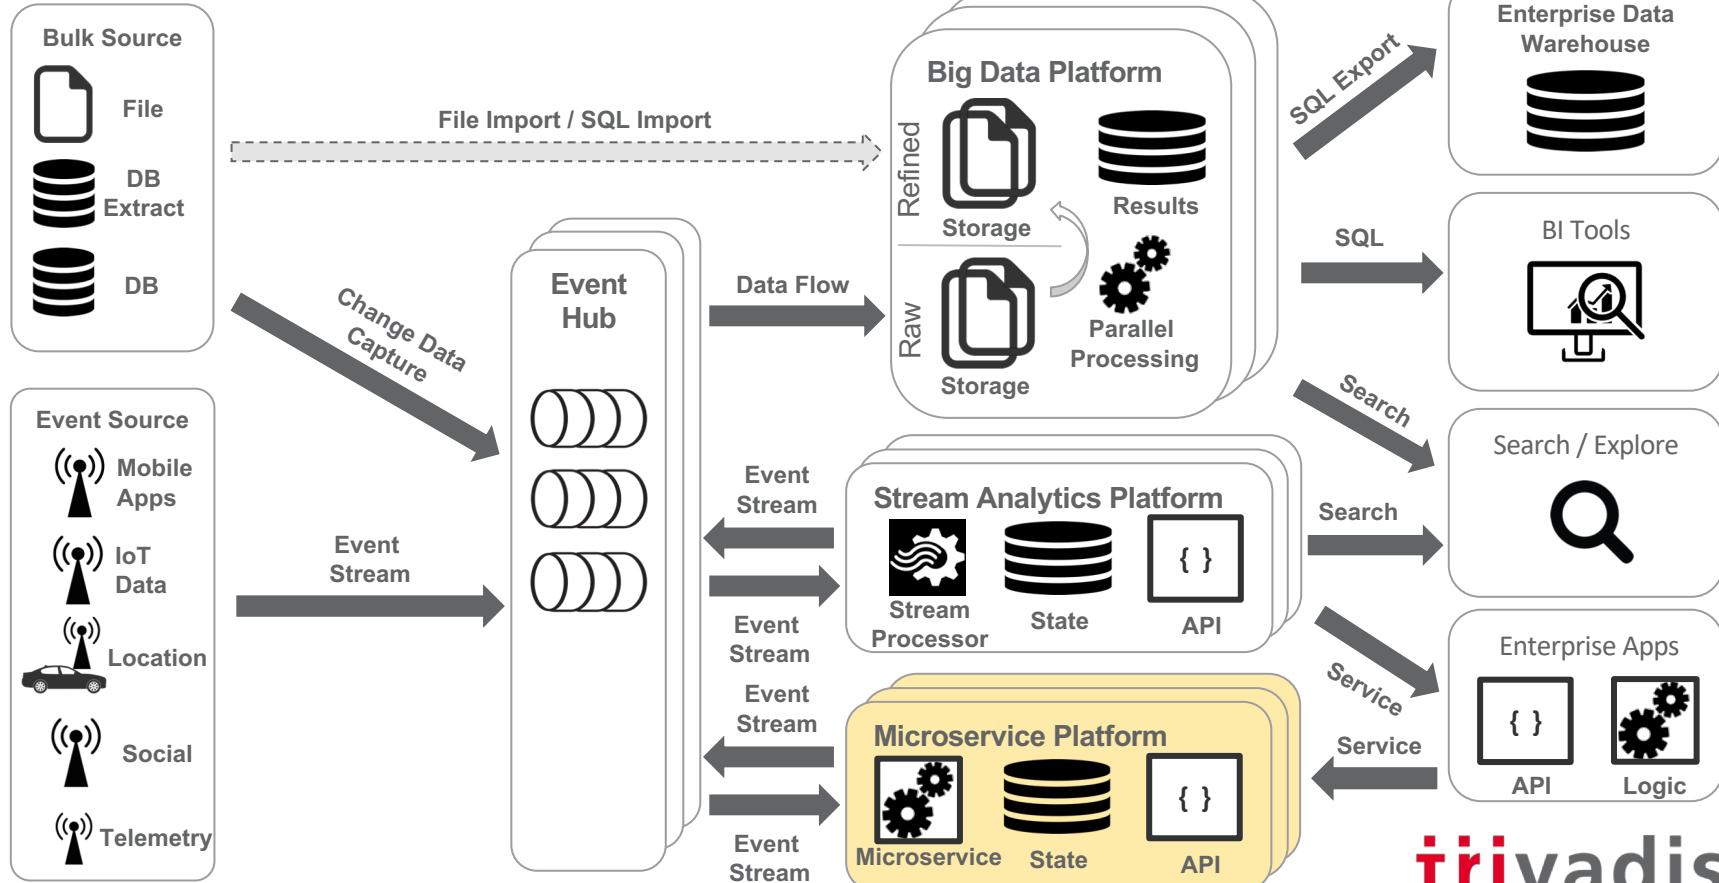
■ Big Data for all historical data analysis



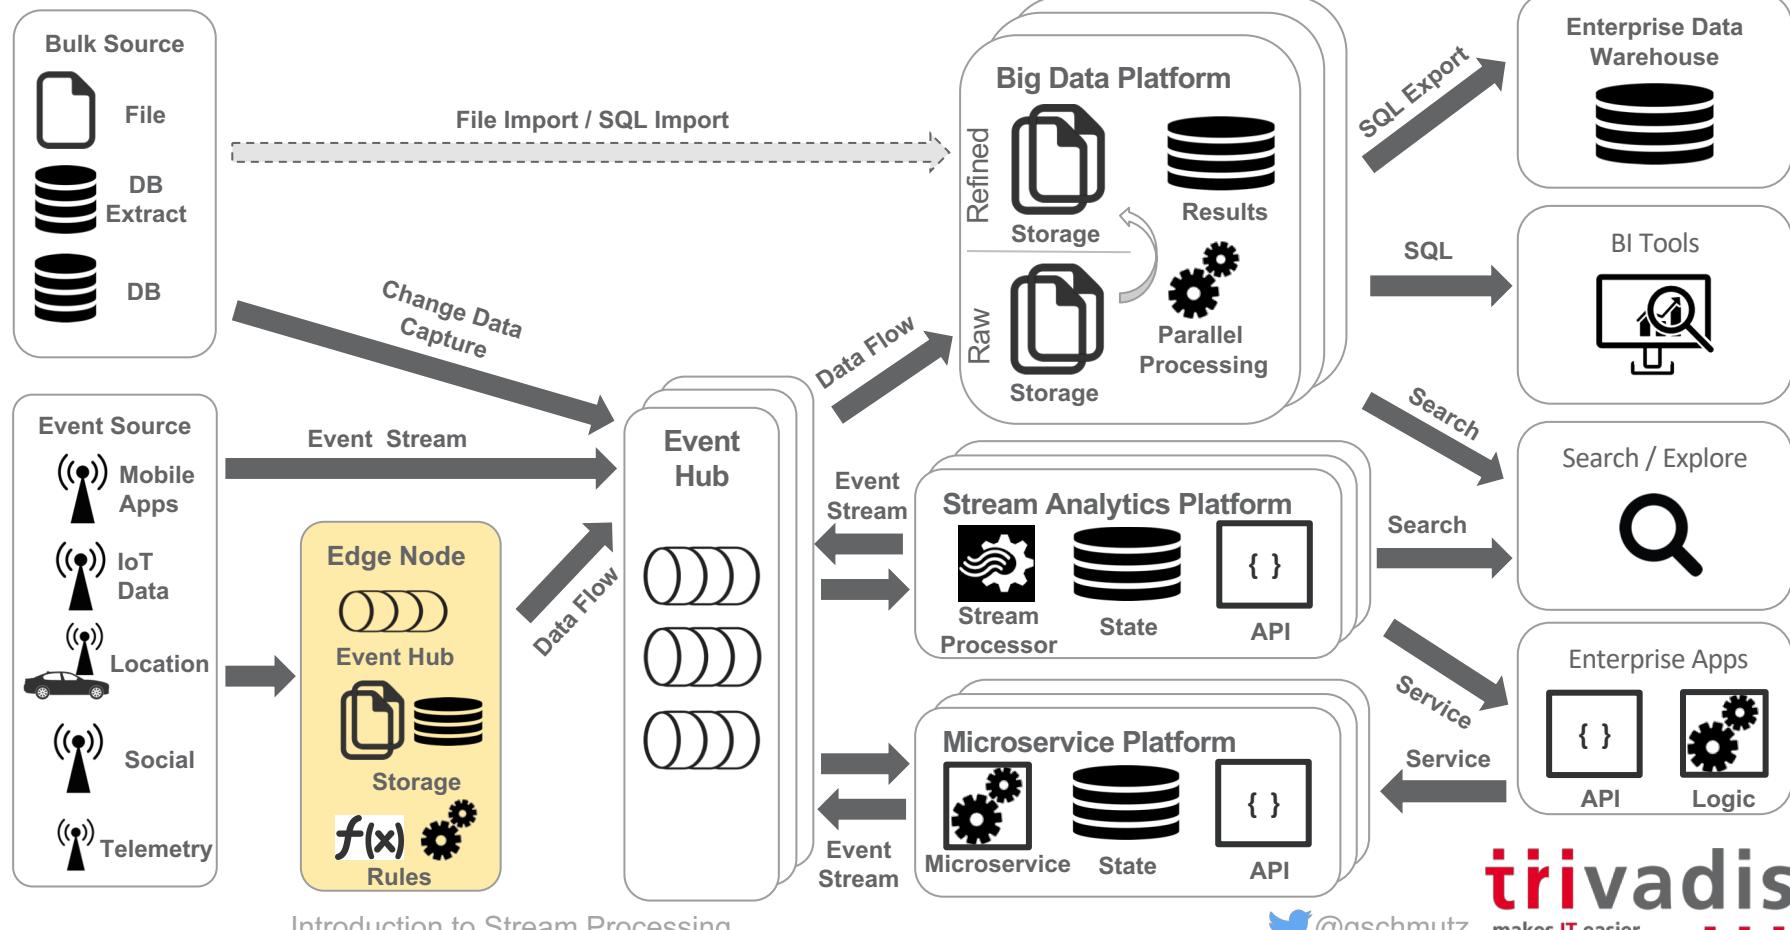
■ Integrate existing systems with lower latency through CDC



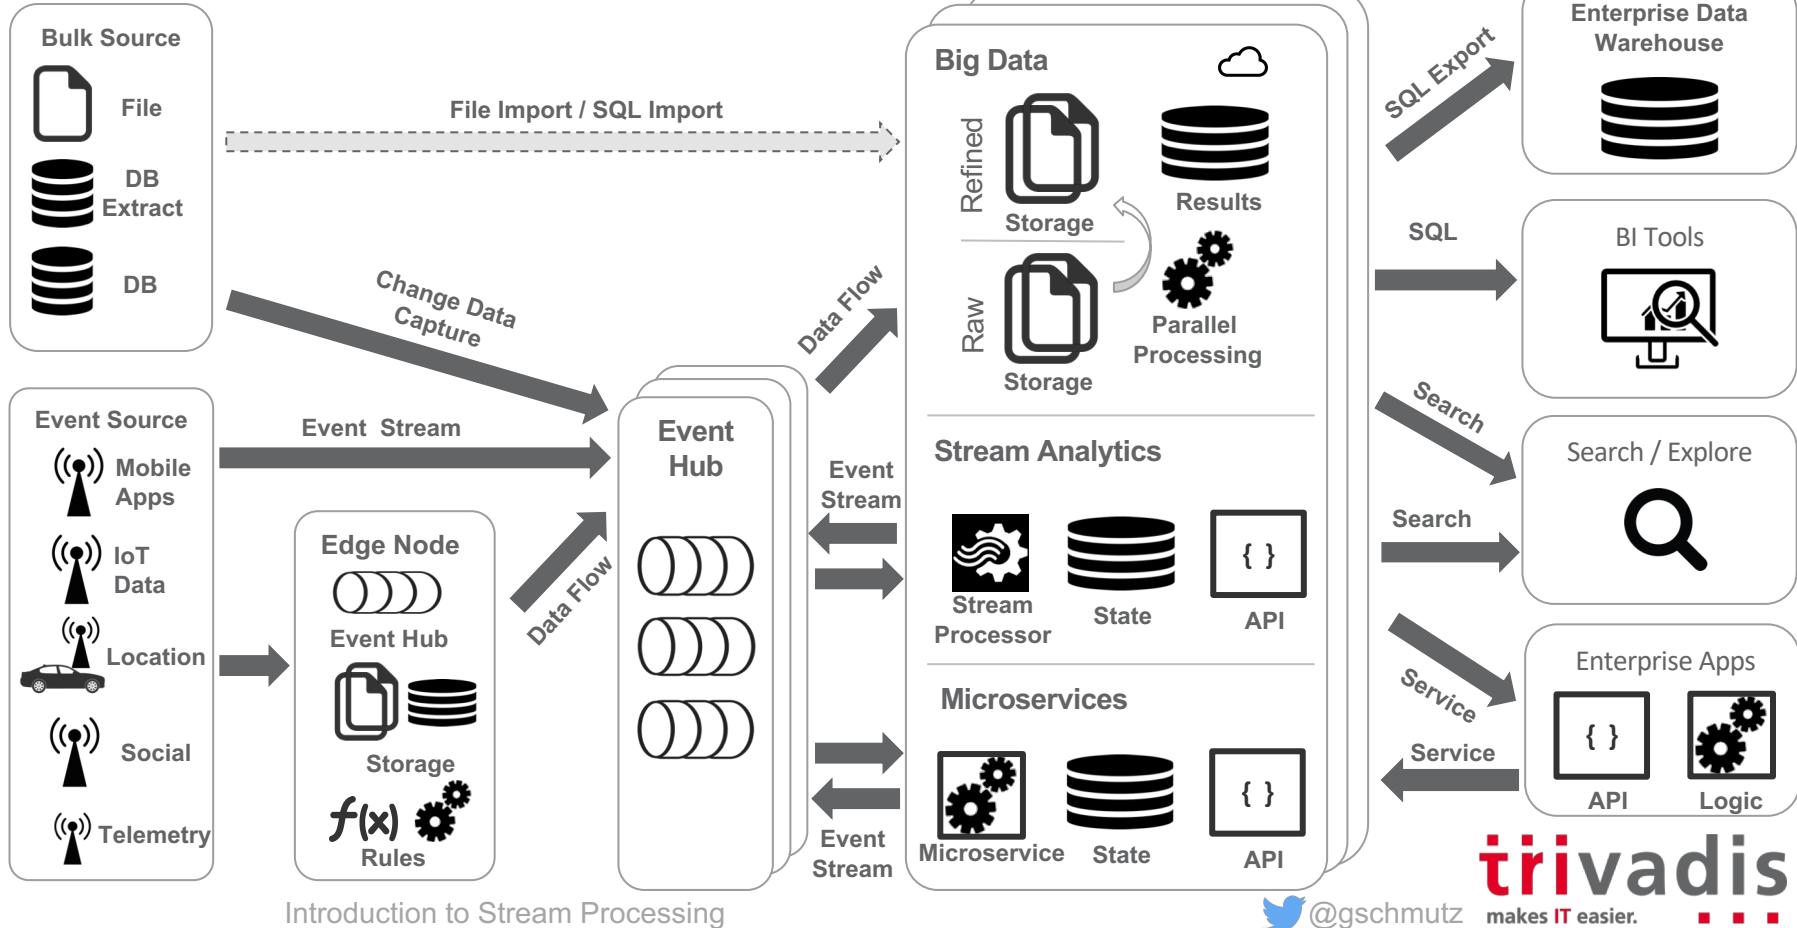
New systems participate in event-oriented fashion



■ Edge computing allows processing close to data sources



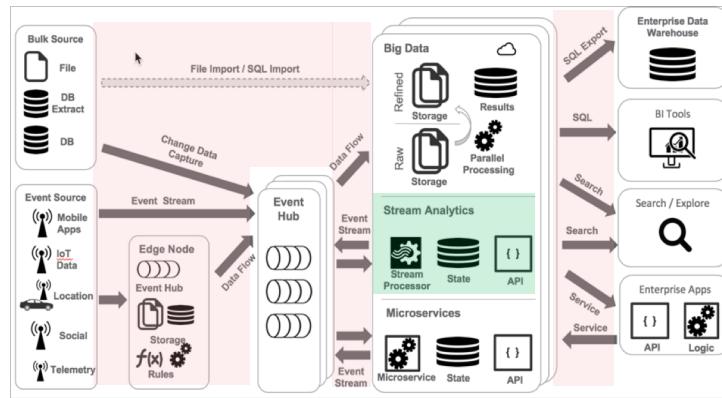
■ Unified Architecture for Modern Data Analytics Solutions



Two Types of Stream Processing (by Gartner)

Stream Data Integration

- focuses on the ingestion and processing of data sources targeting real-time extract-transform-load (ETL) and data integration use cases
- filter and enrich the data



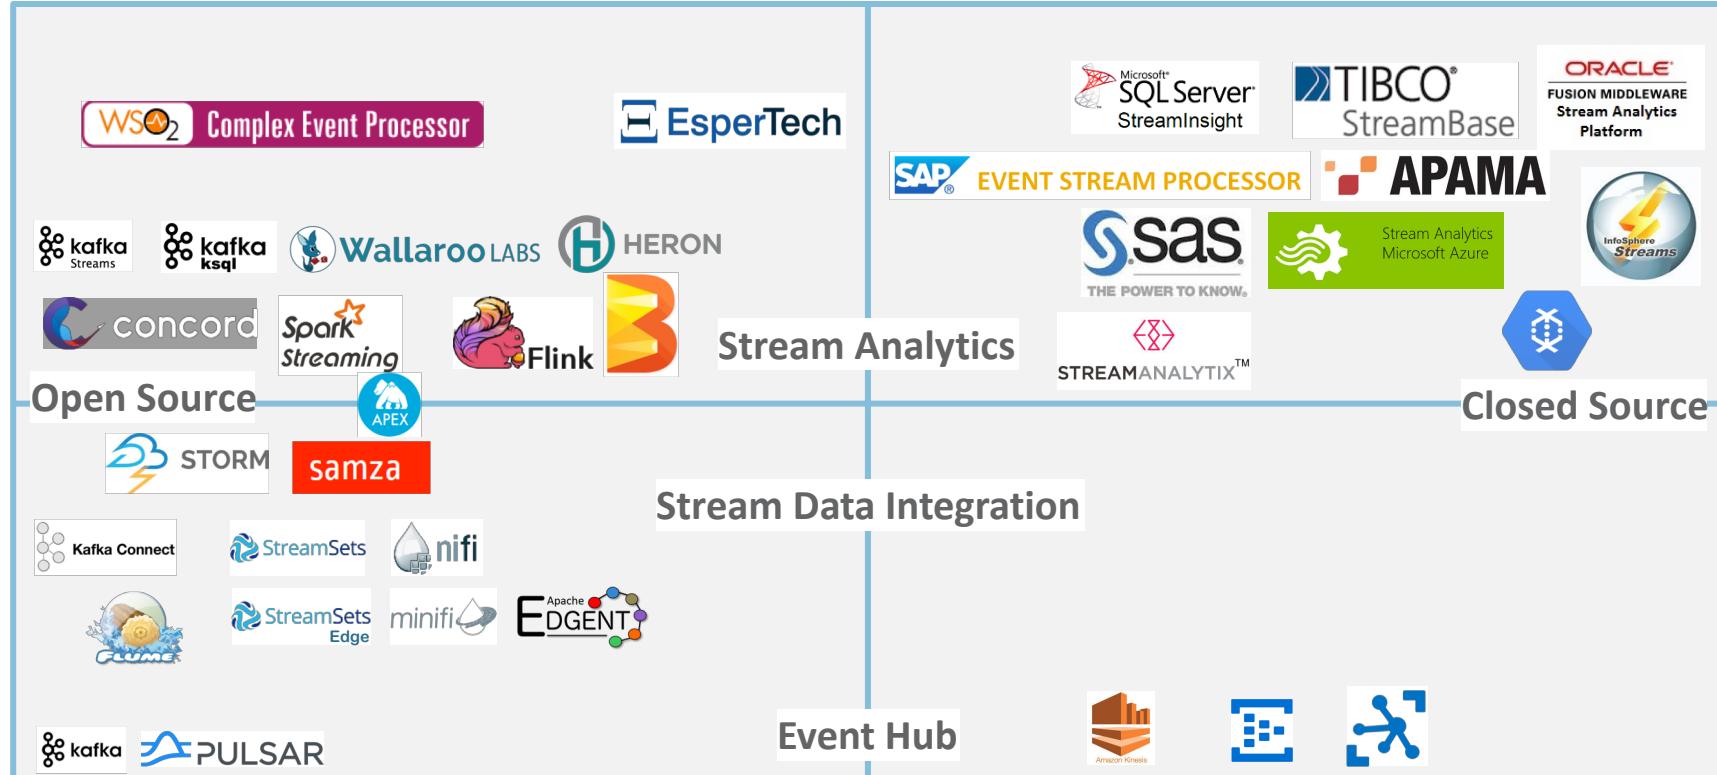
Stream Analytics

- targets analytics use cases
- calculating aggregates and detecting patterns to generate higher-level, more relevant summary information (complex events)
- Complex events may signify threats or opportunities that require a response from the business

Gartner: Market Guide for Event Stream Processing, Nick Heudecker, W. Roy Schulte

Introduction to Stream Processing

■ Stream Processing & Analytics Ecosystem



Source: adapted from Tibco

■ Stream vs. Table / Static

Stream

“History”

an unbounded sequence of structured data (“facts”)

Facts in a stream are **immutable**

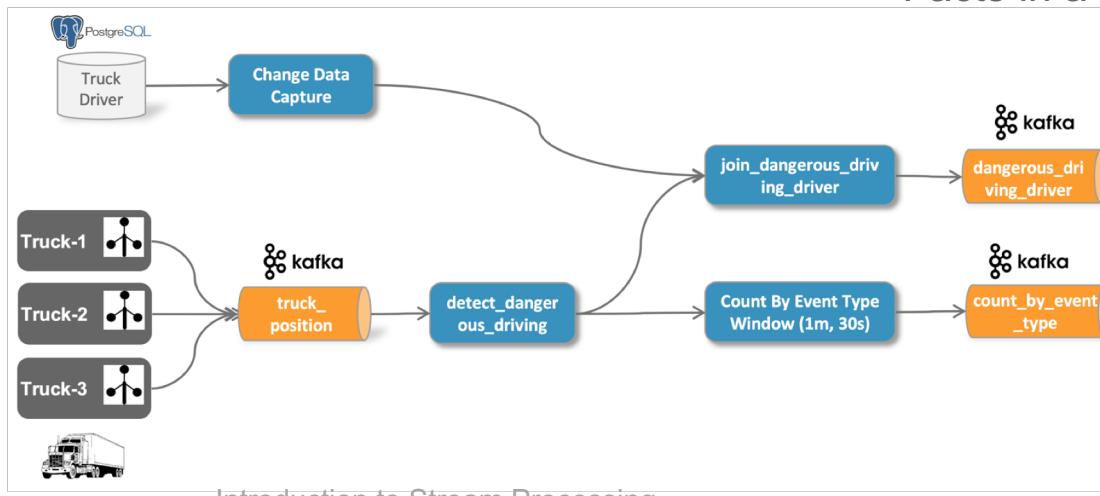
Table / Static

“State”

a view of a stream, or another table, and represents a collection of evolving facts

Latest value for each key in a stream

Facts in a table are **mutable**



Important Capabilities for Stream Processing

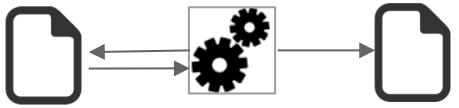
■ Capabilities: Stream Data Integration vs. Stream Analytics

	Stream Data Integration	Stream Analytics
Support for Various Data Sources	yes	partial
Streaming ETL (Transformation/Format Translation, Routing, Validation)	yes	partial
Micro-Batching	yes	partial
Event-at-a-time	yes	yes
Delivery Guarantees	yes	yes
API : GUI-Based API / Declarative API / Programmatic	yes	yes
API: Streaming SQL	-	yes
Event Time vs. Ingestion / Processing Time	-	yes
Windowing	-	yes
Stream-to-Static Joins (Lookup/Enrichment)	partial	yes
Stream-to-Stream Joins	-	yes
State Management	-	yes
Queryable State (aka Interactive Queries)	-	yes
Event Pattern Detection	-	Yes

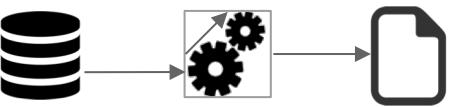


Integrating Data Sources

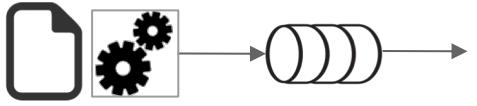
File Polling



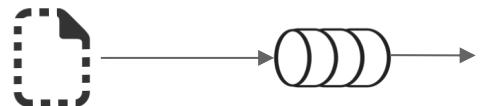
SQL Polling



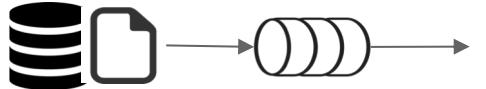
File Stream (File Tailing)



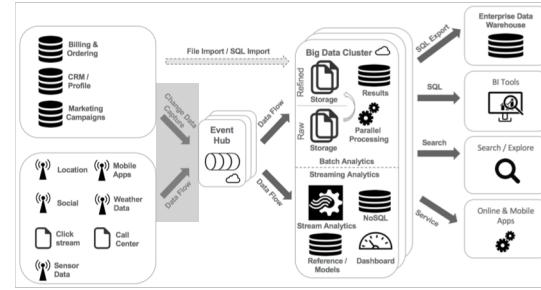
File Stream (Appender)



Change Data Capture (CDC)

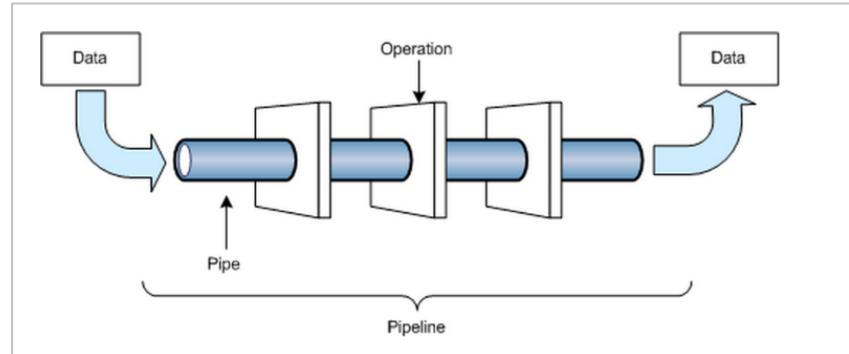


Sensor Stream

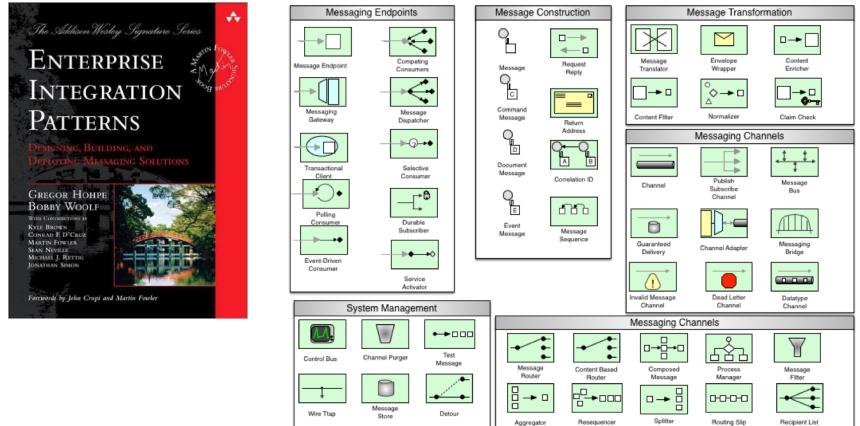
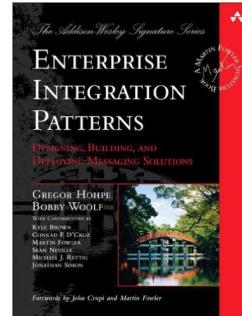


Streaming ETL

- Flow-based "programming"
- Ingest Data from various sources
- Extract – Transform – Load
- High-Throughput, straight-through data flows
- Data Lineage
- Batch- or Stream-Processing
- Visual coding with flow editor
- Event Stream Processing (ESP) **but not** Complex Event Processing (CEP)



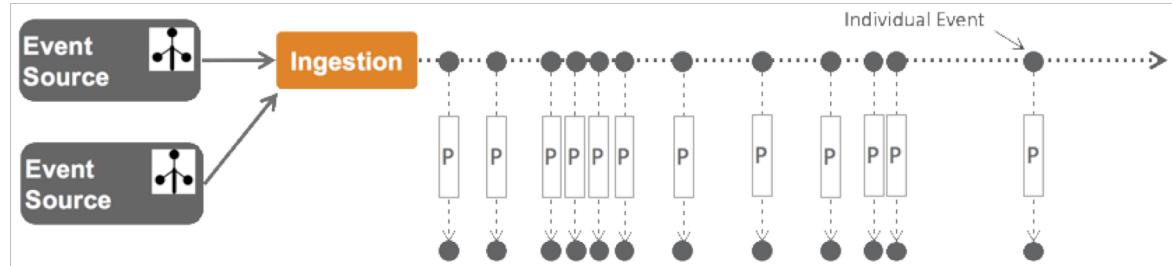
Source: Confluent



■ Event-at-a-time vs. Micro Batch

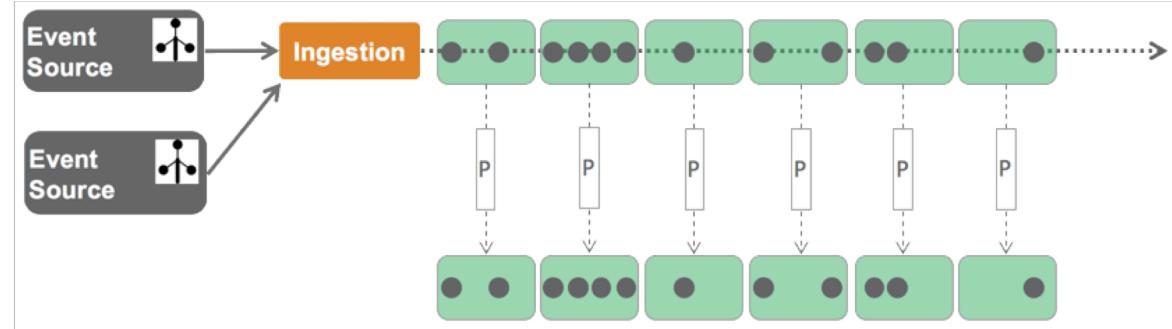
Event-at-a-time Processing

- Events processed as they arrive
- low-latency
- fault tolerance expensive



Micro-Batch Processing

- Splits incoming stream in small batches
- Fault tolerance easier
- Better throughput



■ Delivery Guarantees



At most once (fire-and-forget)

message is sent, but the sender **doesn't care if it's received or lost**. it is *the easiest and most performant behavior to support*.



At least once

retransmission of a message will occur until an acknowledgment is received. Since a delayed acknowledgment from a receiver could be in flight when the sender retransmits, the message may be received **one or more times**.



Exactly once

ensures that a message is received **once and only once**, and is never lost and never repeated. The system must implement whatever mechanisms are required to ensure that a message is received and processed just once

API

GUI-based / Drag-and-Drop

- A graphical way of designing a pipeline
- Often web-based to improve usability



Declarative

- An streaming engine which can be configured in a declarative way
- JSON, YML

```
"config": {  
  "connector.class":  
    "io.confluent.connect.mqtt.MqttSourceConnector",  
  "tasks.max": "1",  
  "mqtt.server.uri": "tcp://mosquitto-1:1883",  
  "mqtt.topics": "truck/+position",  
  "kafka.topic": "truck_position",  
  ...}
```

Programmatic

- Low-level (class) or high-level fluent API
- Higher order function as operators (filter, mapWithState ...)

```
val filteredDf = truckPosDf.  
  where("eventType != 'Normal'")
```

Streaming SQL

- use a stream in a FROM clause
- Extensions supporting windowing, pattern matching, spatial, Operators

```
SELECT * FROM truck_position_s  
WHERE eventType != 'Normal'
```

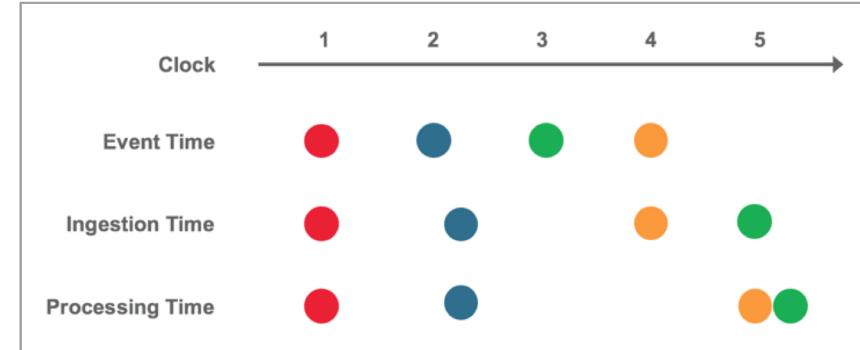
■ Event Time vs. Ingestion / Processing Time

Event time

the time at which events actually *occurred*

Ingestion time / Processing Time

the time at which events are *ingested into / processed by* the system



Not all use cases care about event times but many do

Examples

- characterizing user behavior over time
- most billing applications
- anomaly detection

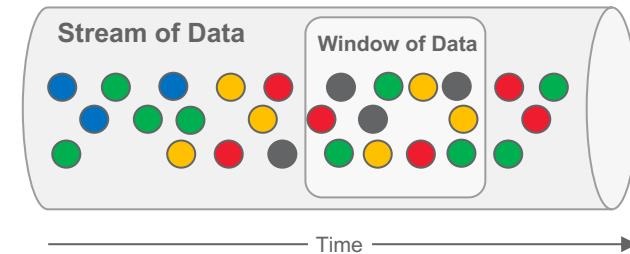
■ Windowing

Computations over events done using
windows of data

Due to size and never-ending nature of it,
it's not feasible to keep entire stream of
data in memory

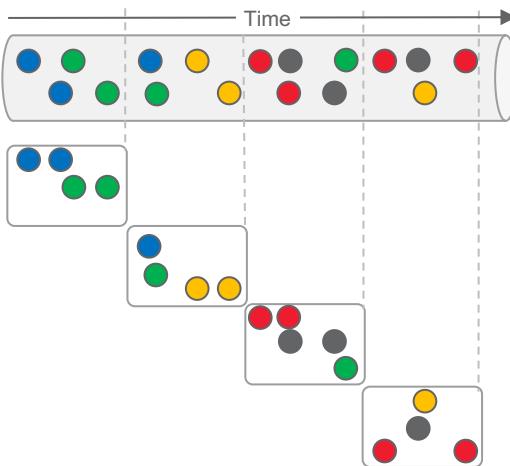
A window of data represents a certain
amount of data where we can perform
computations on

Windows give the power to keep a
working memory and look back at recent
data efficiently

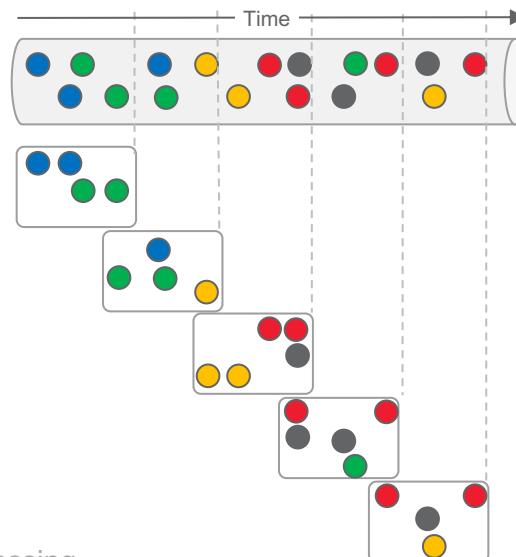


■ Windowing

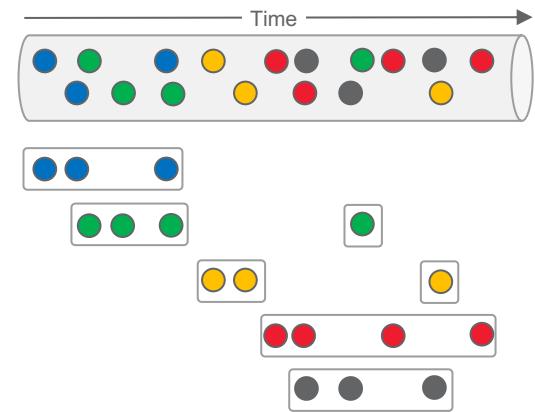
Fixed Window (aka Tumbling Window) - eviction policy always based on the window being full and trigger policy based on either the count of items in the window or time



Sliding Window (aka Hopping Window) - uses eviction and trigger policies that are based on time: *window length* and *sliding interval length*



Session Window – composed of sequences of temporarily related events terminated by a gap of inactivity greater than some timeout

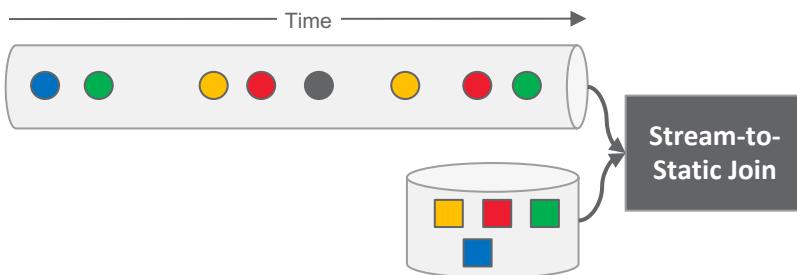


Joining – Stream-to-Static and Stream-to-Stream

Challenges of joining streams

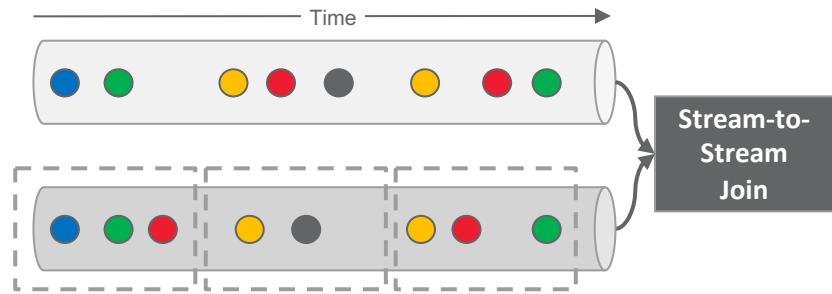
1. Data streams need to be aligned as they come because they have different timestamps
2. since streams are never-ending, the joins must be limited; otherwise join will never end
3. join needs to produce results continuously as there is no end to the data

Stream-to-Static (Table) Join

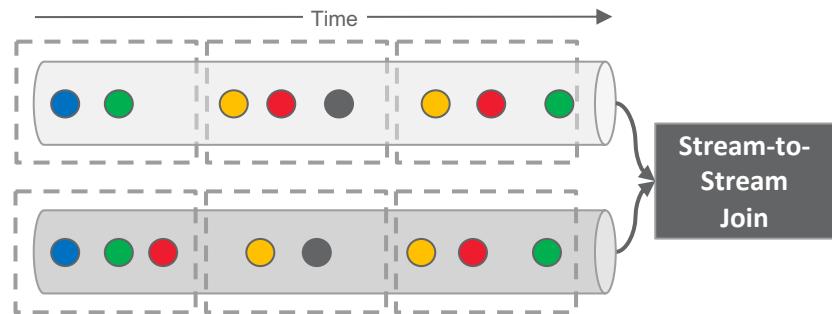


Introduction to Stream Processing

Stream-to-Stream Join (one window join)



Stream-to-Stream Join (two window join)



■ State Management

Necessary if stream processing use case is dependent on previously seen data or external data

Windowing, Joining and Pattern Detection use State Management behind the scenes

State Management services can be made available for custom state handling logic

State needs to be managed as close to the stream processor as possible

Options for State Management

In-Memory

Local,
Embedded
Store

Replicated,
Distributed
Store



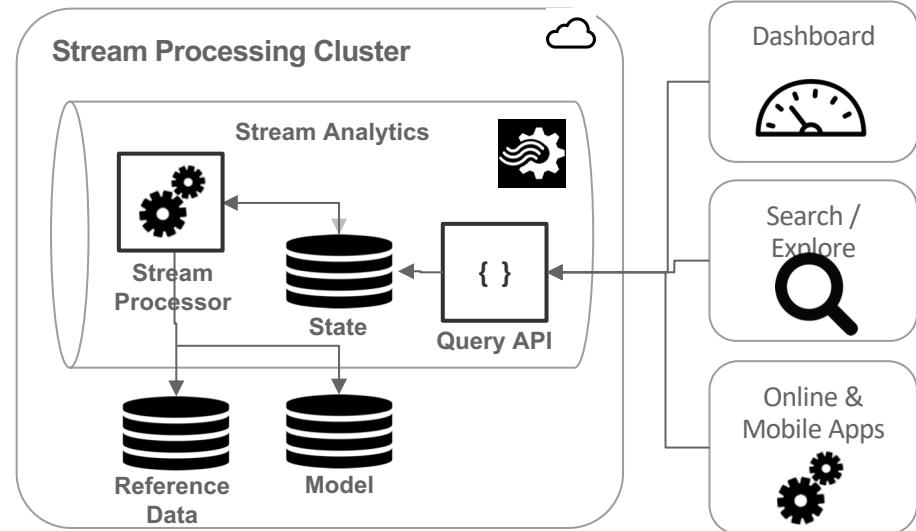
How does it handle failures? If a machine crashes and the/some state is lost?

■ Queryable State (aka. Interactive Queries)

Exposes the state managed by the Stream Analytics solution to the outside world

Allows an application to query the managed state, i.e. to visualize it

For some scenarios, Queryable State can eliminate the need for an external database to keep results



■ Event Pattern Detection

- Streaming Data often contains interesting patterns that only emerge as new streaming data arrives, e.g.
 - Absence Pattern: event A not followed by event B within time window
 - Sequence Pattern: event A followed by event B followed by event C
 - Increasing Pattern: up trend of a value of a certain attribute
 - Decreasing Pattern: down trend of a value of a certain attribute
 - ...
- Pattern operators allow developers to define complex relationships between streaming events

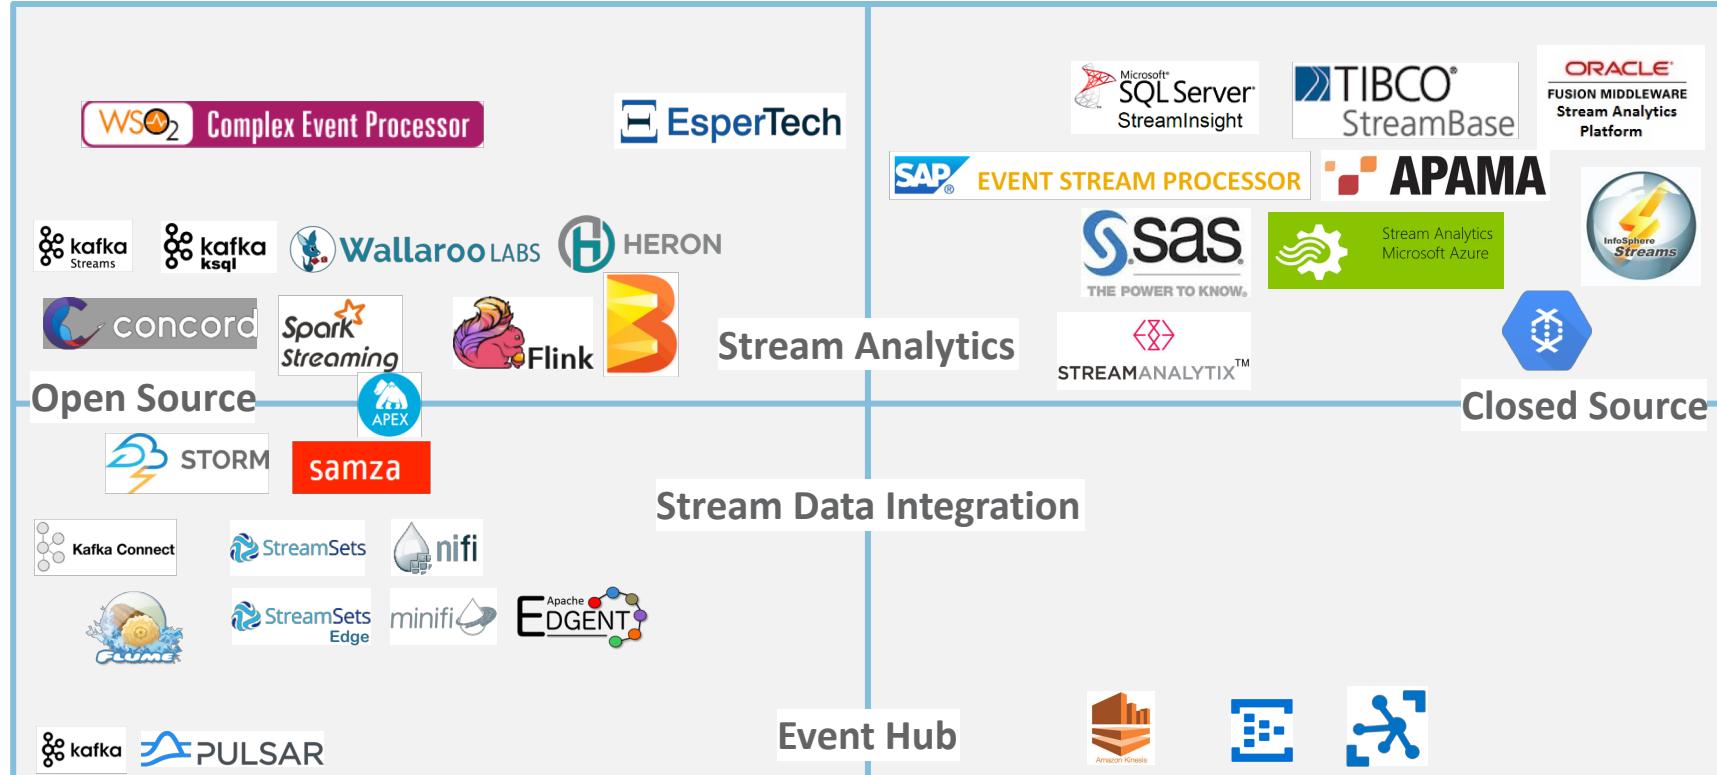
■ Capabilities: Stream Data Integration vs. Stream Analytics

	Stream Data Integration	Stream Analytics
Support for Various Data Sources	yes	-
Streaming ETL (Transformation/Format Translation, Routing, Validation)	yes	partial
Micro-Batching	yes	partial
Event-at-a-time	yes	yes
Delivery Guarantees	yes	yes
API : GUI-Based API / Declarative API / Programmatic	yes	yes
API: Streaming SQL	-	yes
Event Time vs. Ingestion / Processing Time	-	yes
Windowing	-	yes
Stream-to-Static Joins (Lookup/Enrichment)	partial	yes
Stream-to-Stream Joins	-	yes
State Management	-	yes
Queryable State (aka Interactive Queries)	-	yes
Event Pattern Detection	-	Yes



Implementing Stream Processing Solutions

■ Stream Processing & Analytics Ecosystem

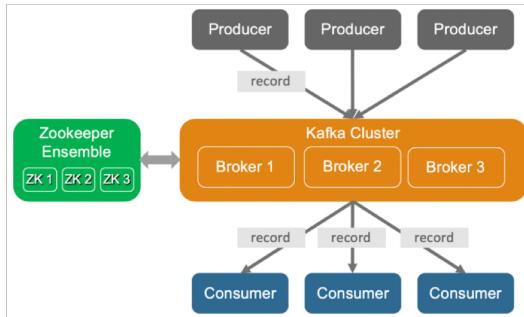


Source: adapted from Tibco

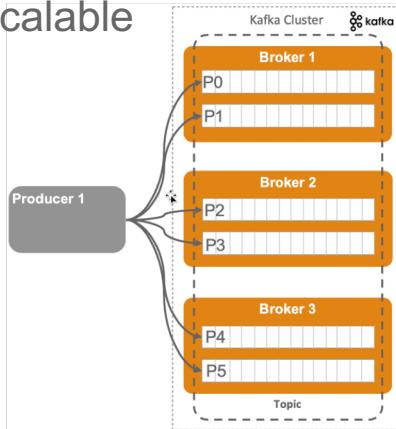
■ Event Hub: Apache Kafka



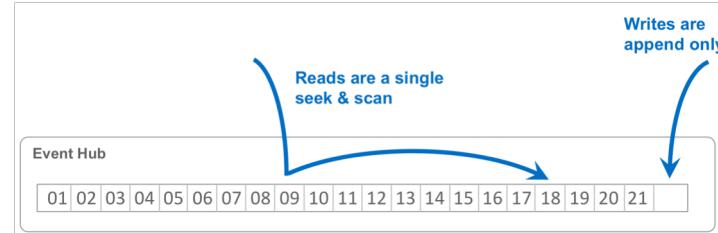
Highly available, Pub/Sub infrastructure



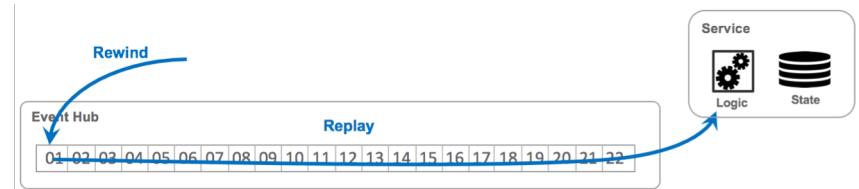
Highly Scalable



Distributed Log at the Core



Logs do not (necessarily) forget

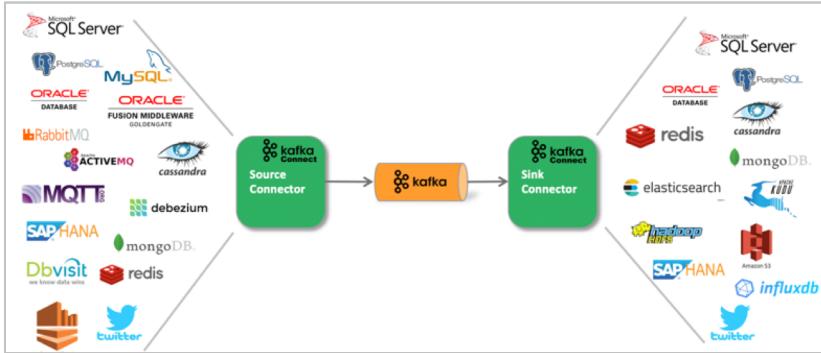


- Never
- Time (TTL) or Size-based
- Log-Compacted based

■ Stream Data Integration: Kafka Connect



declarative style data flows
simplicity - “simple things done simple”
very well integrated with Kafka –
framework is part of Kafka
Single Message Transforms (SMT)



Many connectors available
Implement custom connectors using Java
Supported by Confluent

```
#!/bin/bash
curl -X "POST"
  "http://192.168.69.138:8083/connectors" \
  -H "Content-Type: application/json" \
  -d '${
"name": "mqtt-source",
"config": {
  "connector.class": "...MqttSourceConnector",
  "tasks.max": "1",
  "name": "mqtt-source",
  "mqtt.server.uri": "tcp://mosquitto:1883",
  "mqtt.topics": "truck/+position",
  "kafka.topic": "truck_position",
}
}'
```

■ Stream Data Integration: StreamSets

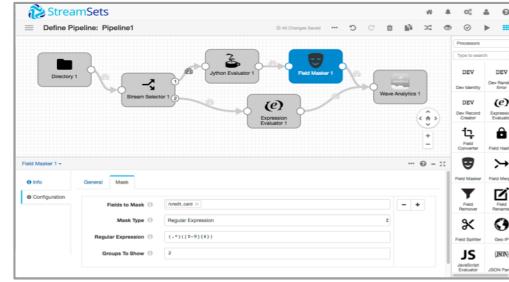
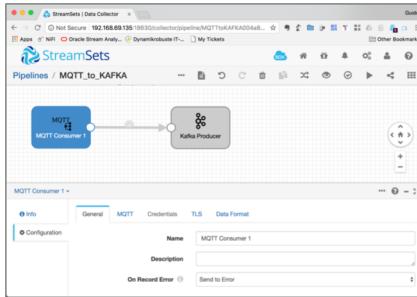


Continuous open source, intent-driven, big data ingest

Visible, record-oriented approach fixes combinatorial explosion

Both stream and batch processing

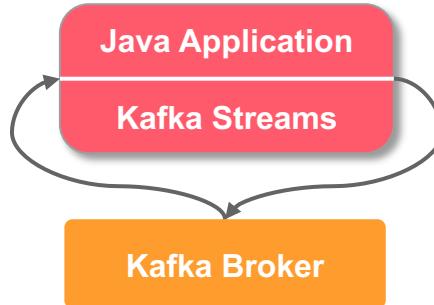
- Standalone, Spark cluster, MapReduce cluster



■ Streaming Analytics: Kafka Streams



Designed as a simple and lightweight library in Apache Kafka
no other dependencies than Kafka
Supports fault-tolerant local state
Supports Windowing (Fixed, Sliding and Session) and Stream-Stream / Stream-Table Joins



Millisecond processing latency, no micro-batching
At-least-once and exactly-once processing guarantees

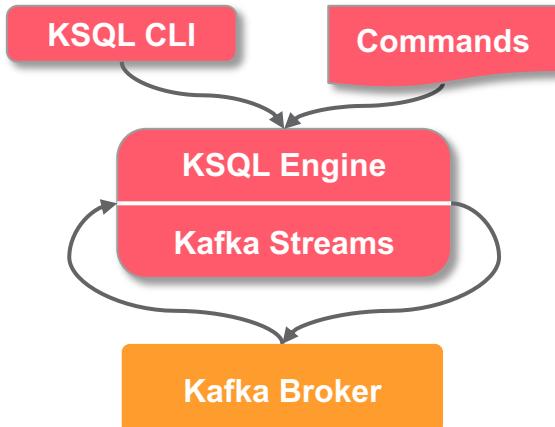
```
KTable<Integer, Customer> customers =  
    builder.stream("customer");  
  
KStream<Integer, Order> orders =  
    builder.stream("order");  
  
KStream<Integer, String> enriched =  
    orders.leftJoin(customers, ...);  
  
joined.to("orderEnriched");
```

■ Streaming Analytics: KSQL

Stream Processing with zero coding
using SQL-like language

Built on top of Kafka Streams

Interactive (CLI) and headless (command
file)



STREAM and TABLE as first-class
citizens

- STREAM = data in motion
- TABLE = collected state of a stream

```
ksql> CREATE STREAM order_s \
      WITH (kafka_topic='order', \
            value_format='AVRO');
```

Message

Stream created

```
ksql> SELECT * FROM order_s \
      WHERE address->country = 'Switzerland';
...
```



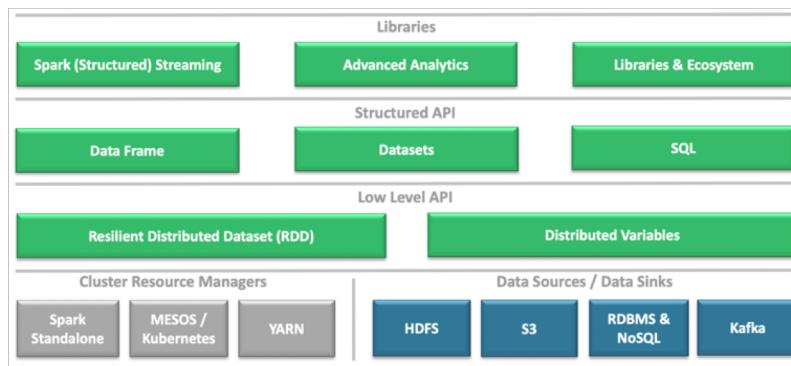
Streaming Analytics: Spark Structured Streaming



2nd generation (1st to be Spark Streaming)

Structured API through DataFrames / Datasets rather than RDDs

Easier code reuse between batch and streaming



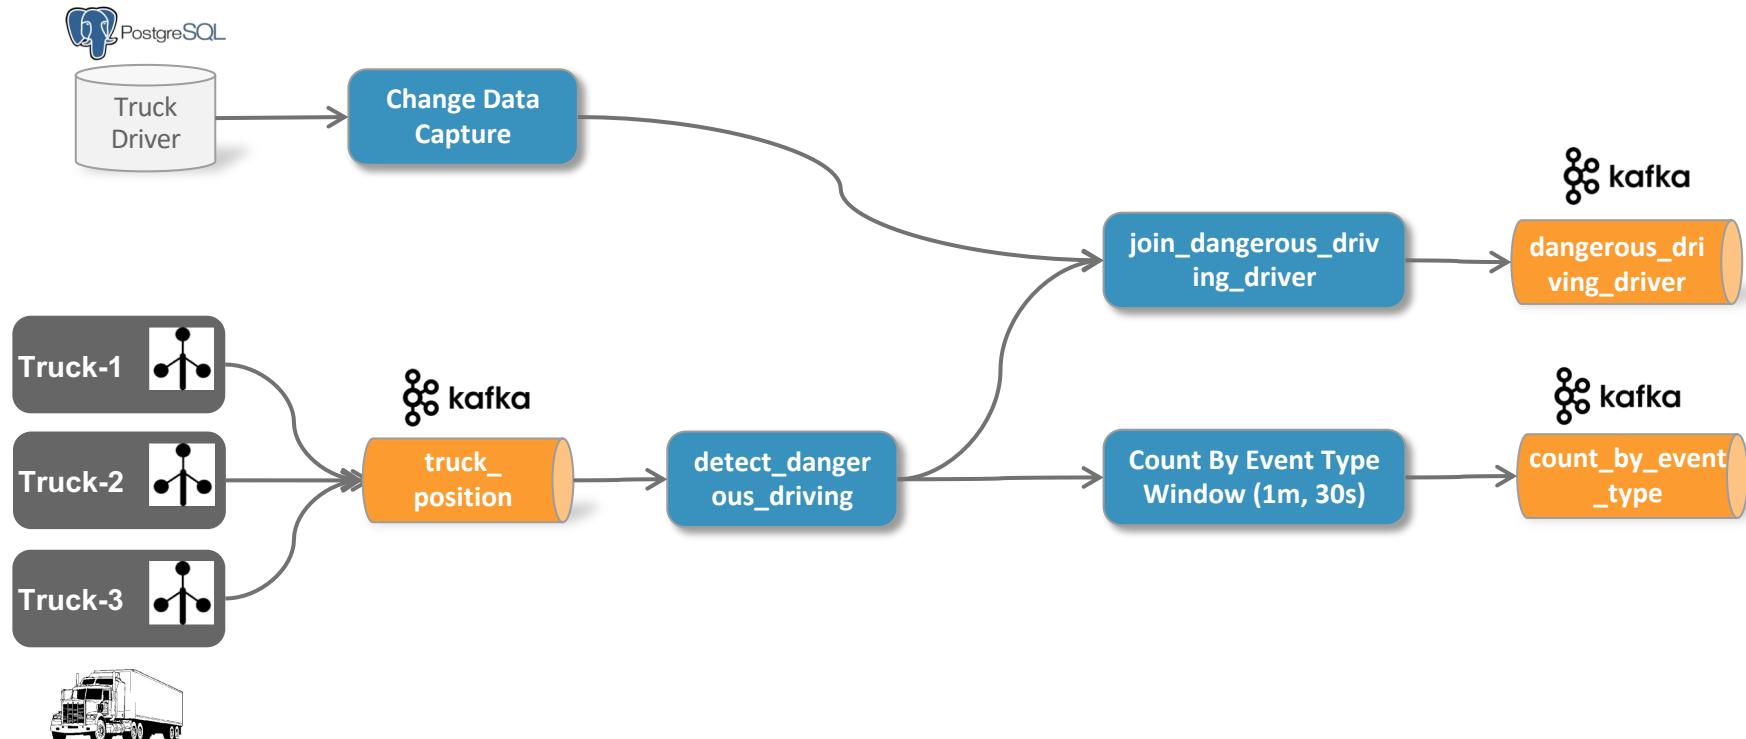
marked production ready in Spark 2.2.0
Support for Java, Scala, Python, R and SQL

```
val orderDF = spark
  .readStream
  .format("kafka")
  .option("kafka.bootstrap.servers",
"broker-1:9092")
  .option("subscribe", "order")
  .load()

val orderFilteredDF =
  orderDF.where(
    "address.country = 'Switzerland'")
```

Demo

■ Sample Use Case



Summary

■ Summary

Stream Processing is the solution for low-latency

Event Hub, Stream Data Integration and Stream Analytics are the main building blocks in your architecture

Kafka is currently the de-facto standard for Event Hub

Various options exists for Stream Data Integration and Stream Analytics

SQL becomes a valid option for implementing Stream Analytics

Still room for improvements (SQL, Event Pattern Detection, Streaming Machine Learning)

Technology on its own won't help you. You need to know how to use it properly.

