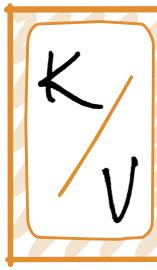


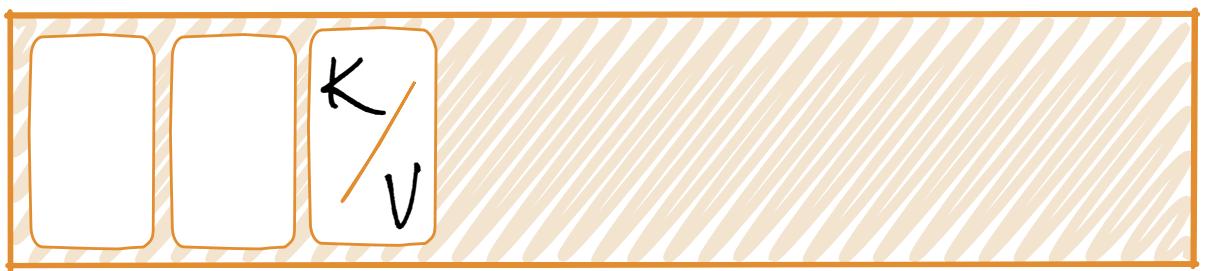
# Architecting Microservices Applications with Instant Analytics

—  
Tim Berglund, Sr. Director, Developer Experience, Confluent  
Rachel Pedreschi, Sr.I Director, Global Field Engineering, Imply Data

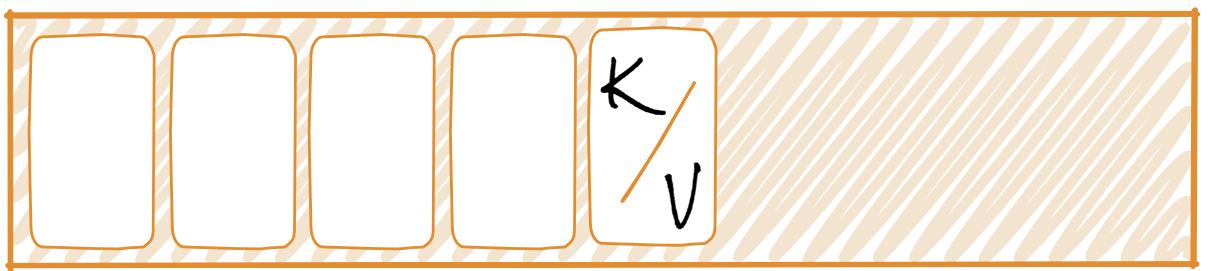
# What the heck is Apache Kafka and Why Should I Care?

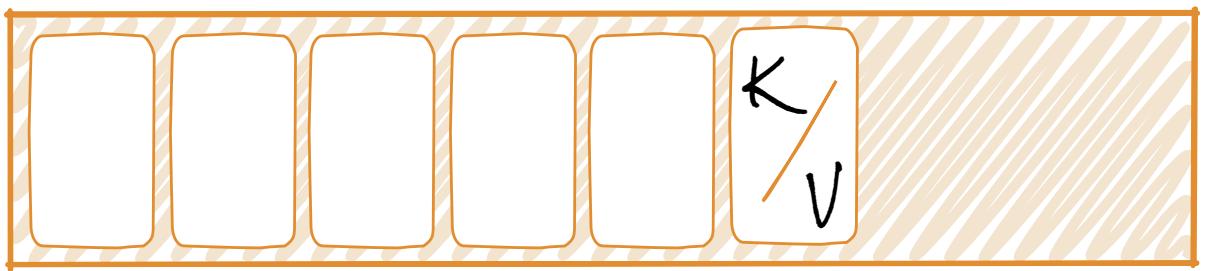


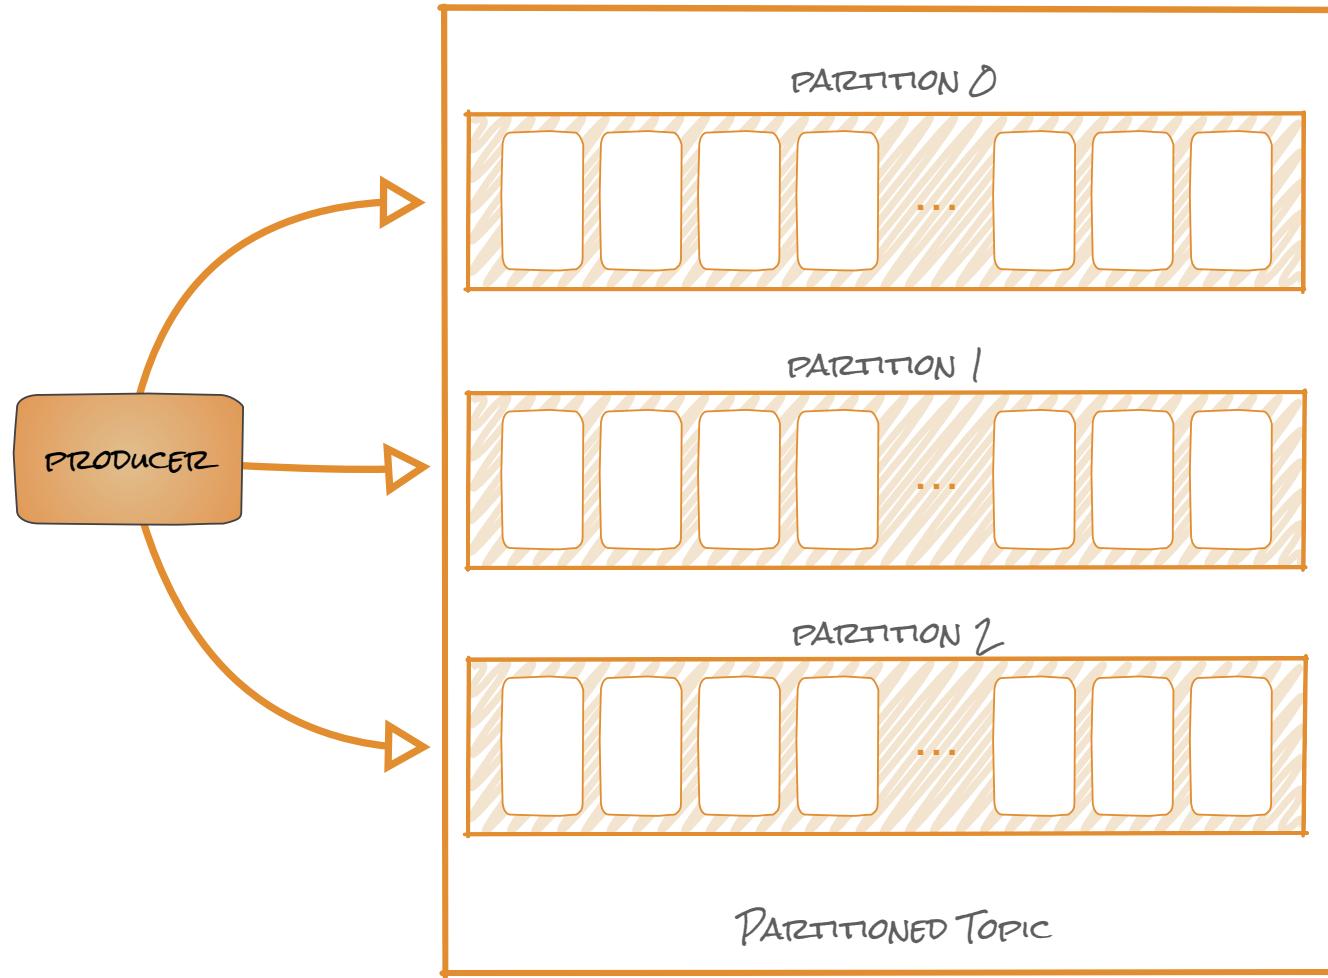


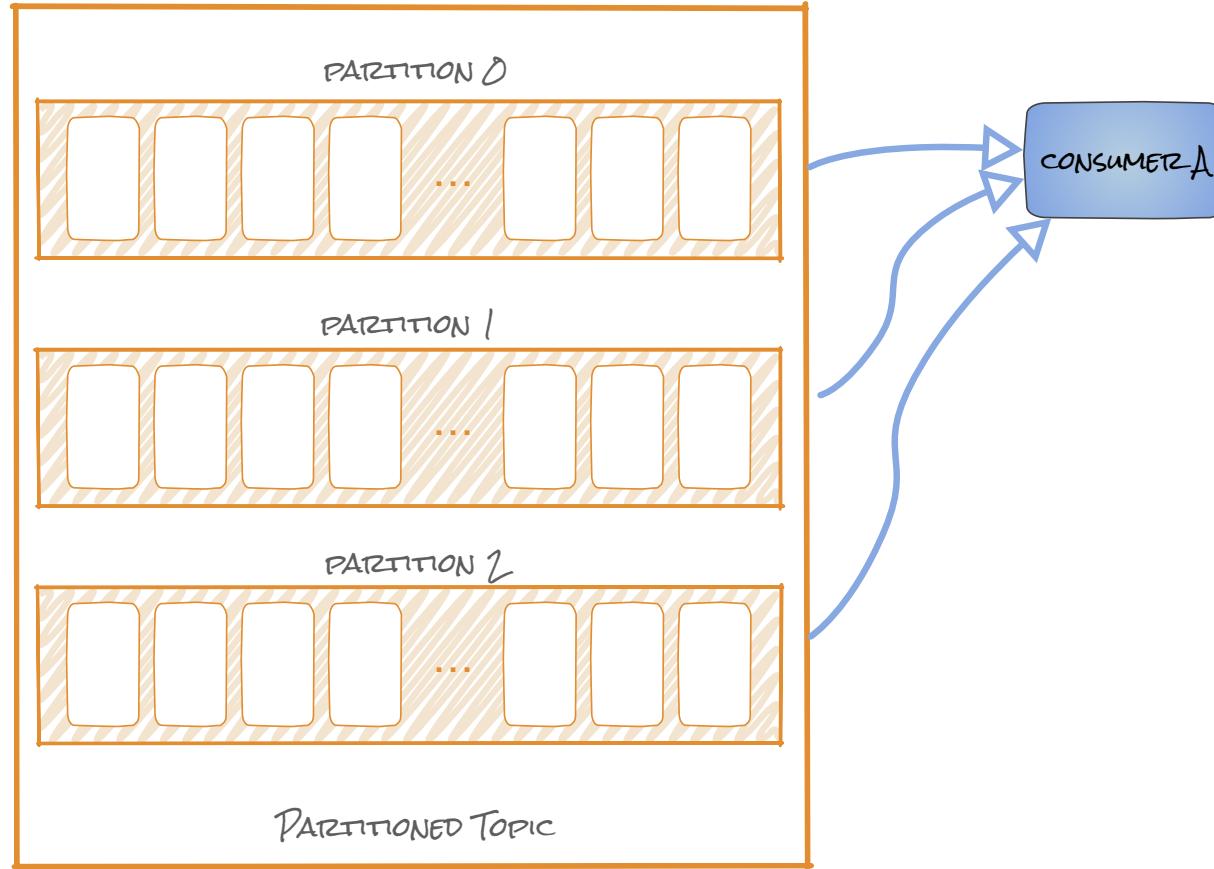


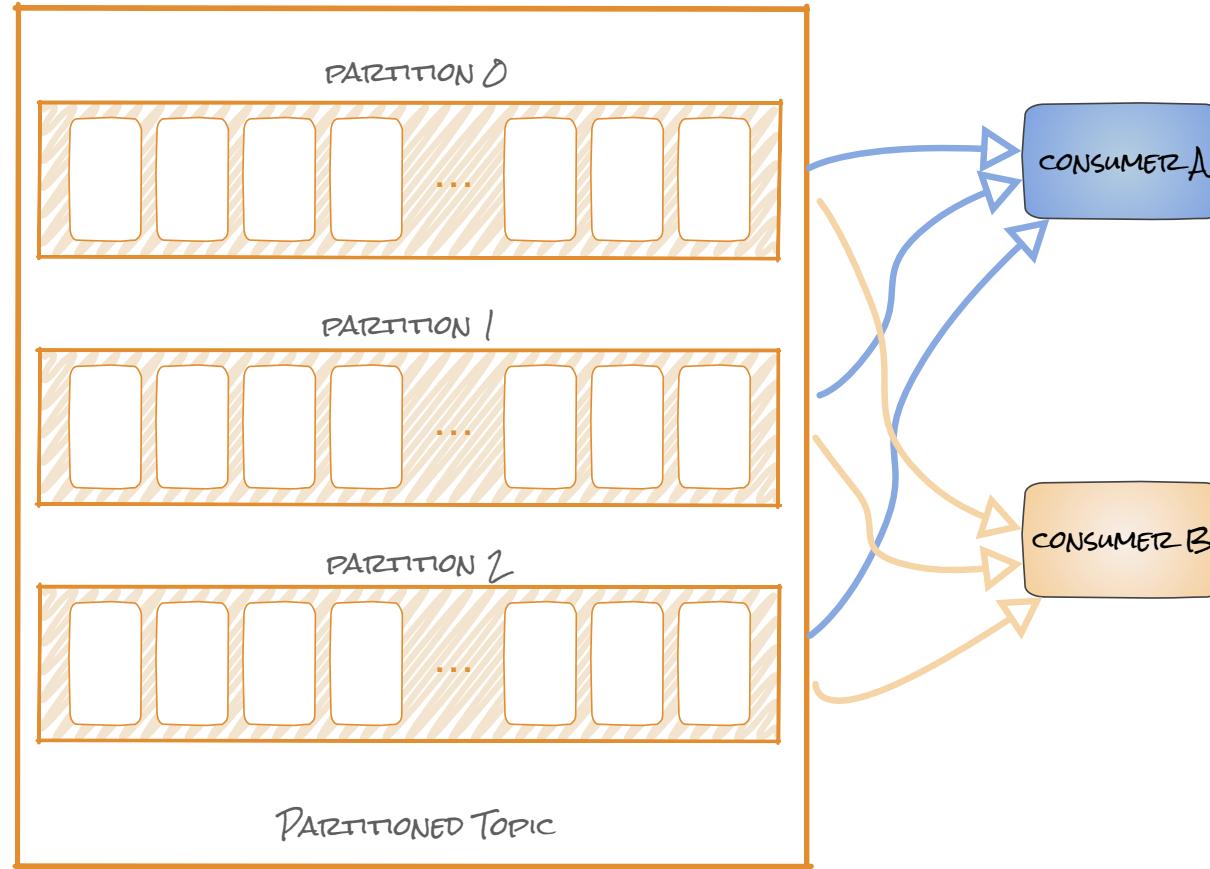


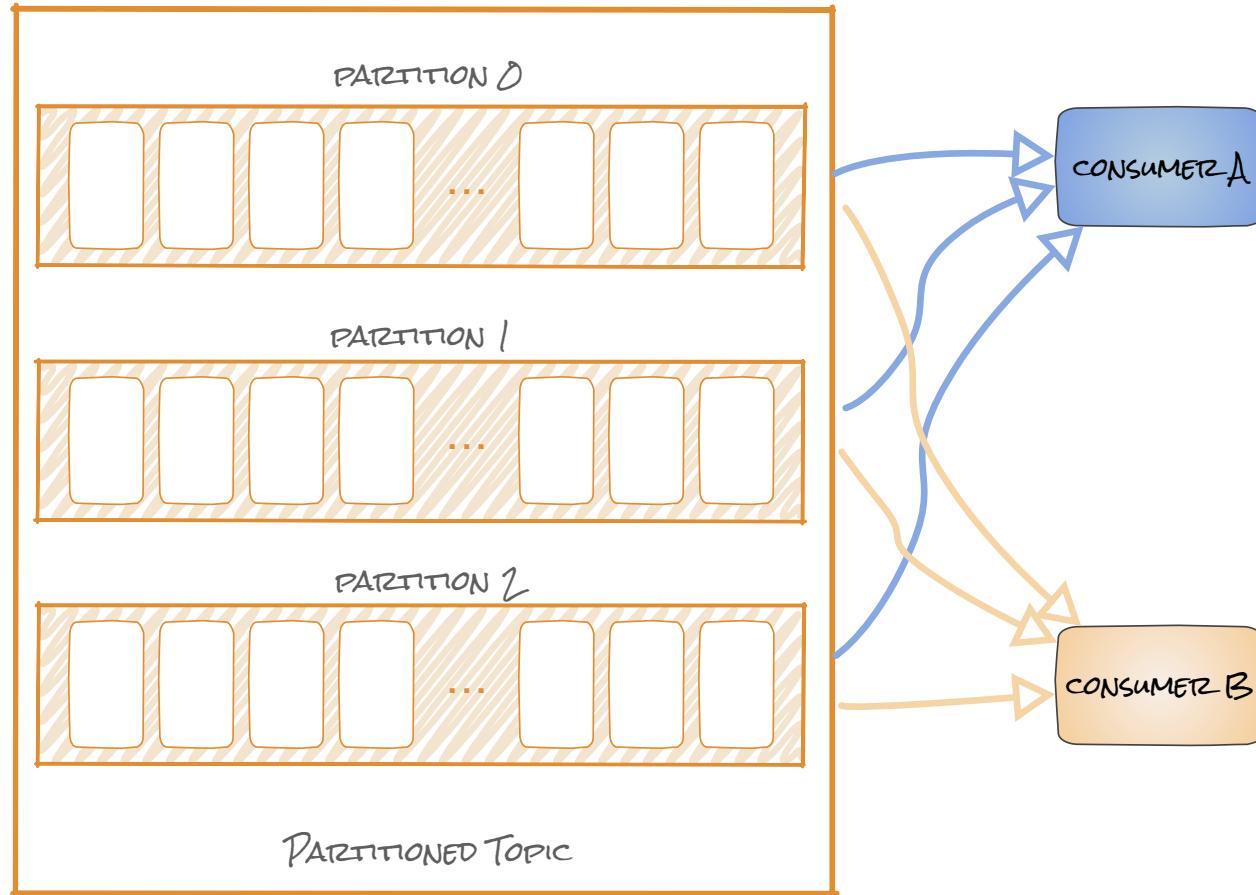


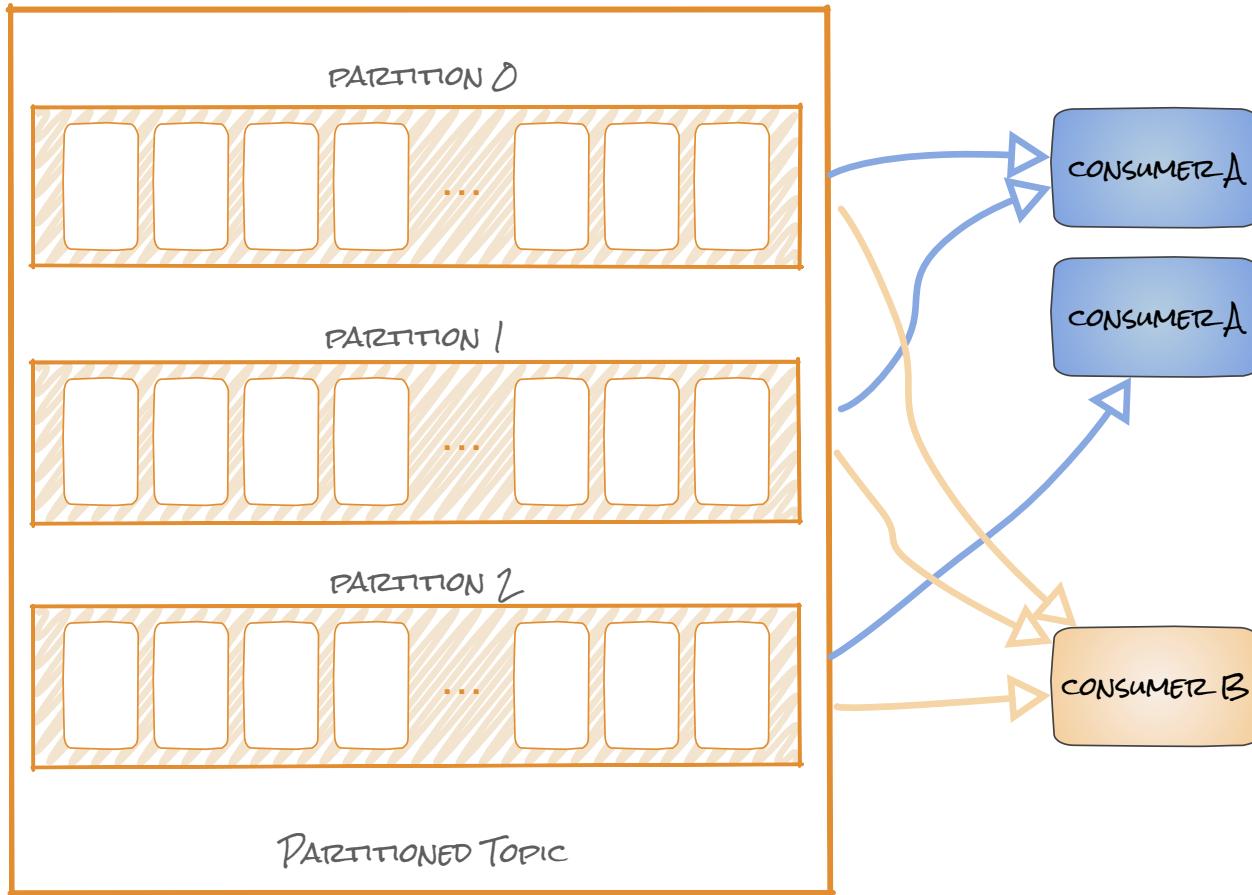


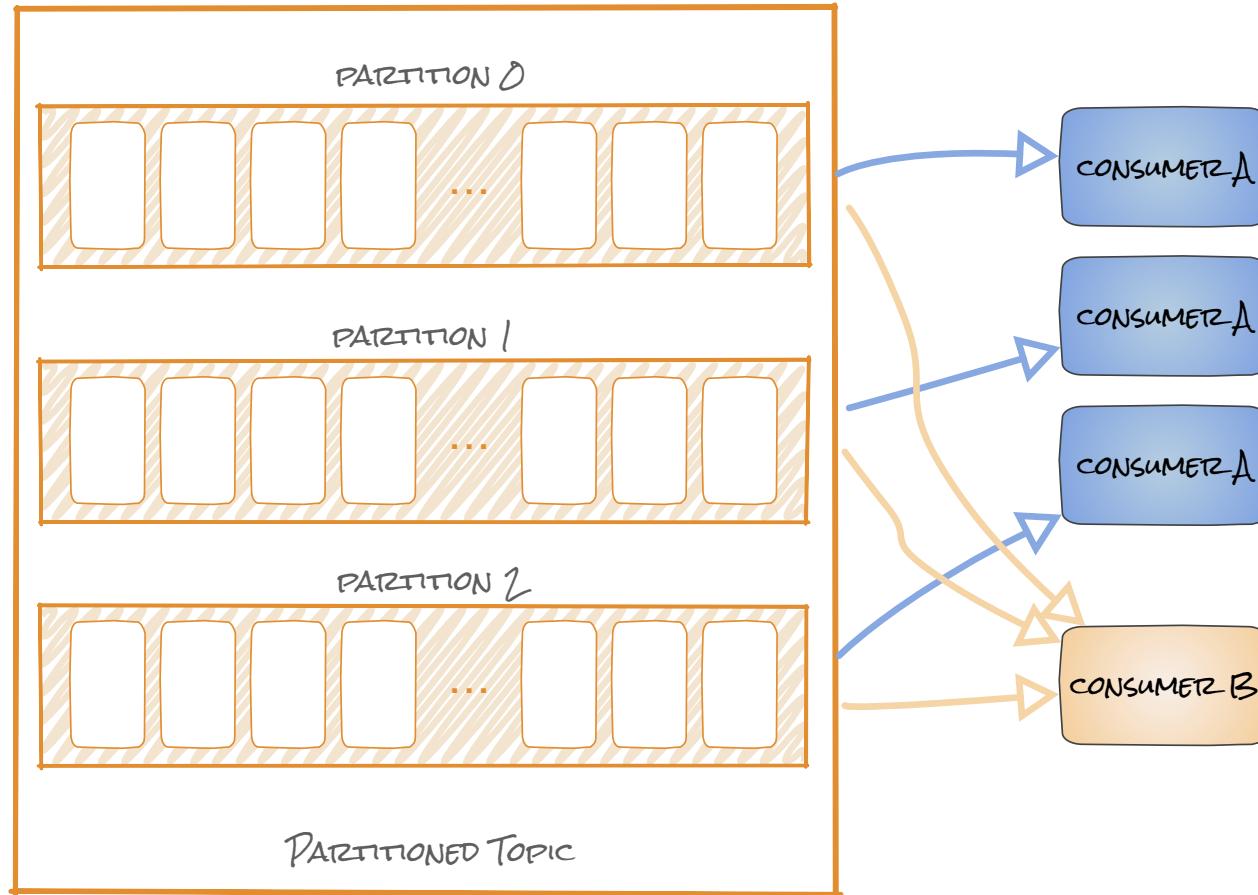


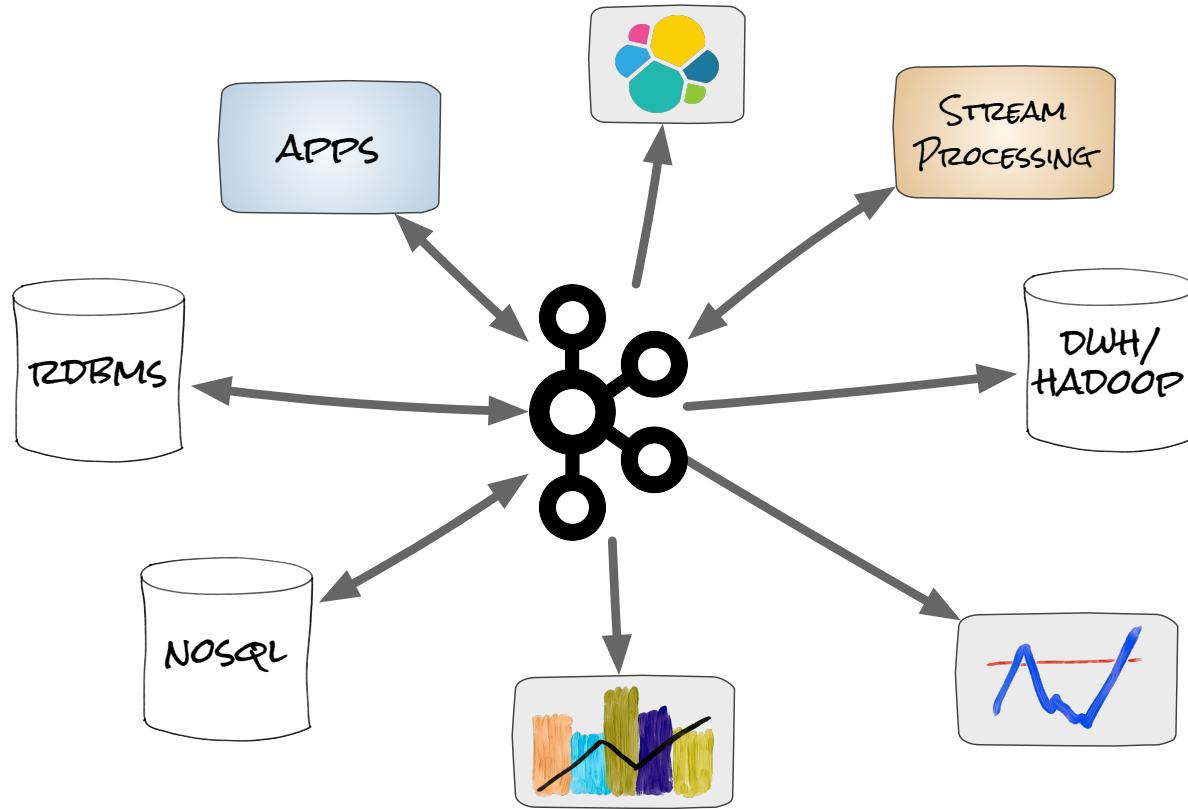


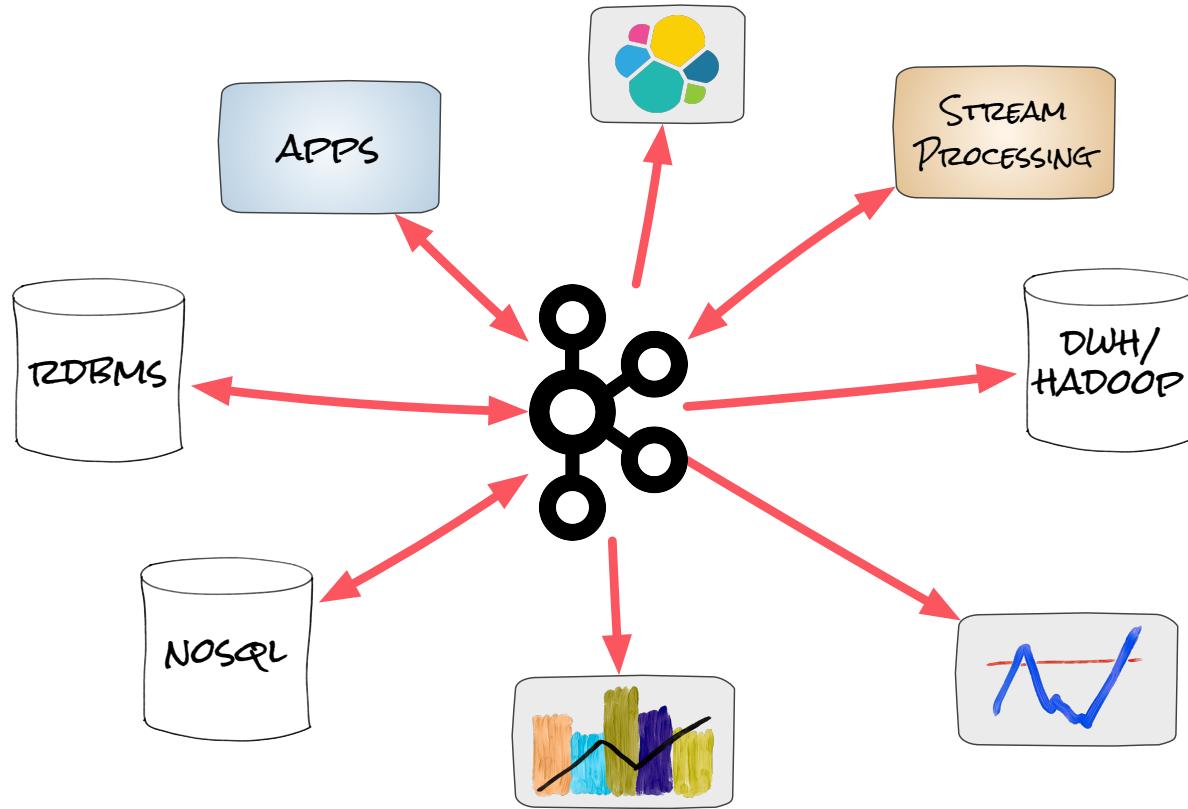


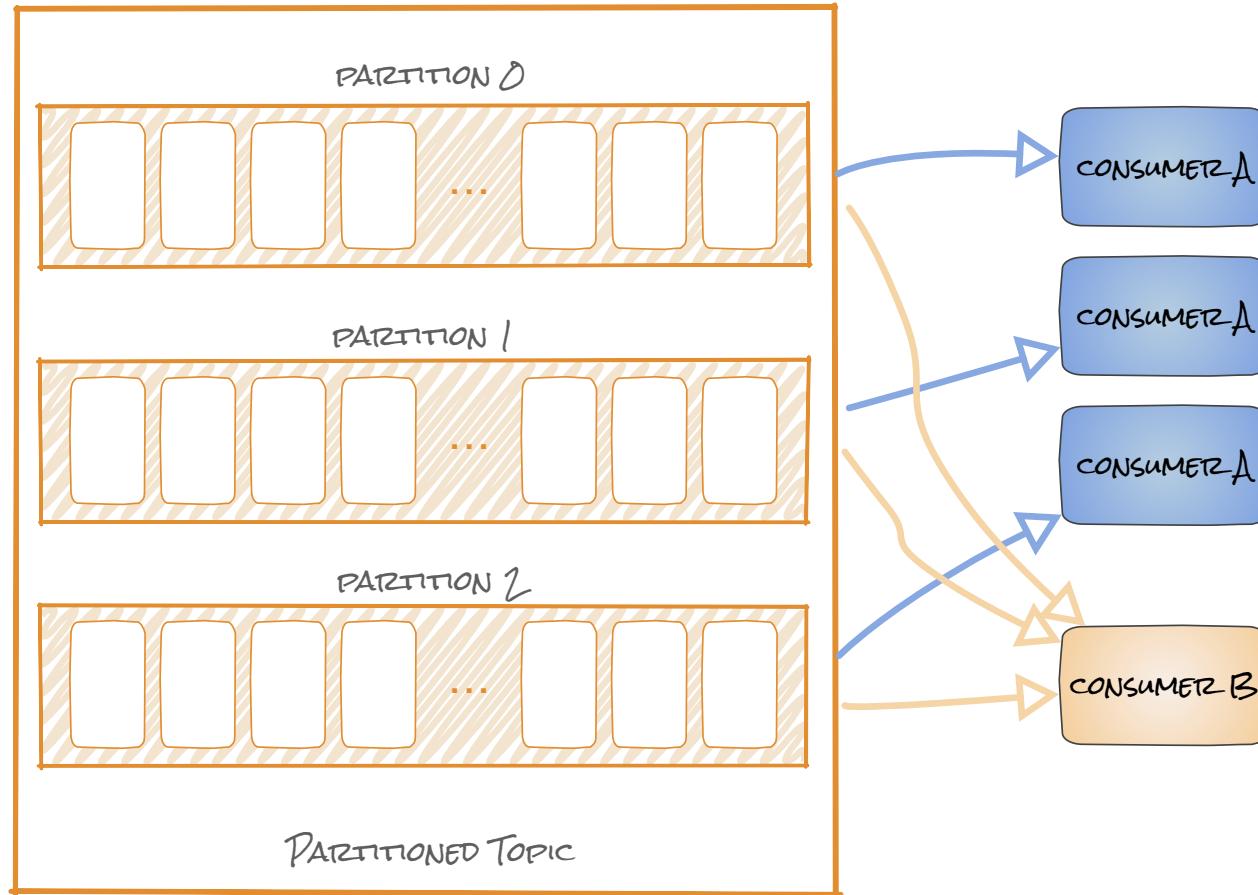


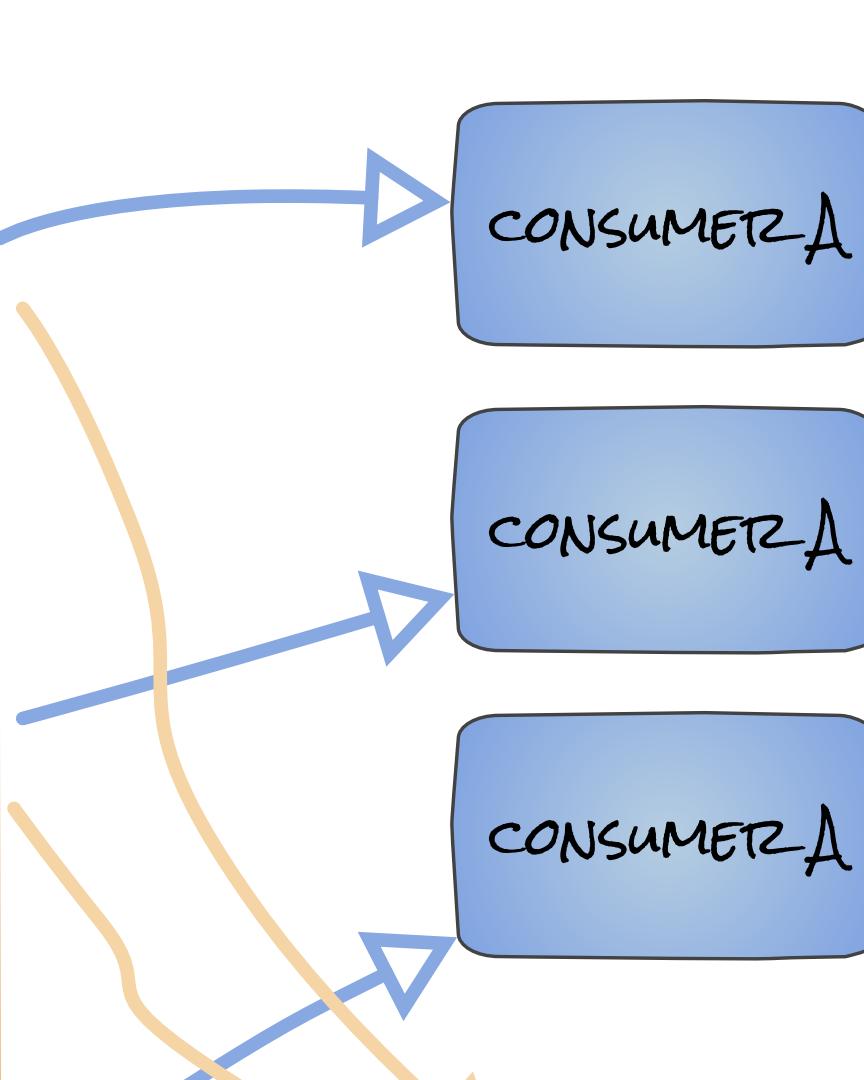
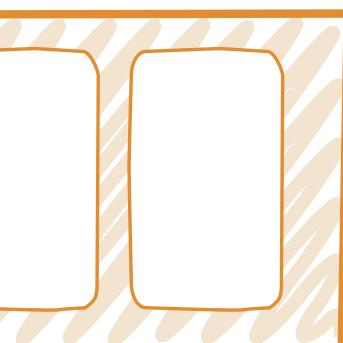
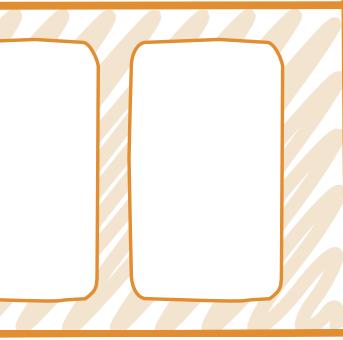


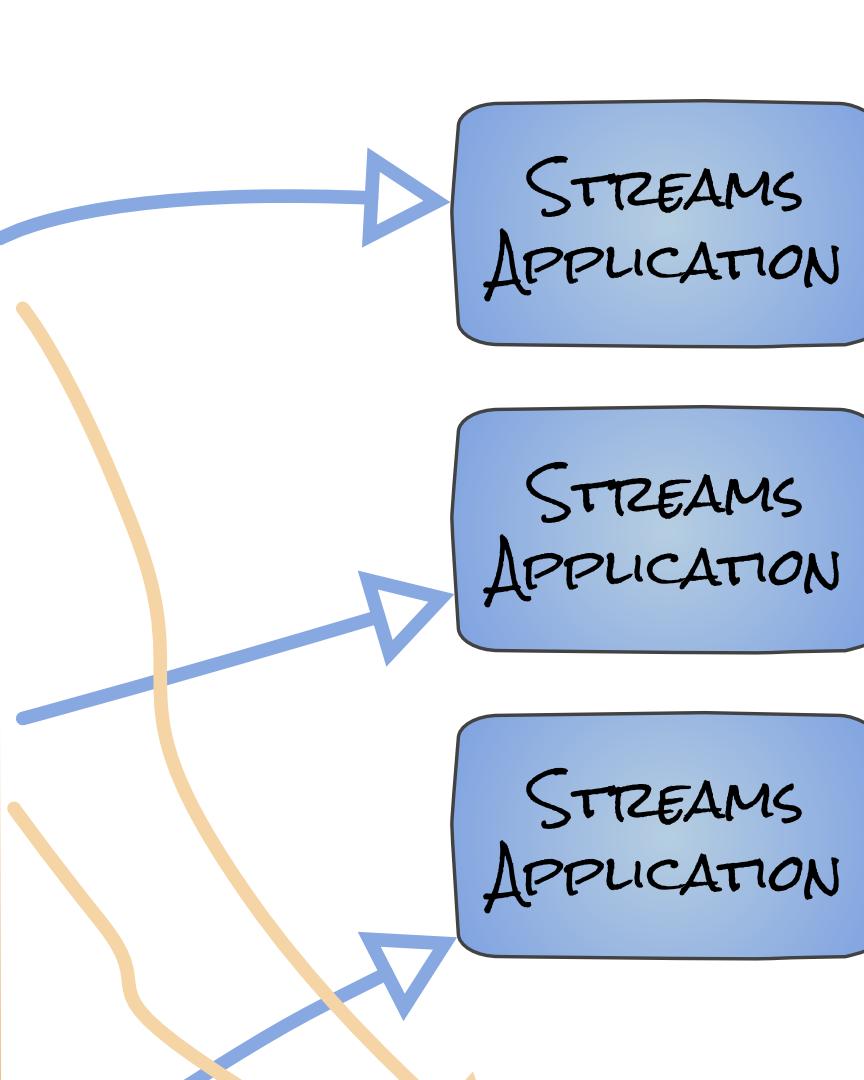
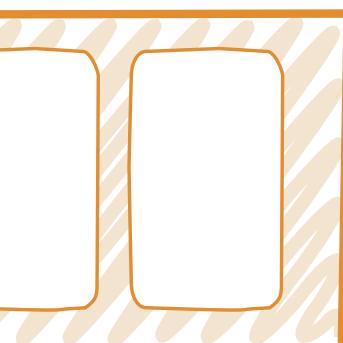
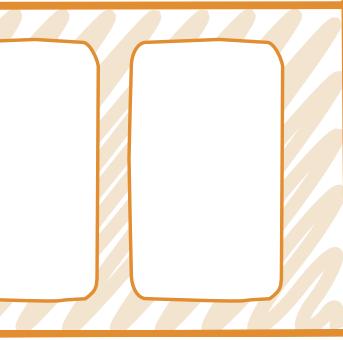












```
public static void main(String args[]) {  
    Properties streamsConfiguration = getProperties(SCHEMA_REGISTRY_URL);  
    final Map<String, String> serdeConfig =  
        Collections.singletonMap(AbstractKafkaAvroSerDeConfig.SCHEMA_REGISTRY_URL_CONFIG,  
            SCHEMA_REGISTRY_URL);  
    final SpecificAvroSerde<Movie> movieSerde = getMovieAvroSerde(serdeConfig);  
    final SpecificAvroSerde<Rating> ratingSerde = getRatingAvroSerde(serdeConfig);  
    final SpecificAvroSerde<RatedMovie> ratedMovieSerde = new SpecificAvroSerde<>();  
    ratingSerde.configure(serdeConfig, false);  
    StreamsBuilder builder = new StreamsBuilder();  
    KTable<Long, Double> ratingAverage = getRatingAverageTable(builder);  
    getRatedMoviesTable(builder, ratingAverage, movieSerde);  
    Topology topology = builder.build();  
    KafkaStreams streams = new KafkaStreams(topology, streamsConfiguration);  
    Runtime.getRuntime().addShutdownHook(new Thread(streams::close));  
    streams.start();  
}  
  
private static SpecificAvroSerde<Rating> getRatingAvroSerde(Map<String, String> serdeConfig) {  
    final SpecificAvroSerde<Rating> ratingSerde = new SpecificAvroSerde<>();  
    ratingSerde.configure(serdeConfig, false);  
    return ratingSerde;  
}  
  
public static SpecificAvroSerde<Movie> getMovieAvroSerde(Map<String, String> serdeConfig) {
```

```
}

public static SpecificAvroSerde<Movie> getMovieAvroSerde(Map<String, String> serdeConfig) {
    final SpecificAvroSerde<Movie> movieSerde = new SpecificAvroSerde<>();
    movieSerde.configure(serdeConfig, false);
    return movieSerde;
}

public static KTable<Long, String> getRatedMoviesTable(StreamsBuilder builder,
                                                       KTable<Long, Double> ratingAverage,
                                                       SpecificAvroSerde<Movie> movieSerde) {

    builder.stream("raw-movies", Consumed.with(Serdes.Long(), Serdes.String()))
        .mapValues(Parser::parseMovie)
        .map((key, movie) -> new KeyValue<>(movie.getId(), movie))
        .to("movies", Produced.with(Serdes.Long(), movieSerde));

    KTable<Long, Movie> movies = builder.table("movies",
                                                Materialized
                                                    .<Long, Movie, KeyValueStore<Bytes, byte[]>>as(
                                                        "movies-store")
                                                    .withValueSerde(movieSerde)
                                                    .withKeySerde(Serdes.Long())
    );
}

KTable<Long, String> ratedMovies = ratingAverage
    .join(movies, (avg, movie) -> movie.getTitle() + "=" + avg);

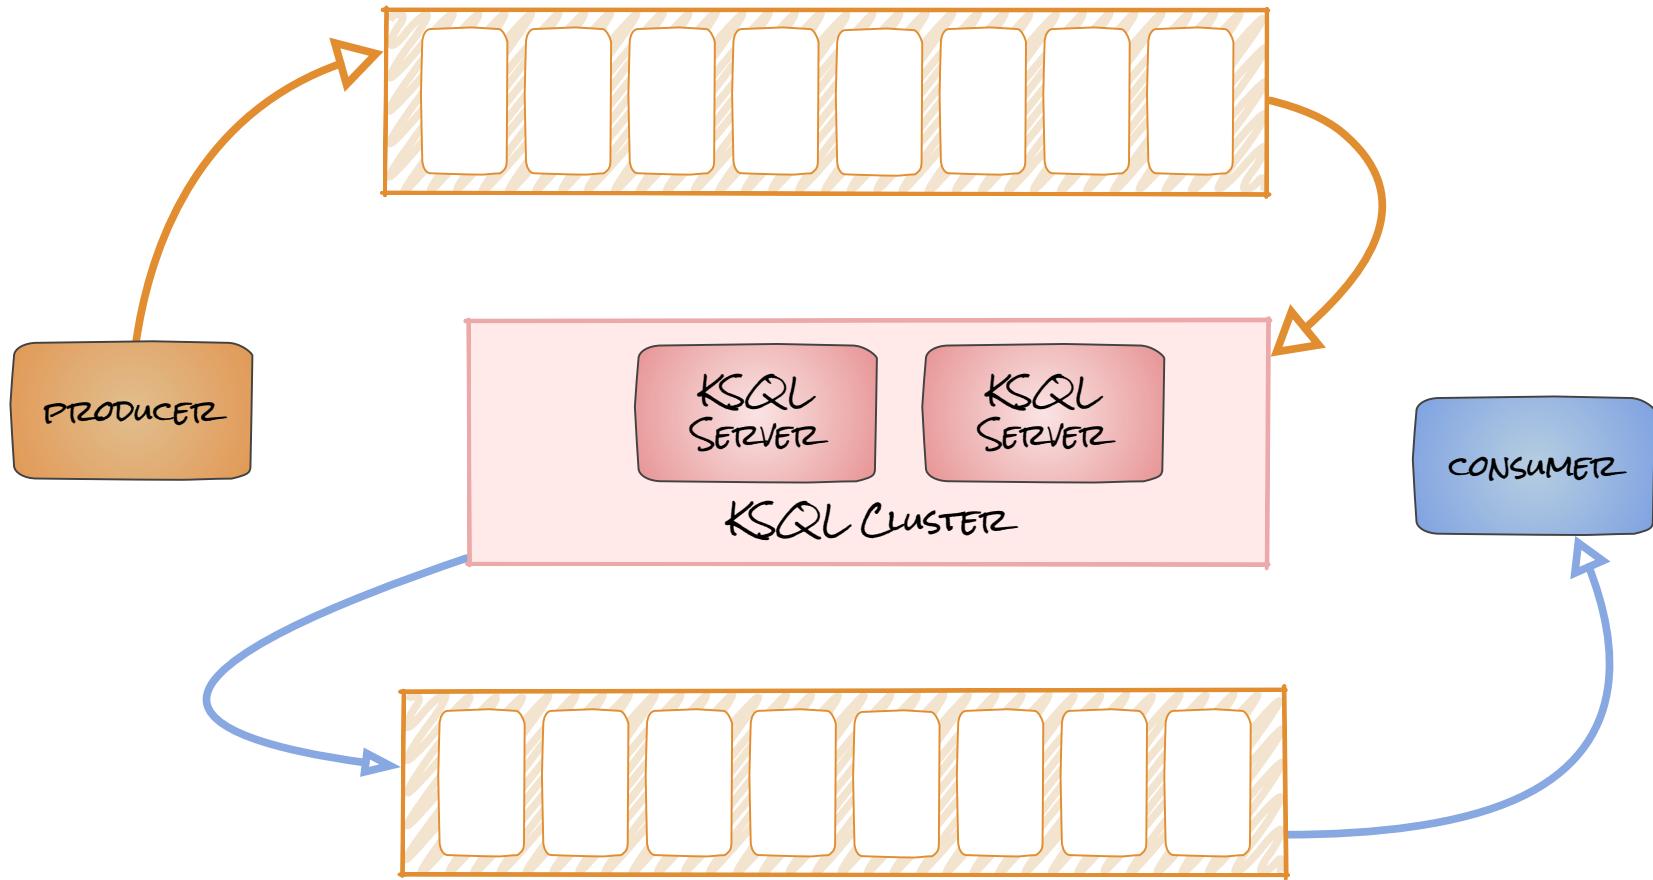
ratedMovies.toStream().to("rated-movies", Produced.with(Serdes.Long(), Serdes.String()));
return ratedMovies;
}

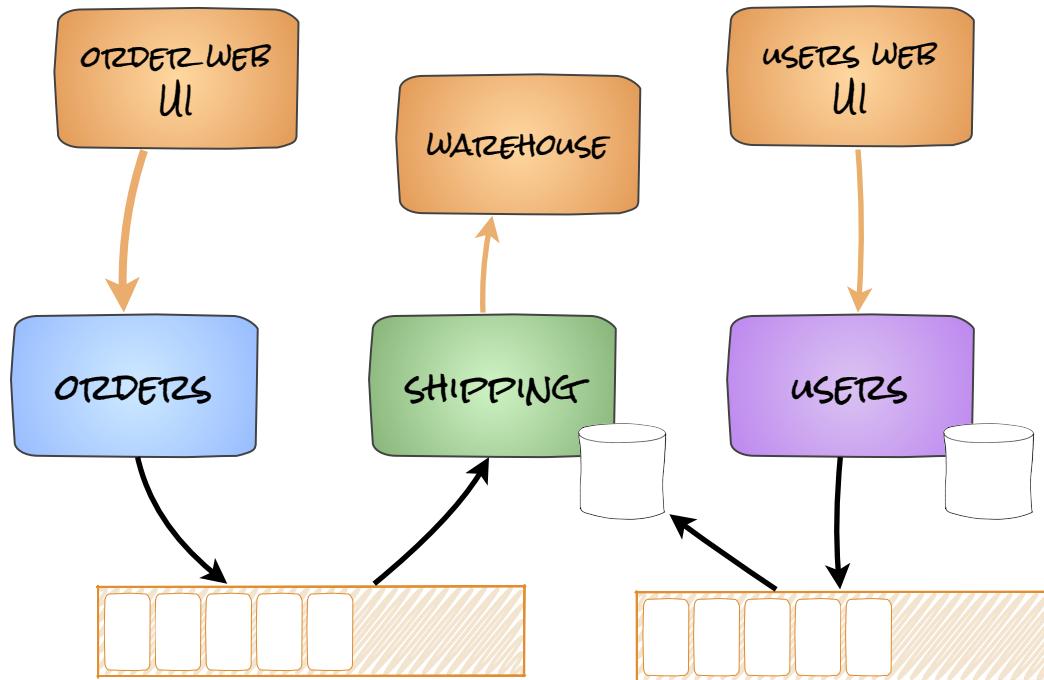
public static KTable<Long, Double> getRatingAverageTable(StreamsBuilder builder) {
    KStream<Long, String> rawRatings = builder.stream("raw-ratings",
                                                       Consumed.with(Serdes.Long(),
                                                         Serdes.String()));
    KStream<Long, Rating> ratings = rawRatings.mapValues(Parser::parseRating)
```

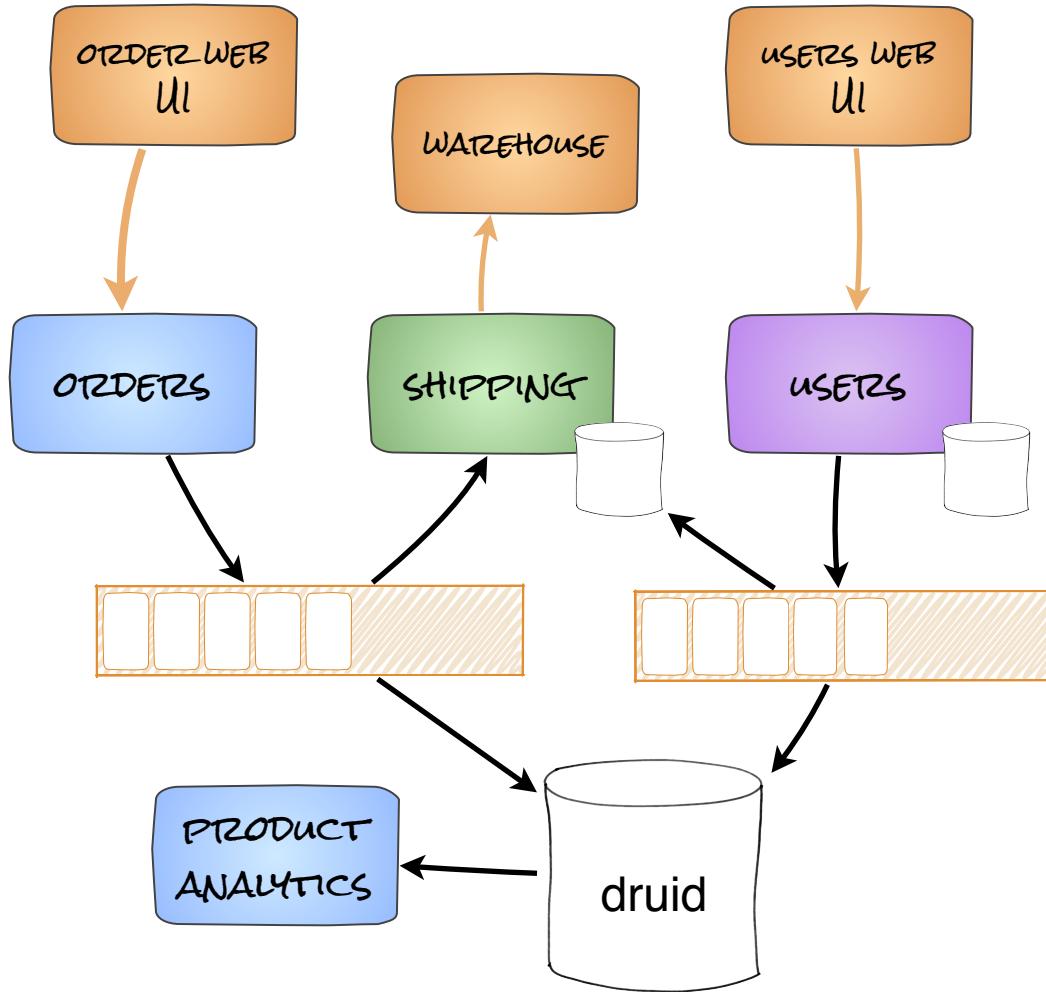
```
    return ratedMovies;
}

public static KTable<Long, Double> getRatingAverageTable(StreamsBuilder builder) {
    KStream<Long, String> rawRatings = builder.stream("raw-ratings",
        Consumed.with(Serdes.Long(),
                      Serdes.String()));
    KStream<Long, Rating> ratings = rawRatings.mapValues(Parser::parseRating)
        .map((key, rating) -> new KeyValue<>(rating.getMovieId(), rating));
    KStream<Long, Double> numericRatings = ratings.mapValues(Rating::getRating);
    KGroupedStream<Long, Double> ratingsById = numericRatings.groupByKey();
    KTable<Long, Long> ratingCounts = ratingsById.count();
    KTable<Long, Double> ratingSums = ratingsById.reduce((v1, v2) -> v1 + v2);
    KTable<Long, Double> ratingAverage = ratingSums.join(ratingCounts,
        (sum, count) -> sum / count.doubleValue(),
        Materialized.as("average-ratings"));
    ratingAverage.toStream()
        /*.peek((key, value) -> { // debug only
            System.out.println("key = " + key + ", value = " + value);
        }*/)
        .to("average-ratings");
    return ratingAverage;
}
```

```
CREATE TABLE movie_ratings AS
    SELECT title,
           SUM(rating)/COUNT(rating) AS avg_rating,
           COUNT(rating) AS num_ratings
    FROM ratings
        LEFT OUTER JOIN movies
        ON ratings.movie_id = movies.movie_id
GROUP BY title;
```



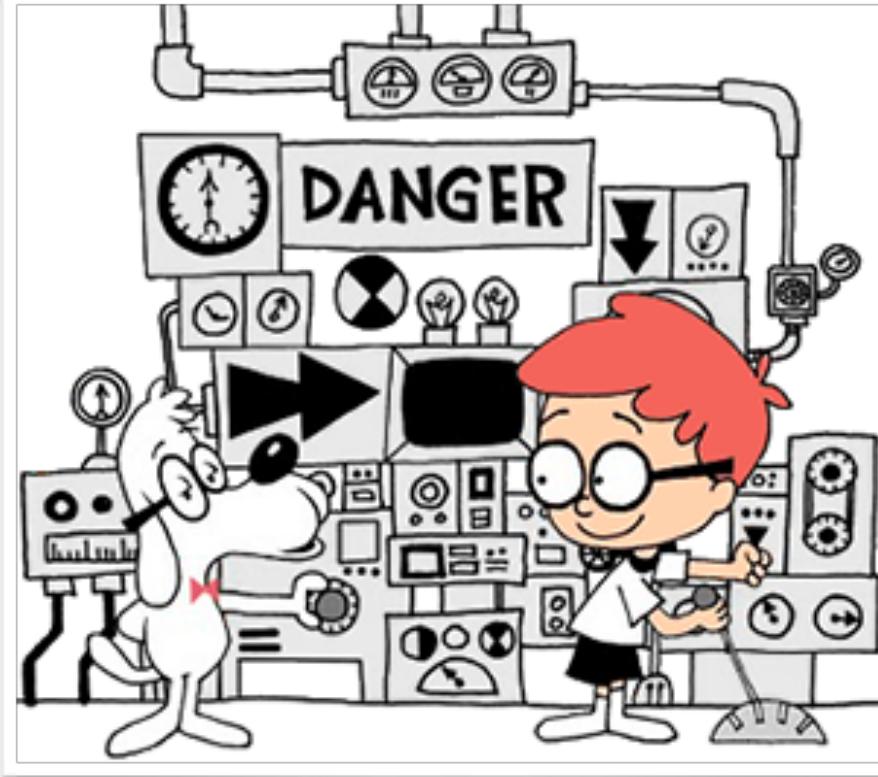




# What the heck is Apache Druid and Why Should I Care?

Rachel

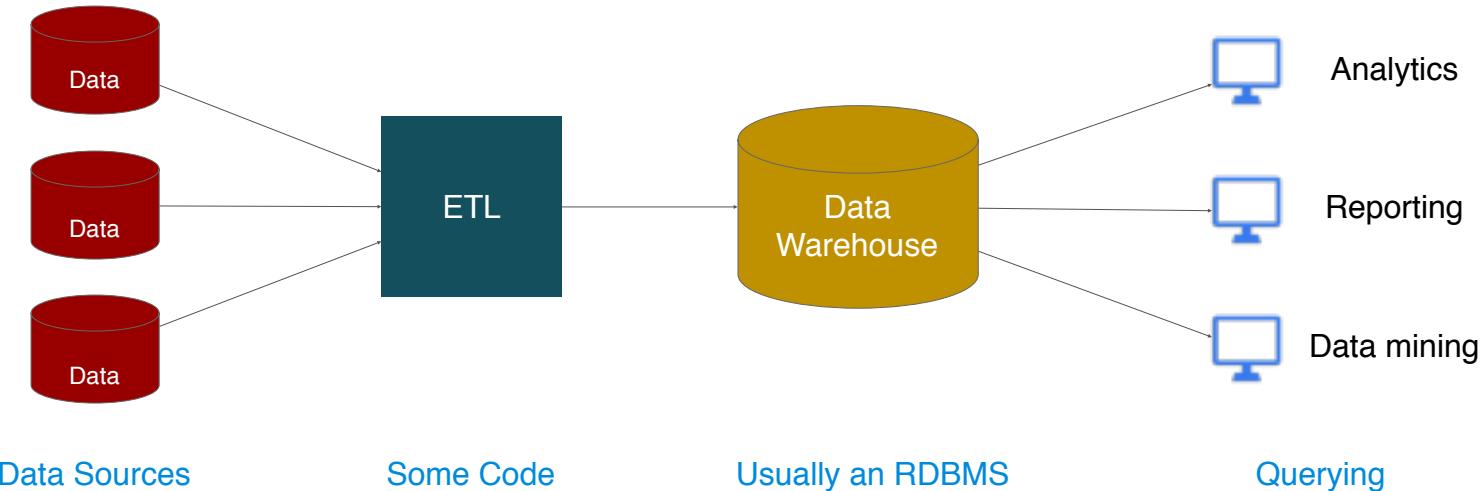
Tim

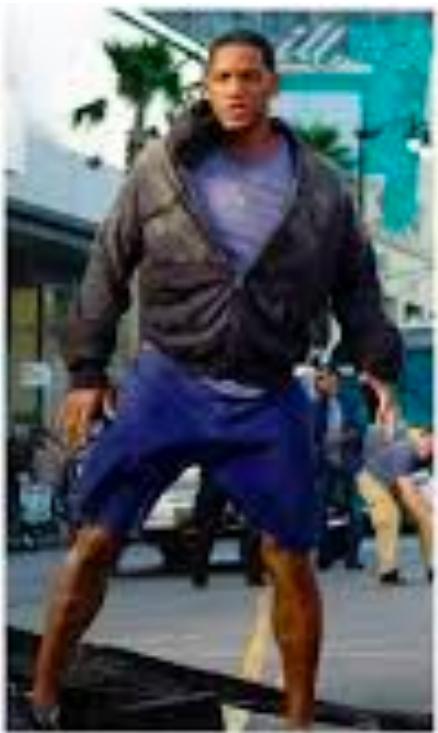


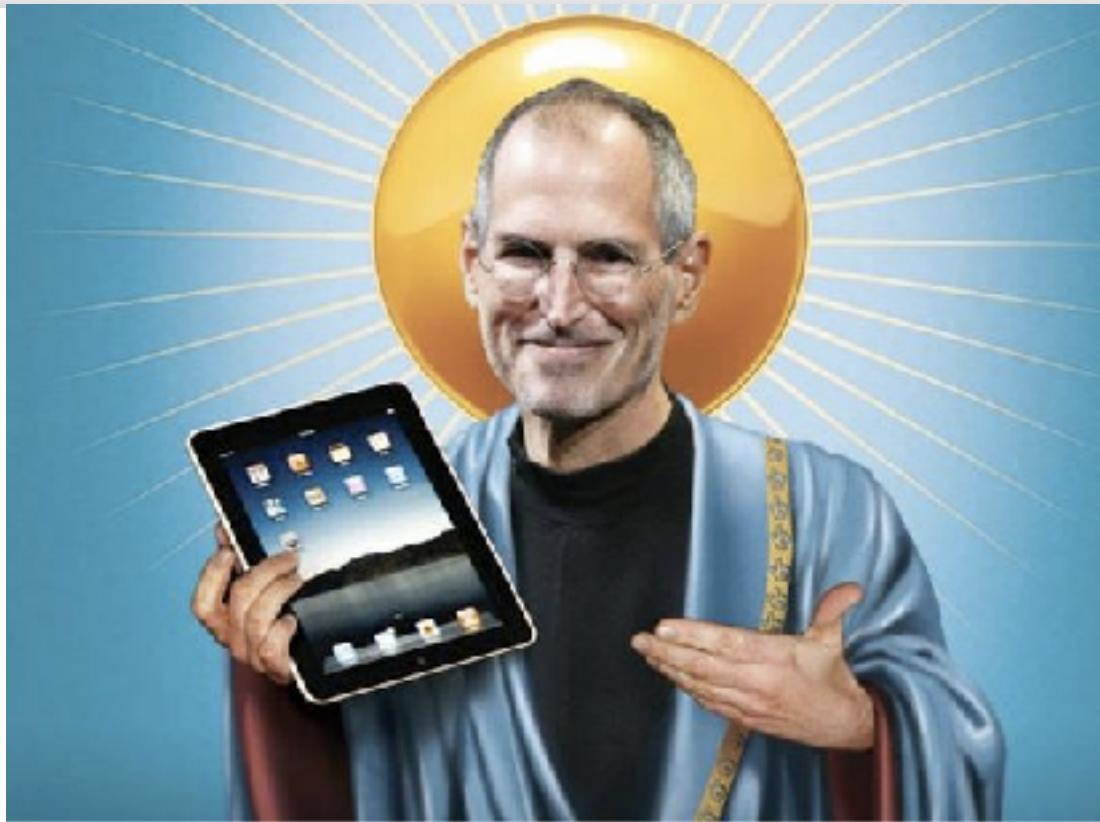


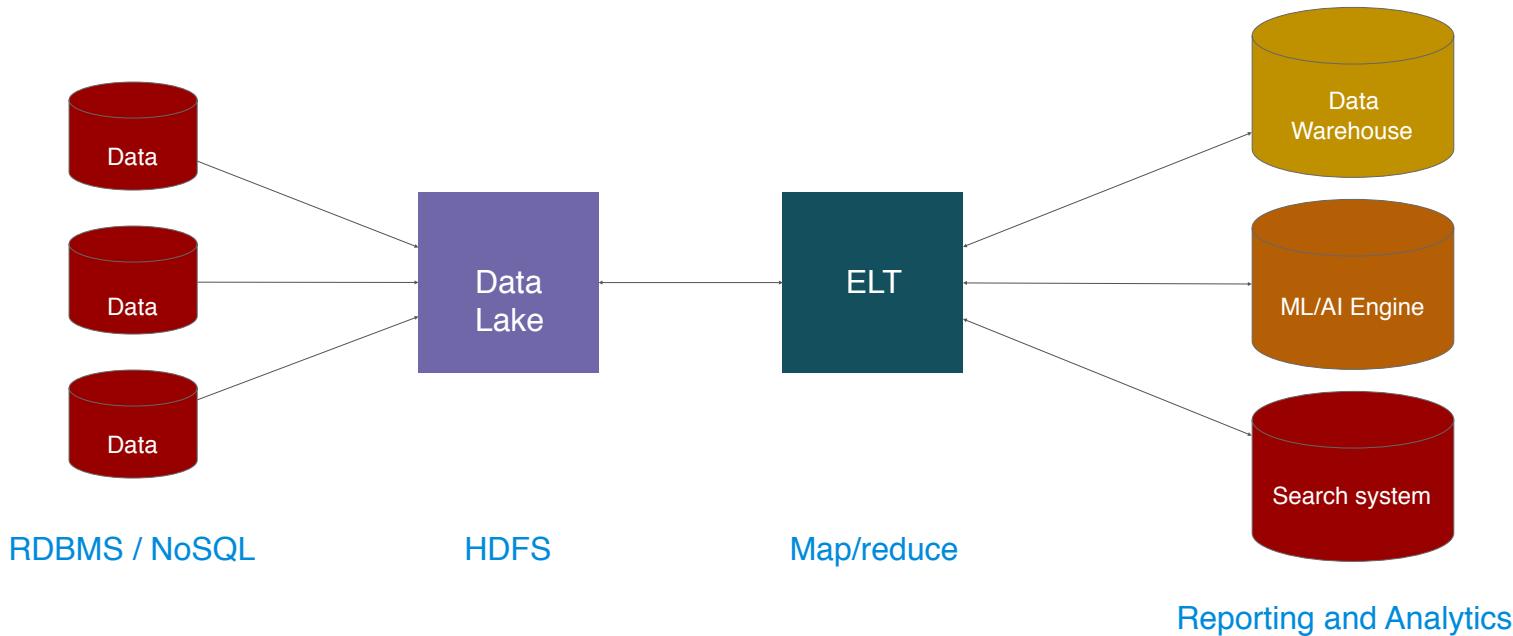




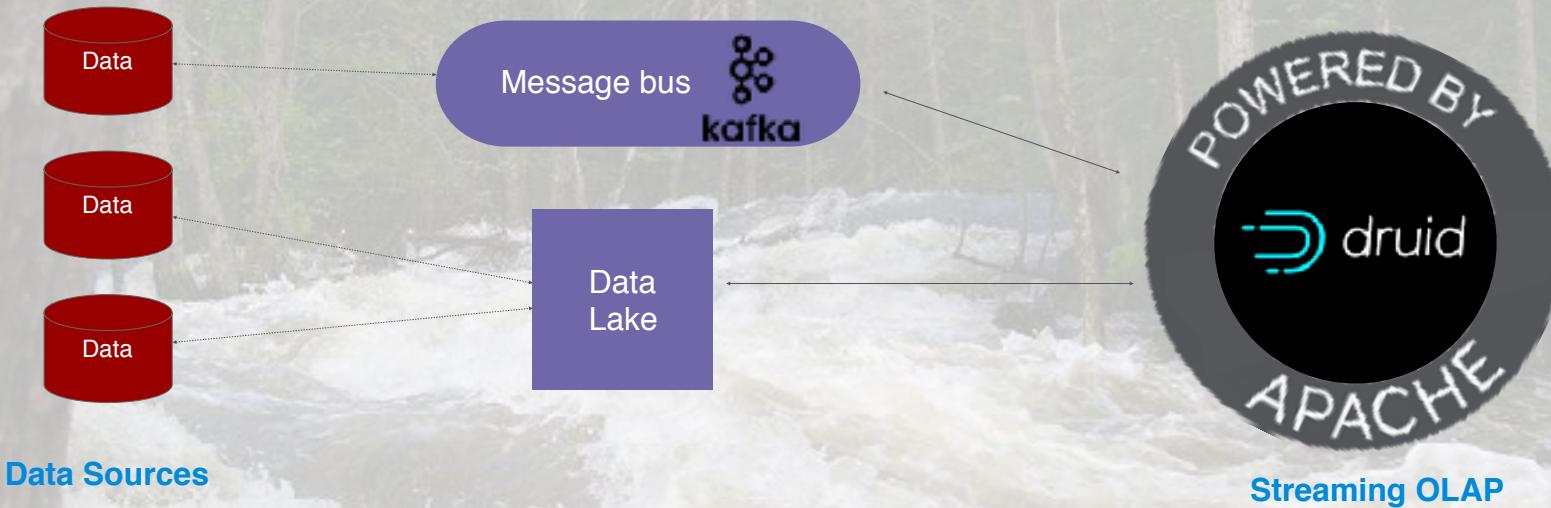














**DevOps Borat** @DEVOPS\_BORAT · 23 Mar 2013

In startup we have great of capability for churn out solution. Please send problem, we are pay good money.

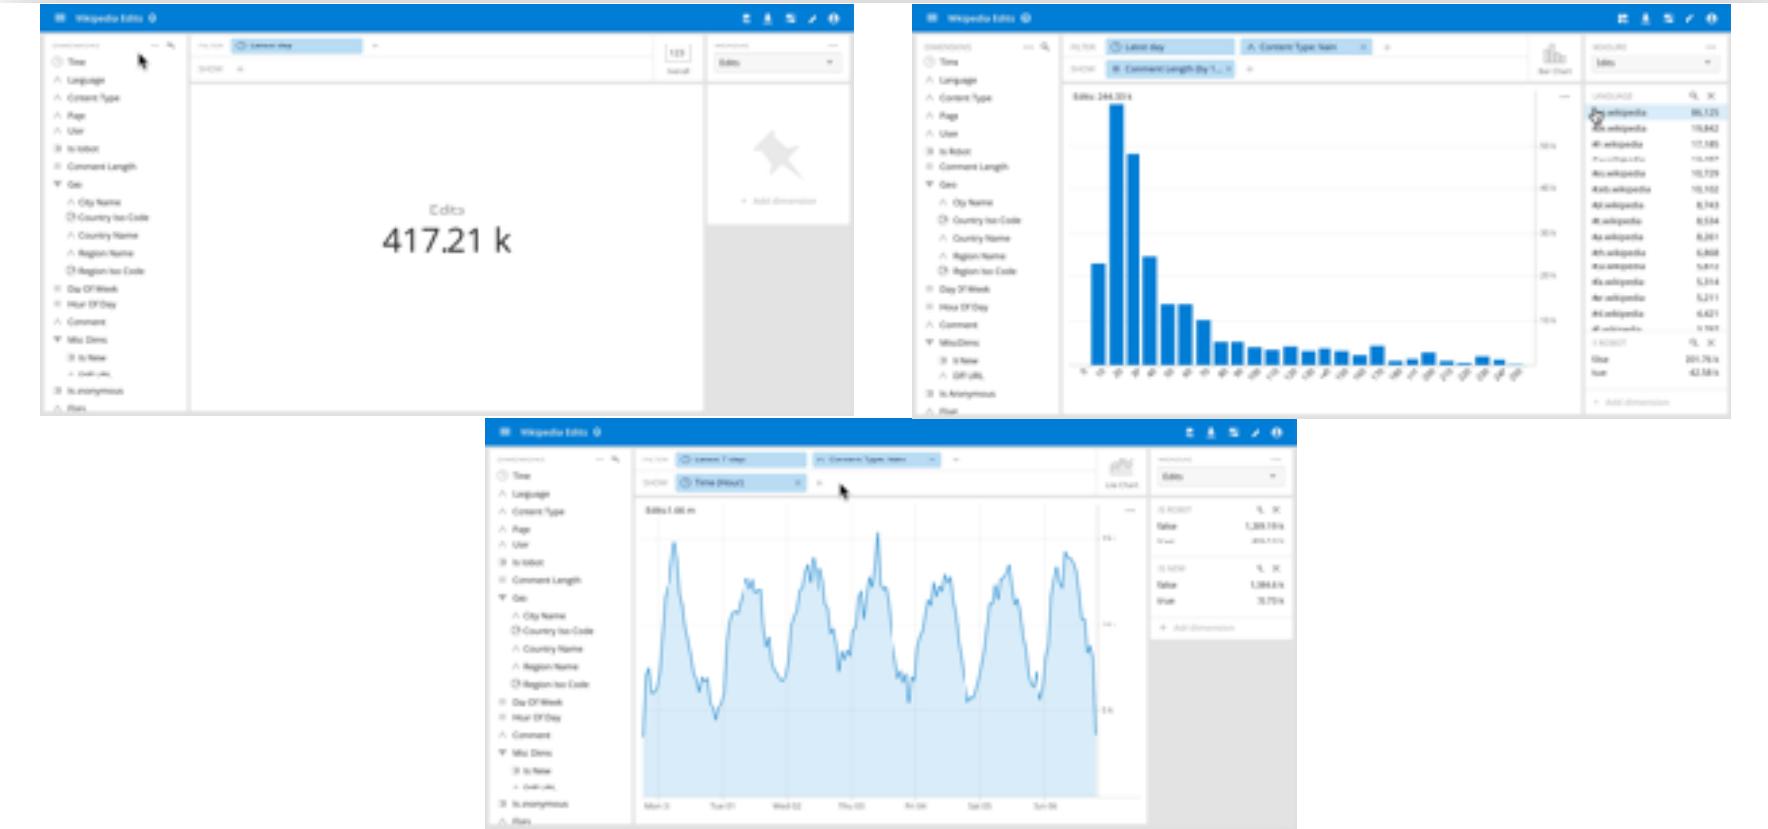


71



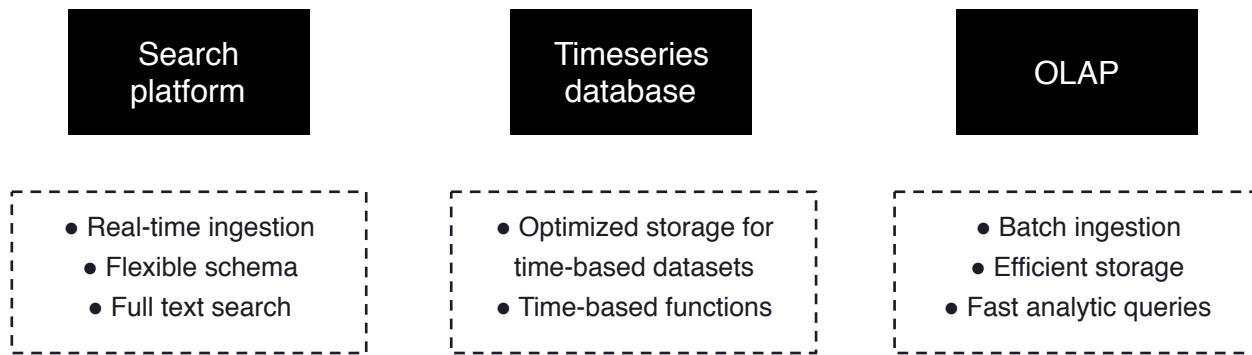
28

...



# Typical Big Data++ Challenges

- Scale: when data is large, we need a lot of servers
- Speed: aiming for sub-second response time
- Complexity: too much fine grain to precompute
- High dimensionality: 10s or 100s of dimensions
- Concurrency: many users and tenants
- Freshness: load from streams



Search  
platform

Timeseries  
database

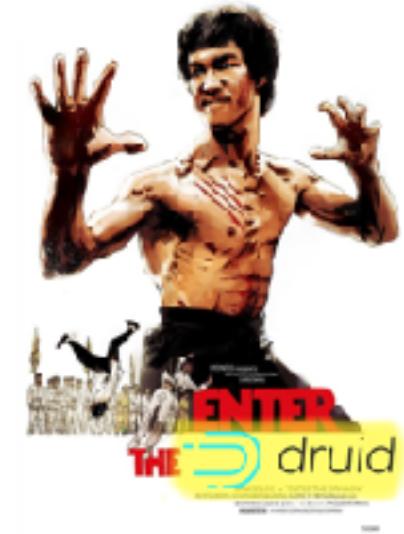
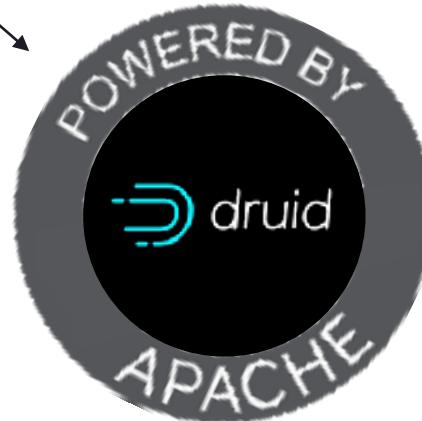
OLAP

- Real-time ingestion
- Flexible schema
- Full text search

- Optimized storage for time-based datasets
- Time-based functions

- Batch ingestion
- Efficient storage
- Fast analytic queries

high performance  
analytics **database** for  
event-driven data



# Druid Use Cases in the Wild

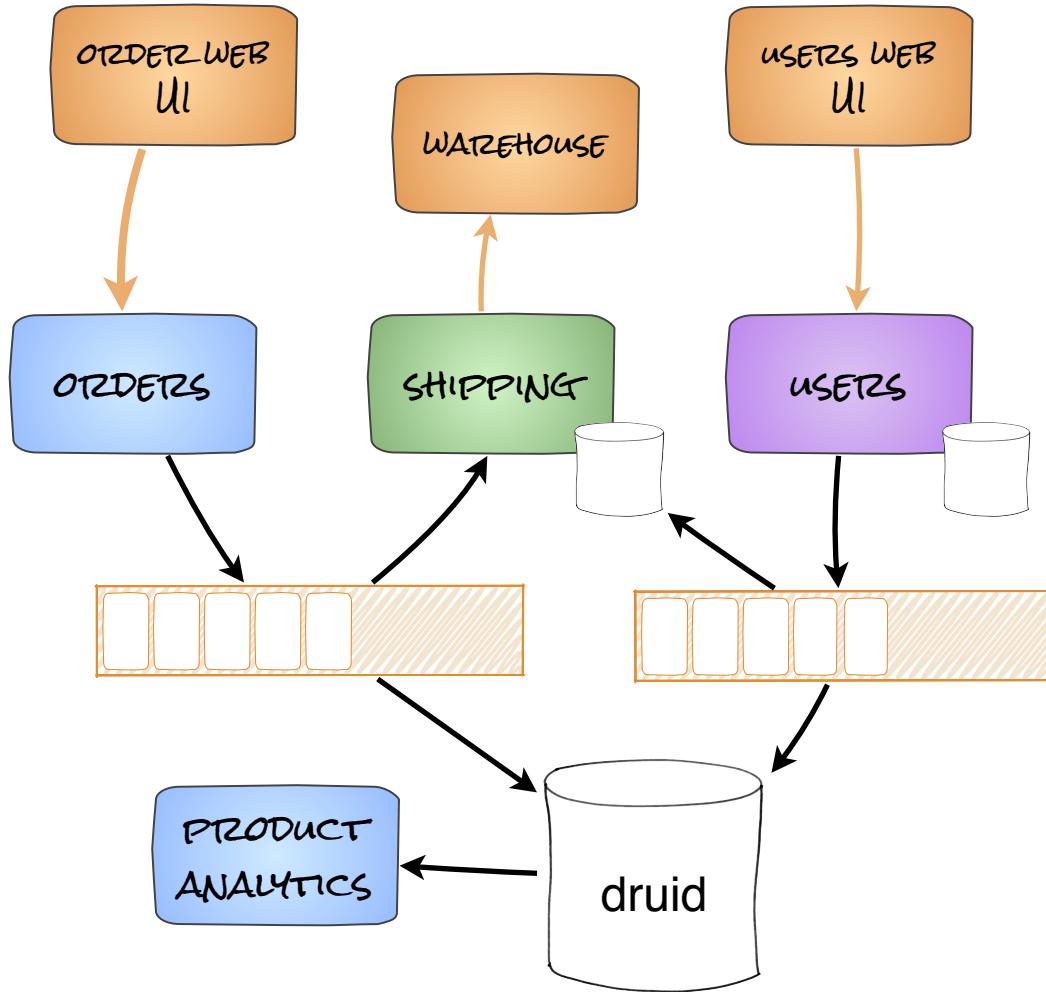
- - 1. Digital Advertising - Publishers, Advertisers, Exchanges
  - 2. User Event Analytics- Clickstream, QoS, Usage
  - 3. Network Telemetry
  - 4. Lots and Lots of Data- IoT, Product Analytics, Fraud
-

# Gratuitous Customer Quote

—  
“The performance is great ... some of the tables that we have internally in Druid have **billions and billions of events** in them, and we’re scanning them in **under a second.**”



*Source: <https://www.infoworld.com/article/2949168/hadoop/yahoo-struts-its-hadoop-stuff.html>*





# You can Druid too!

—  
Druid community site: <https://druid.apache.org/>

Imply distribution: <https://imply.io/get-started>

Contribute: <https://github.com/apache/druid>

—

