

Modernizing Infrastructures for Fast Data

Spark, Kafka, Cassandra, Reactive Platform and Mesos

by Dean Wampler, Ph.D. (@deanwampler)

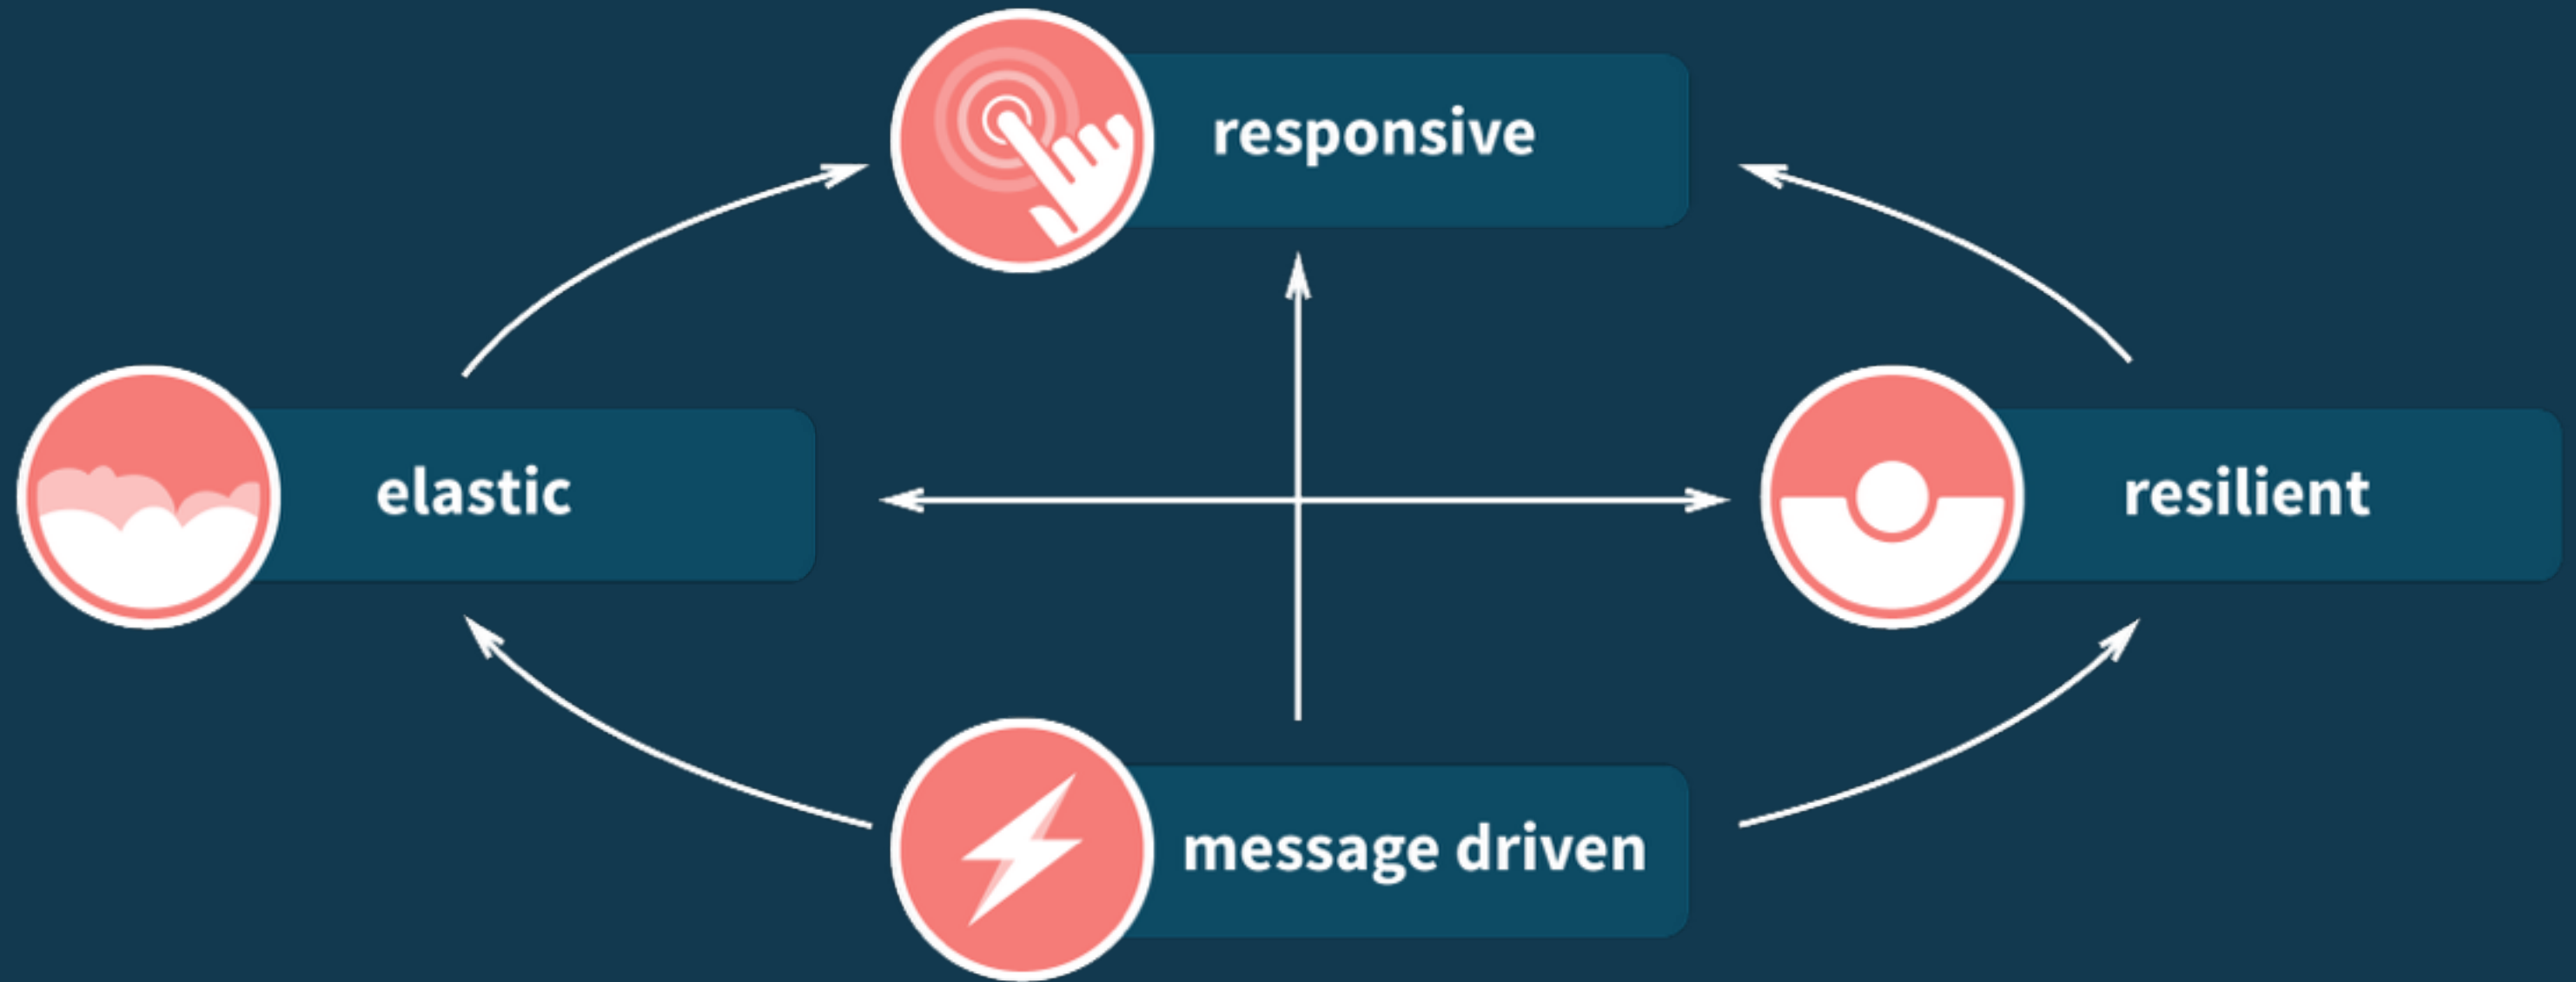


Outline


- Reactive Enterprise Architectures: The Lightbend Perspective
- Big Data and the Emergence of Apache Spark
- An Architecture for Fast Data

Reactive Enterprise Applications: The Lightbend Perspective

Reactive Manifesto



The Lightbend Reactive Platform



Media



Social



Technology



Education



IoT



Online Services



Retail



Finance



Big Data and the Emergence of Apache Spark

Distributed compute frameworks: MapReduce

- Distribution computation over that data.

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

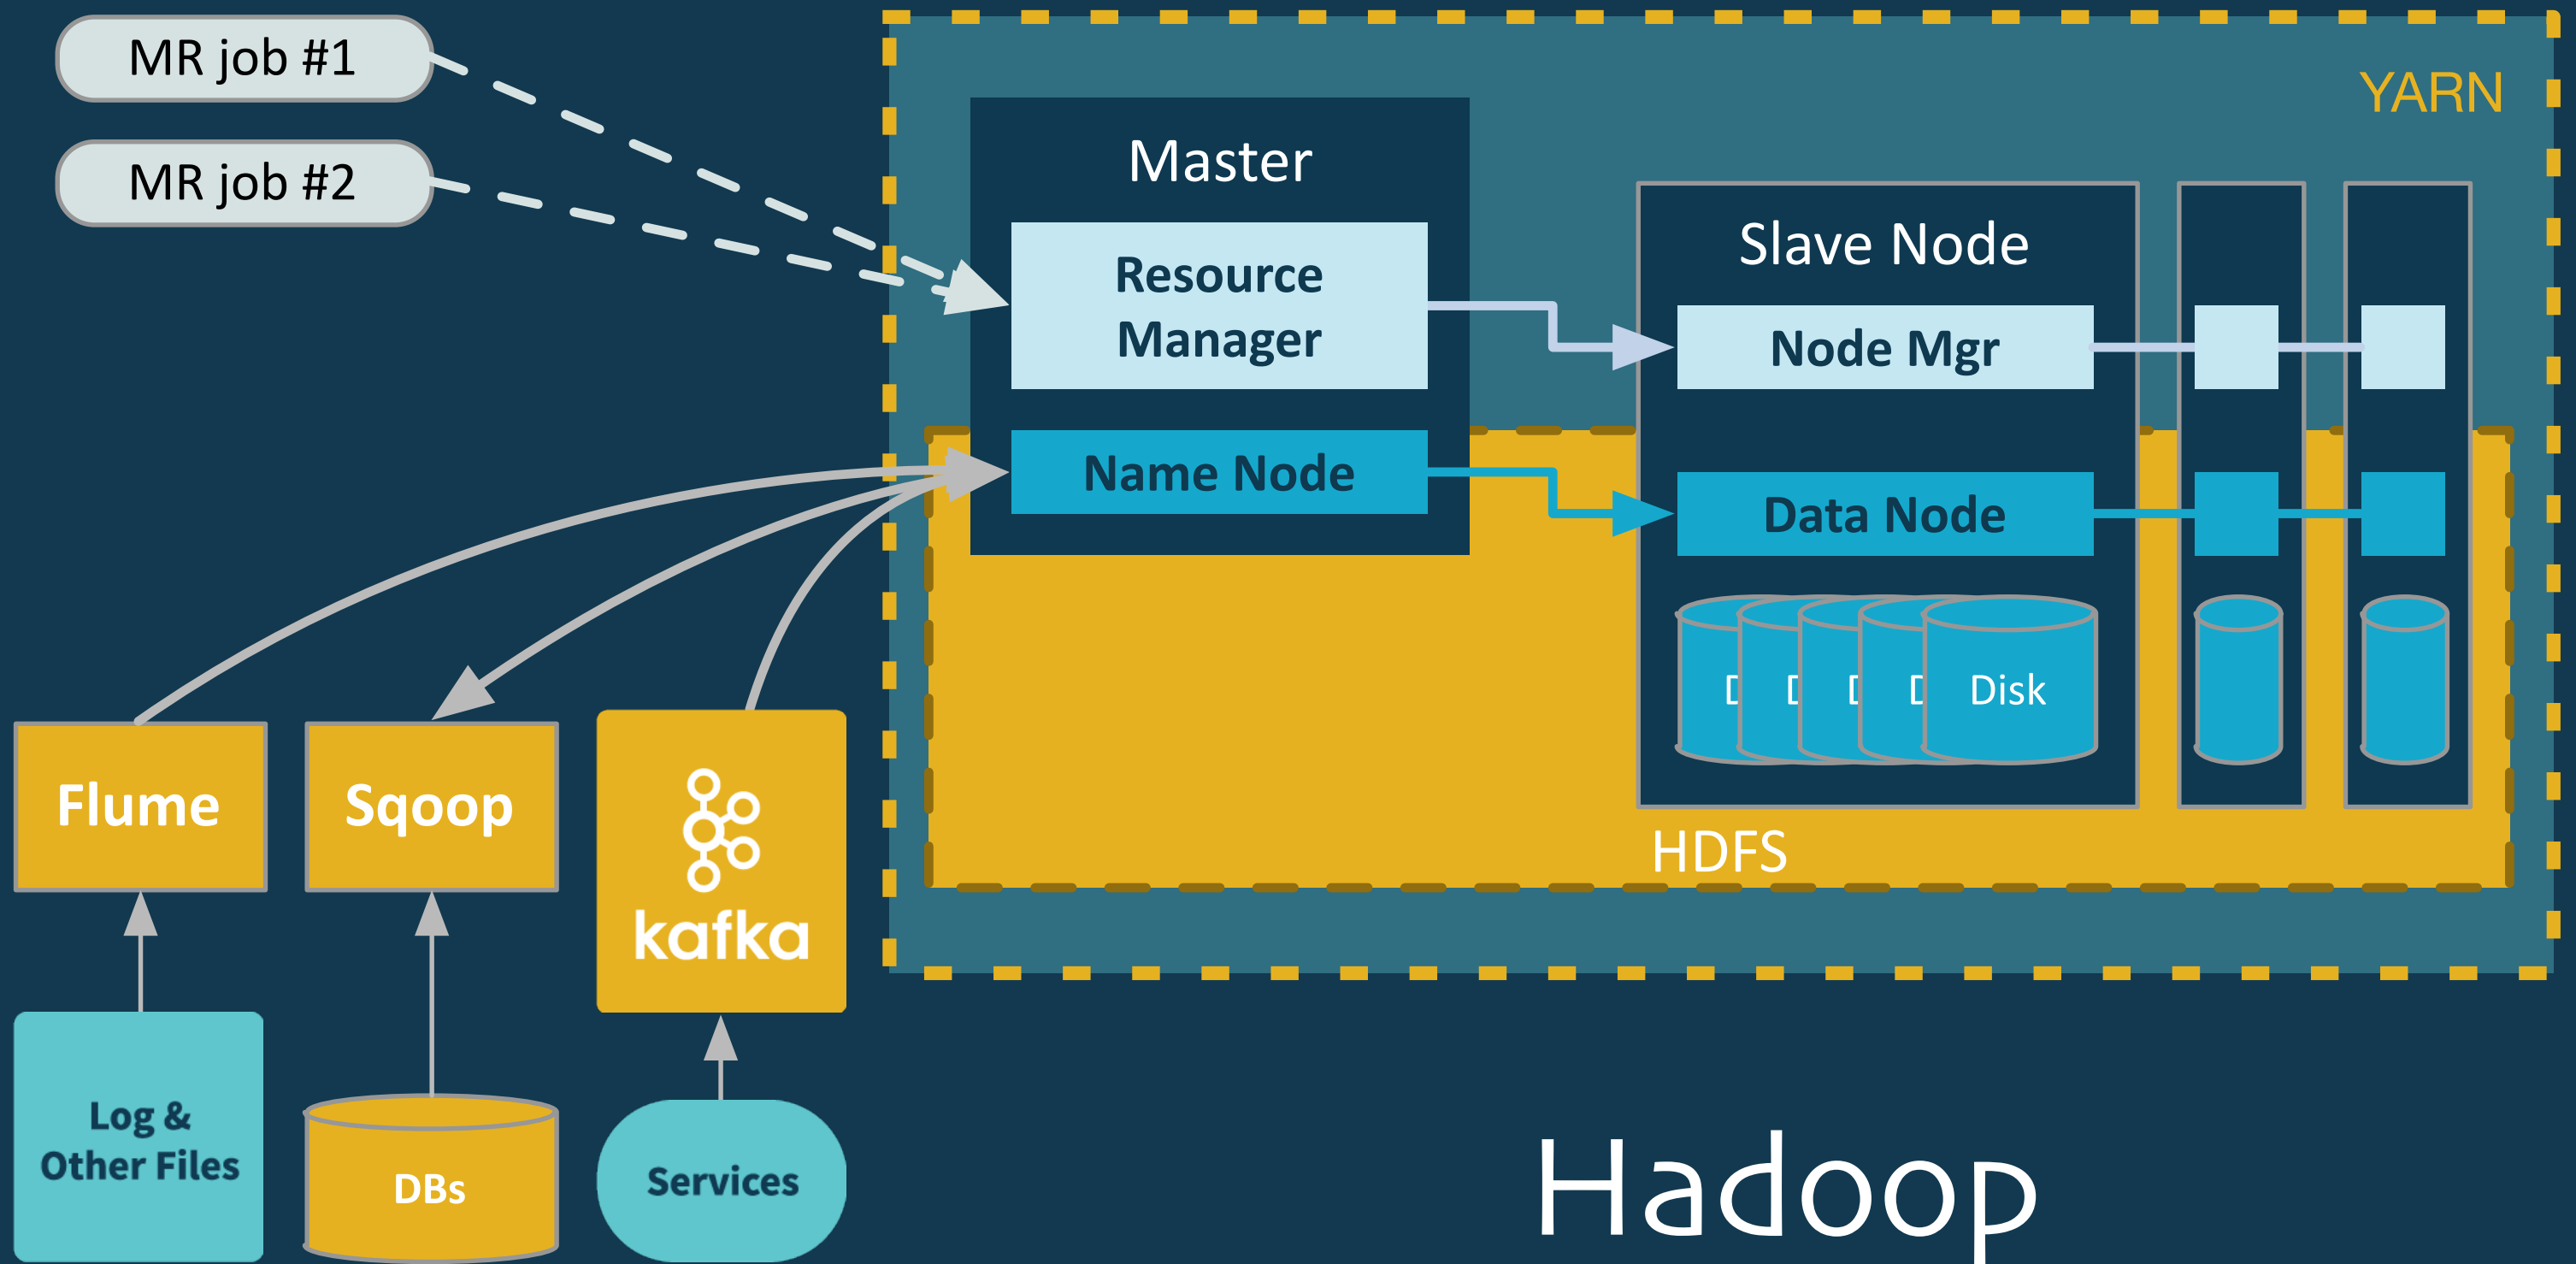
Google, Inc.



Abstract

MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a key/value pair to generate a set of intermediate key/value pairs, and a *reduce* function that merges all intermediate

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across hundreds or thousands of machines in order to finish in a reasonable amount of time. The issues of how to parallelize the computation, distribute the data, and handle



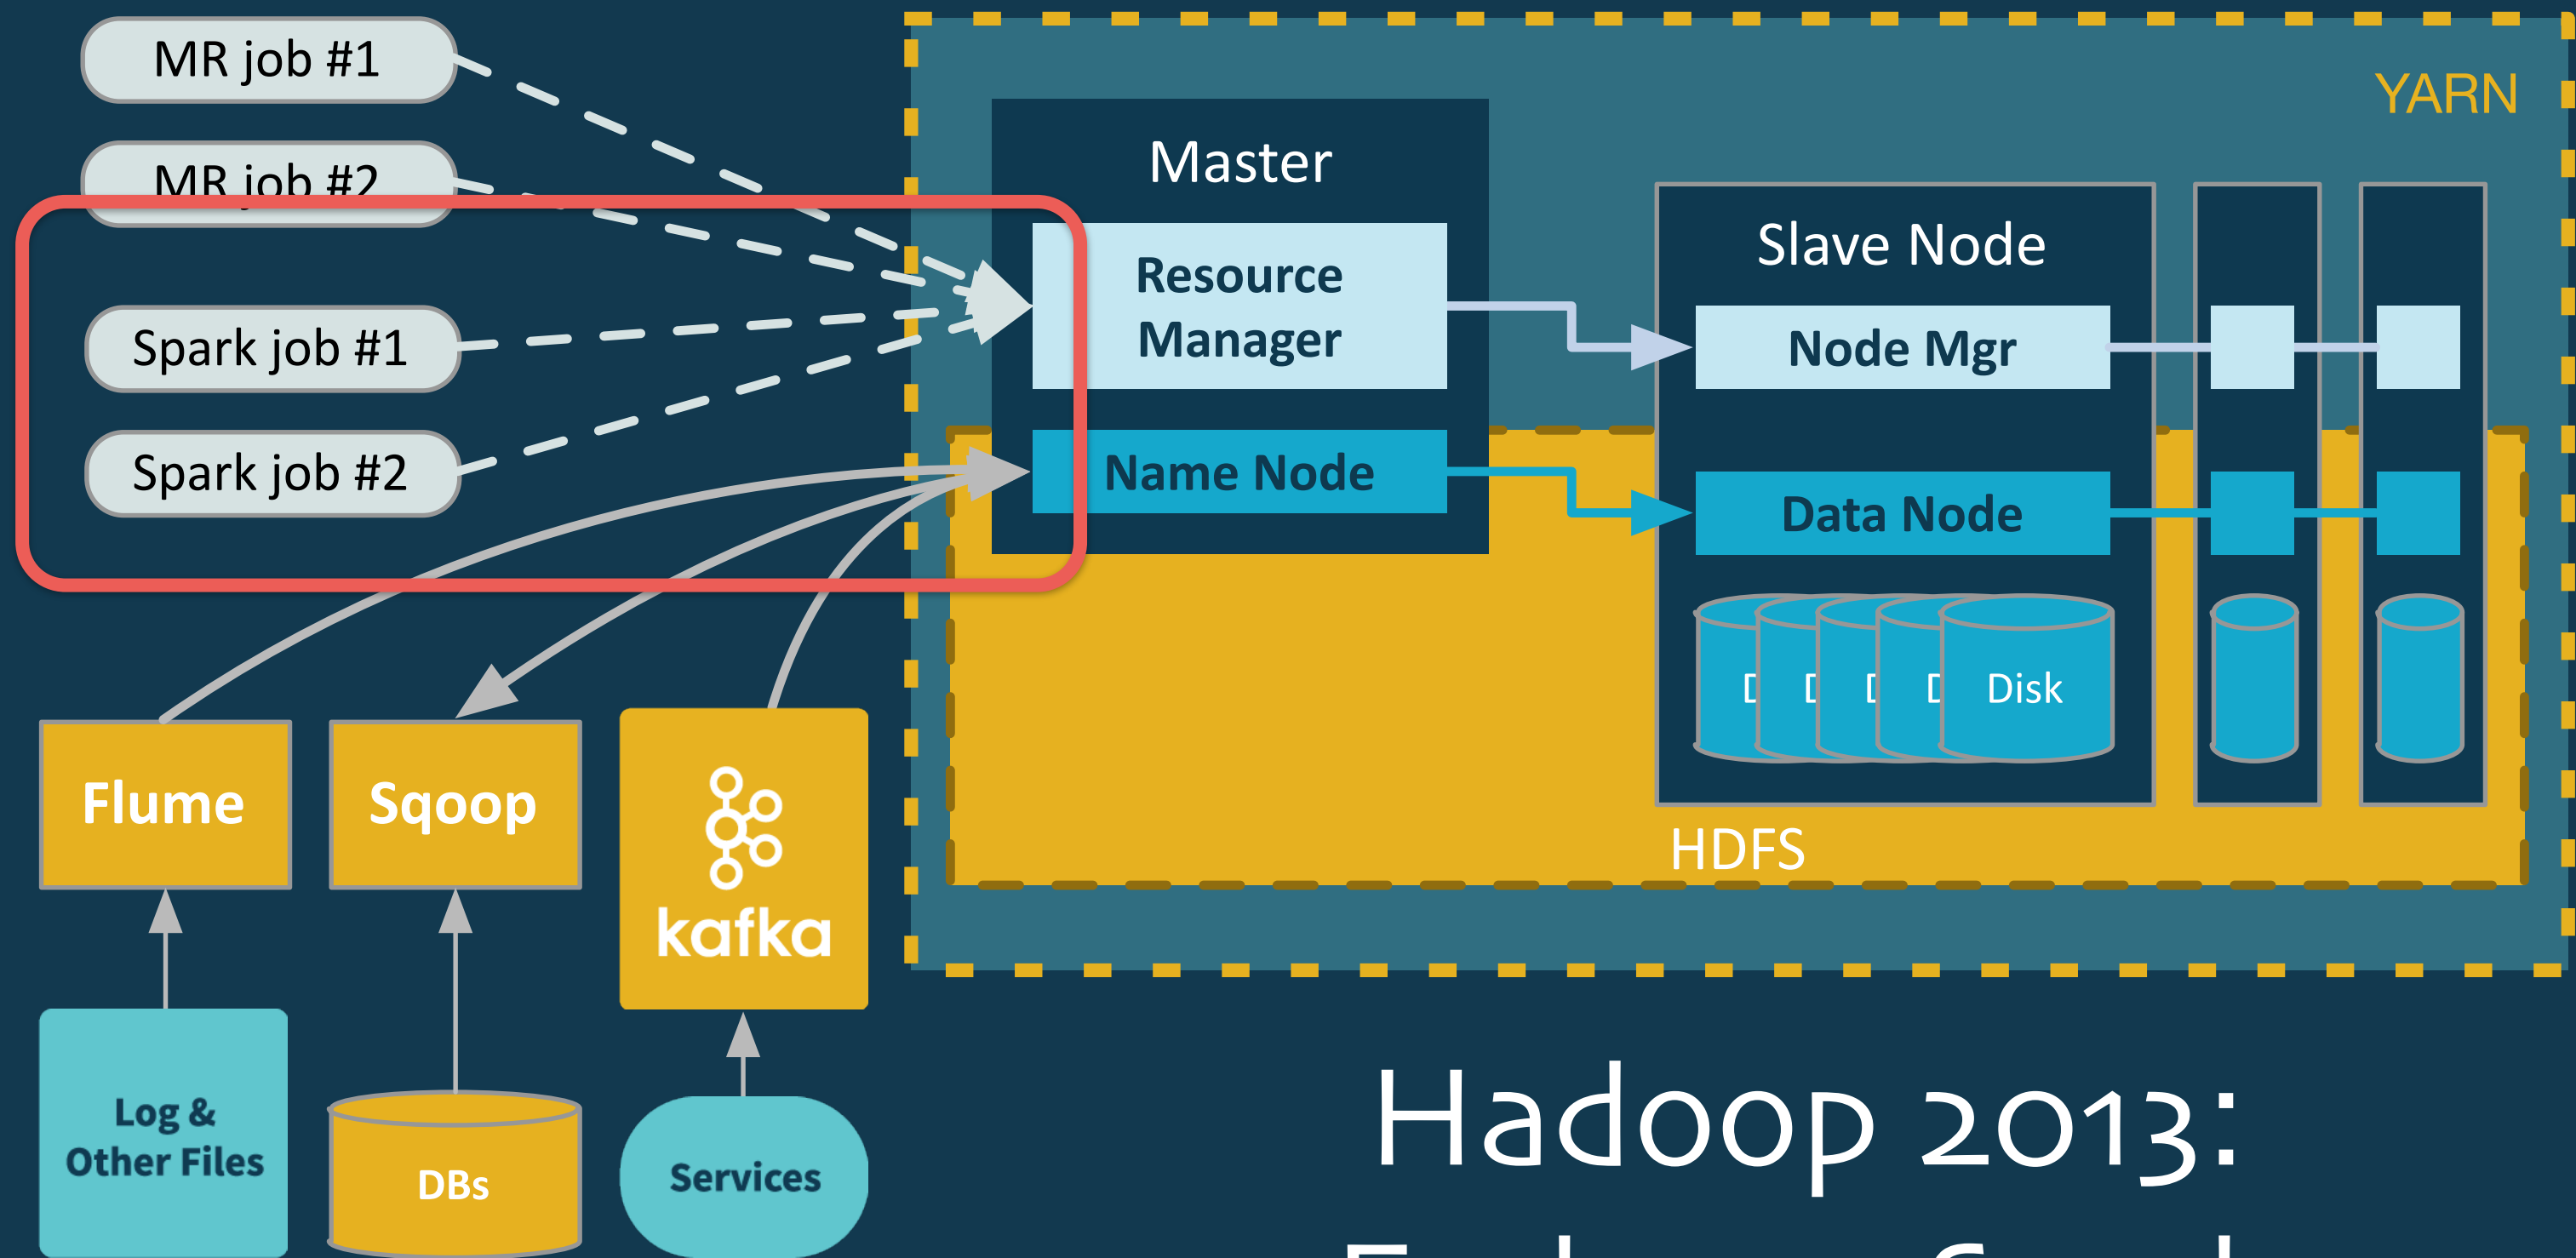
Hadoop Strengths

- Lowest CapEx system for Big Data.
- Excellent for ingesting and integrating diverse datasets.
- Flexible: from classic analytics (aggregations and data warehousing) to machine learning.

Hadoop Weaknesses

- Complex administration.
- YARN can't manage all distributed services.
- MapReduce:
 - Has poor performance.
 - A difficult programming model.
 - Doesn't support stream processing.

Why Apache Spark?




Hadoop 2013: Embrace Spark

Spark vs. MapReduce Performance

Iterative algorithm used in machine learning

100x better for many algorithms.



Spark: Major Performance Improvements

2013 Record:
Hadoop

2100 machines

72 minutes



2014 Record:
Spark

207 machines

23 minutes




Sort 100TB



One of the Fastest Growing OS Projects

Contributors / Month to Spark



Modules

SQL/DataFrames
(Structured Data)

Spark Streaming
(~Real Time)

MLlib
(Machine Learning)

GraphX
(Graphs)

Spark RDD
(Core)



The Core - Resilient Distributed Datasets

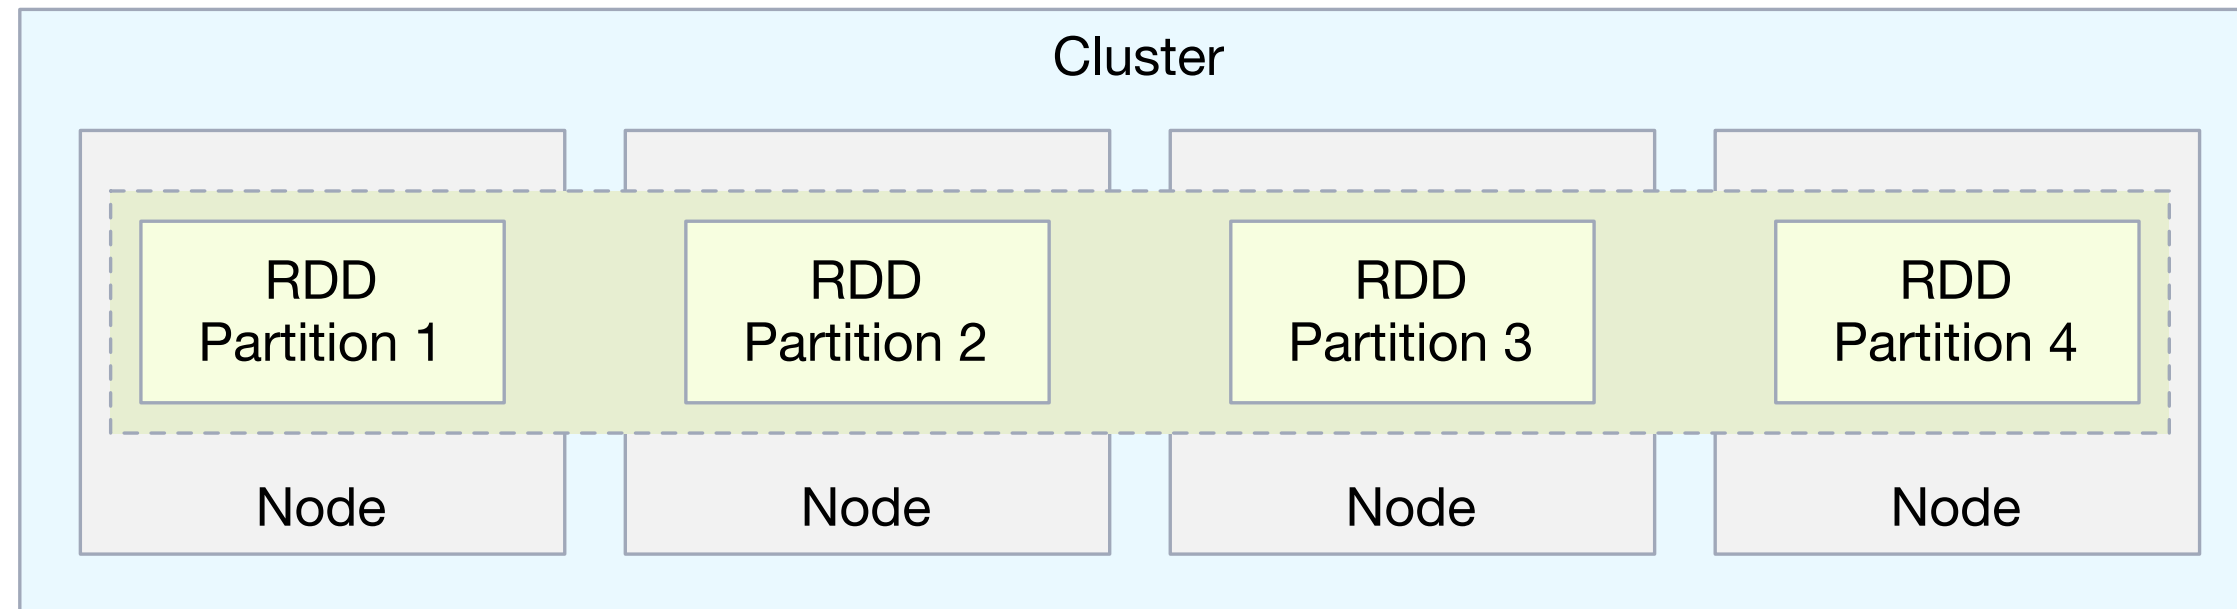
SQL/DataFrames
(Structured Data)

Spark Streaming
(~Real Time)

MLlib
(Machine Learning)

GraphX
(Graphs)

Spark RDD
(Core)



"Inverted Index" in Spark

textFile

map

flatMap

reduceByKey

map

groupByKey

map

saveAsTextFile

```
sparkContext.textFile("/path/to/input")
  .map { line =>
    val array = line.split(",", 2)
    (array(0), array(1))
  }.flatMap {
    case (id, contents) => toWords(contents).map(w=>((w, id), 1))
  }.reduceByKey {
    (count1, count2) => count1 + count2
  }.map {
    case ((word, path), n) => (word, (path, n))
  }.groupByKey
  .map {
    case (word, list) => (word, sortByCount(list))
  }.saveAsTextFile("/path/to/output")
```

SQL queries and a “DataFrame” DSL

SQL/DataFrames
(Structured Data)

Spark Streaming
(~Real Time)

MLlib
(Machine Learning)

GraphX
(Graphs)

Spark RDD
(Core)

- For data with a fixed schema...
 - Write SQL queries (currently a subset of HiveQL).
 - Use equivalent Python-inspired **DataFrame** API.



Use SQL or the Idiomatic DataFrame API

```
# SQL:
sqlContext.sql("""
    SELECT state, age, COUNT(*) AS cnt
    FROM people
    GROUP BY state, age
    ORDER BY cnt DESC, state ASC, age ASC
    """)

// DataFrame (Scala):
people.state($"state", $"age")
    .groupBy($"state", $"age").count()
    .orderBy($"count".desc, $"state".asc, $"age".asc)
```


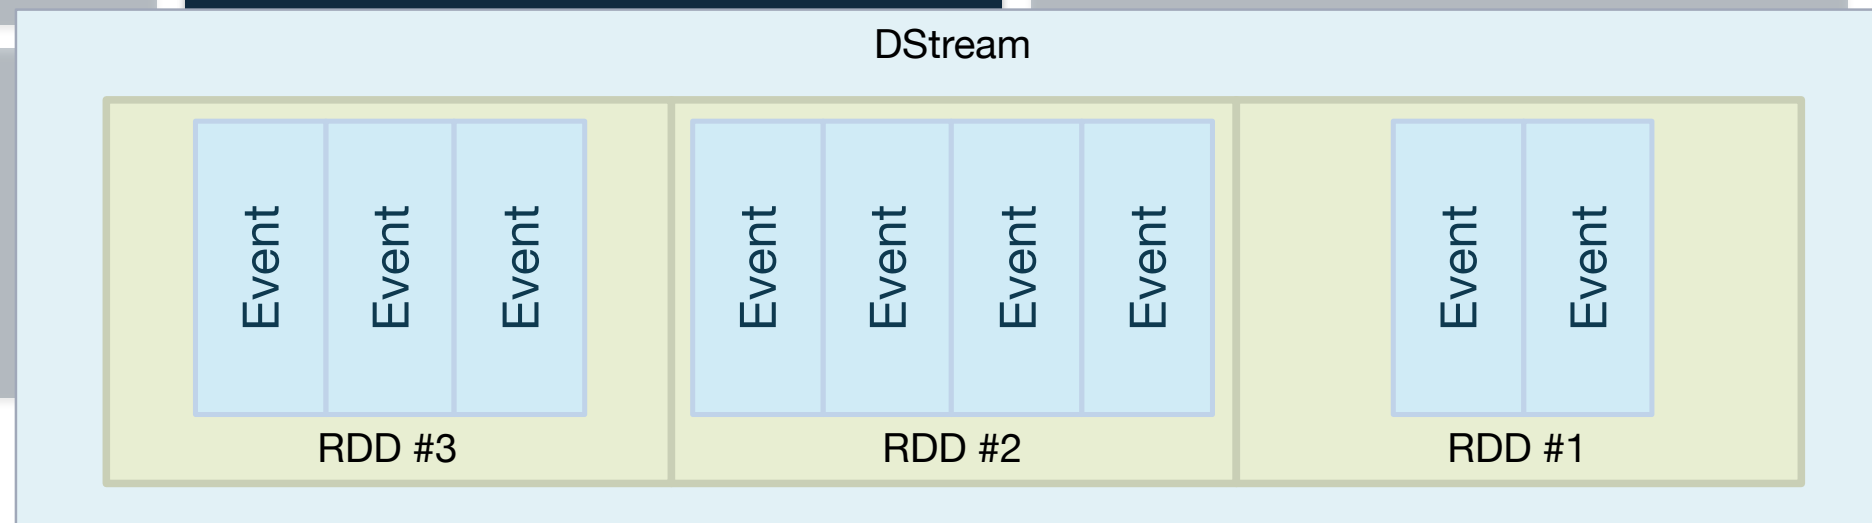

Spark Streaming: “Mini-batch” Processing

SQL/DataFrames
(Structured Data)

Spark Streaming
(~Real Time)

MLlib
(Machine Learning)

GraphX
(Graphs)



Streaming Inverted Index

```
val kafkaBrokers = "host1:port1,host2:port2,..."
val kafkaTopics  = Set("topic1", "topic2", ...)

val sparkConf = new SparkConf().setAppName("...")
val ssc = new StreamingContext(sparkConf, Seconds(2))

// Create direct kafka stream with kafkaBrokers and kafkaTopics
val kafkaParams = Map[String, String]("metadata.broker.list" -> kafkaBrokers)
val messages =
    KafkaUtils.createDirectStream[String,String,StringDecoder,StringDecoder](
        ssc, kafkaParams, kafkaTopics)

messages.flatMap {case (topic,text) => toWords(text).map(w=>((w,topic),1L))}
    .reduceByKey (_ + _)
    .map {case ((word, topic), n) => (word, (path, n))}
    .groupByKey
    .map {case (word, list) => (word, sortByCount(list))}
```



```
val ssc = new StreamingContext(sparkConf, Seconds(2))

// Create direct kafka stream with kafkaBrokers and kafkaTopics
val kafkaParams = Map[String, String]("metadata.broker.list" -> kafkaBrokers)
val messages =
    KafkaUtils.createDirectStream[String, String, StringDecoder, StringDecoder](
        ssc, kafkaParams, kafkaTopics)

messages.flatMap {case (topic, text) => toWords(text).map(w=>((w, topic), 1L))}
    .reduceByKey (_ + _)
    .map {case ((word, topic), n) => (word, (path, n))}
    .groupByKey
    .map {case (word, list) => (word, sortByCount(list))}
    .saveAsTextFiles("/path/to/output")

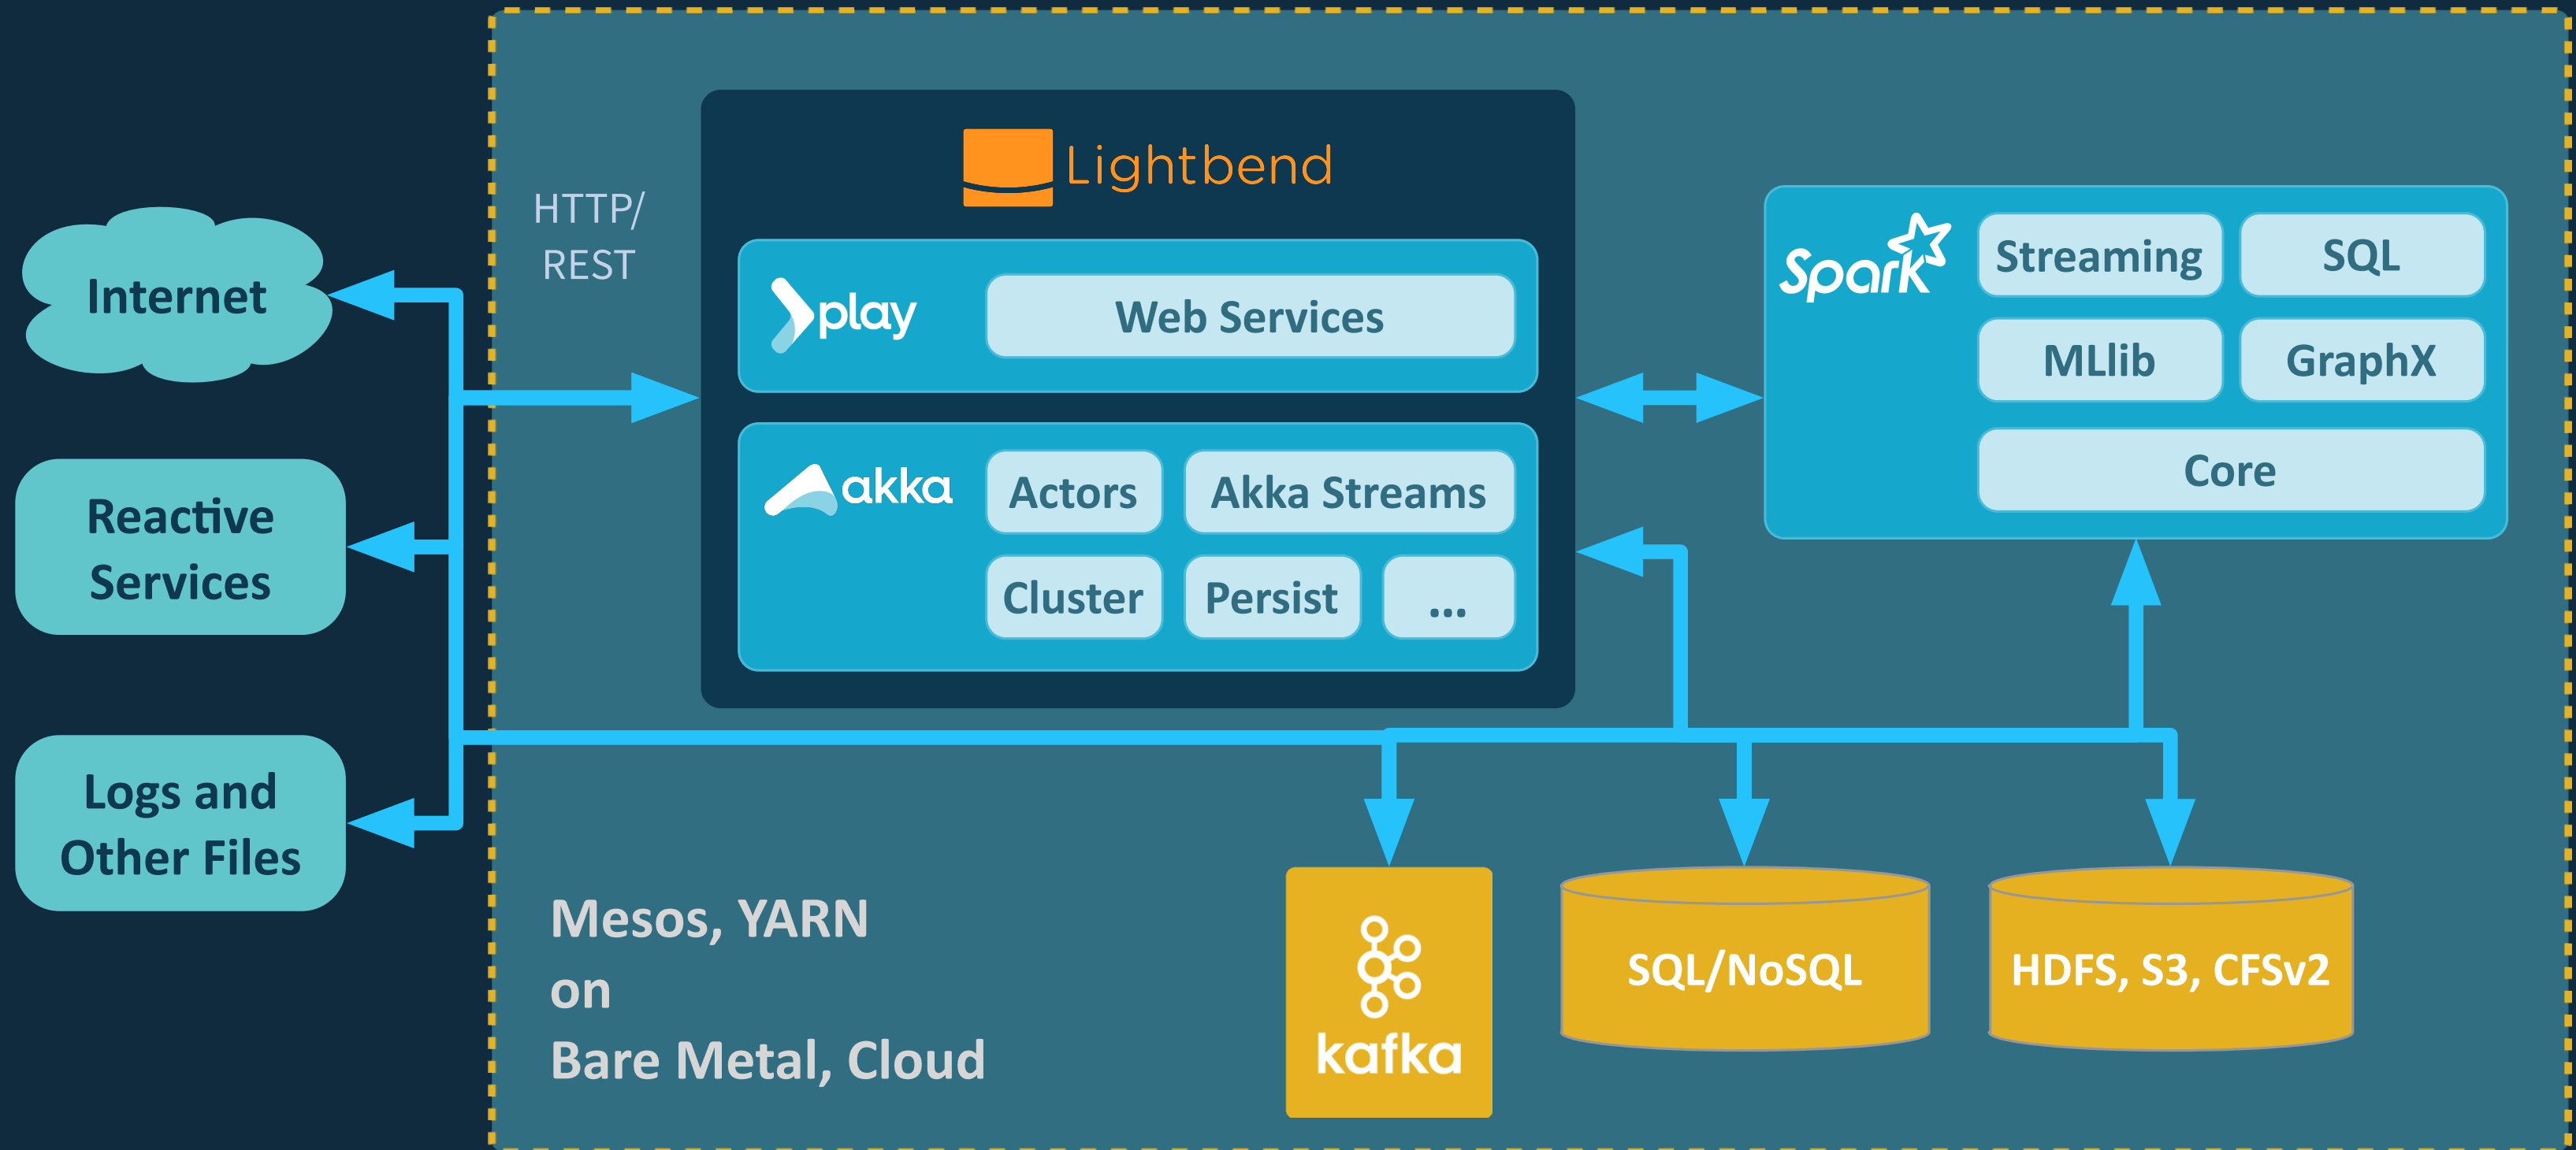
ssc.start()
ssc.awaitTermination()
```

An Architecture for Fast Data

Fast as in Streaming. Why?

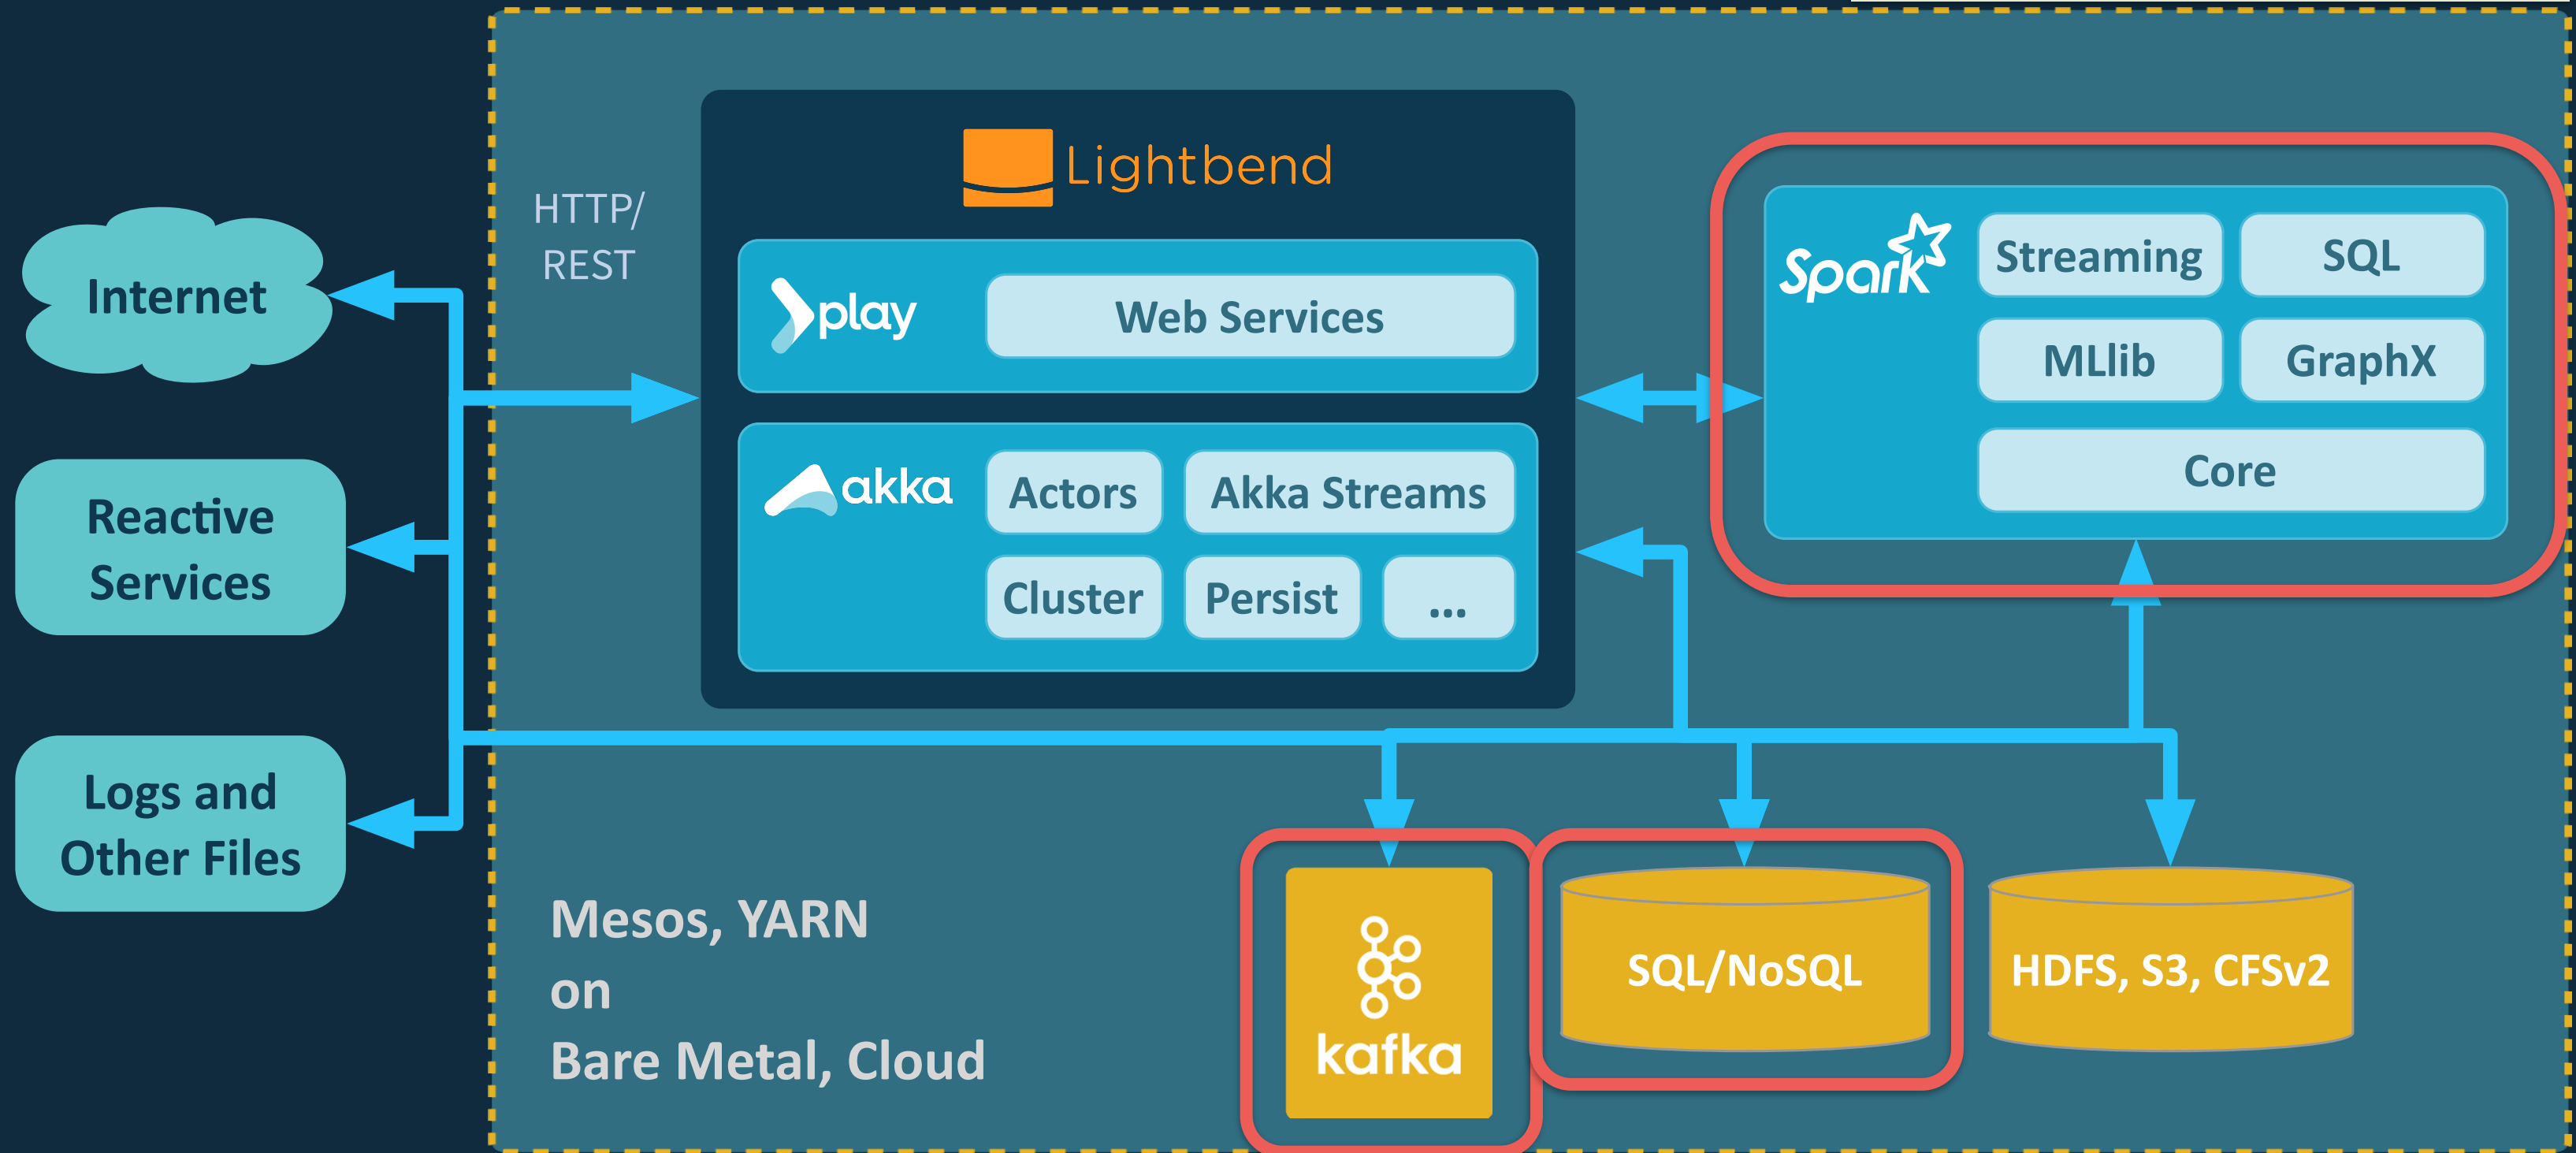
- Update a search engine in real time as web page or documents change.
- Train a SPAM filter with every email.
- Detect anomalies as they happen through processing of logs and monitoring data.

Fast Data Architecture



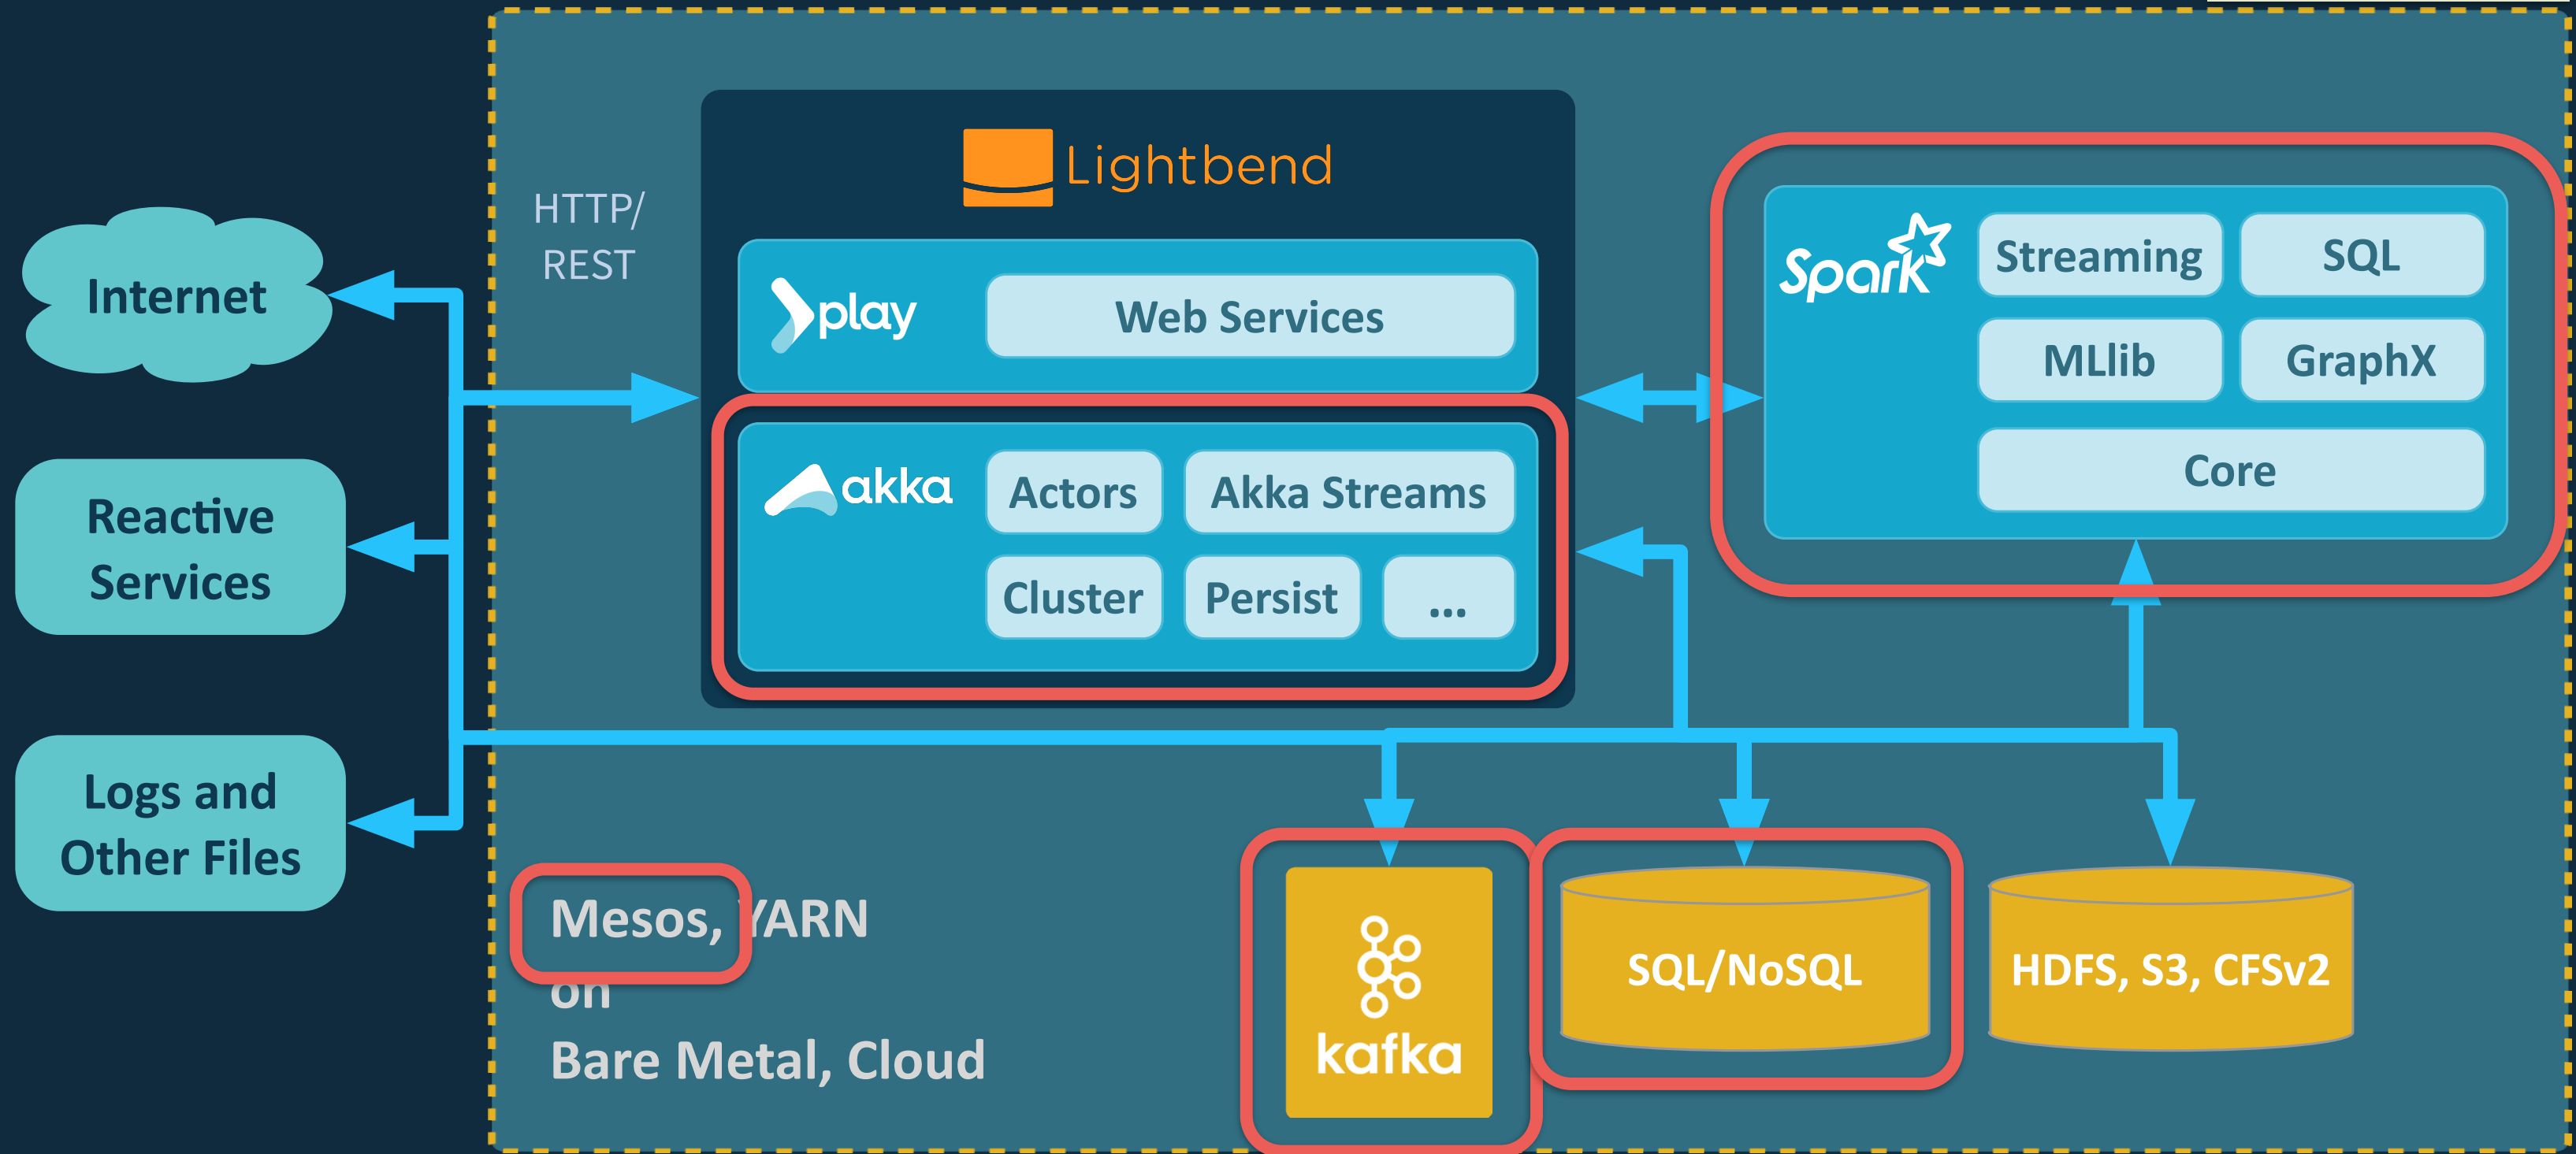
Fast Data Architecture

Core of Spark, Kafka,
and Cassandra



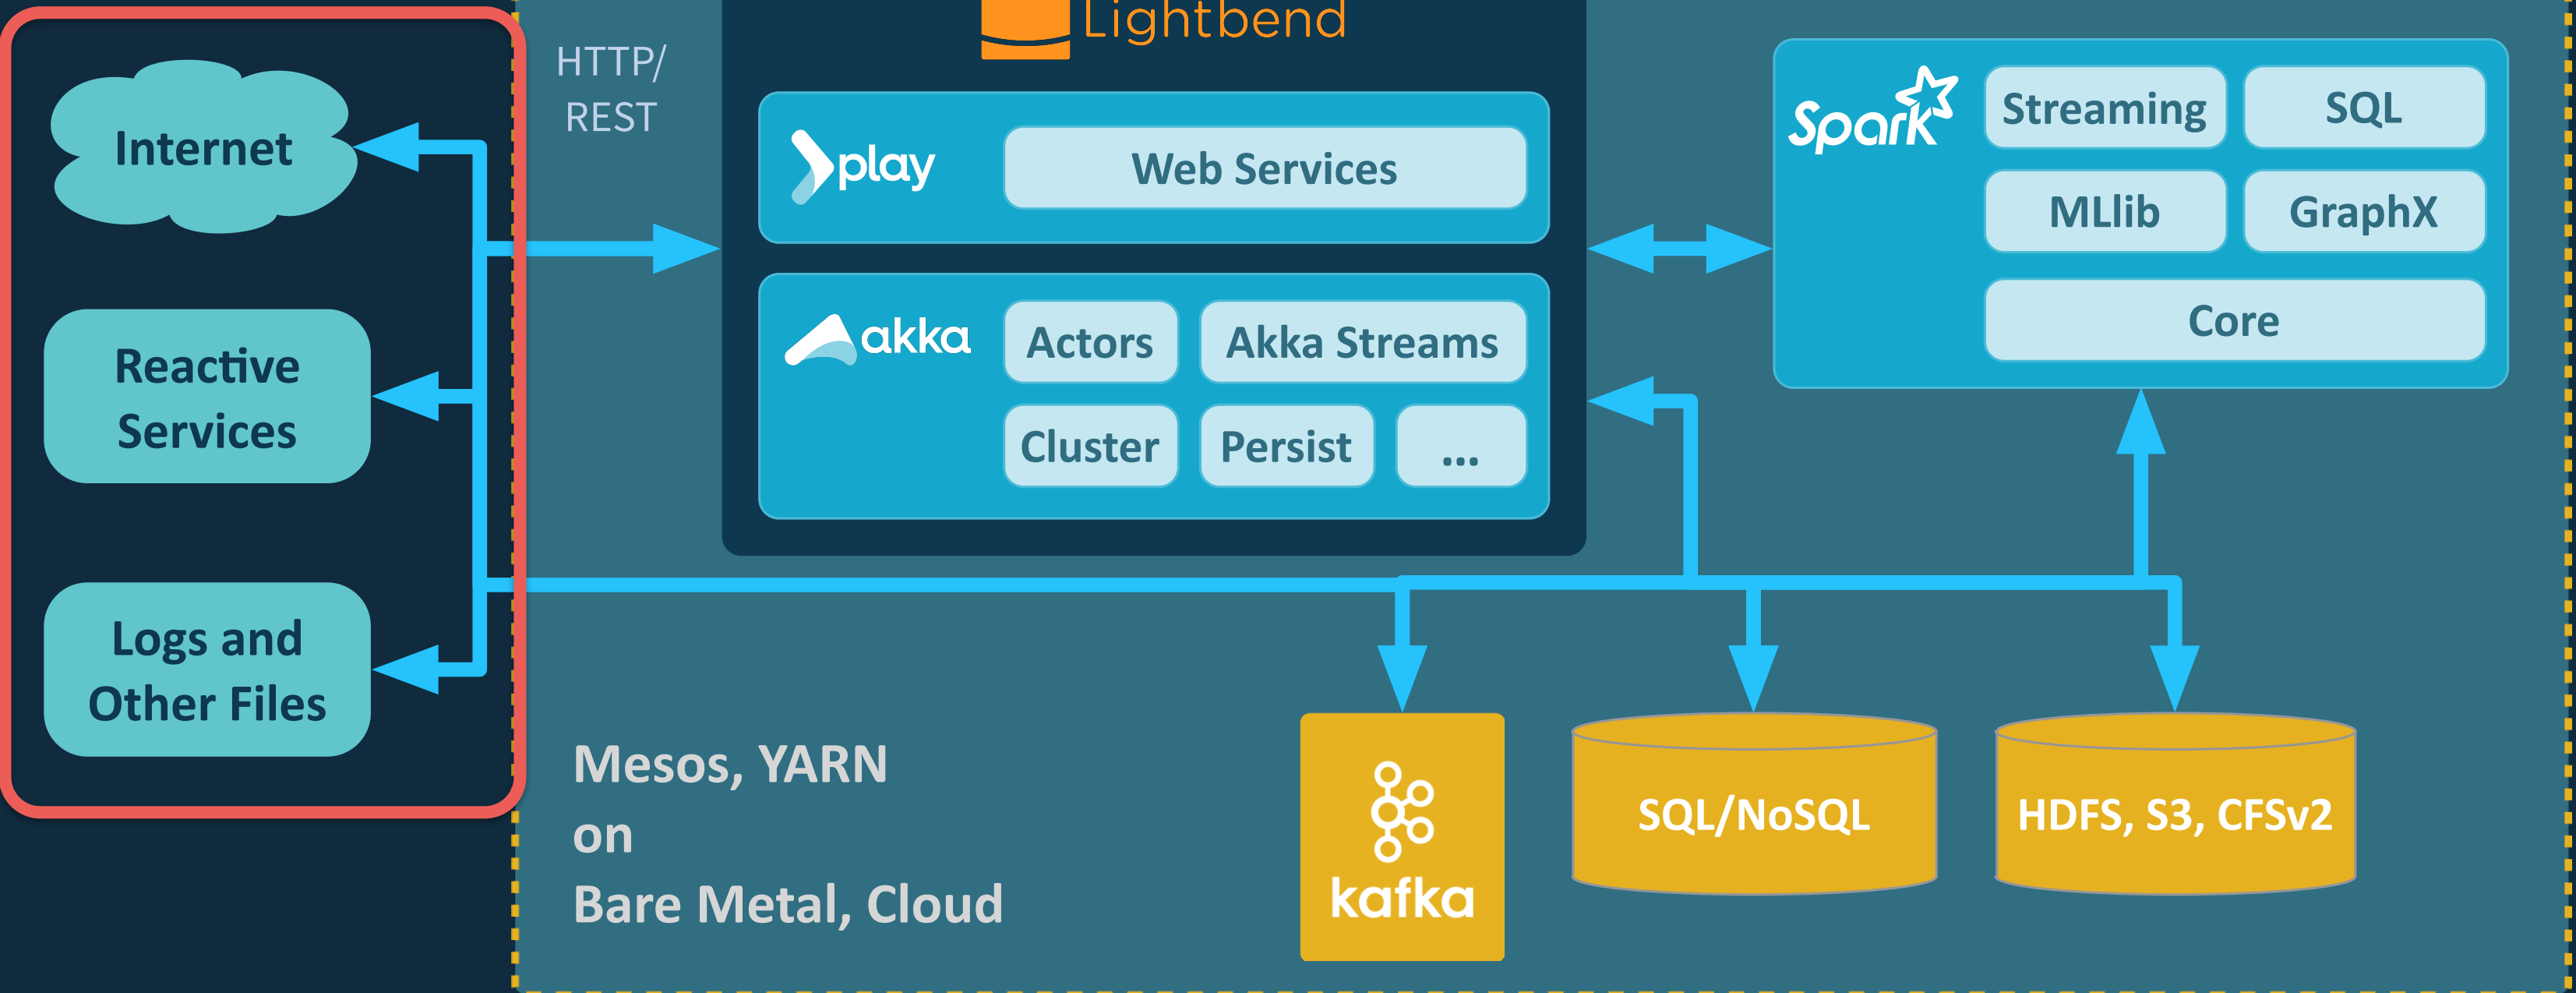
Fast Data Architecture

“SMACK”
Stack

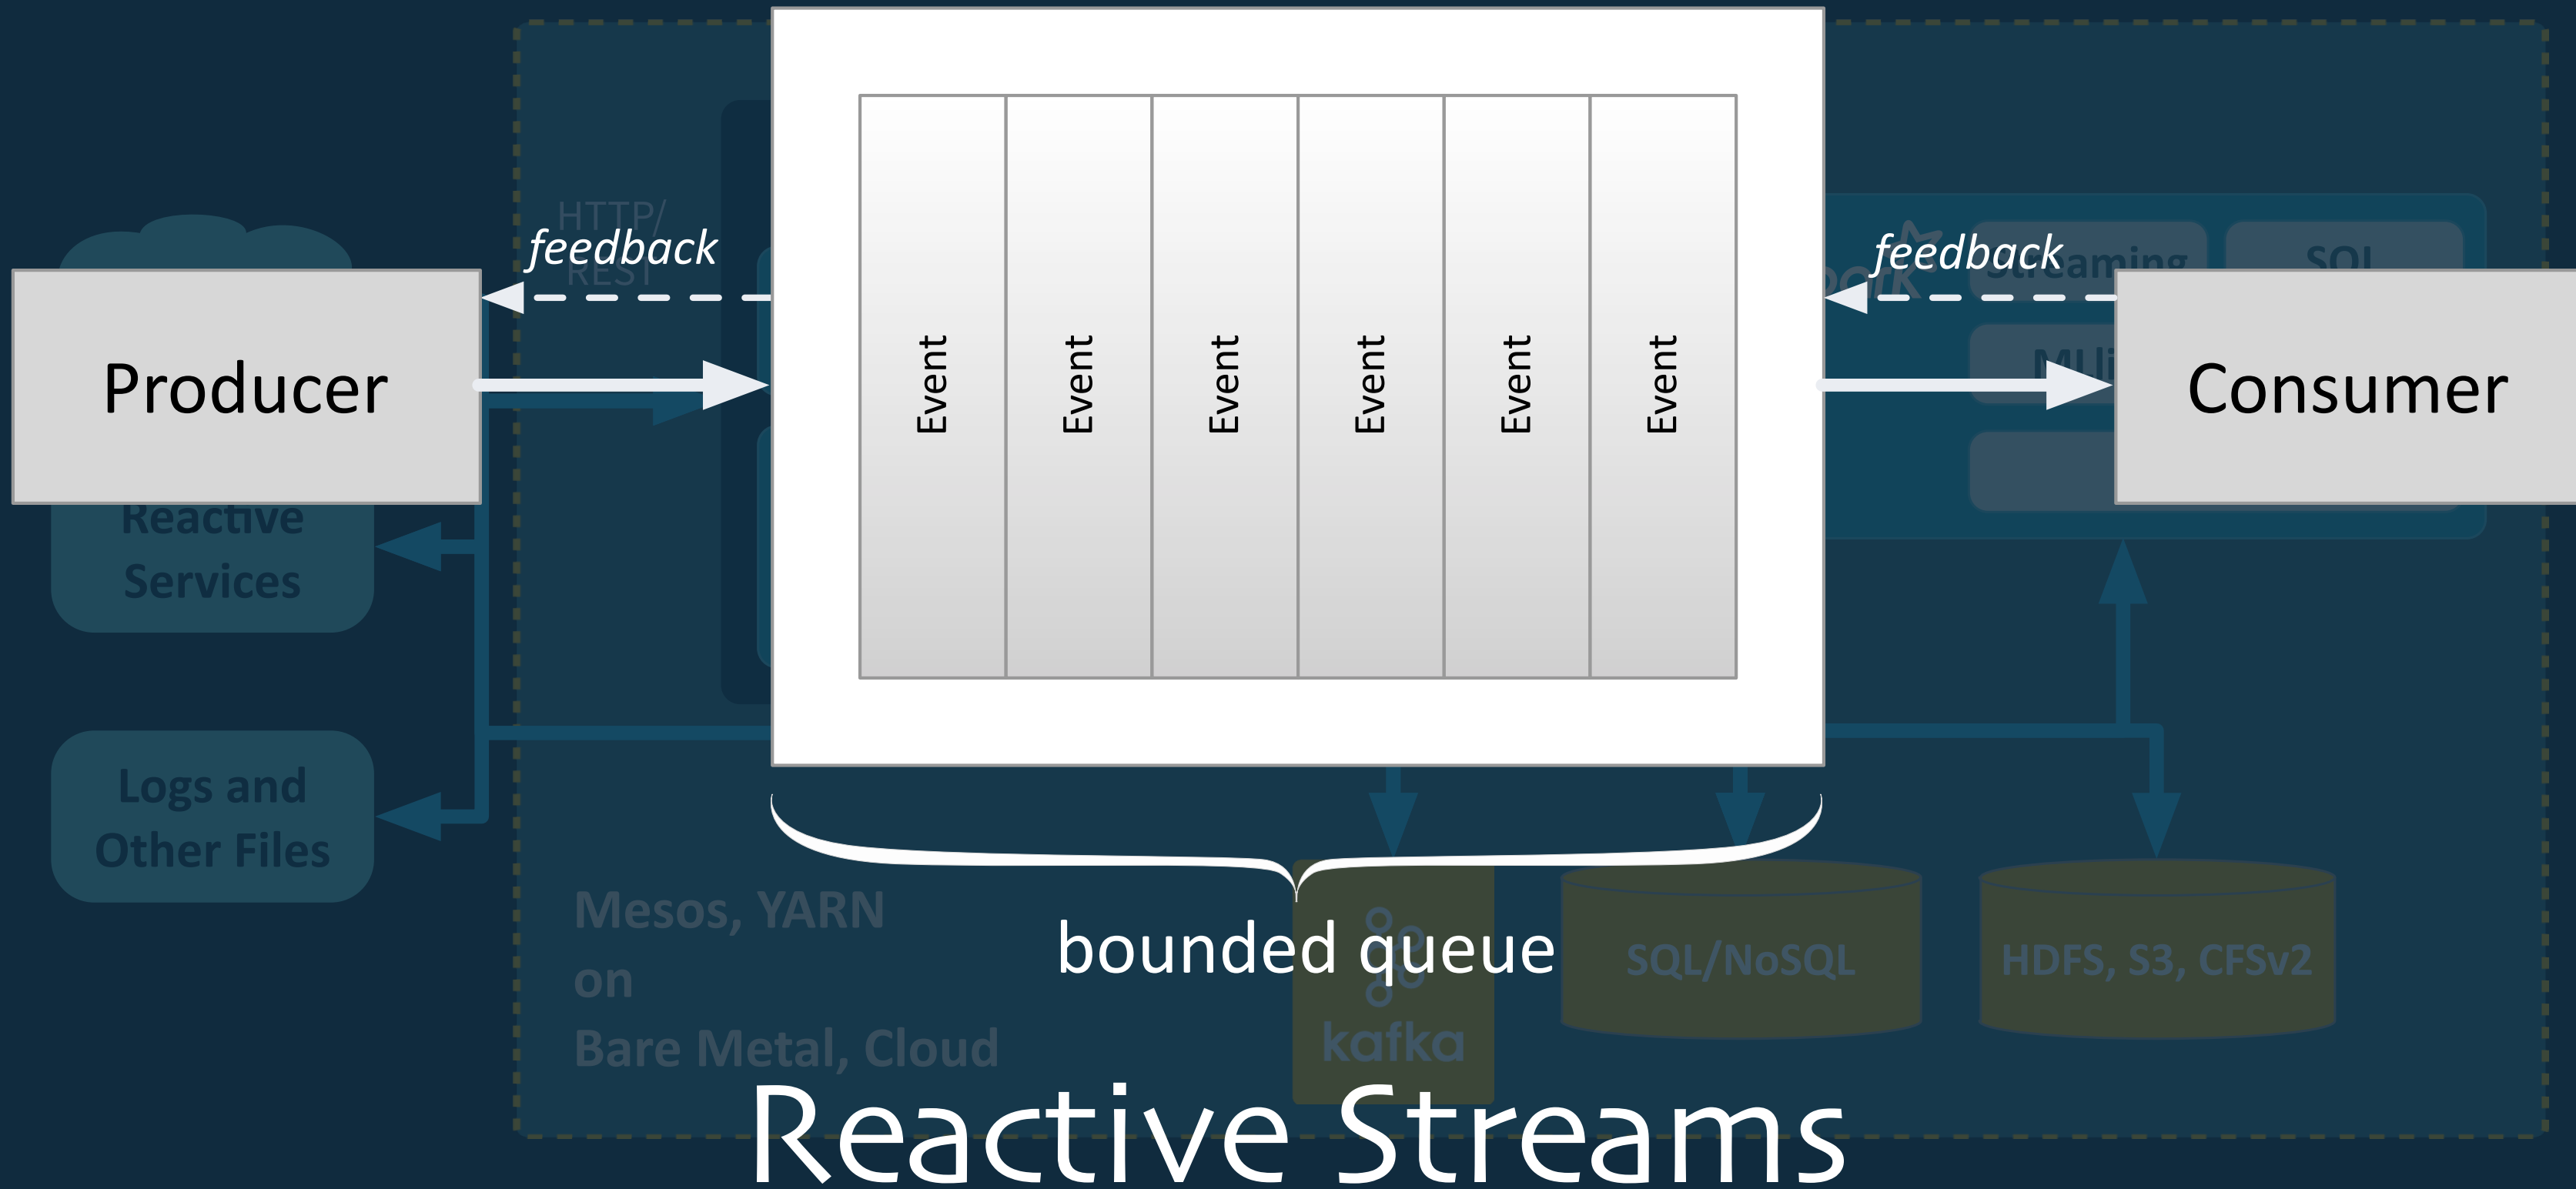


Fast Data Architecture

Data
Sources

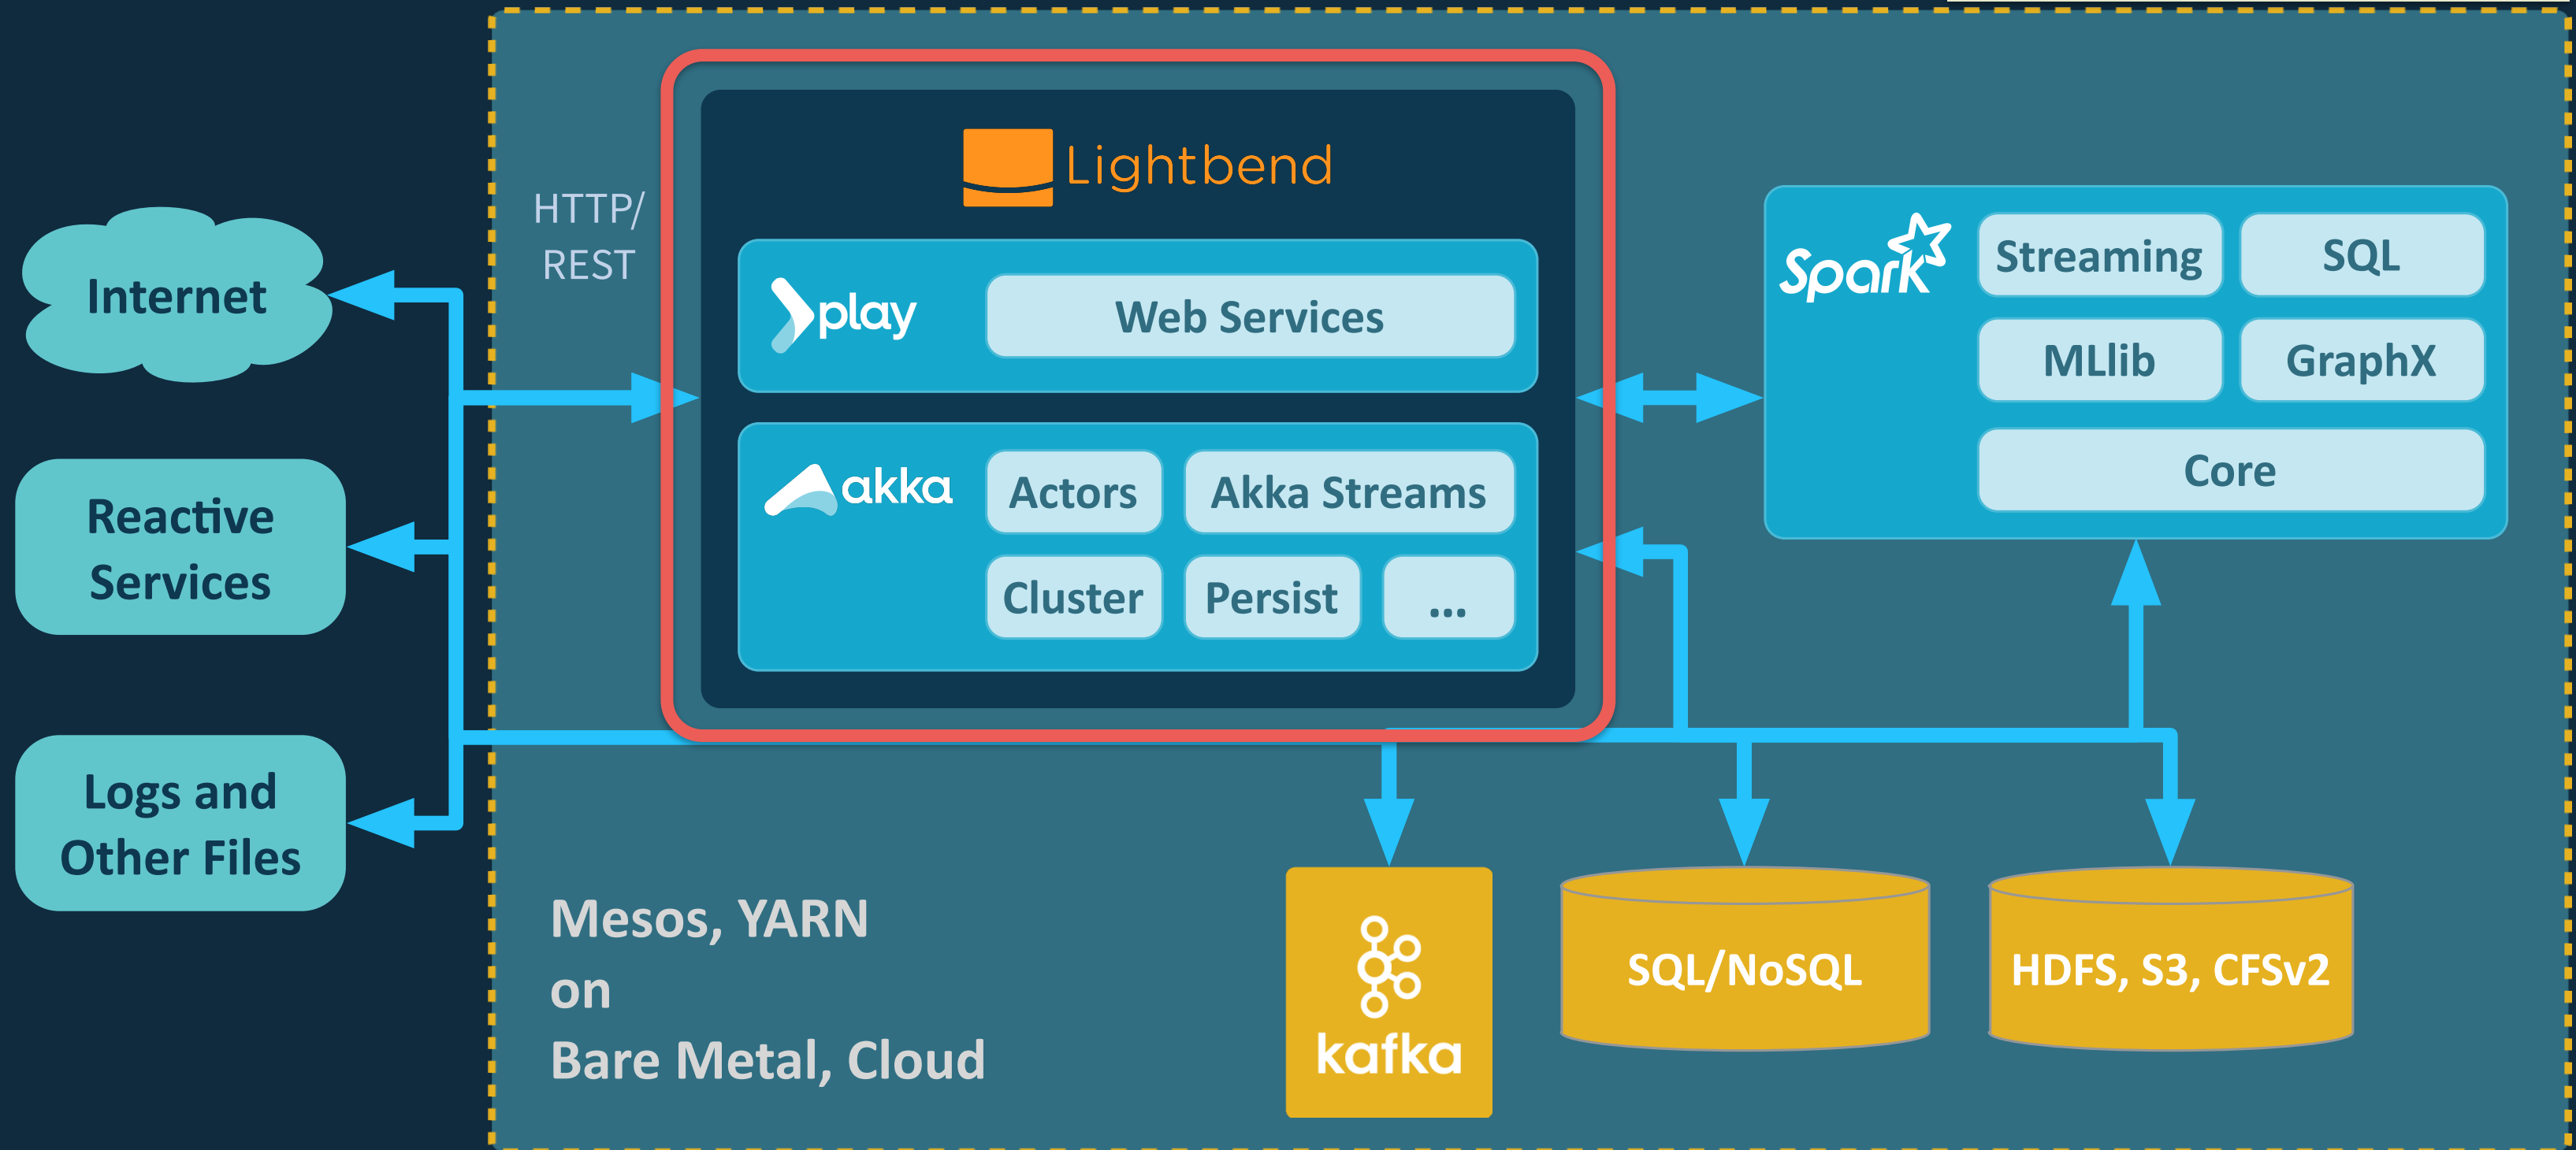


Fast Data Architecture



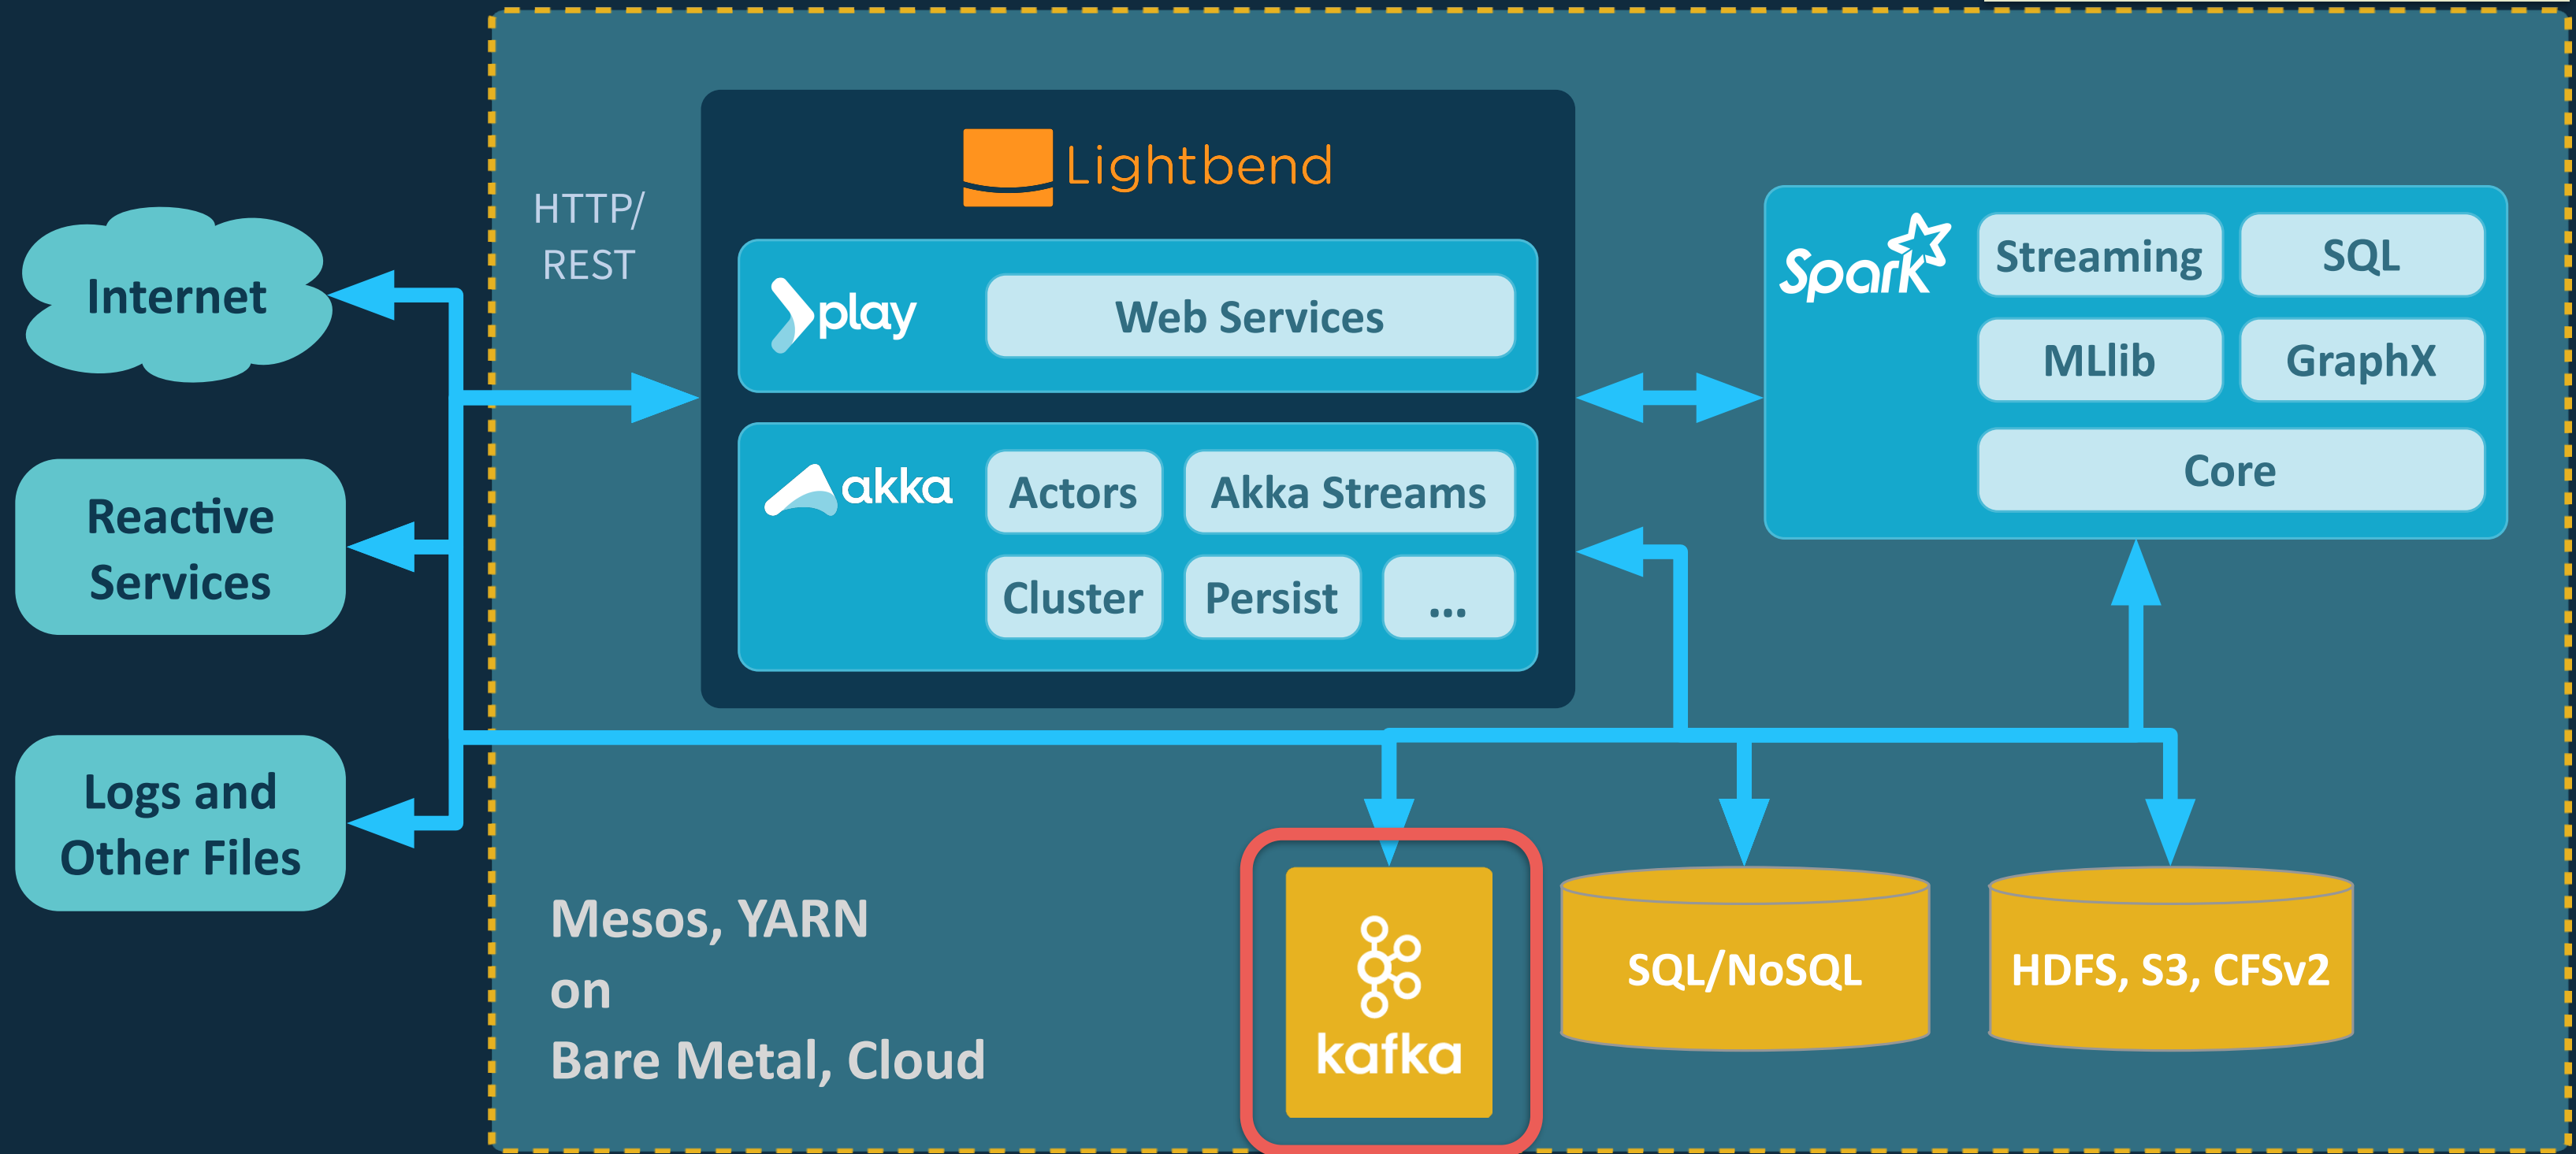
Fast Data Architecture

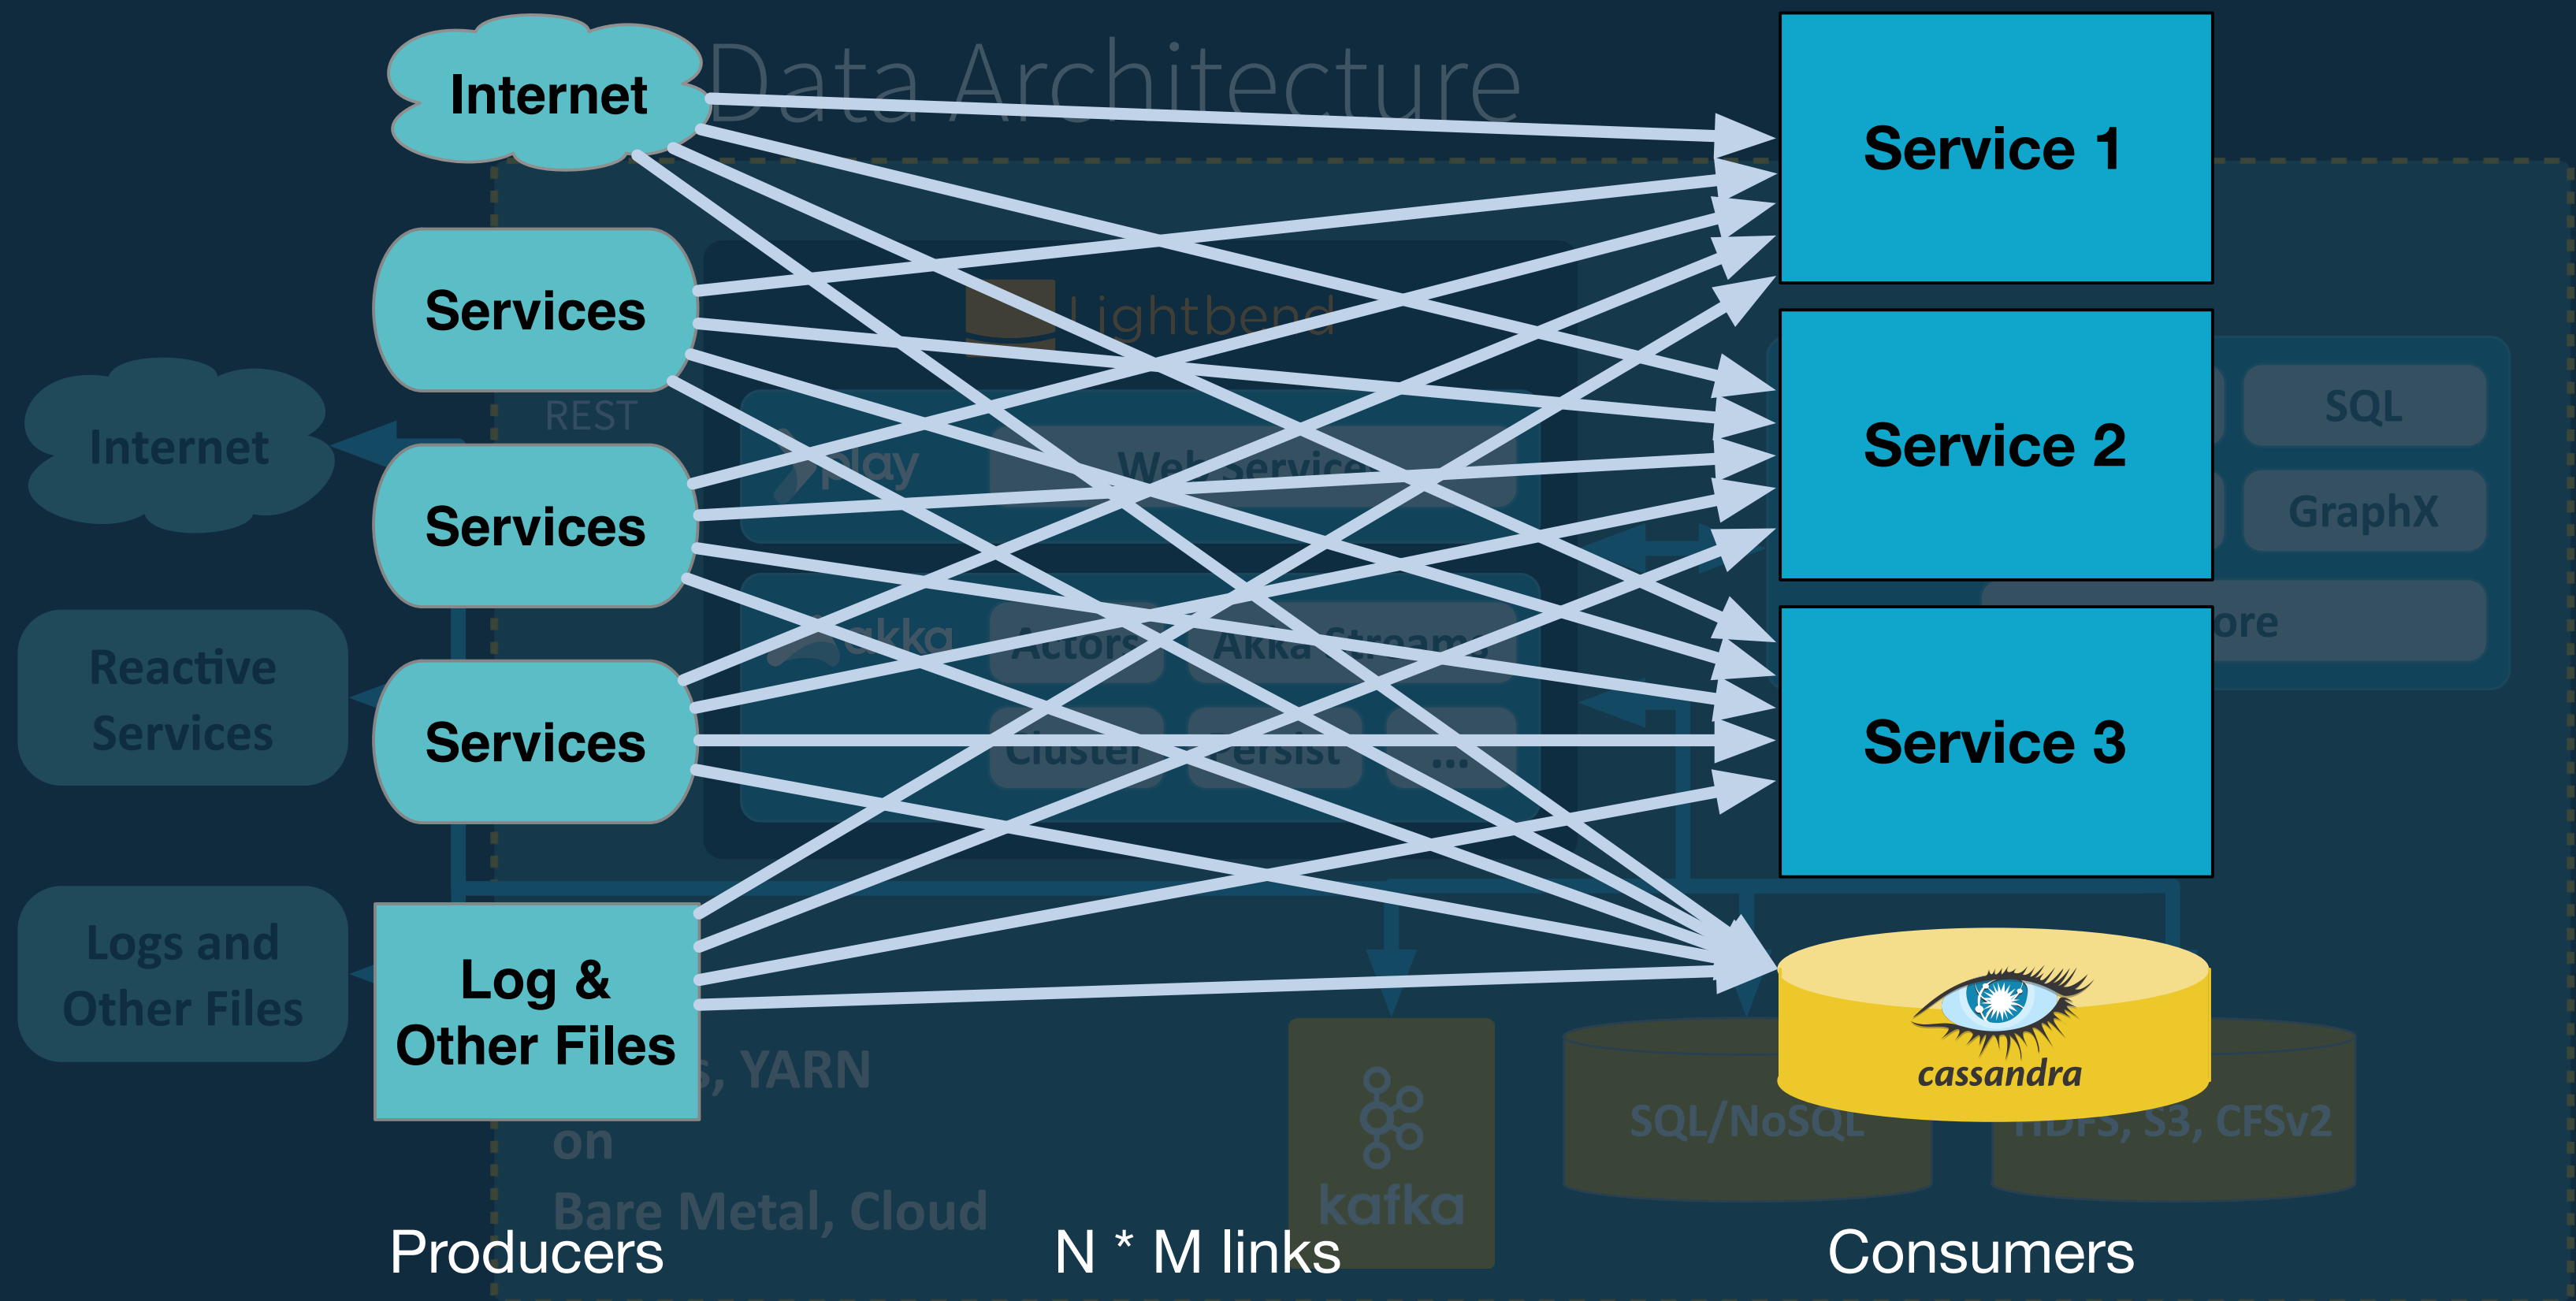
Lightbend Reactive Platform



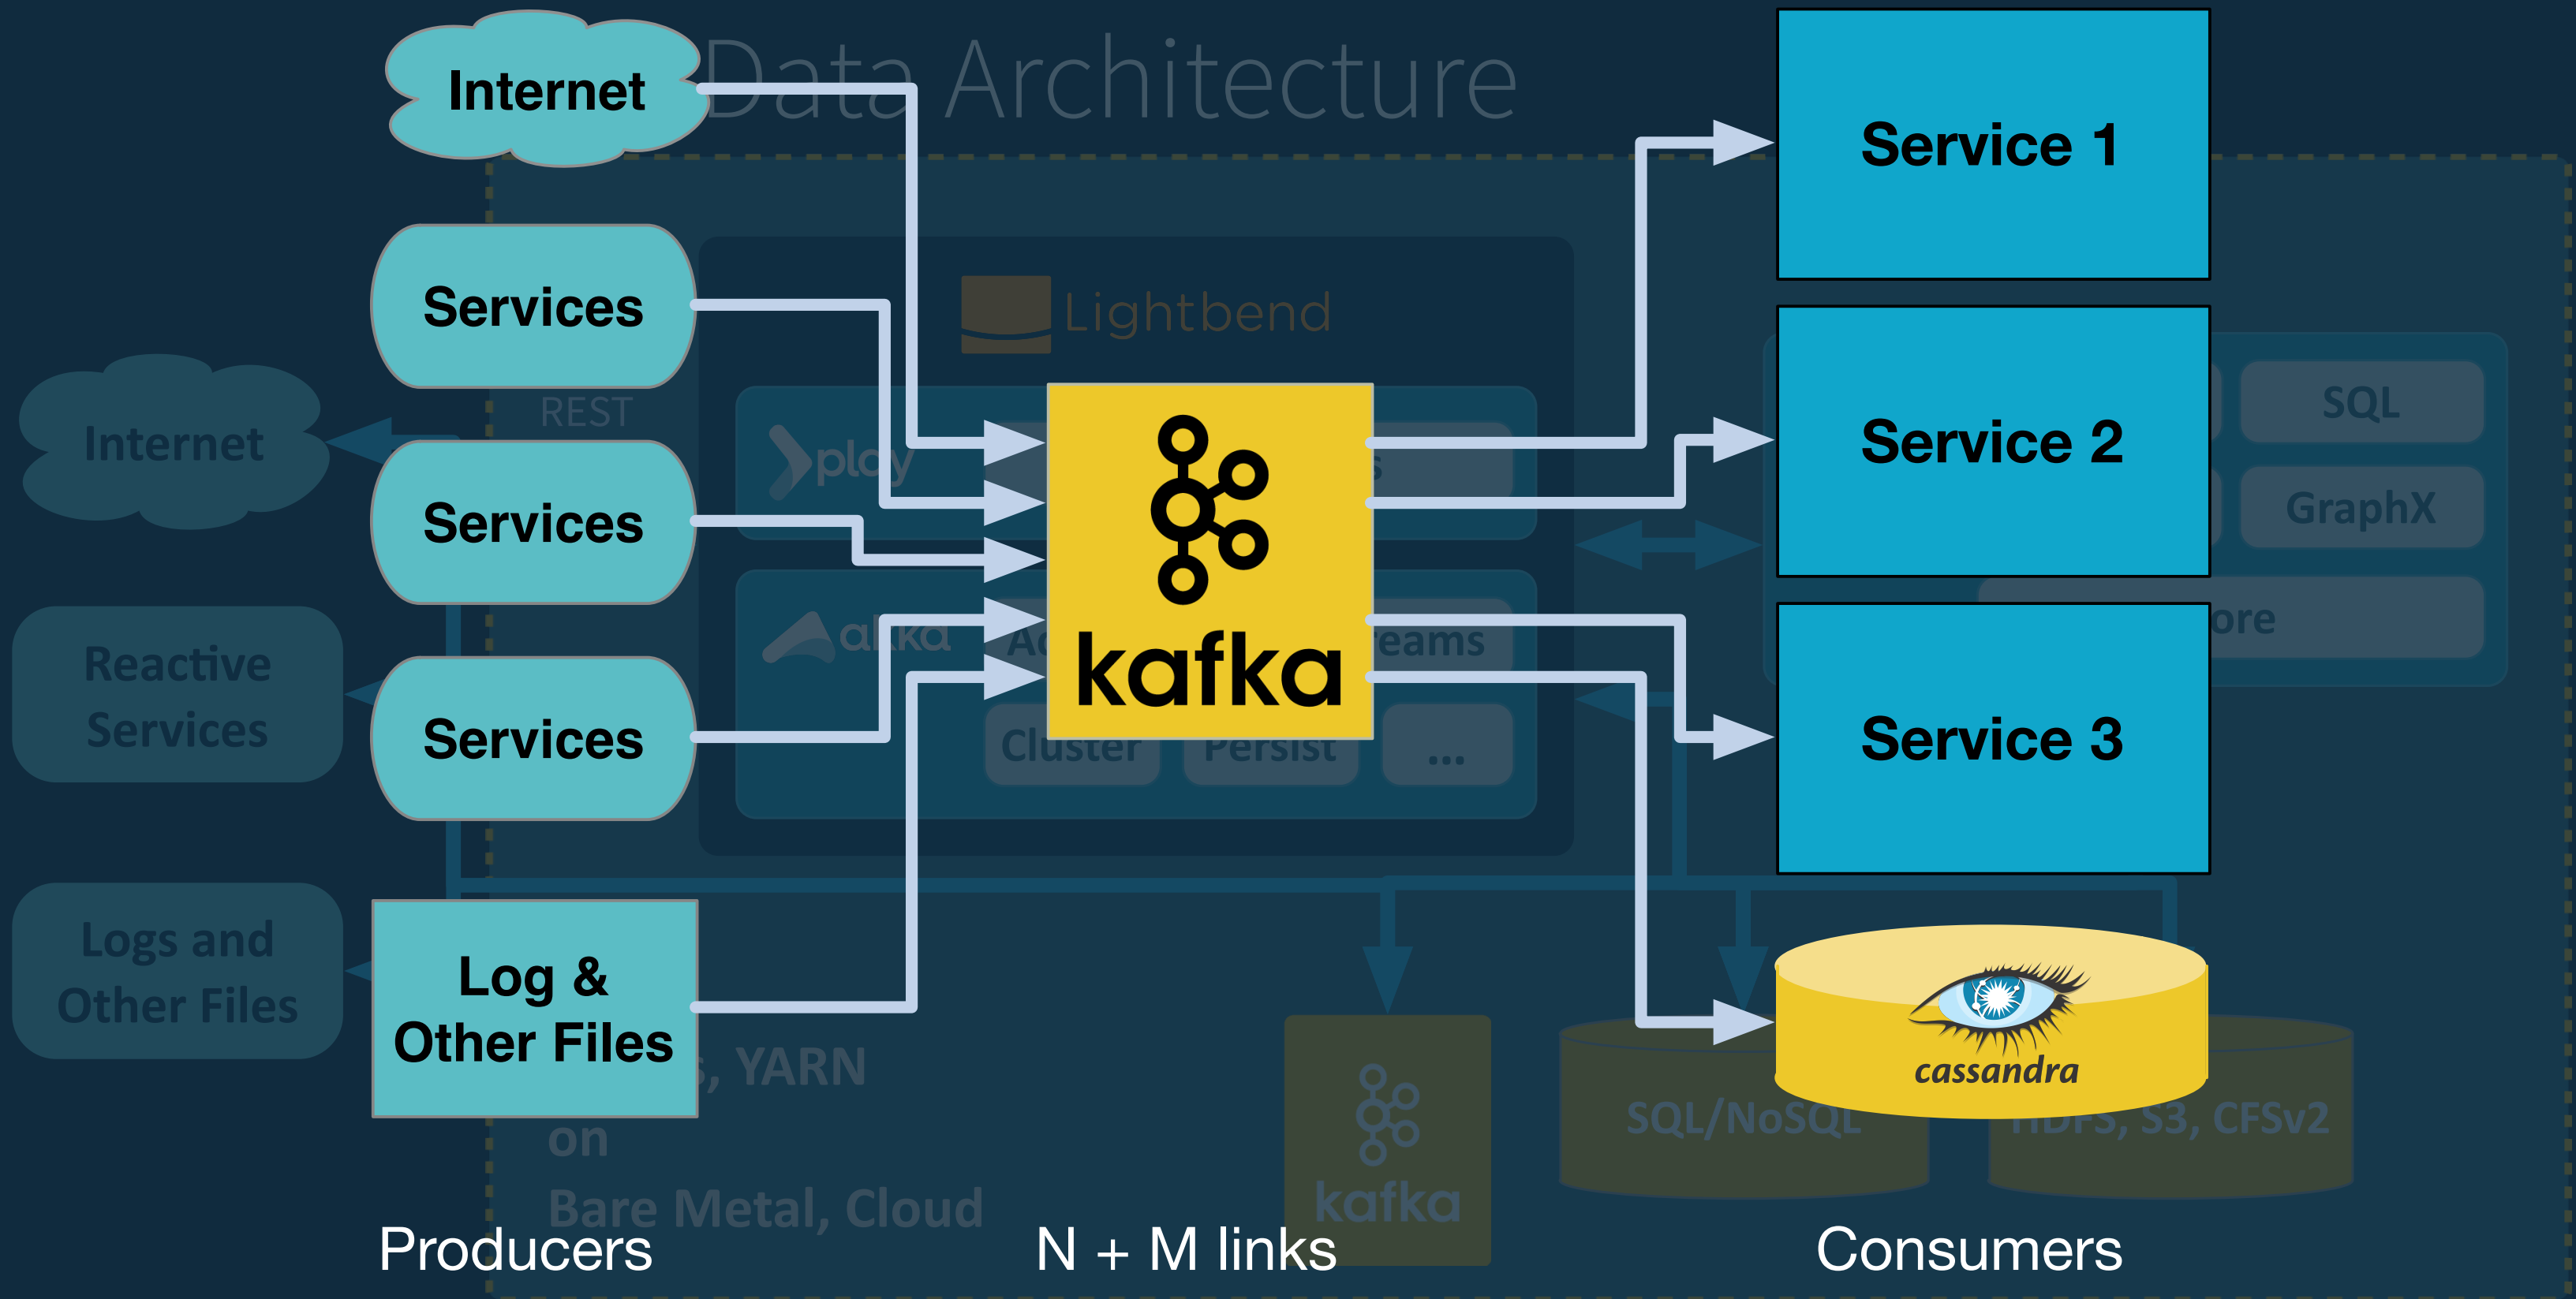
Fast Data Architecture

Kafka for Stream
Storage



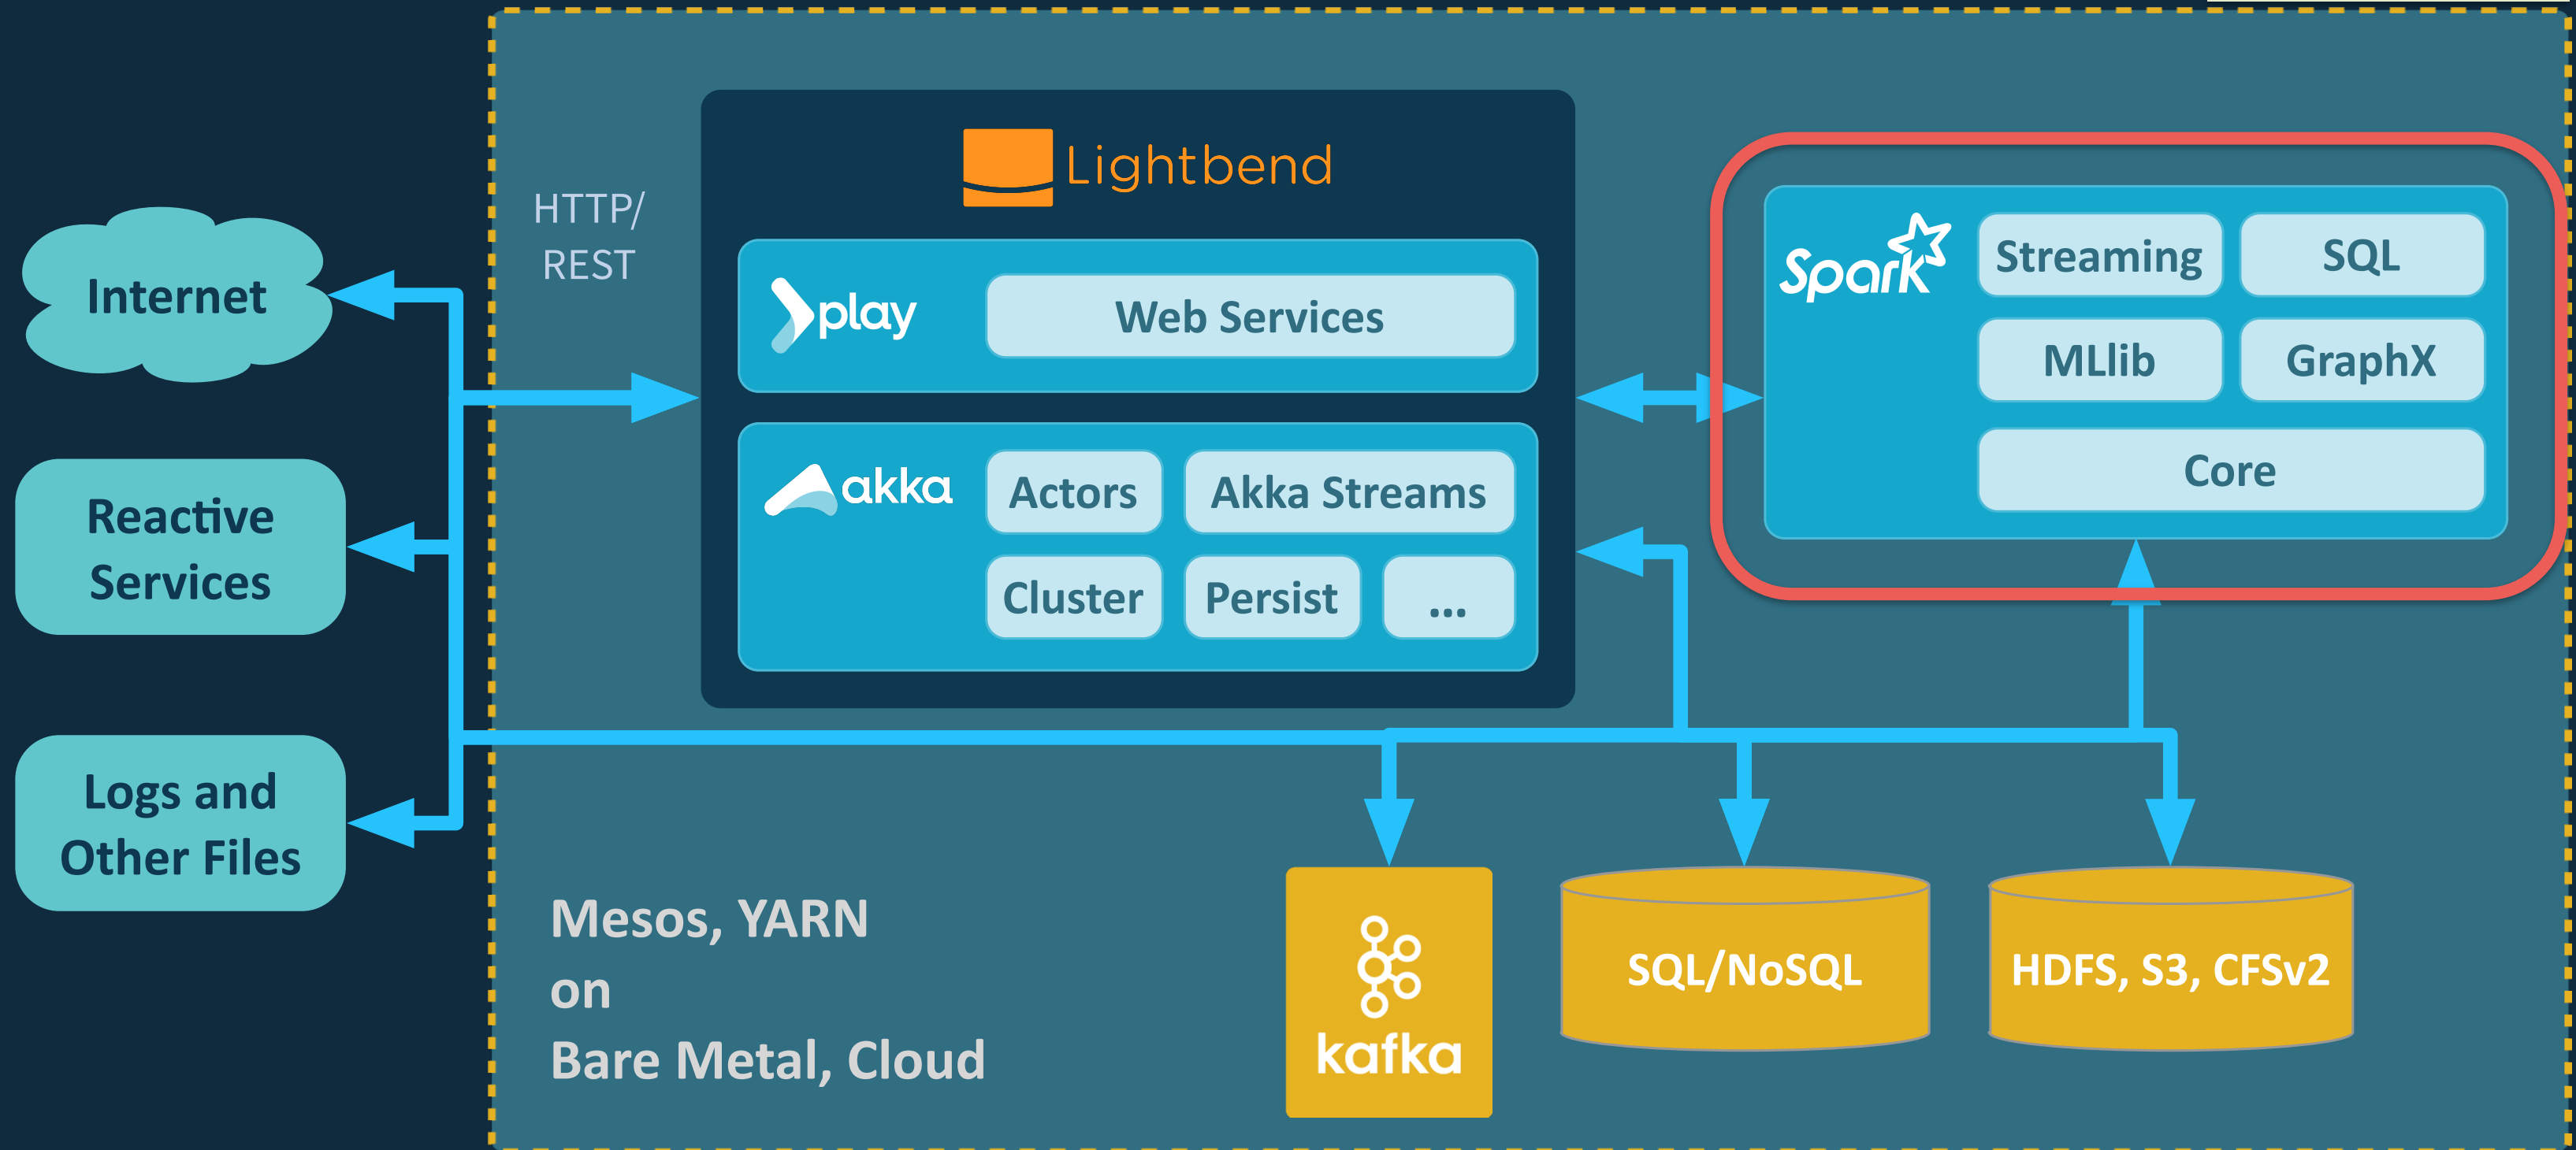


Data Architecture



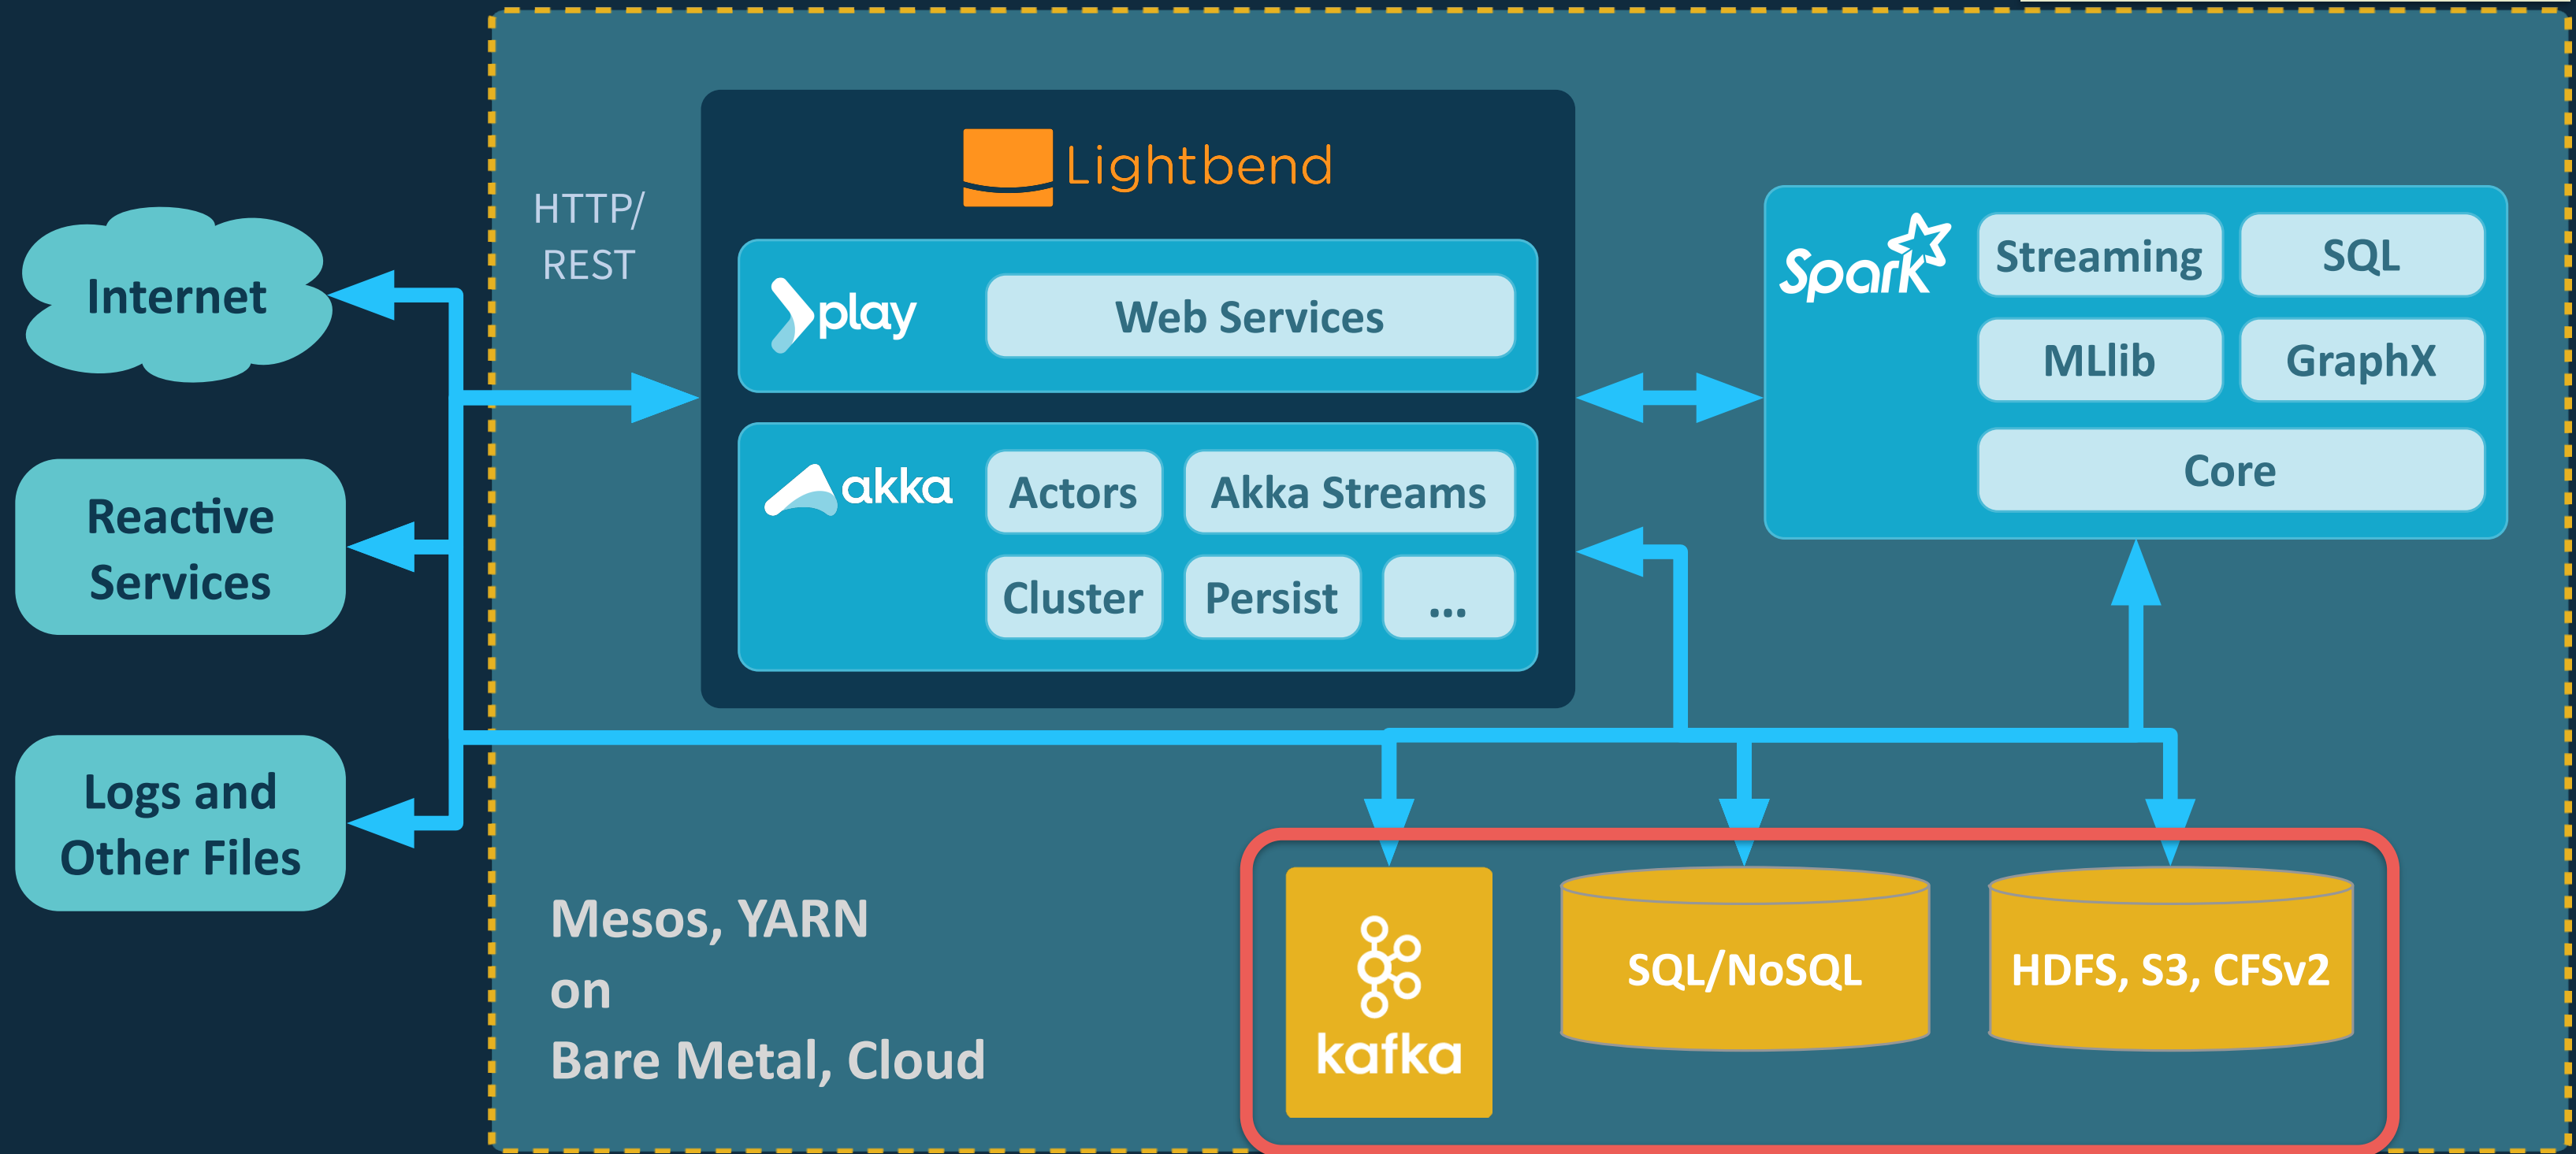
Fast Data Architecture

Minibatch
Processing



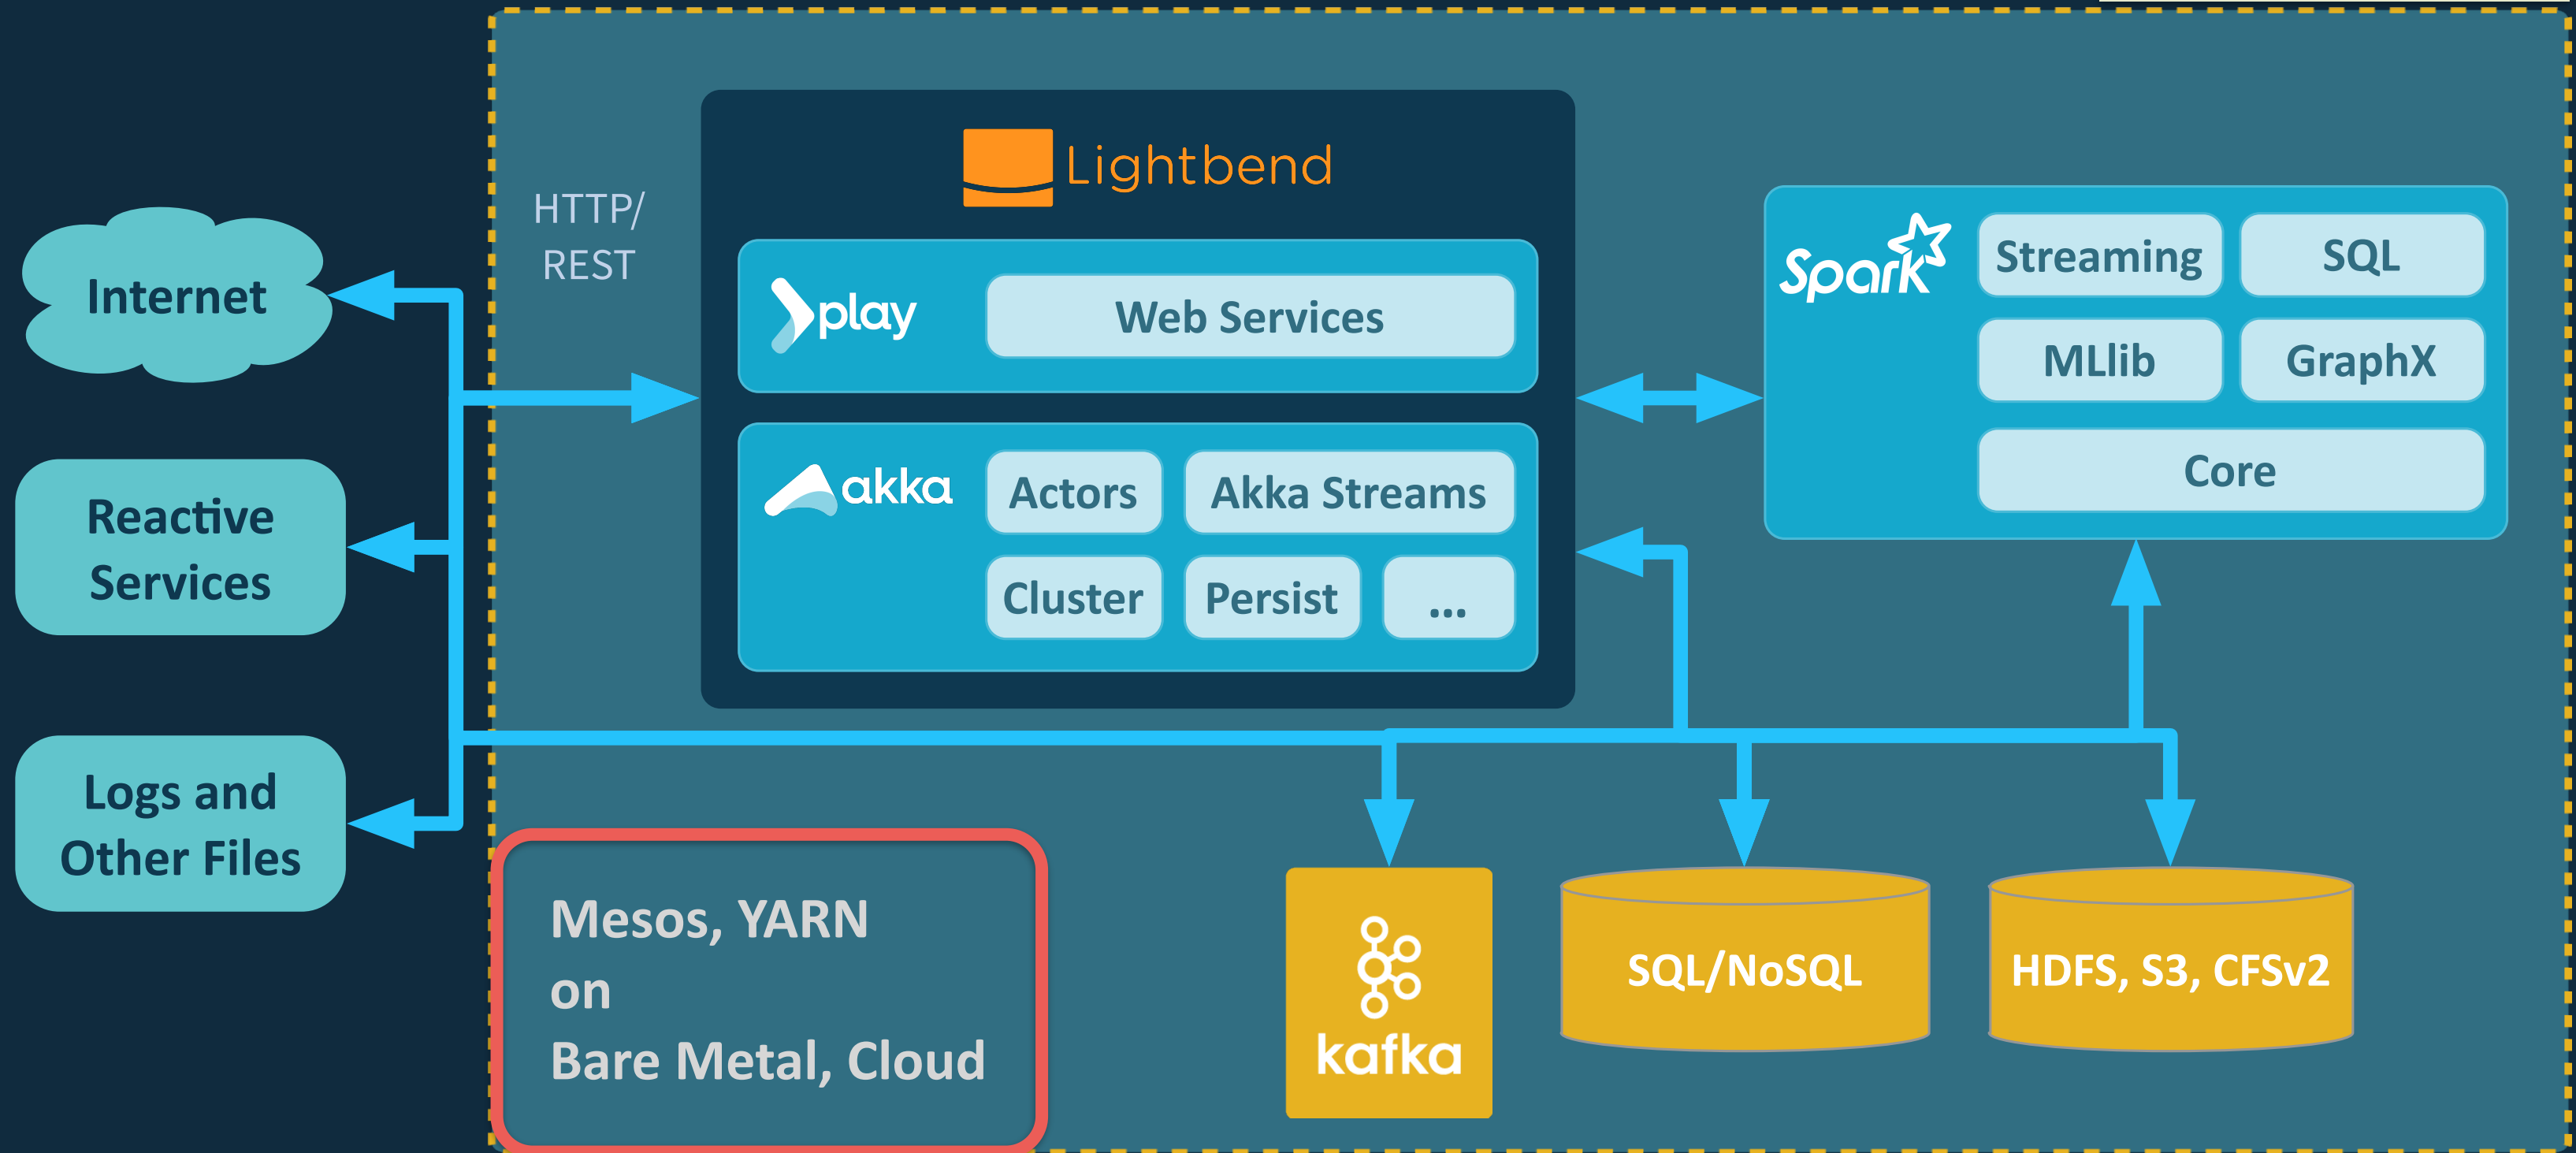
Fast Data Architecture

Short and Long-term Storage



Fast Data Architecture

Infrastructure



- Next Steps

- Learn - Fast Data: Big Data Evolved
- Watch - Using Spark, Kafka, Cassandra and Akka on Mesos for Real-Time Personalization
- Review - Spark success stories by Lightbend clients

