



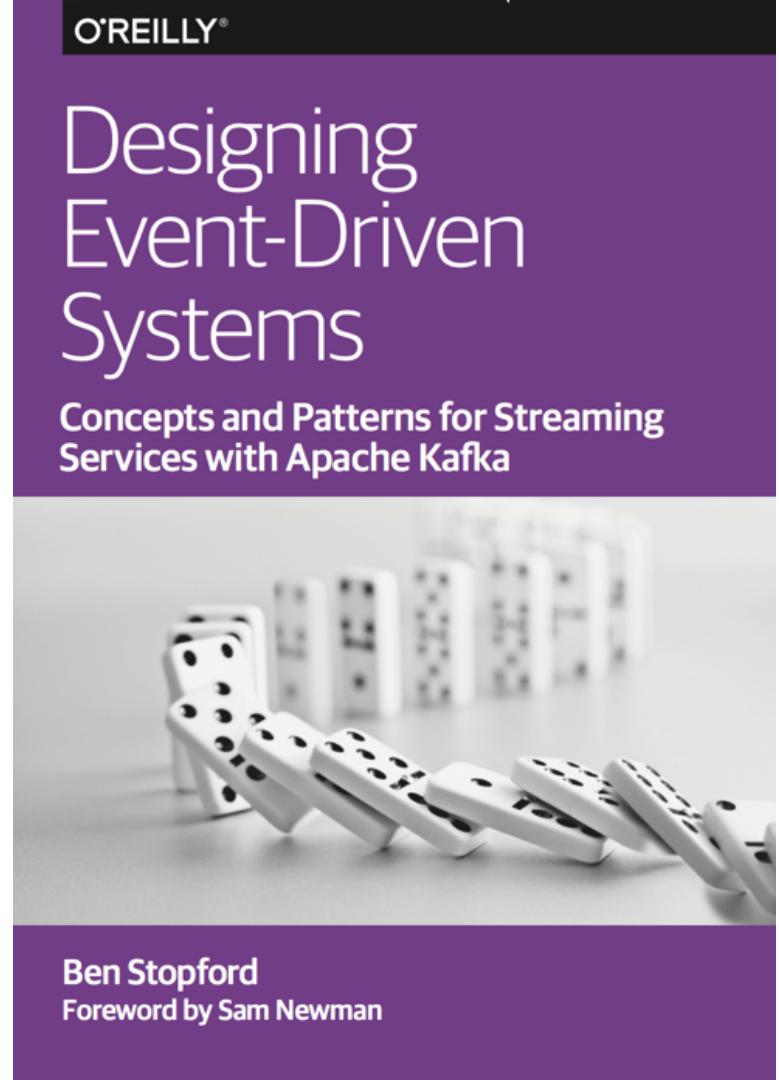
Building Event Driven Services with Apache Kafka and Kafka Streams

Ben Stopford

@benstopford

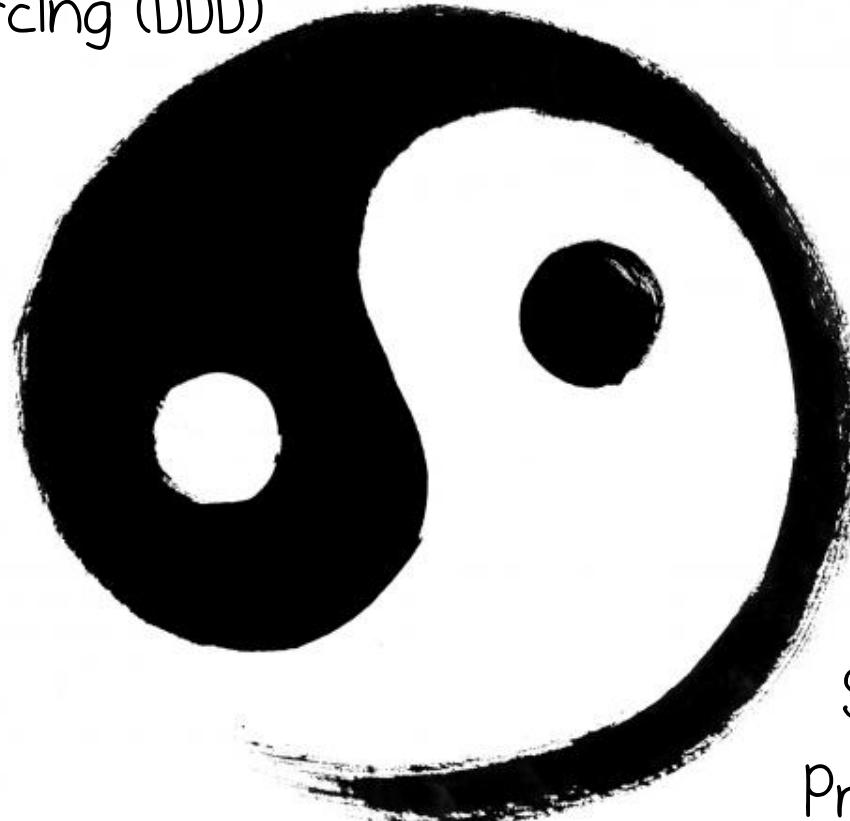
There is a book!

<http://bit.ly/designing-event-driven-systems>



Event Driven Architectures

Event Sourcing (DDD)



Stream
Processing

Today's ecosystems get pretty big

- 2.2 trillion messages per day (6 Petabytes)
- Up to 400 Microservices per cluster.
- 20-200 Brokers per cluster



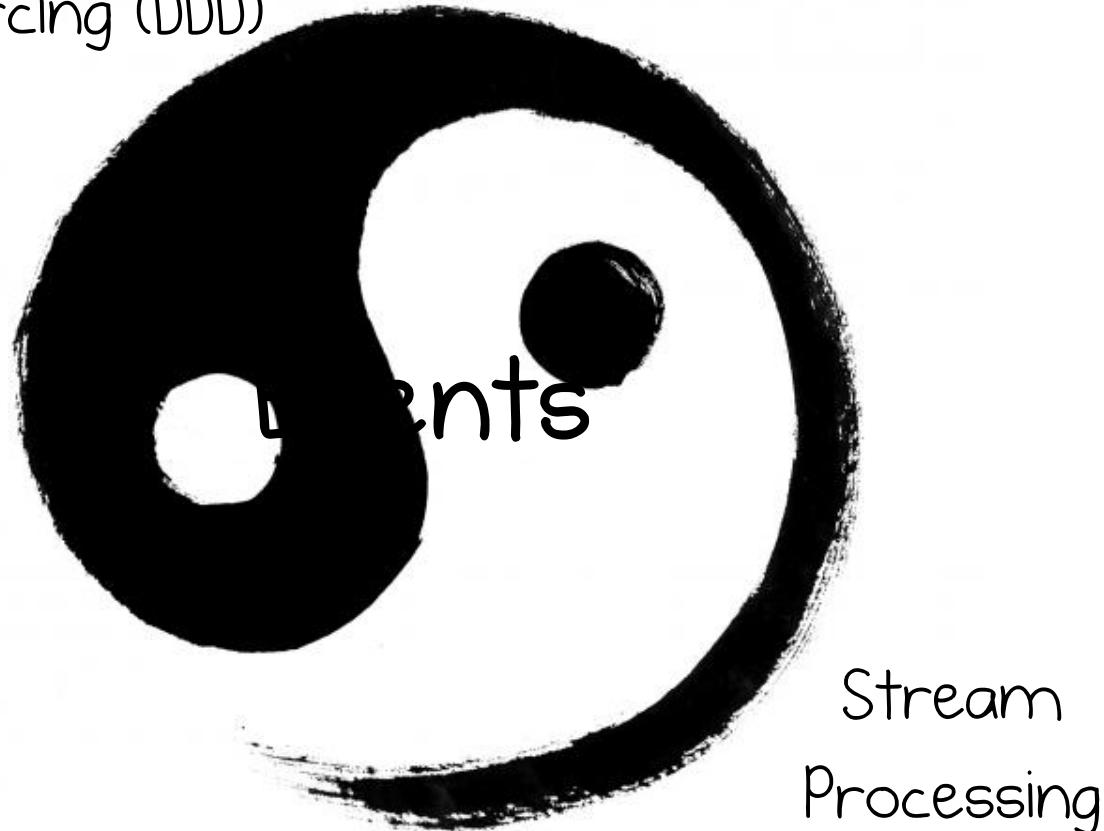
Today's ecosystems get pretty big

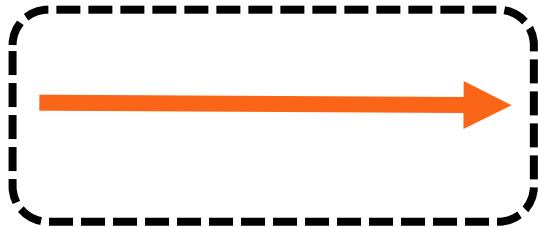


- 1 billion messages per day
- 20,000 messages per second
- 100 teams

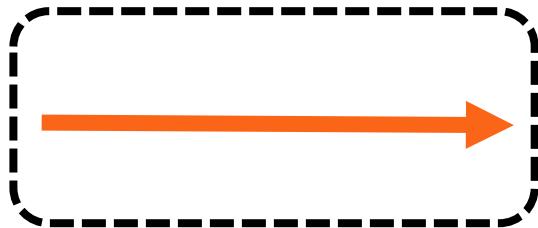
Event Driven Architectures

Event Sourcing (DDD)

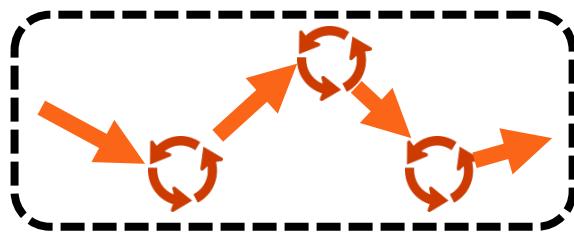




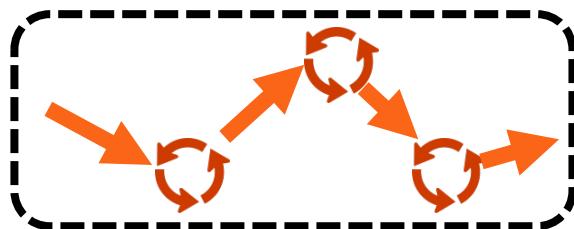
Notification



Data
replication

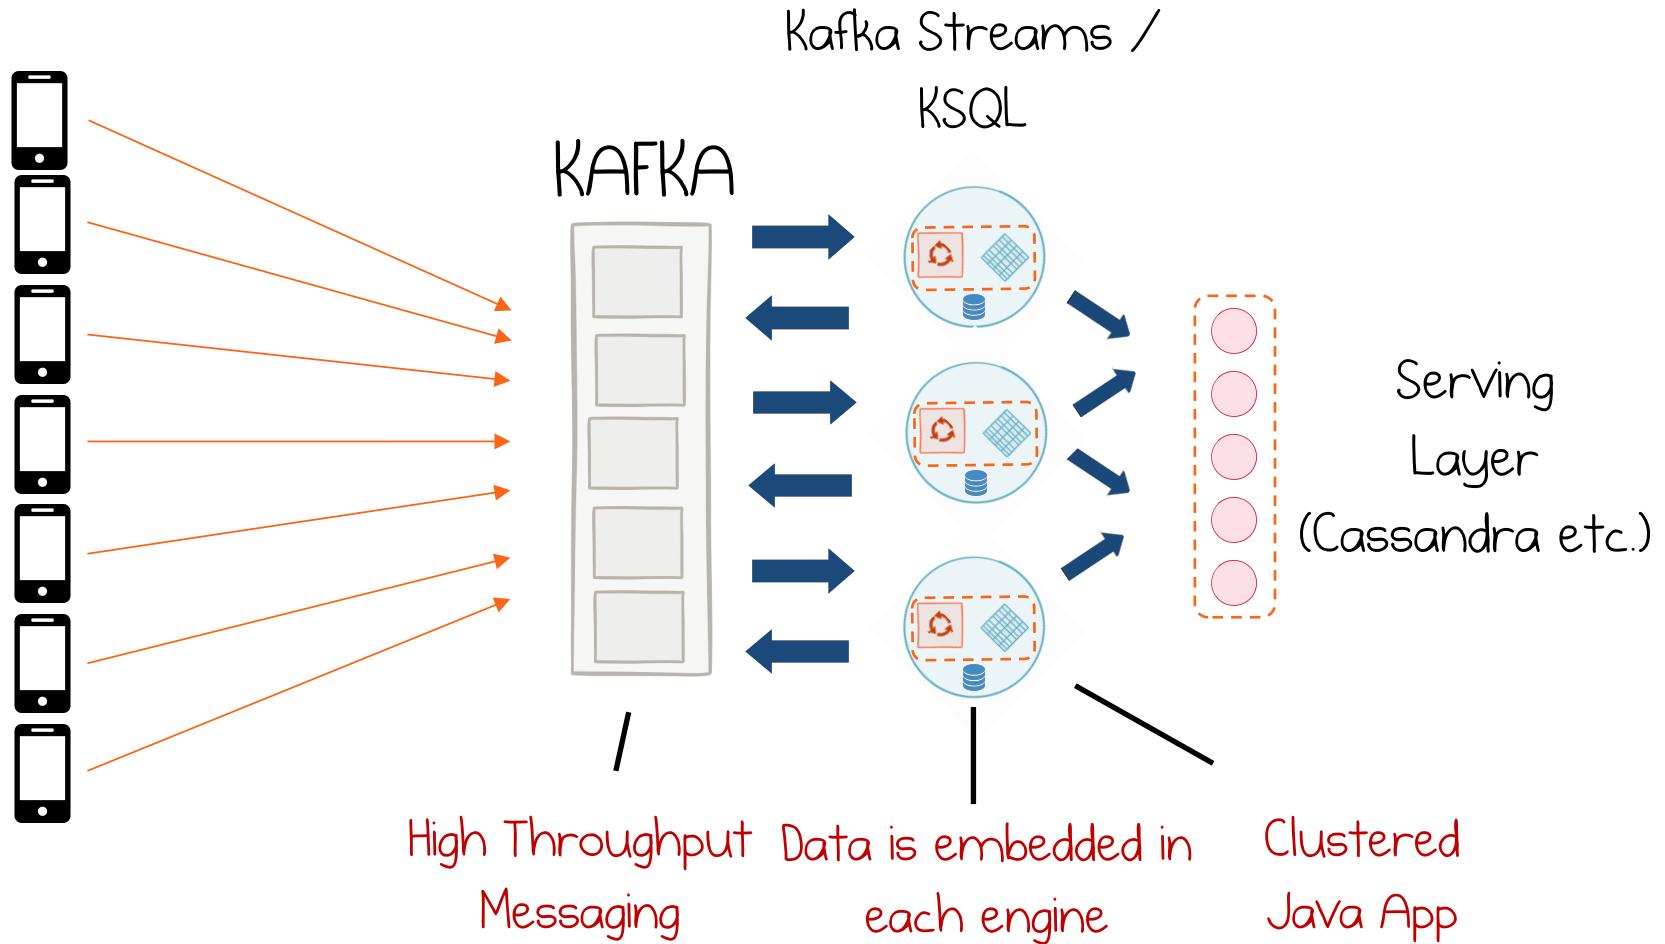


Notification

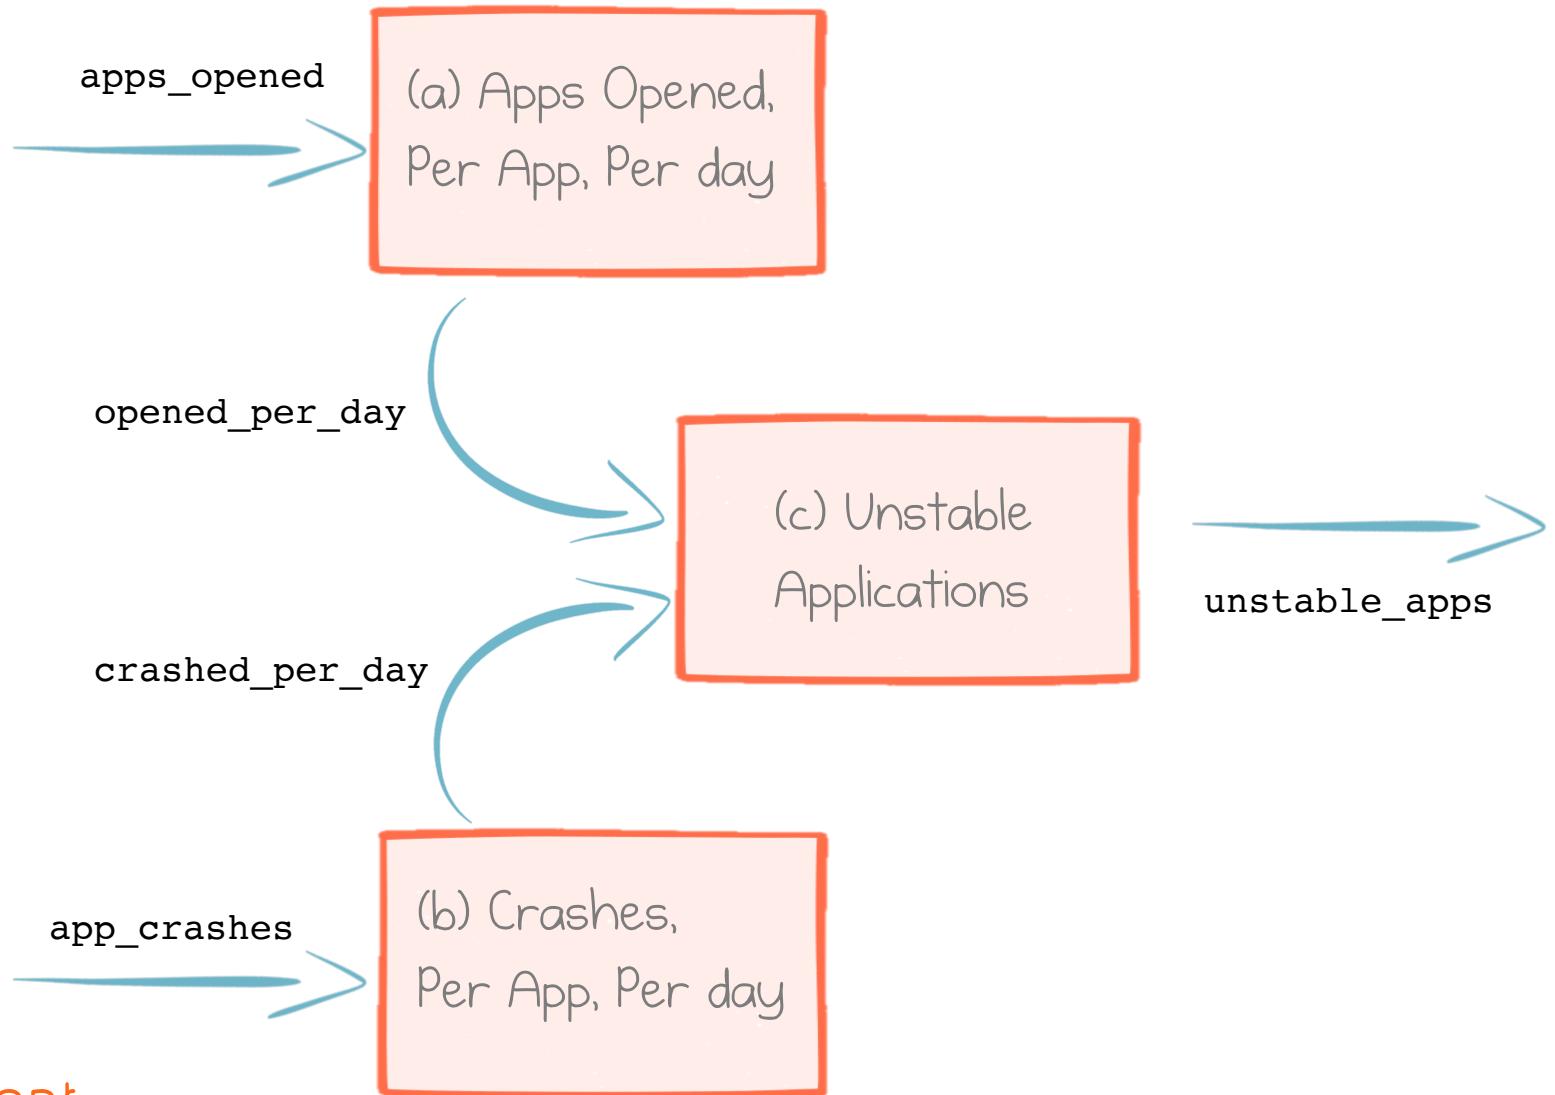


Data
replication

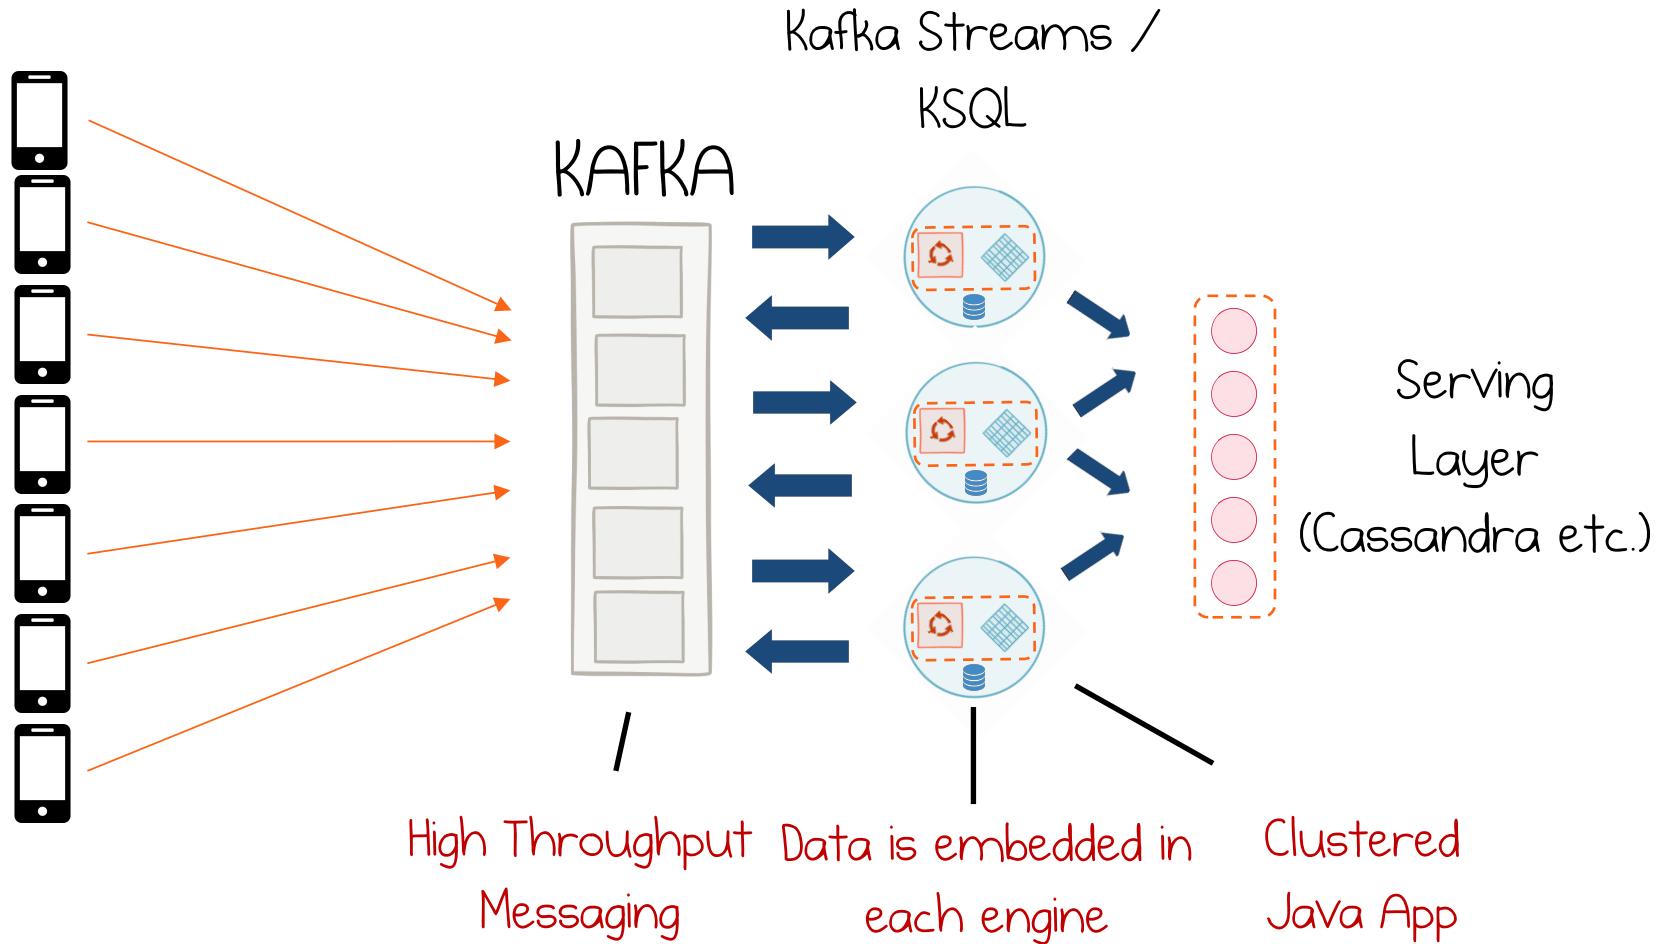
Streaming Platforms



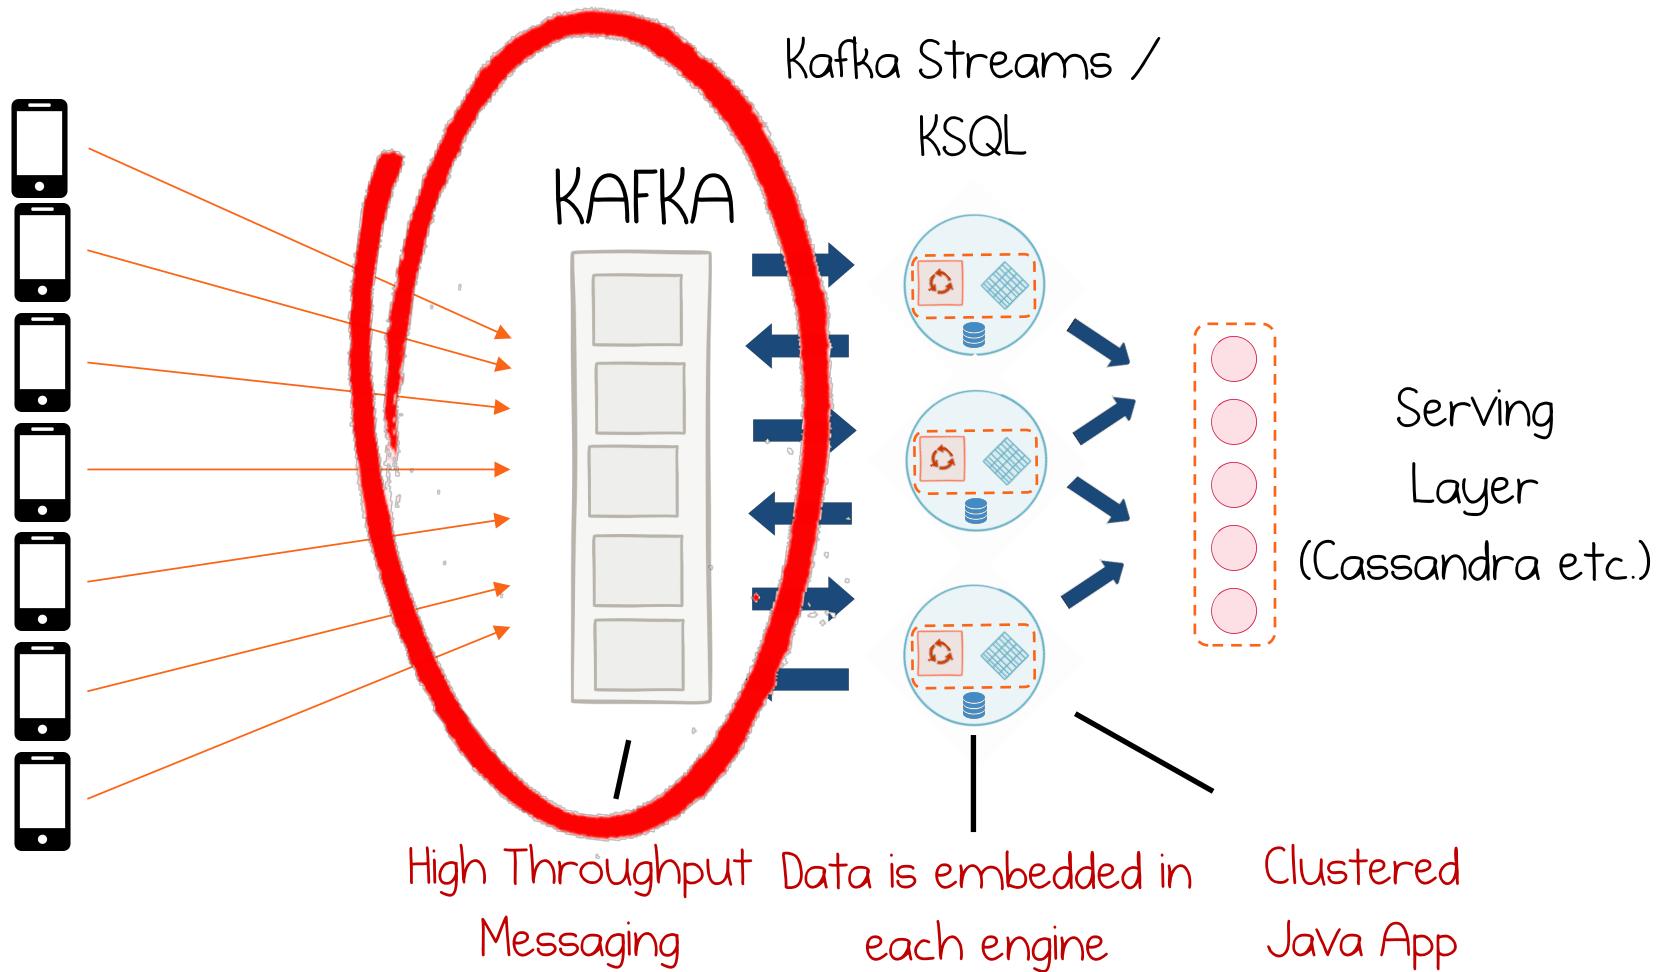
Streaming Pipeline



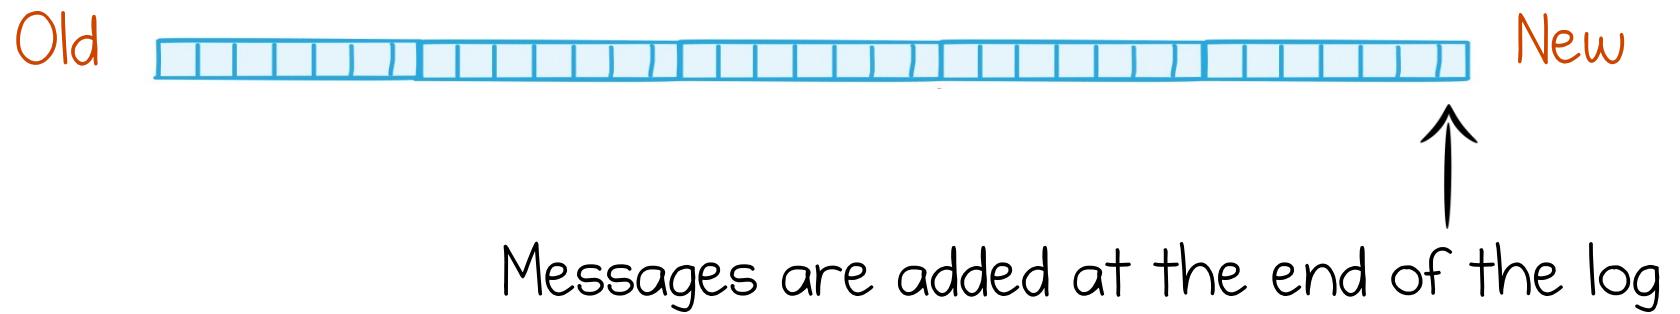
Streaming Platforms



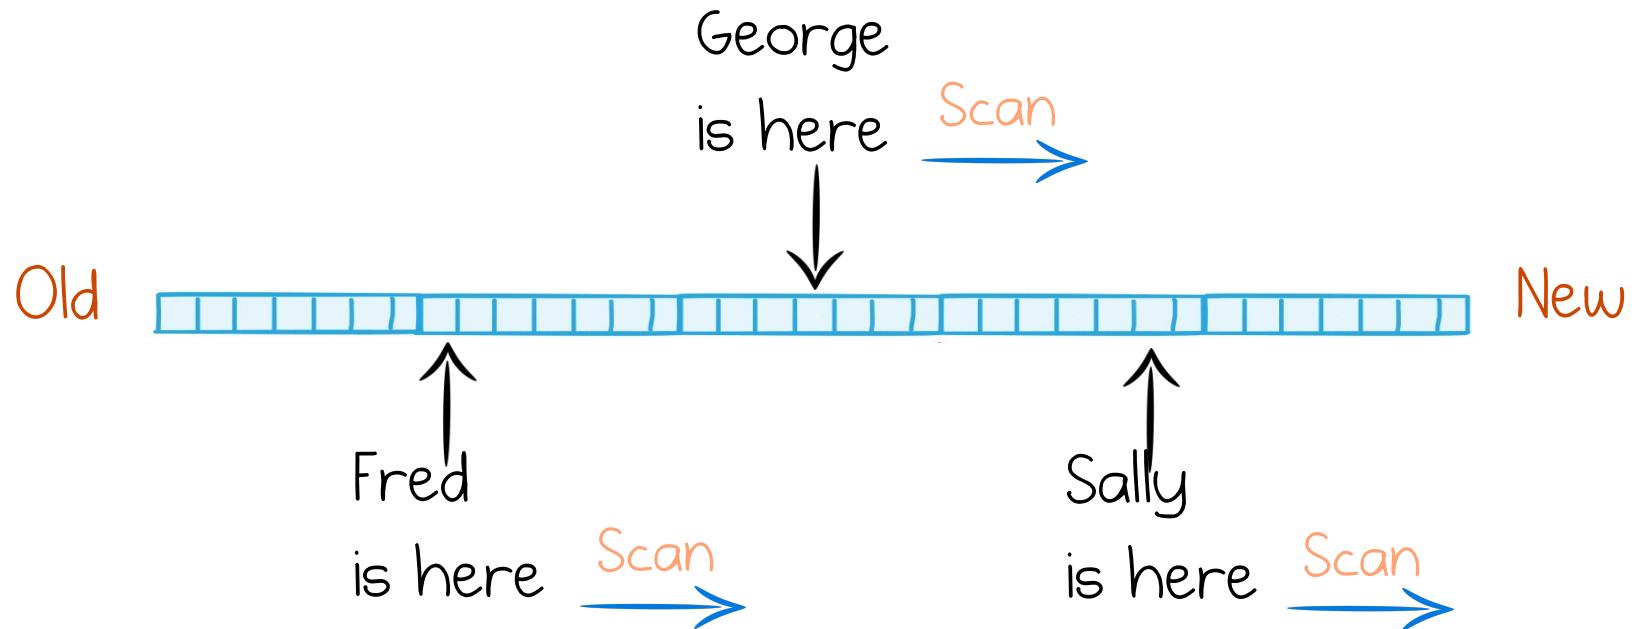
Streaming Platforms



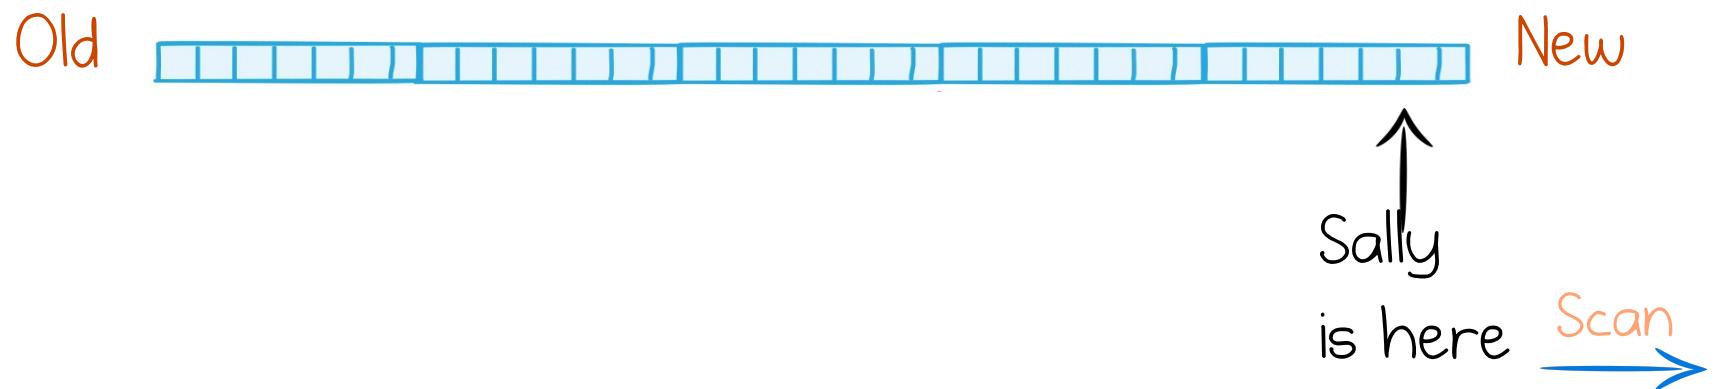
An event log is a simple idea



Readers have a position all of their own

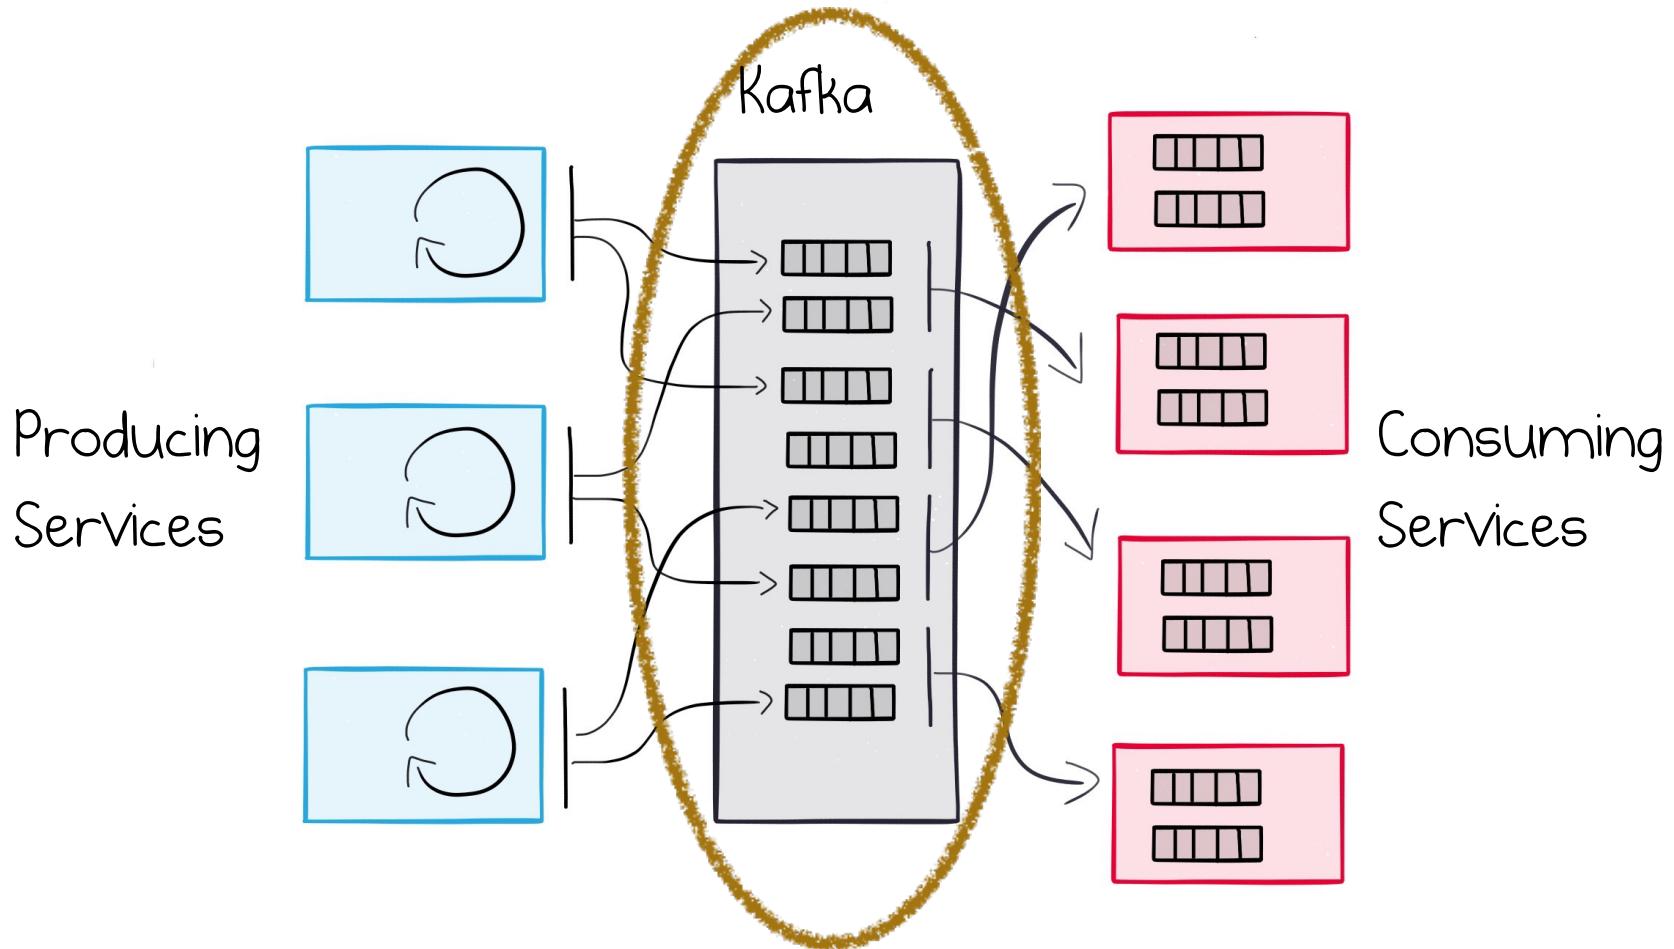


You can rewind and replay, just like TiVo!

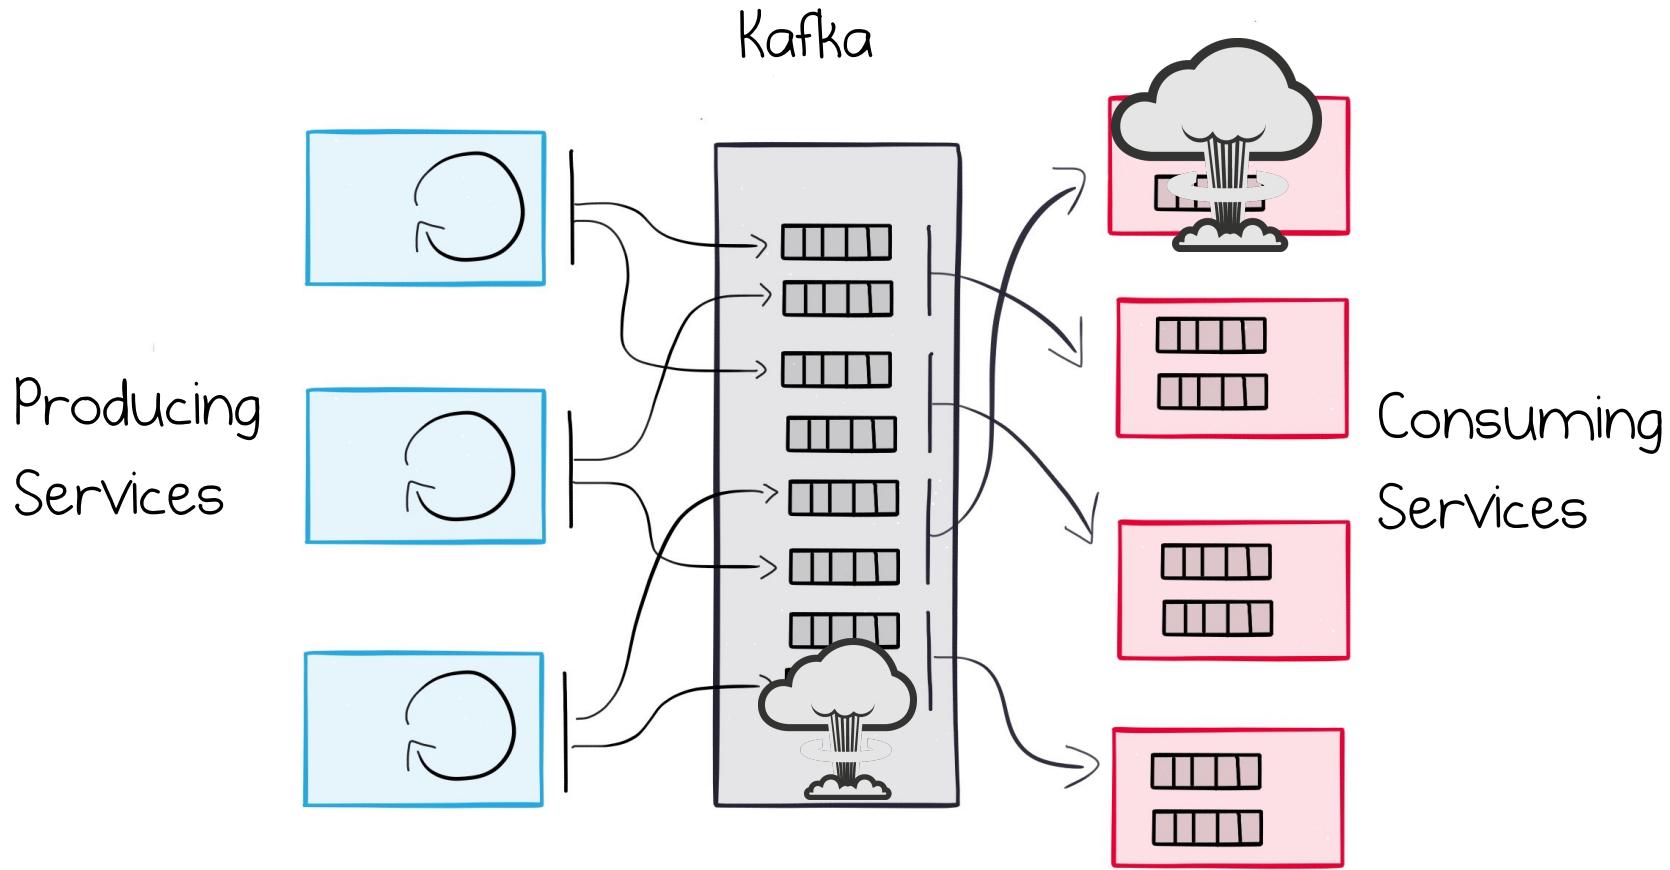


The hard part: Tying it all together!

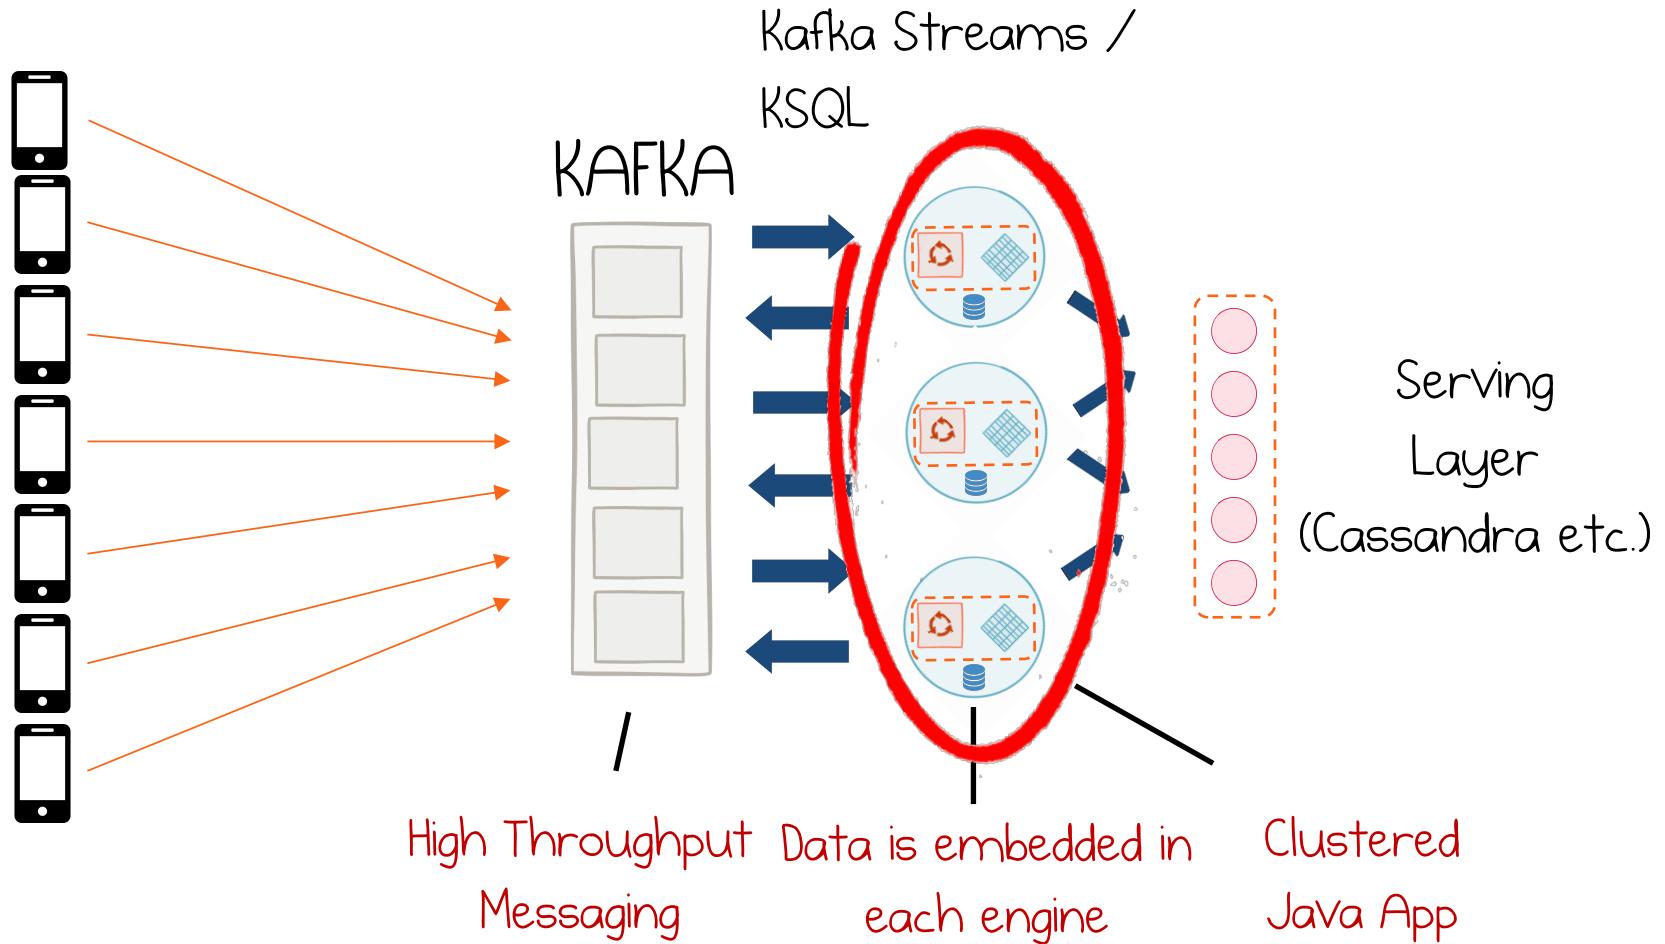
Many "logs" over many machines



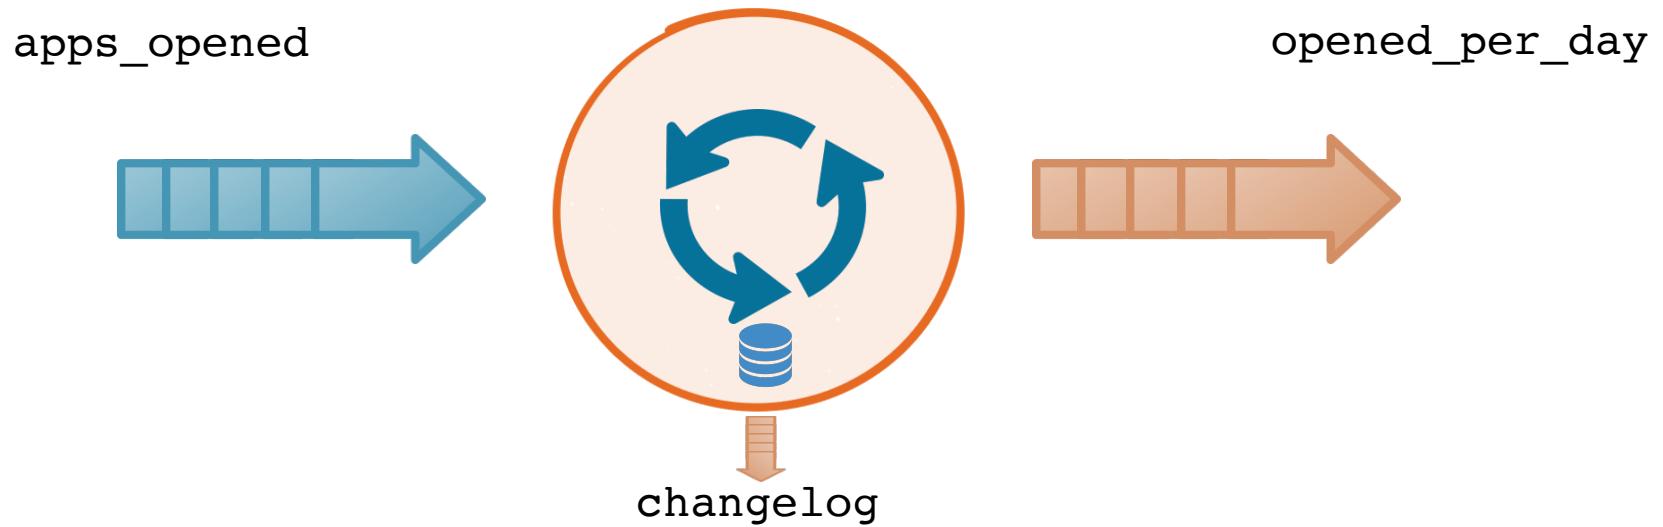
Resistant to Failure

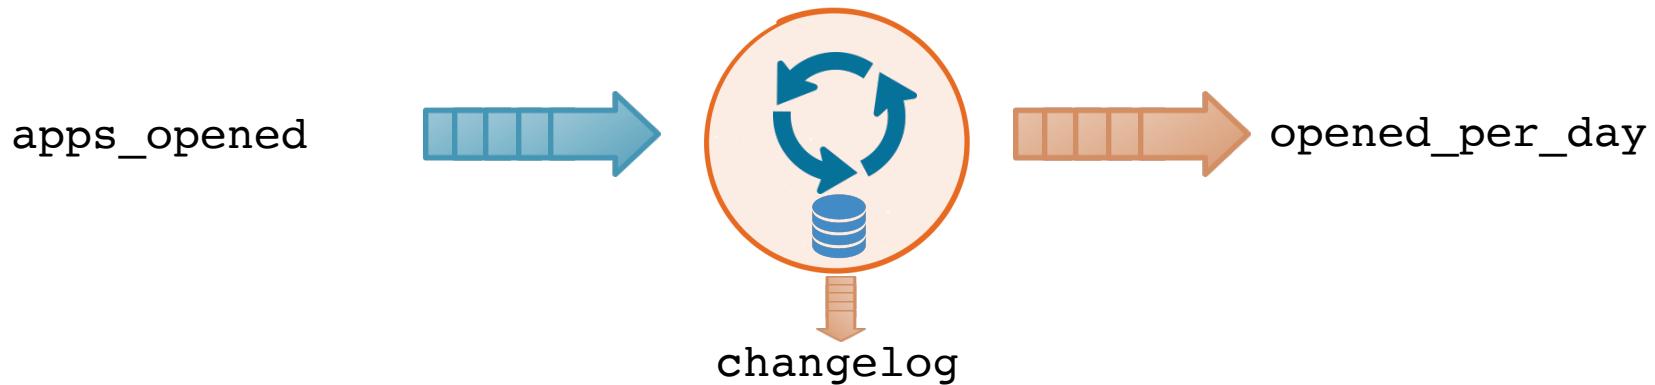


Streaming Platforms

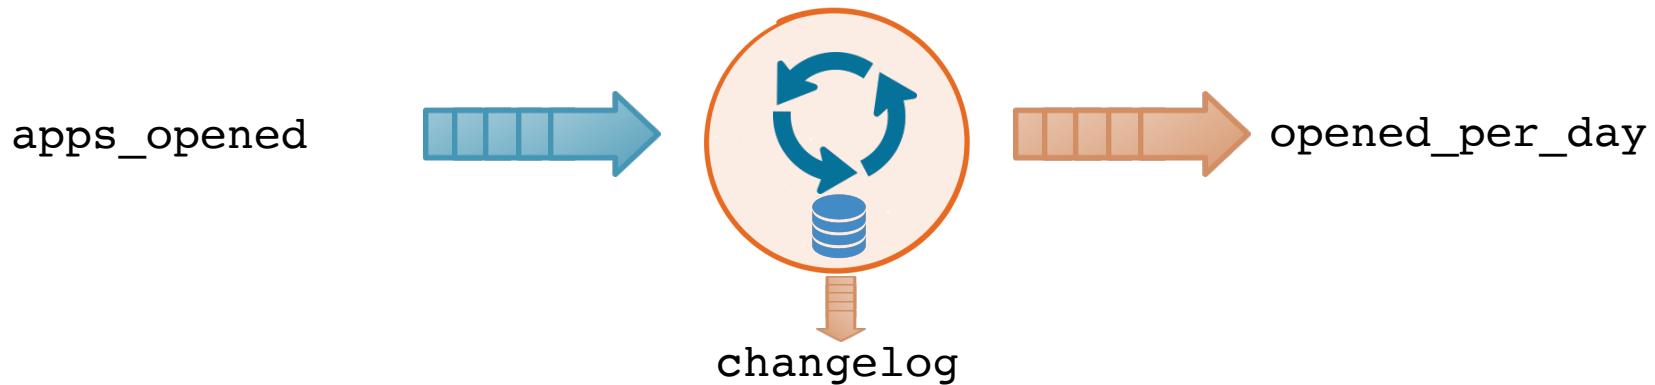


Streaming Example

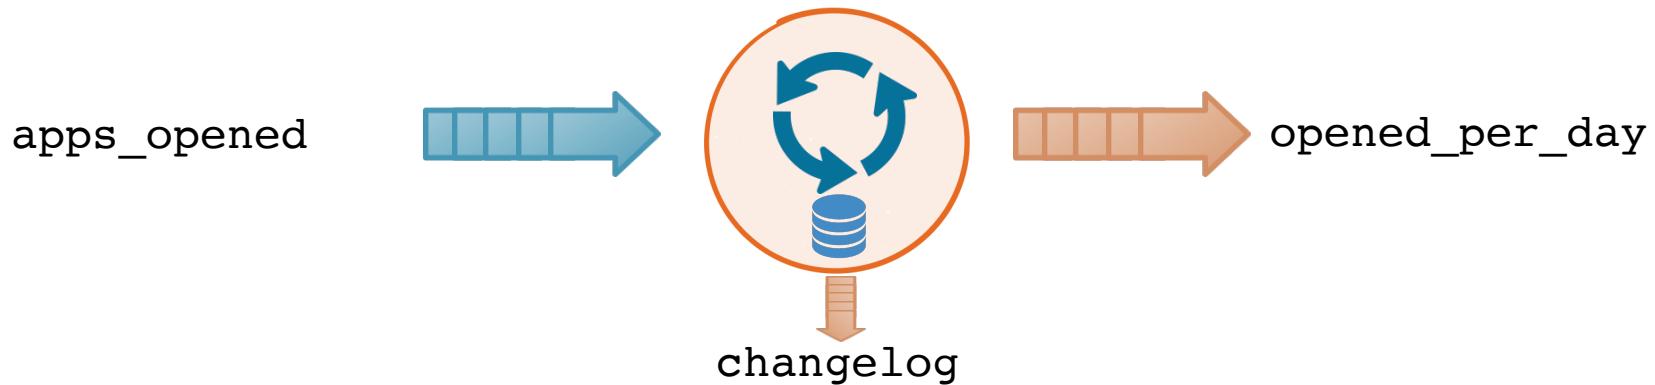




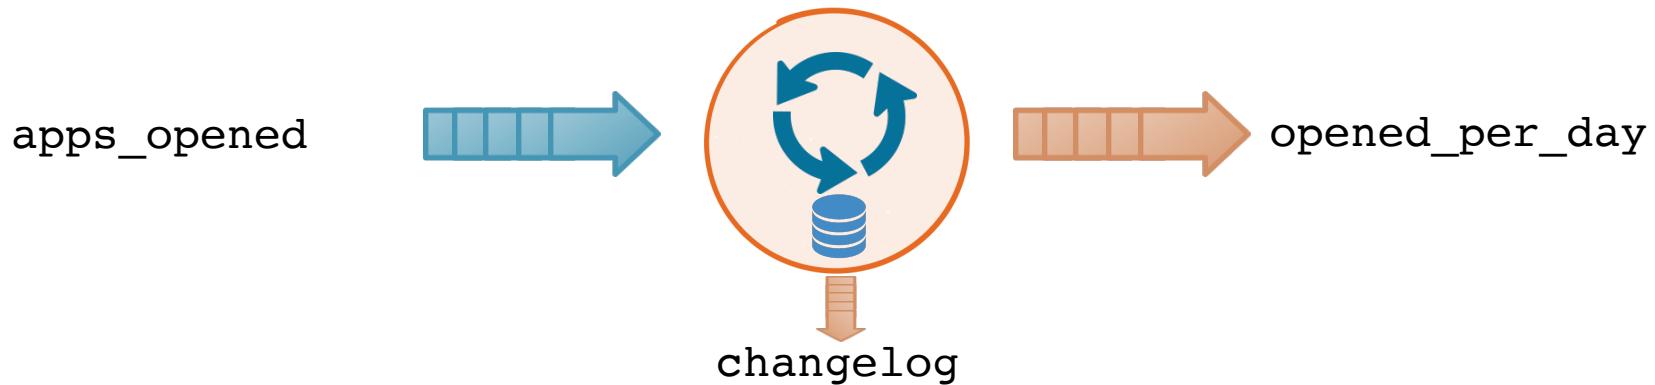
```
CREATE TABLE opened_per_day AS  
SELECT app_id, count(*)  
FROM apps_opened  
WINDOW TUMBLING (SIZE 1 DAY)  
GROUP BY app_id;
```



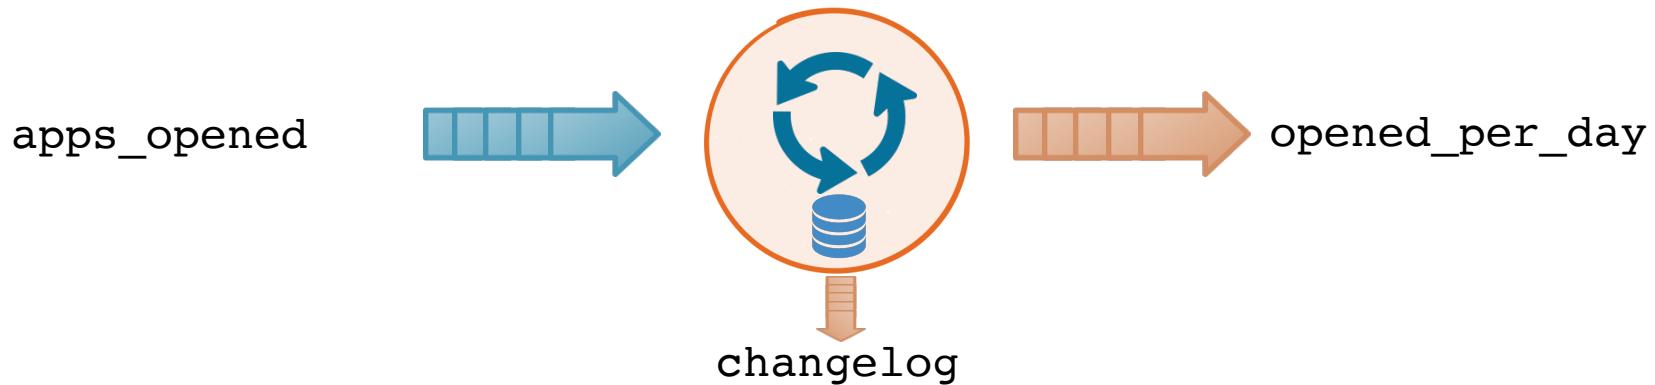
```
CREATE TABLE opened per day AS  
SELECT app_id, count(*)  
FROM apps_opened  
WINDOW TUMBLING (SIZE 1 DAY)  
GROUP BY app_id;
```



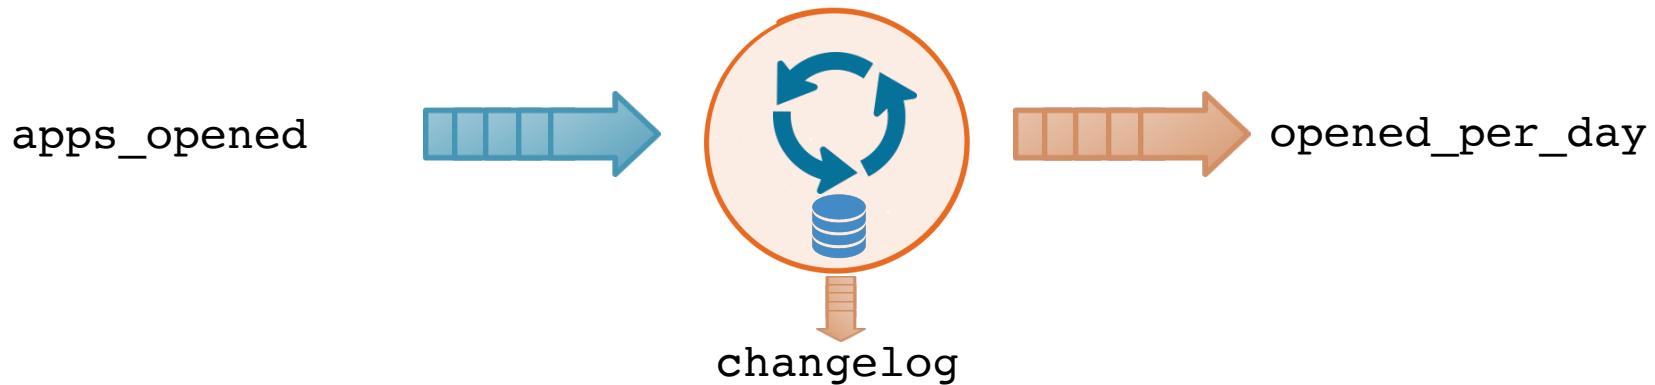
```
CREATE TABLE opened_per_day AS  
SELECT app_id, count(*)  
FROM apps_opened  
WINDOW TUMBLING (SIZE 1 DAY)  
GROUP BY app_id;
```



```
CREATE TABLE opened_per_day AS
SELECT app_id, count(*)
FROM apps opened
WINDOW TUMBLING (SIZE 1 DAY)
GROUP BY app_id;
```



```
CREATE TABLE opened_per_day AS  
SELECT app_id, count(*)  
FROM apps_opened  
WINDOW TUMBLING (SIZE 1 DAY)  
GROUP BY app_id;
```

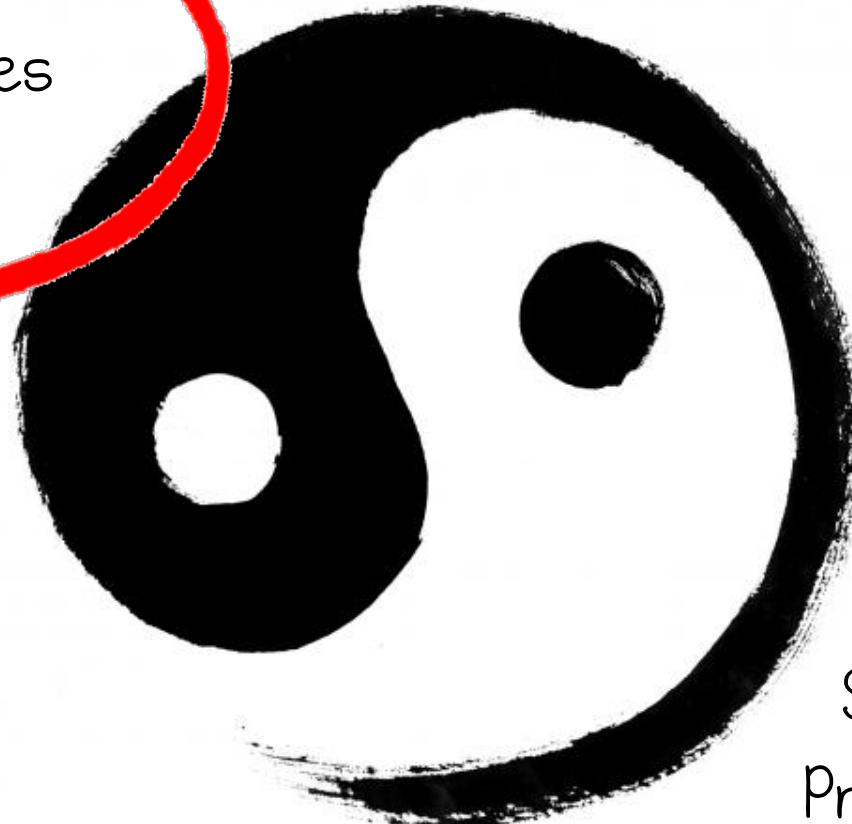


```
CREATE TABLE opened_per_day AS  
SELECT app_id, count(*)  
FROM apps_opened  
WINDOW TUMBLING (SIZE 1 DAY)  
GROUP BY app_id;
```

Streaming is manipulating events in flight,
at scale.

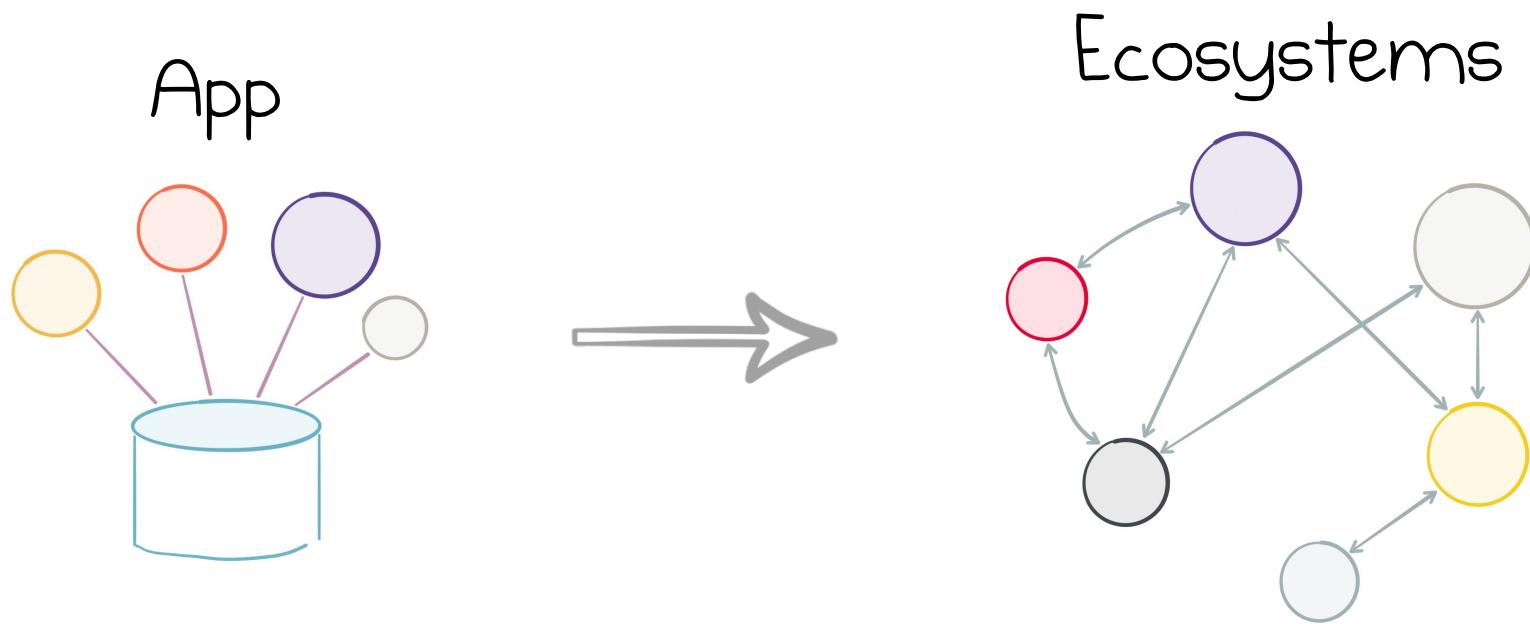


Event Driven
Architectures

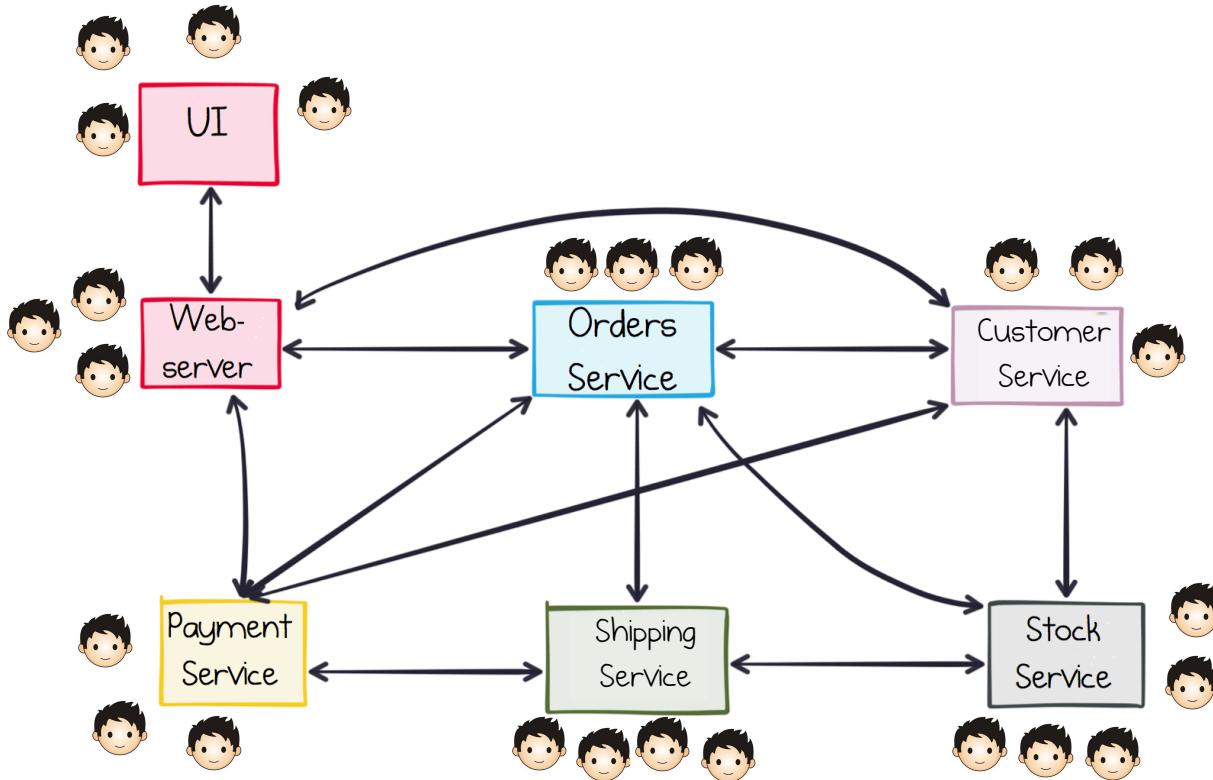


Stream
Processing

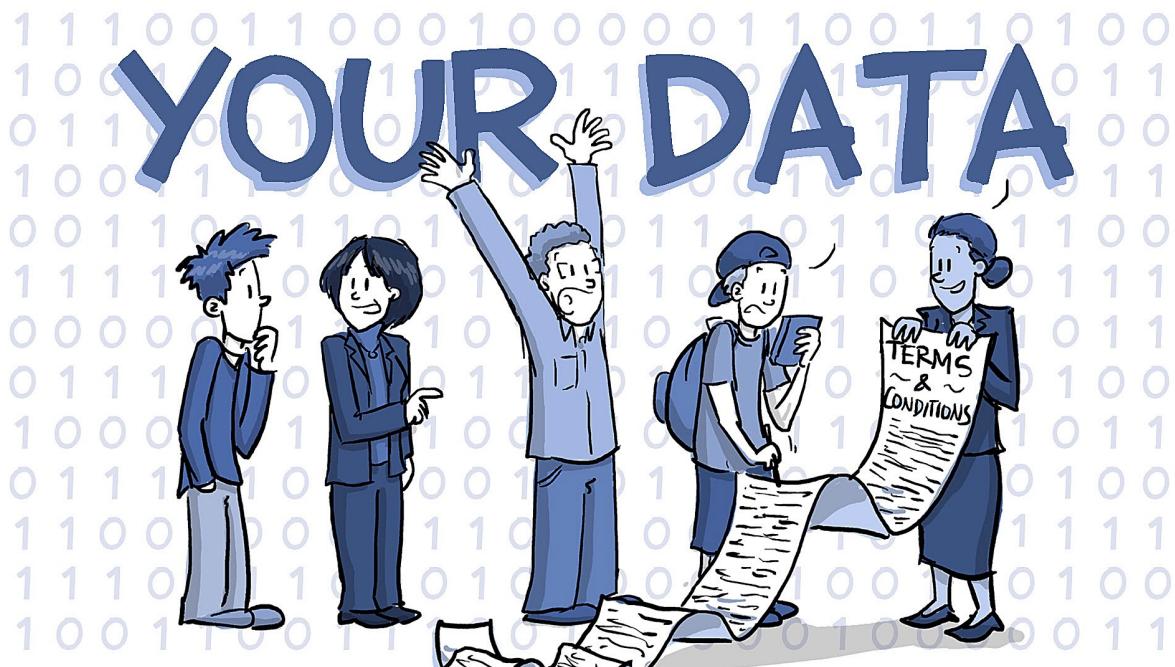
Increasingly we build ecosystems



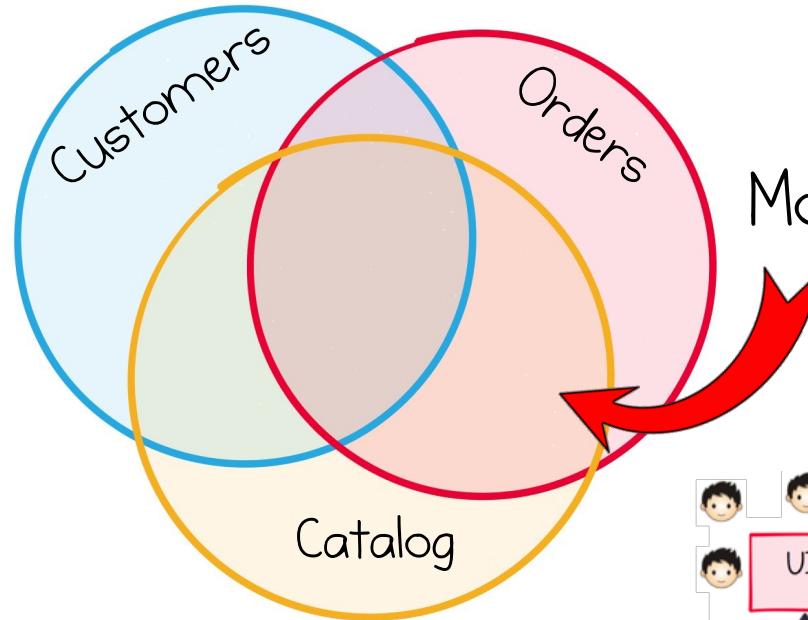
SOA / Microservices / EDA



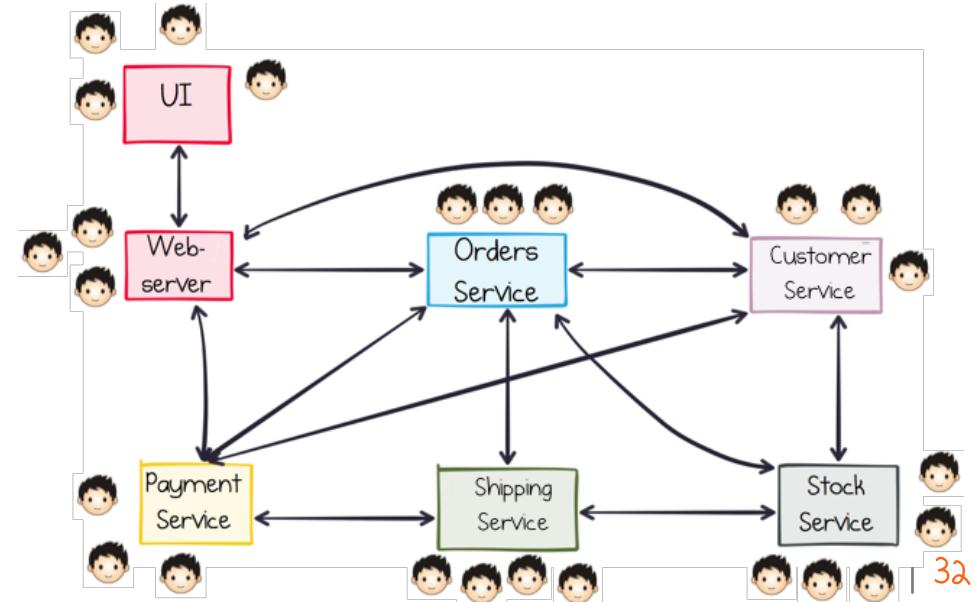
The Problem is DATA



Most services share the same core facts.



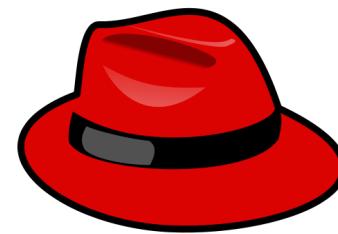
Most services live
in here



Events have two hats

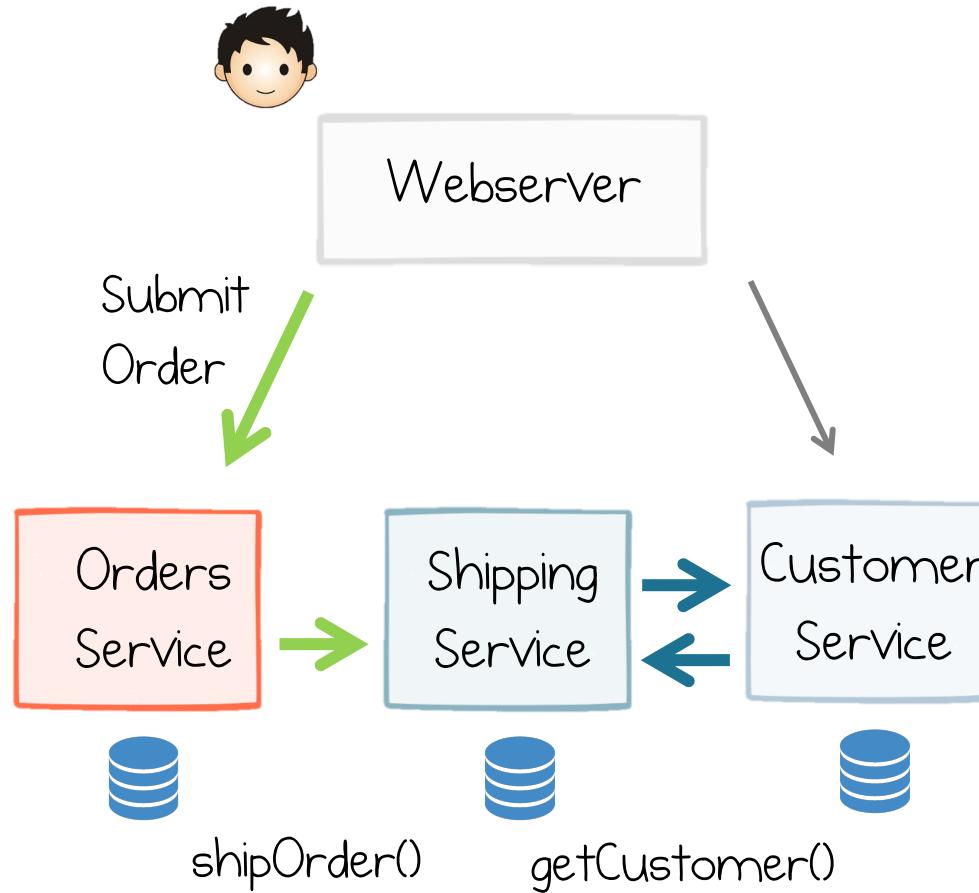


Notification



Data
replication

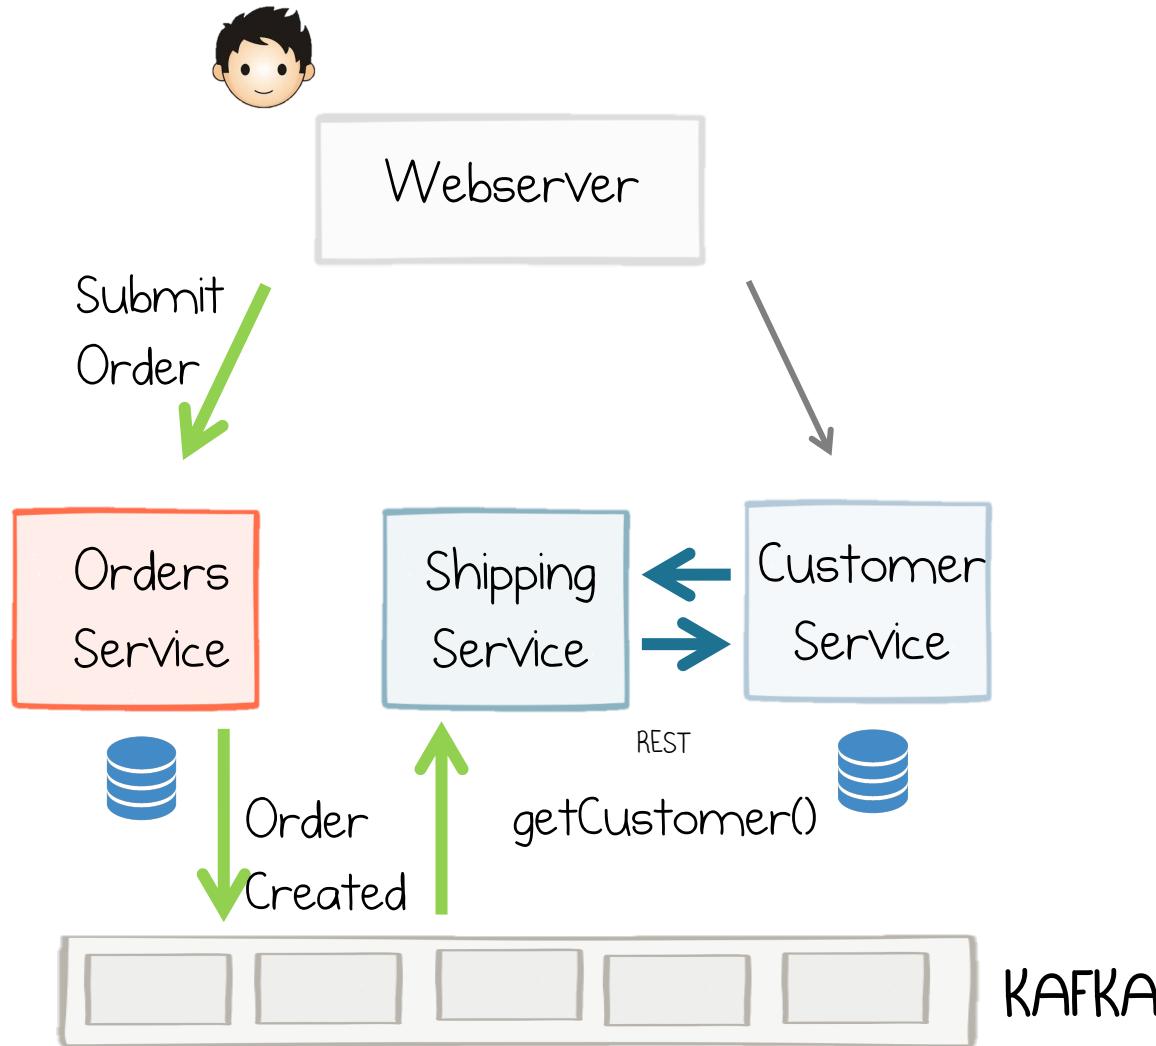
Buying an iPad (with REST/RPC)



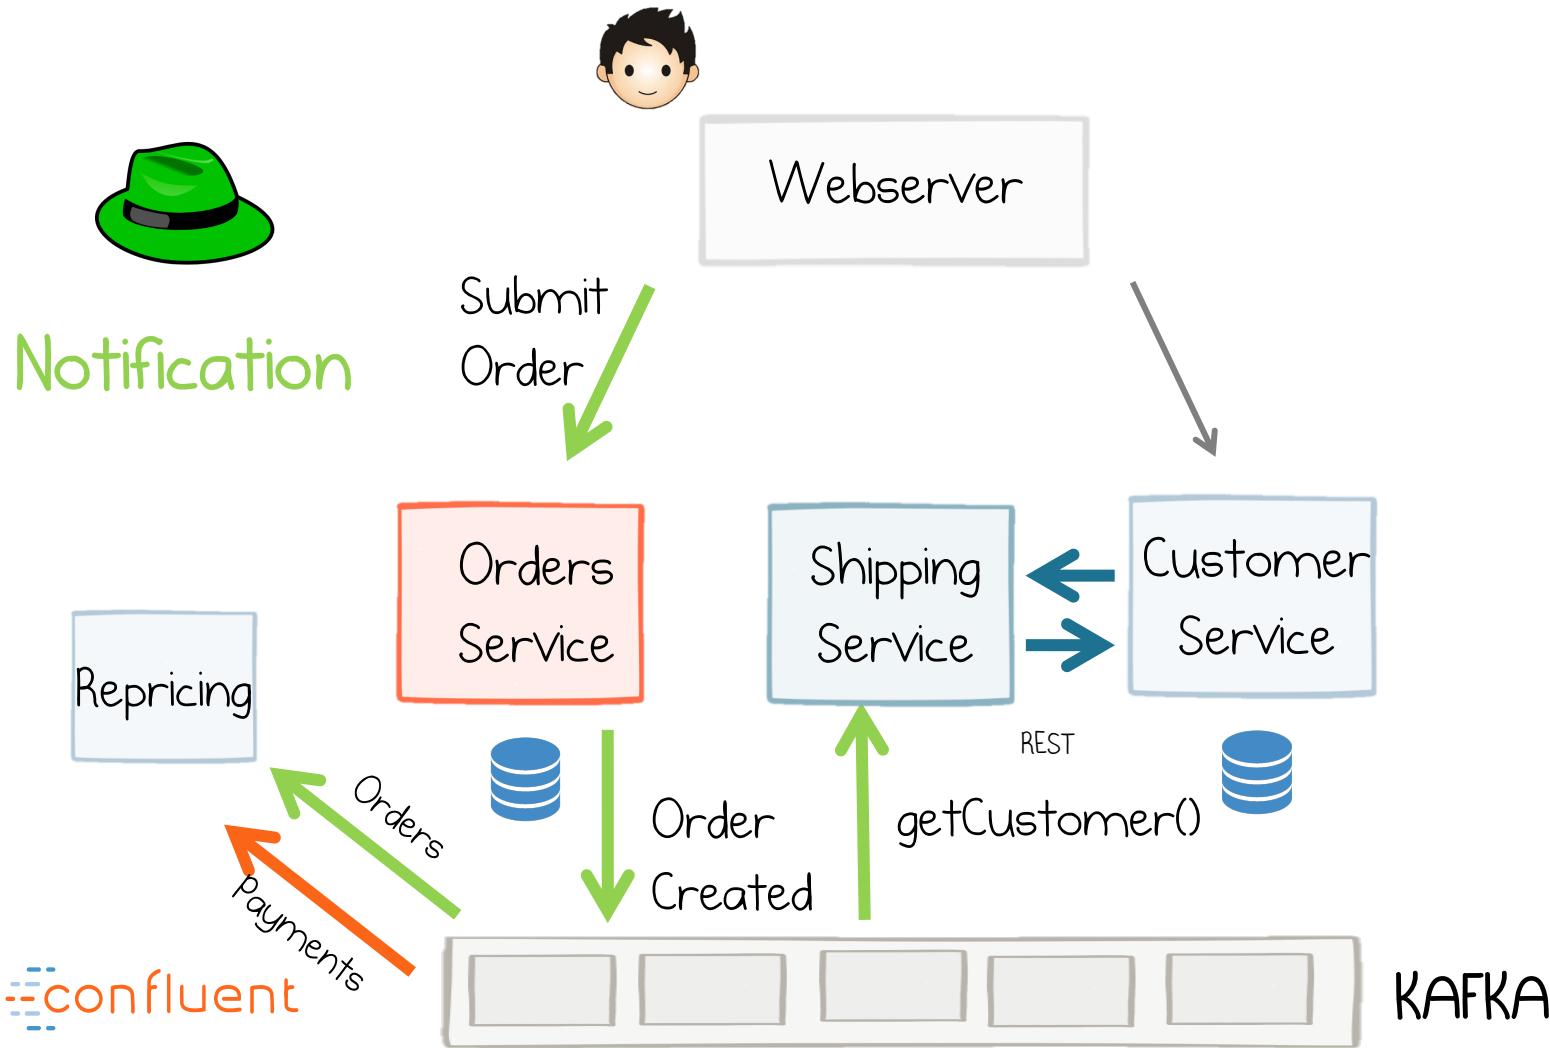
Events for Notification Only



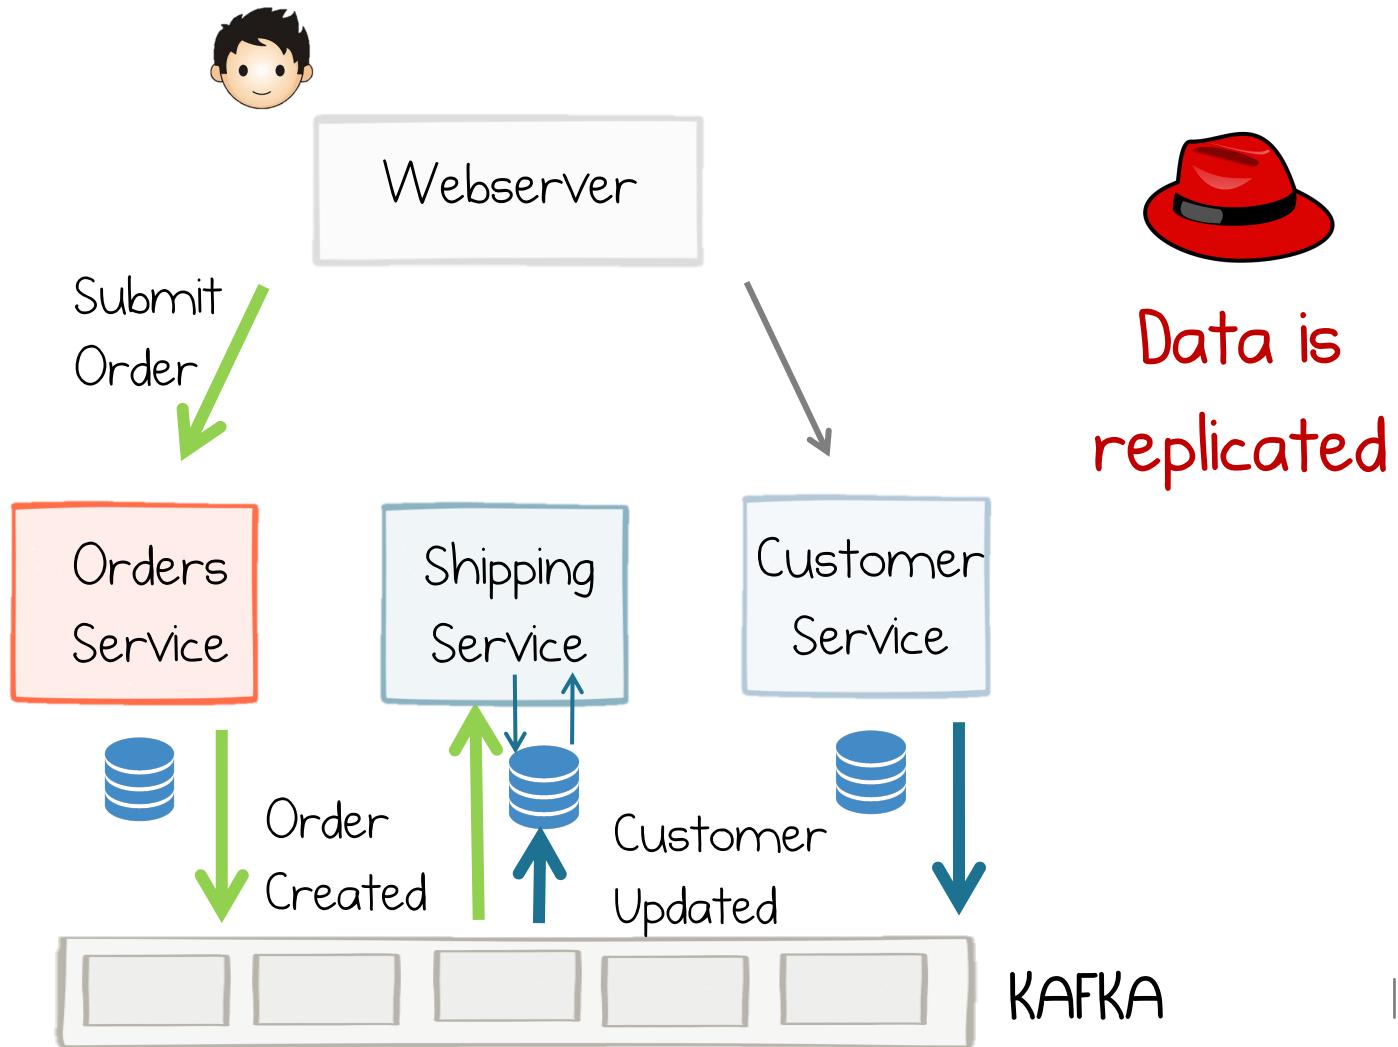
Notification



Pluggability



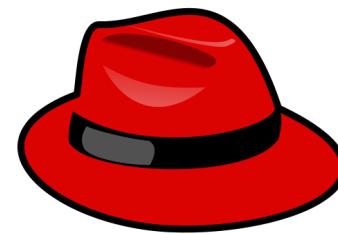
Events for Data Locality



Events have two hats



Notification



Data
replication

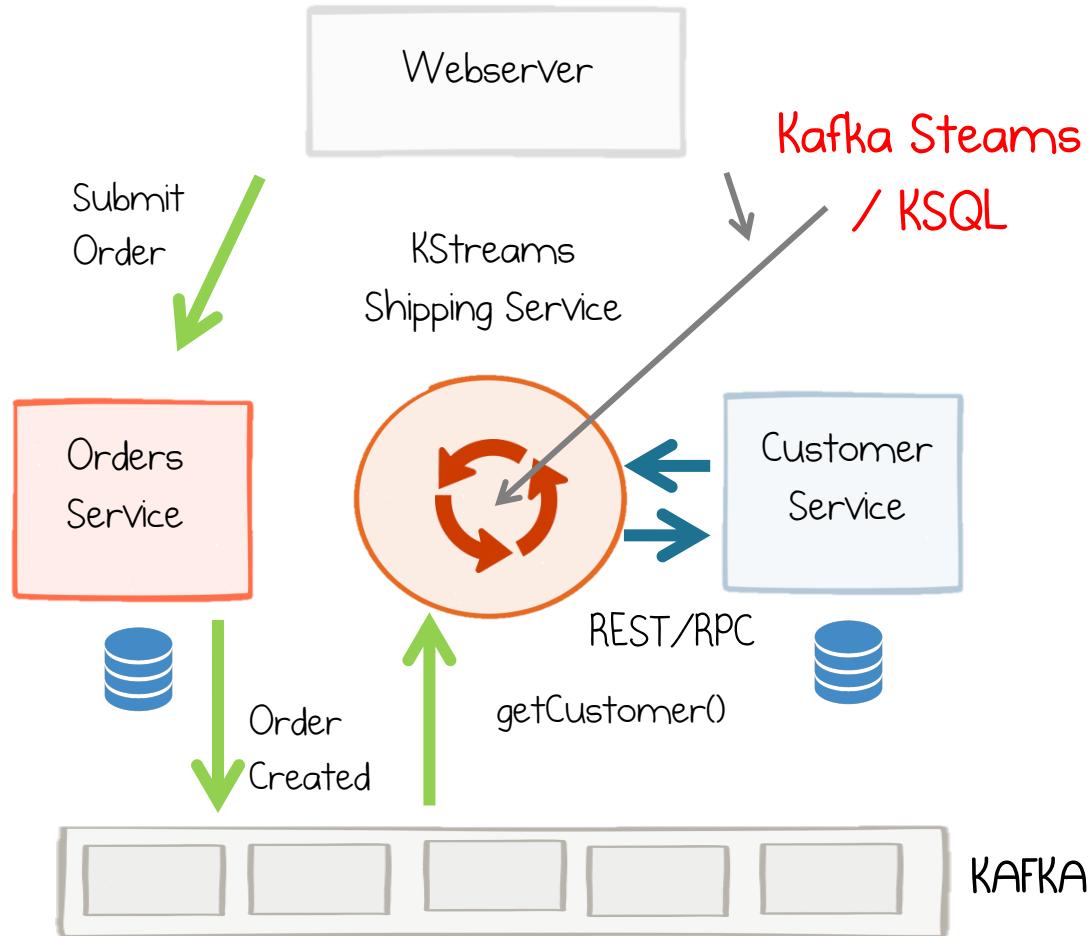
Stateless / Stateful Stream Processing

Relates to these hats

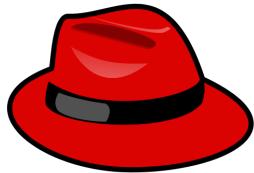
Stateless Stream Processing



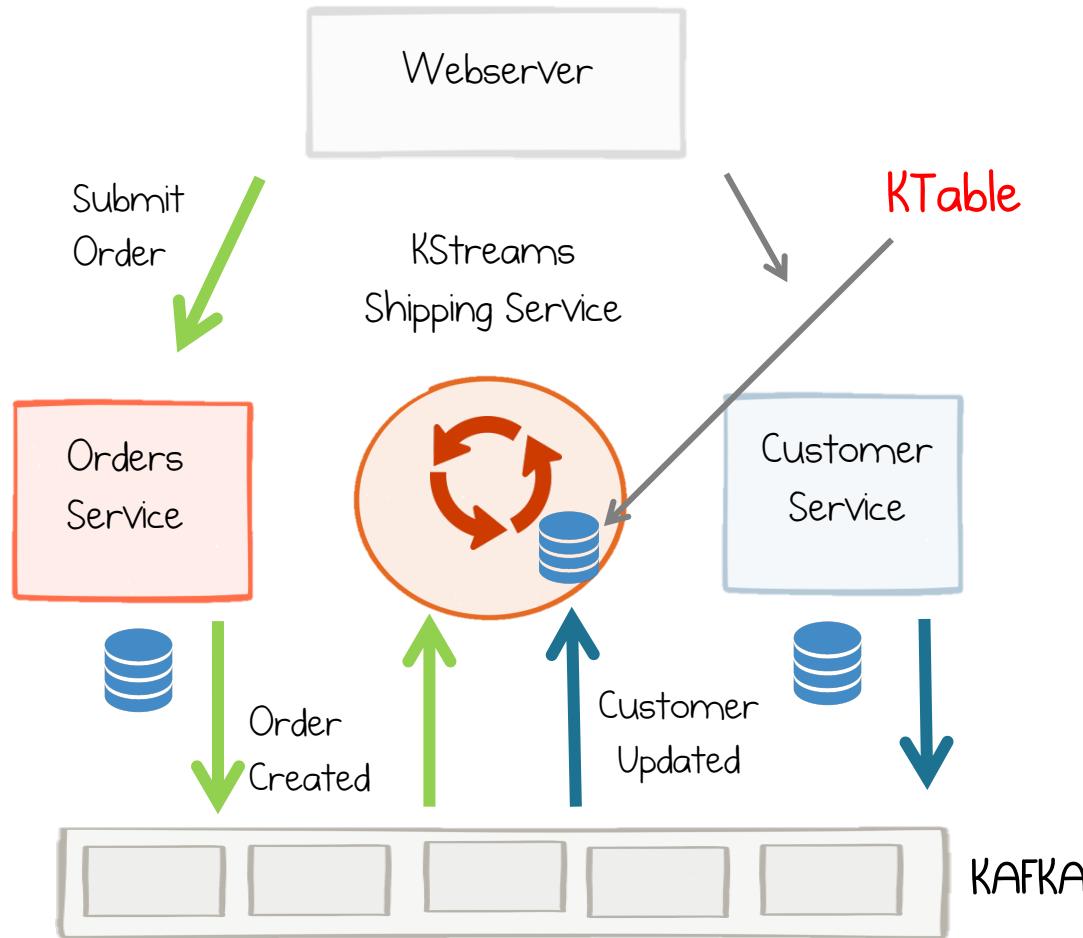
Notification



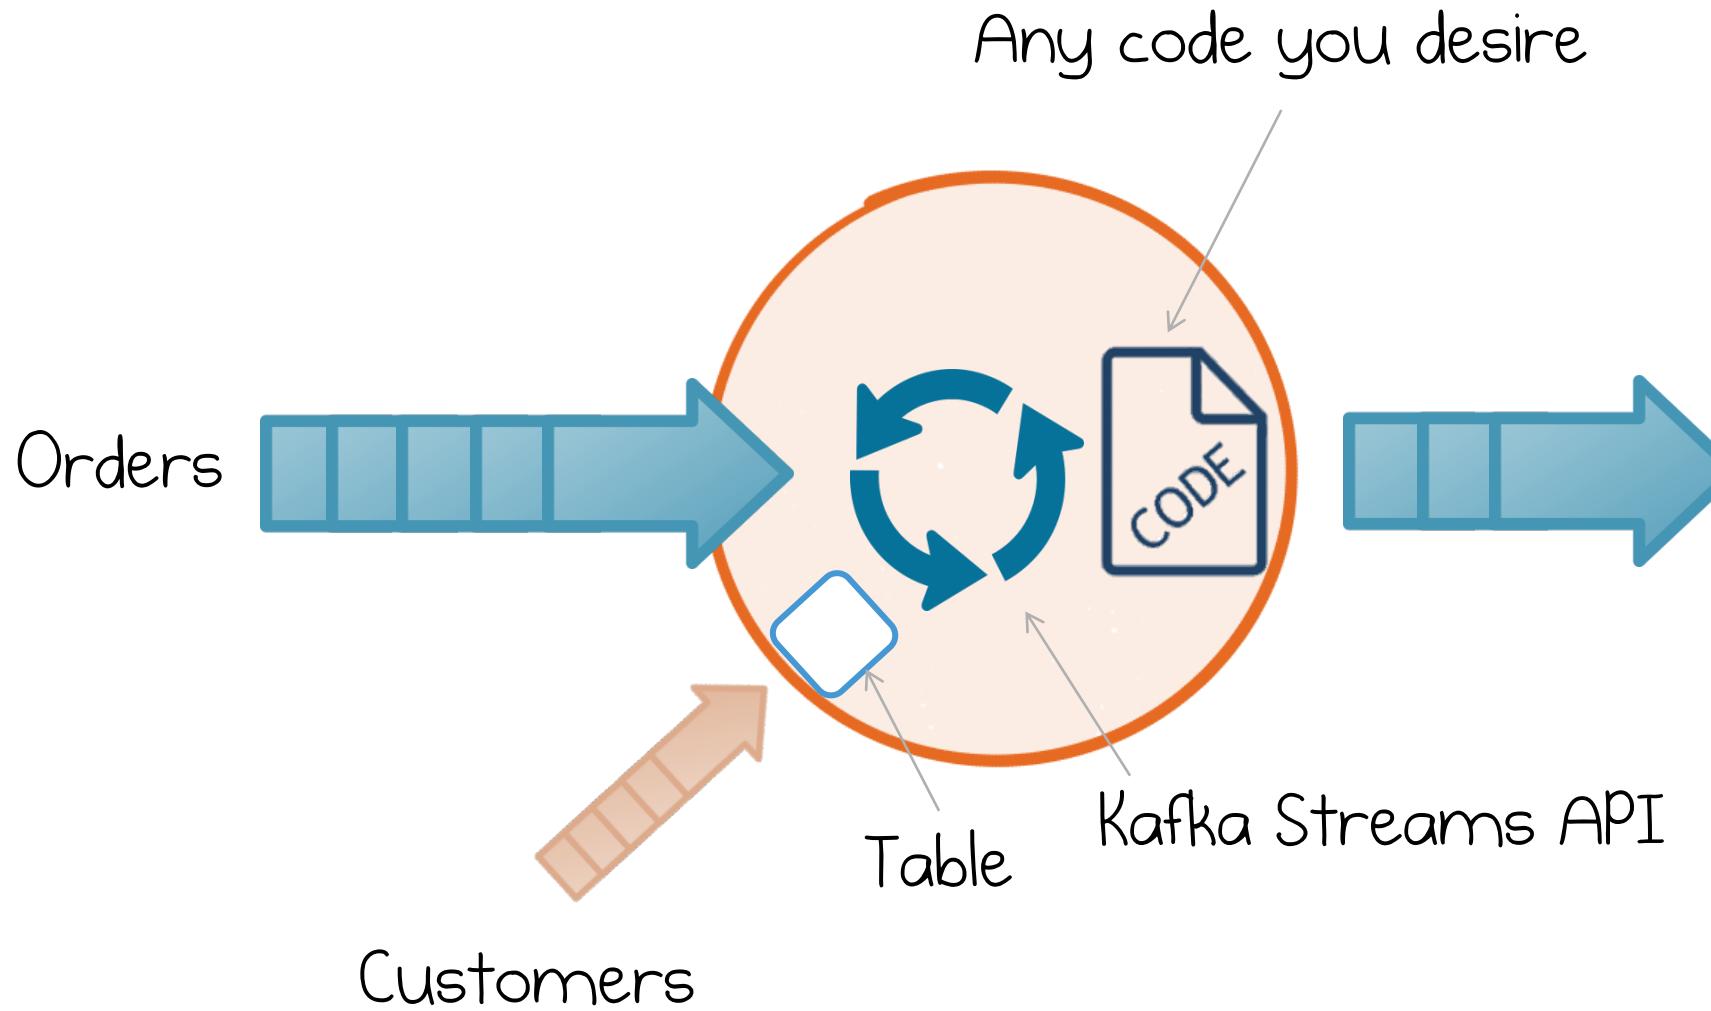
Stateful Stream Processing



Data
replication



KSQL ~ KStreams

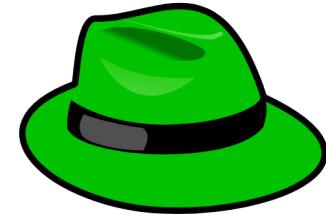


It's just java

```
builder.stream("Orders")
    .join("Customers", ...)
    .transform((key, value) ->
        {//Any code your heart desires!})
)
.to("Shipments");
```

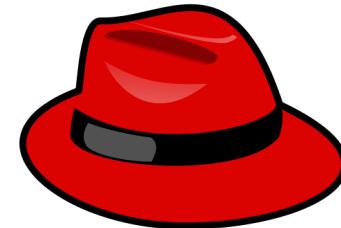
Streaming can be stateful or stateless

1. Joining & Operating on Streams



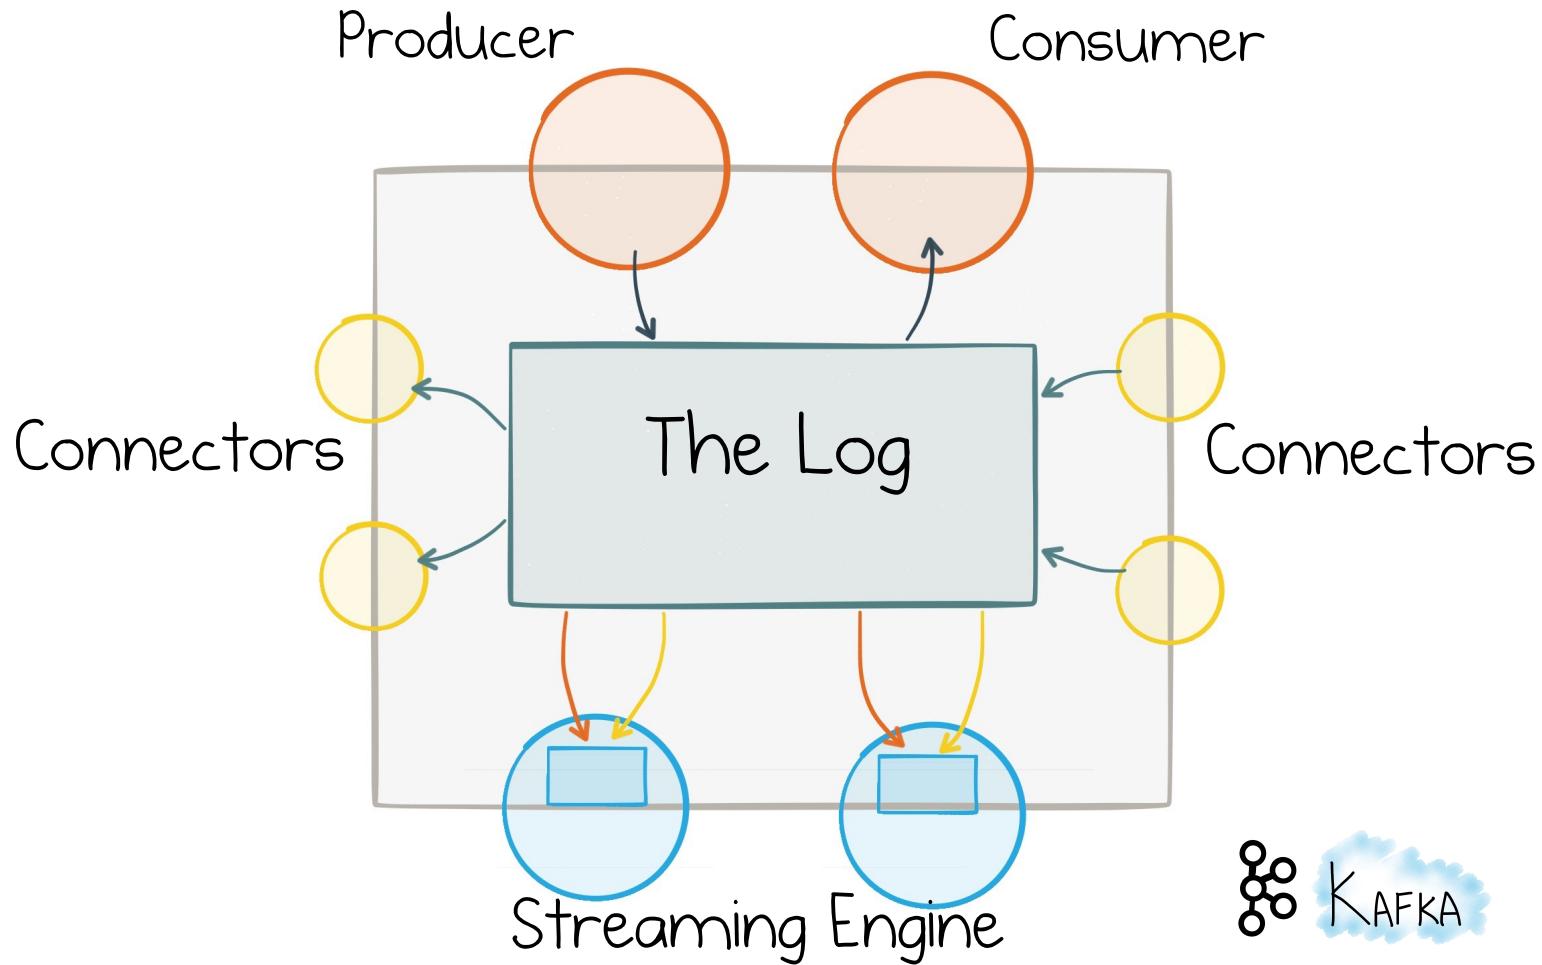
On Notification

2. Joining & Operating on Materialized Tables



Data Replication

Streaming Platform



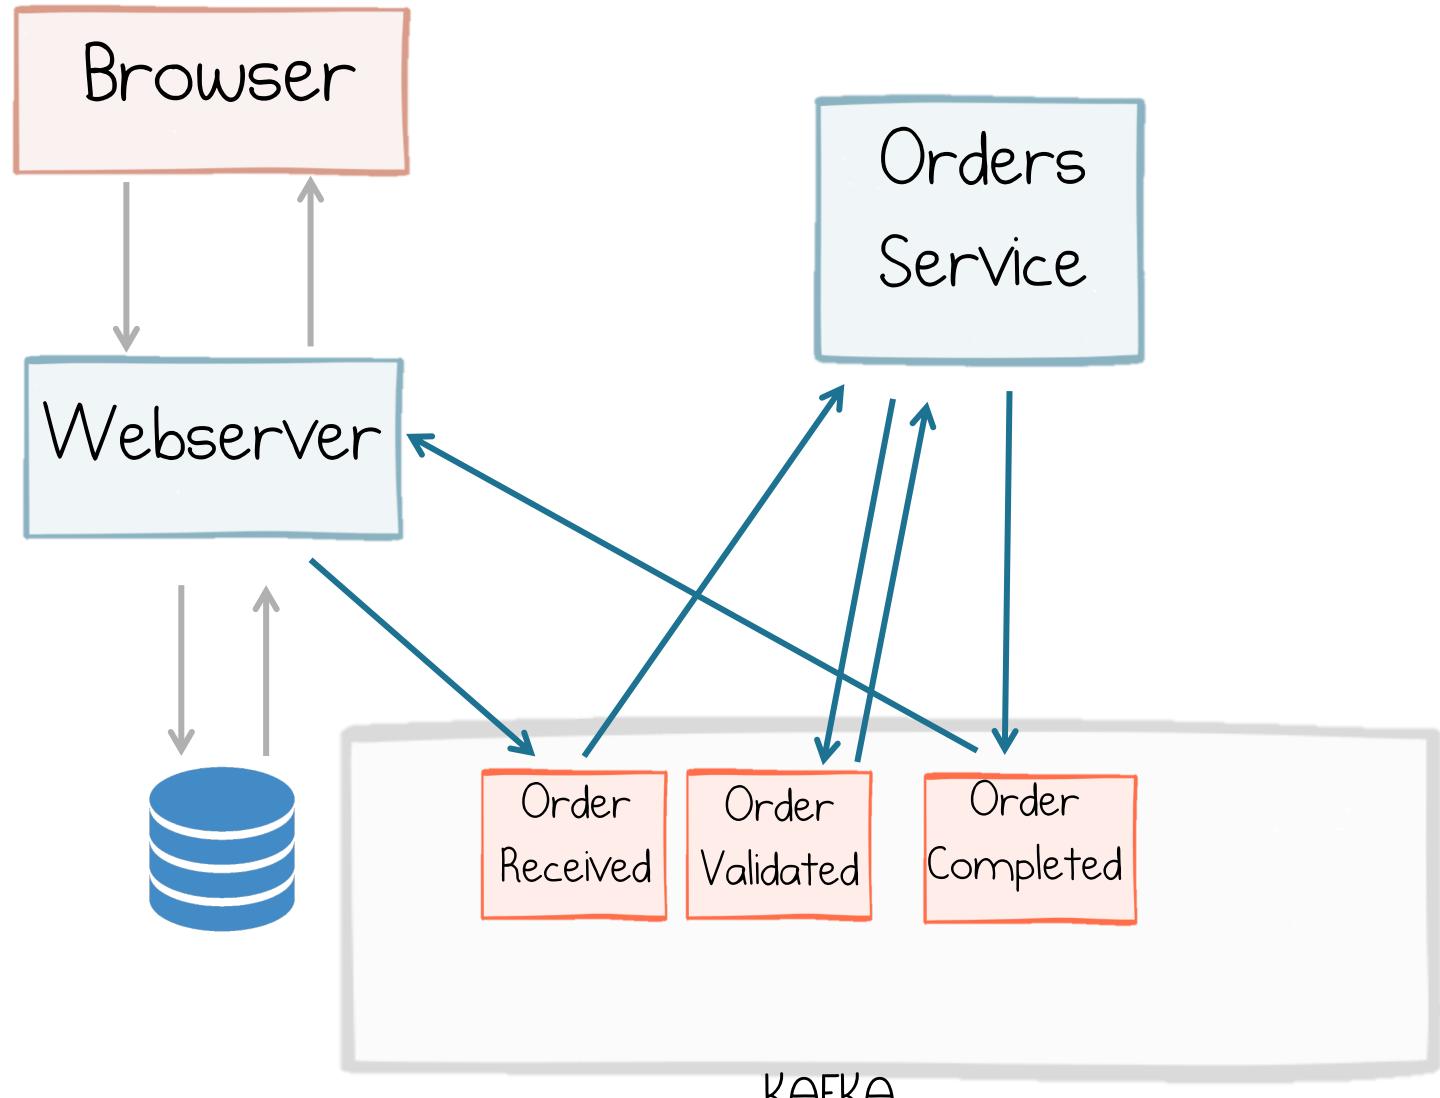
Event Driven Example

I. Use events to decouple and to collaborate

Event Collaboration

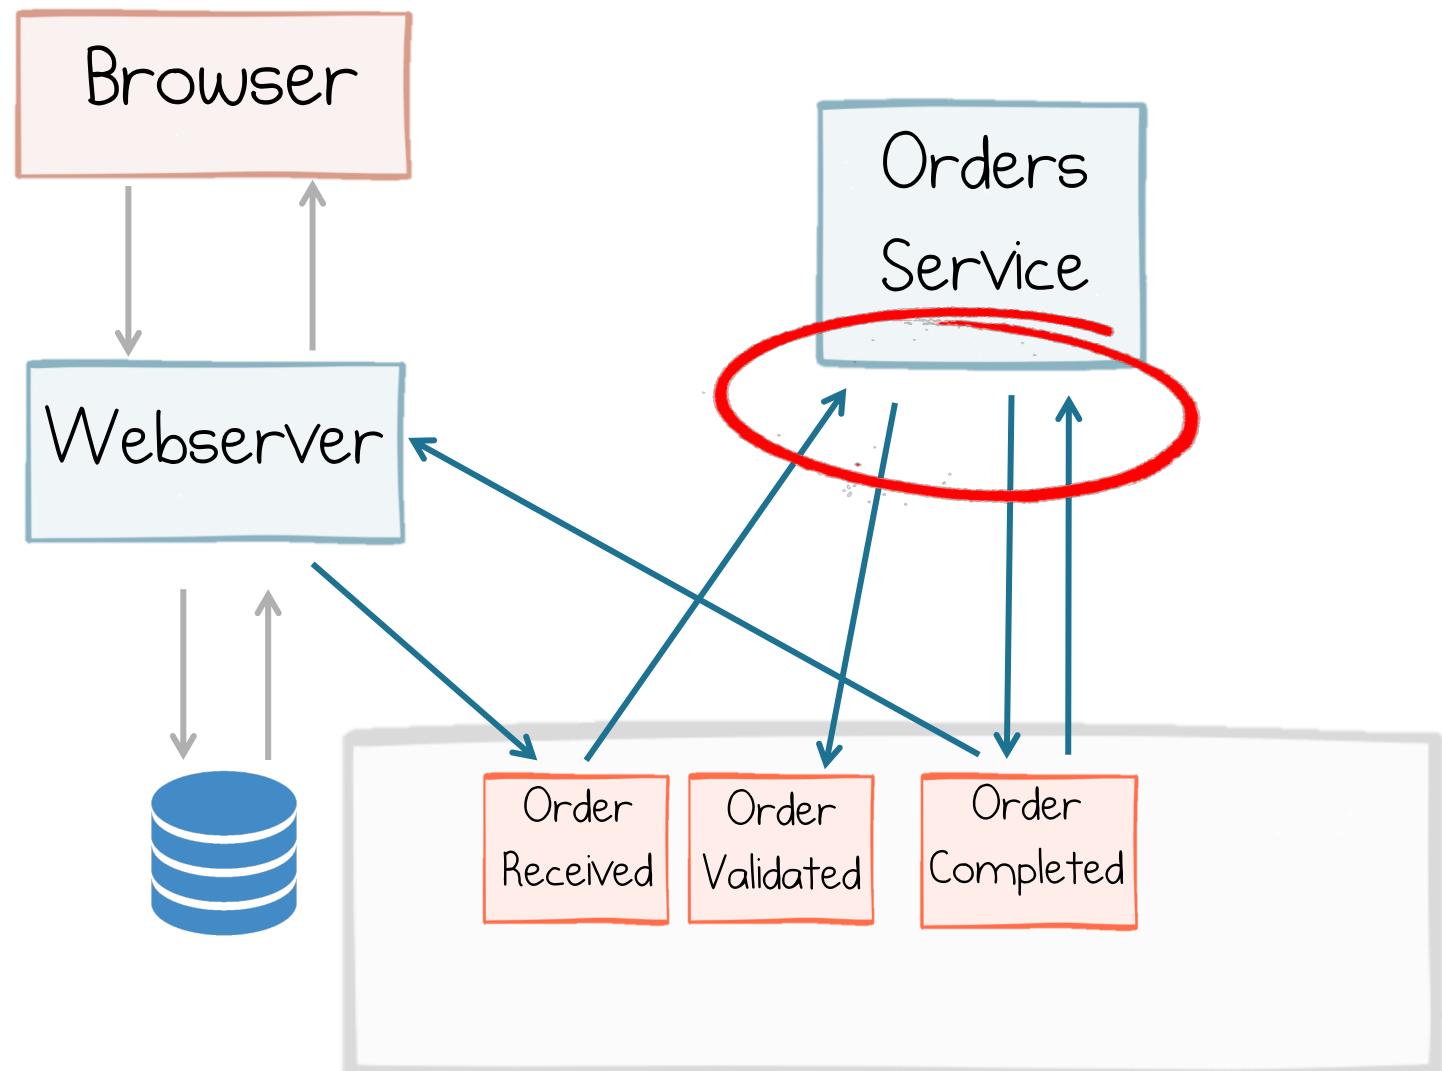


Notification

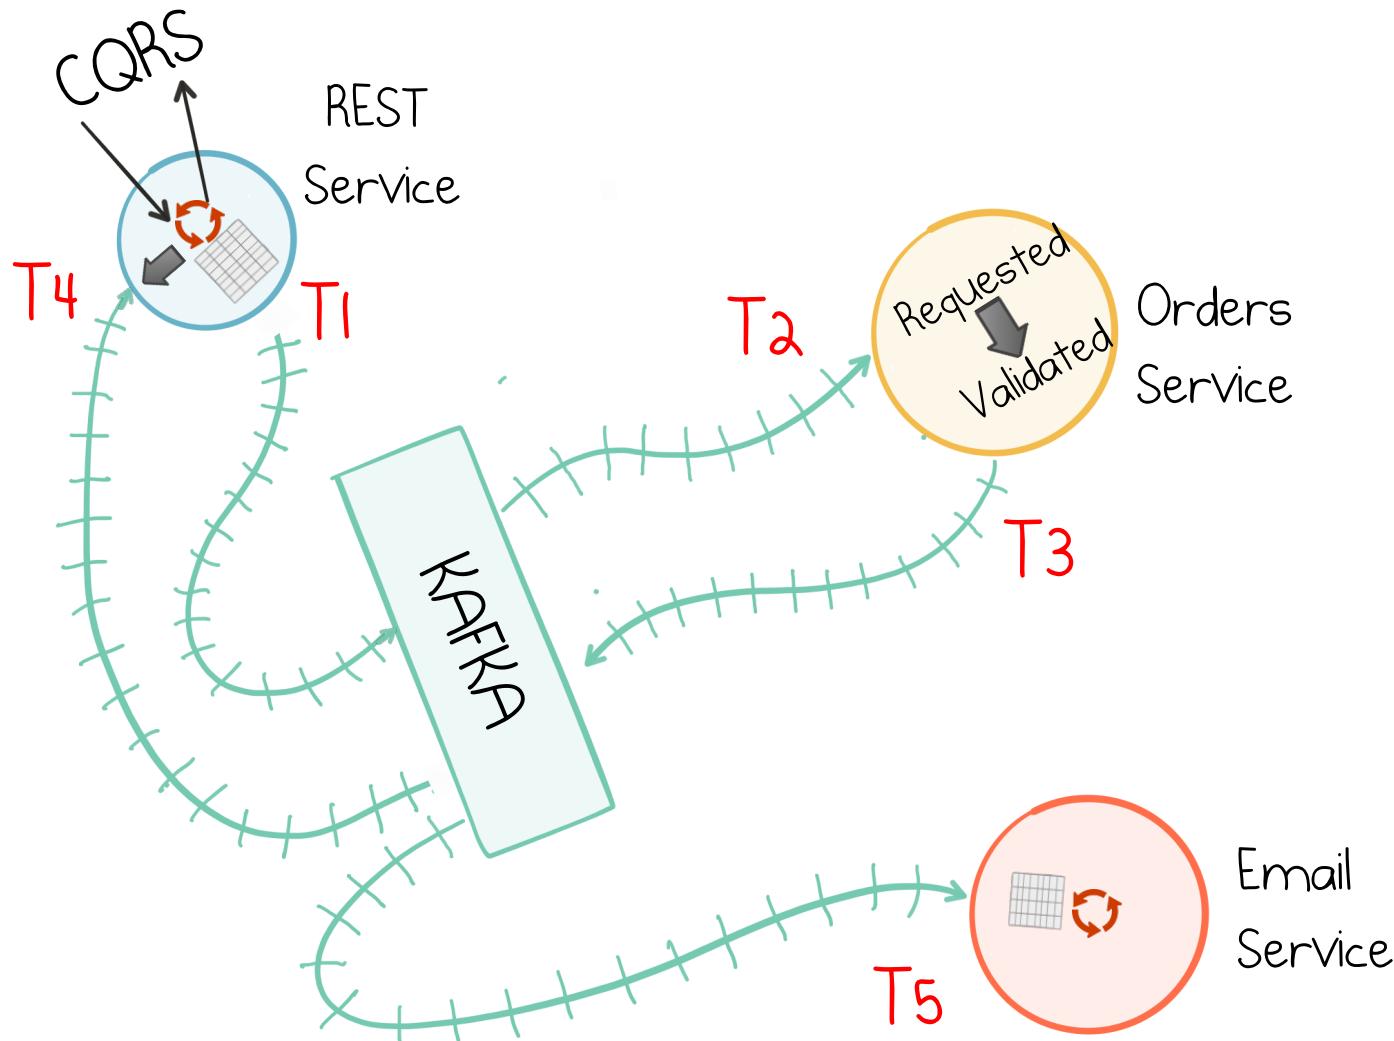


2. Use the Single Writer Principal

State changes to a topic owned by one service

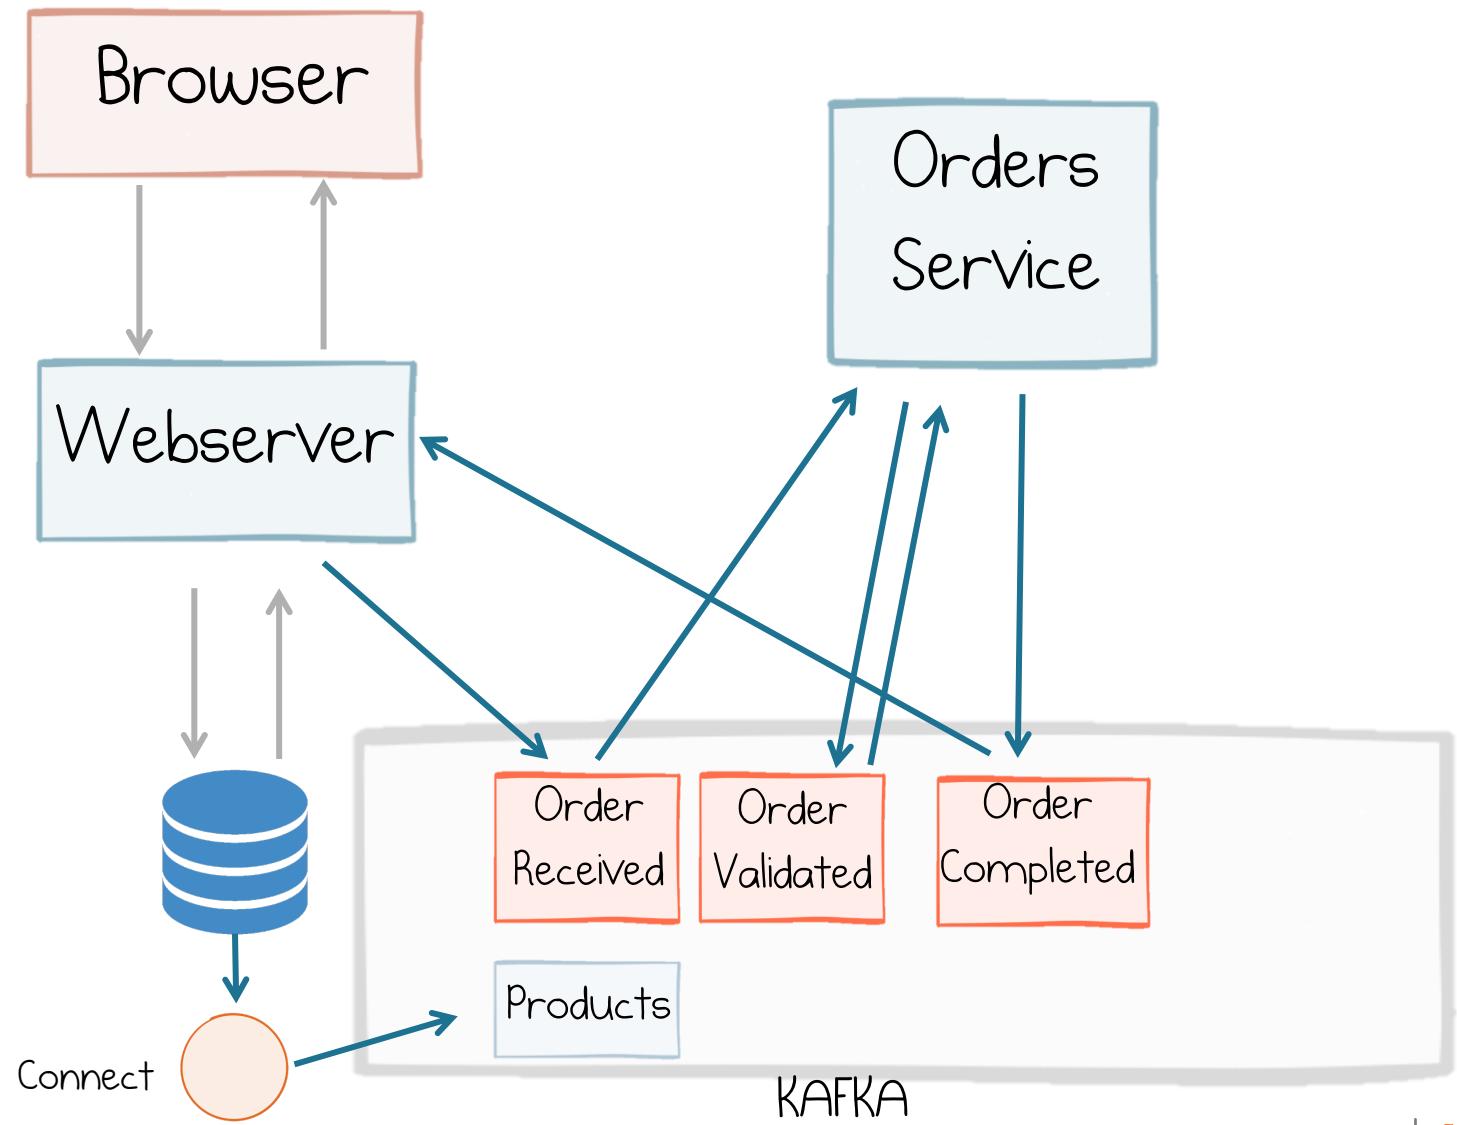


Local consistency points in the absence of Global Consistency



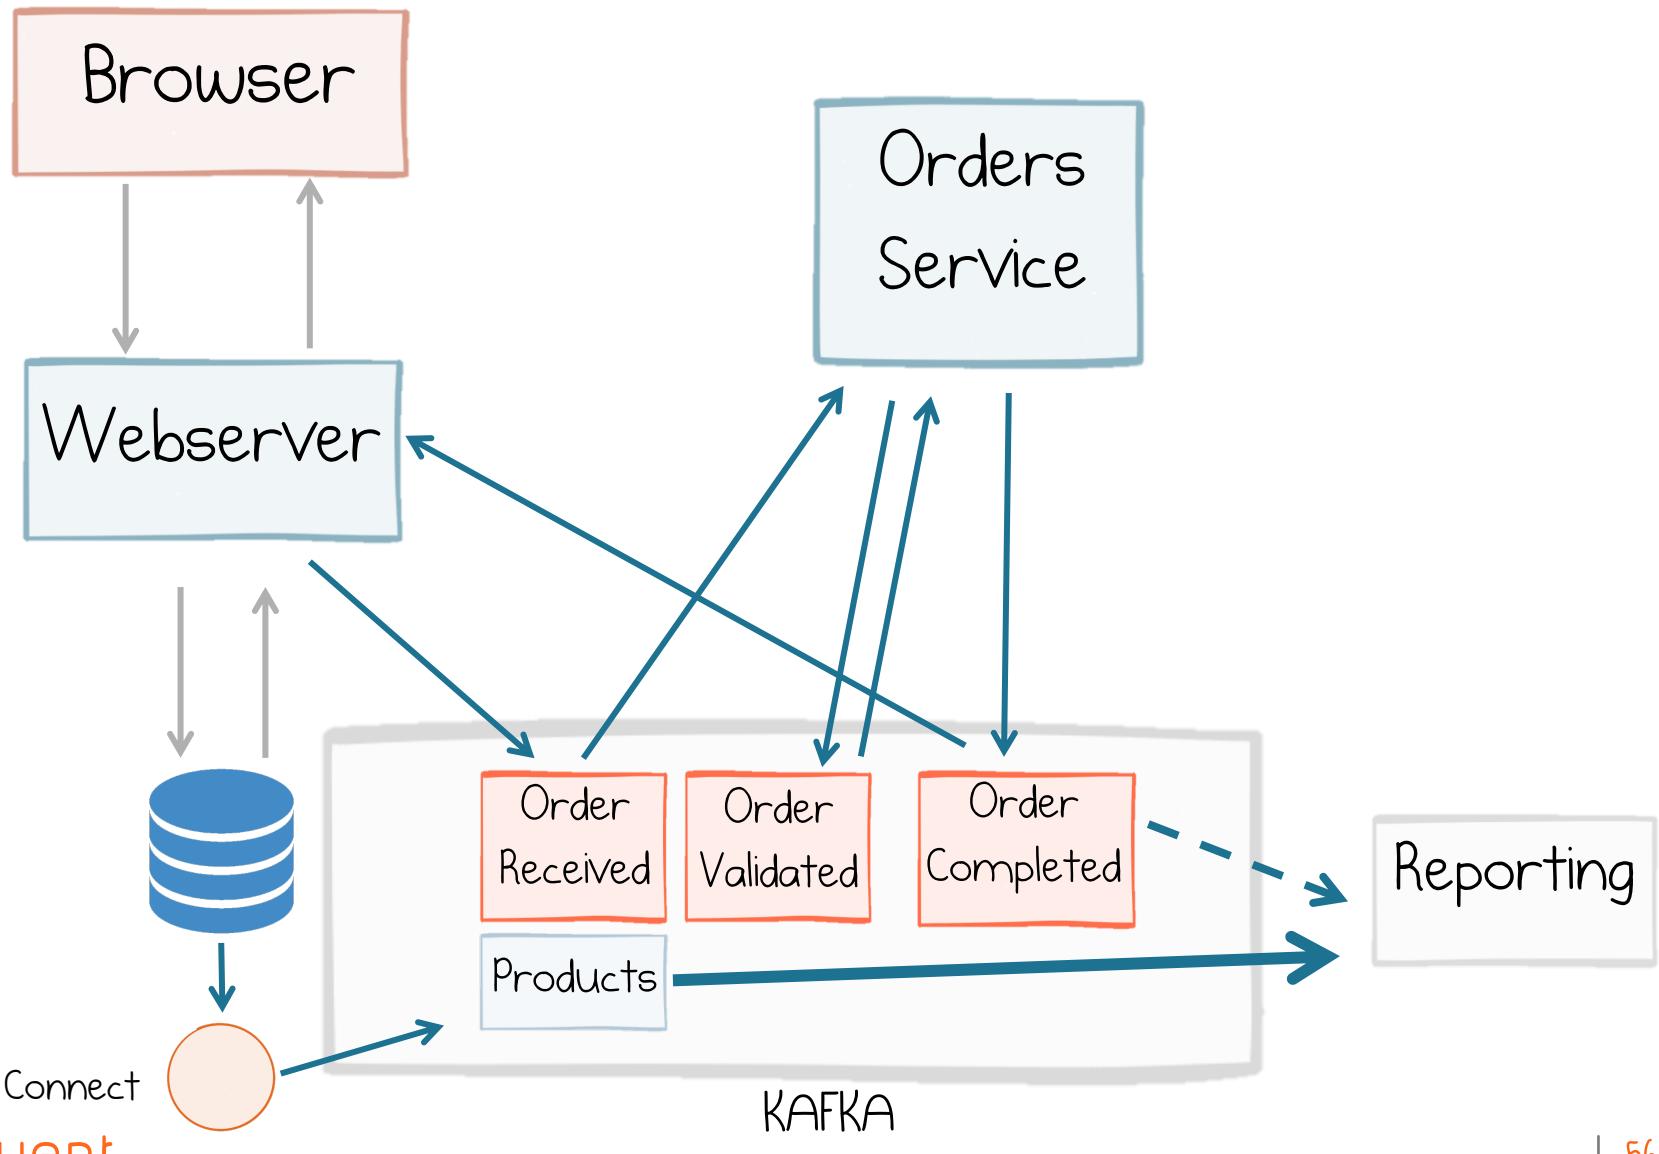
3. Convert legacy databases to Events (with CDC)

Make Legacy Datasets Available via the Log

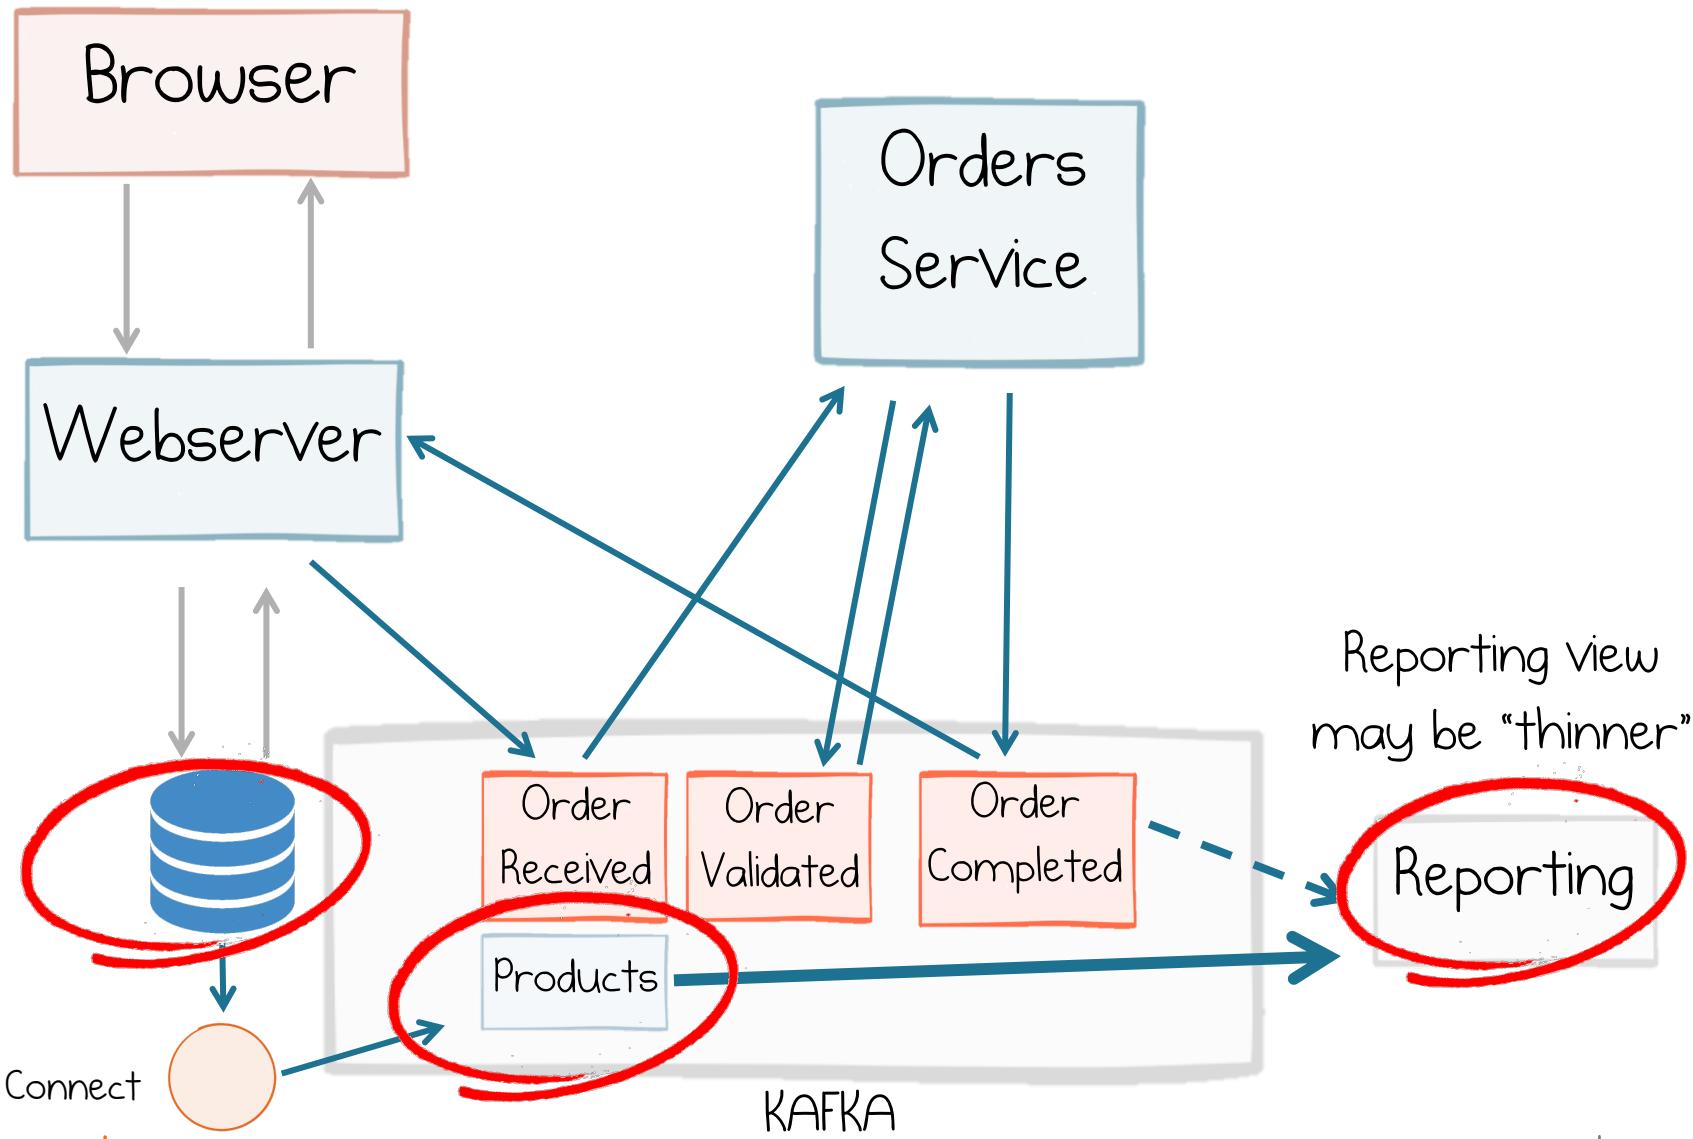


4. Use Kafka as an Event Store

Shared Source of Truth

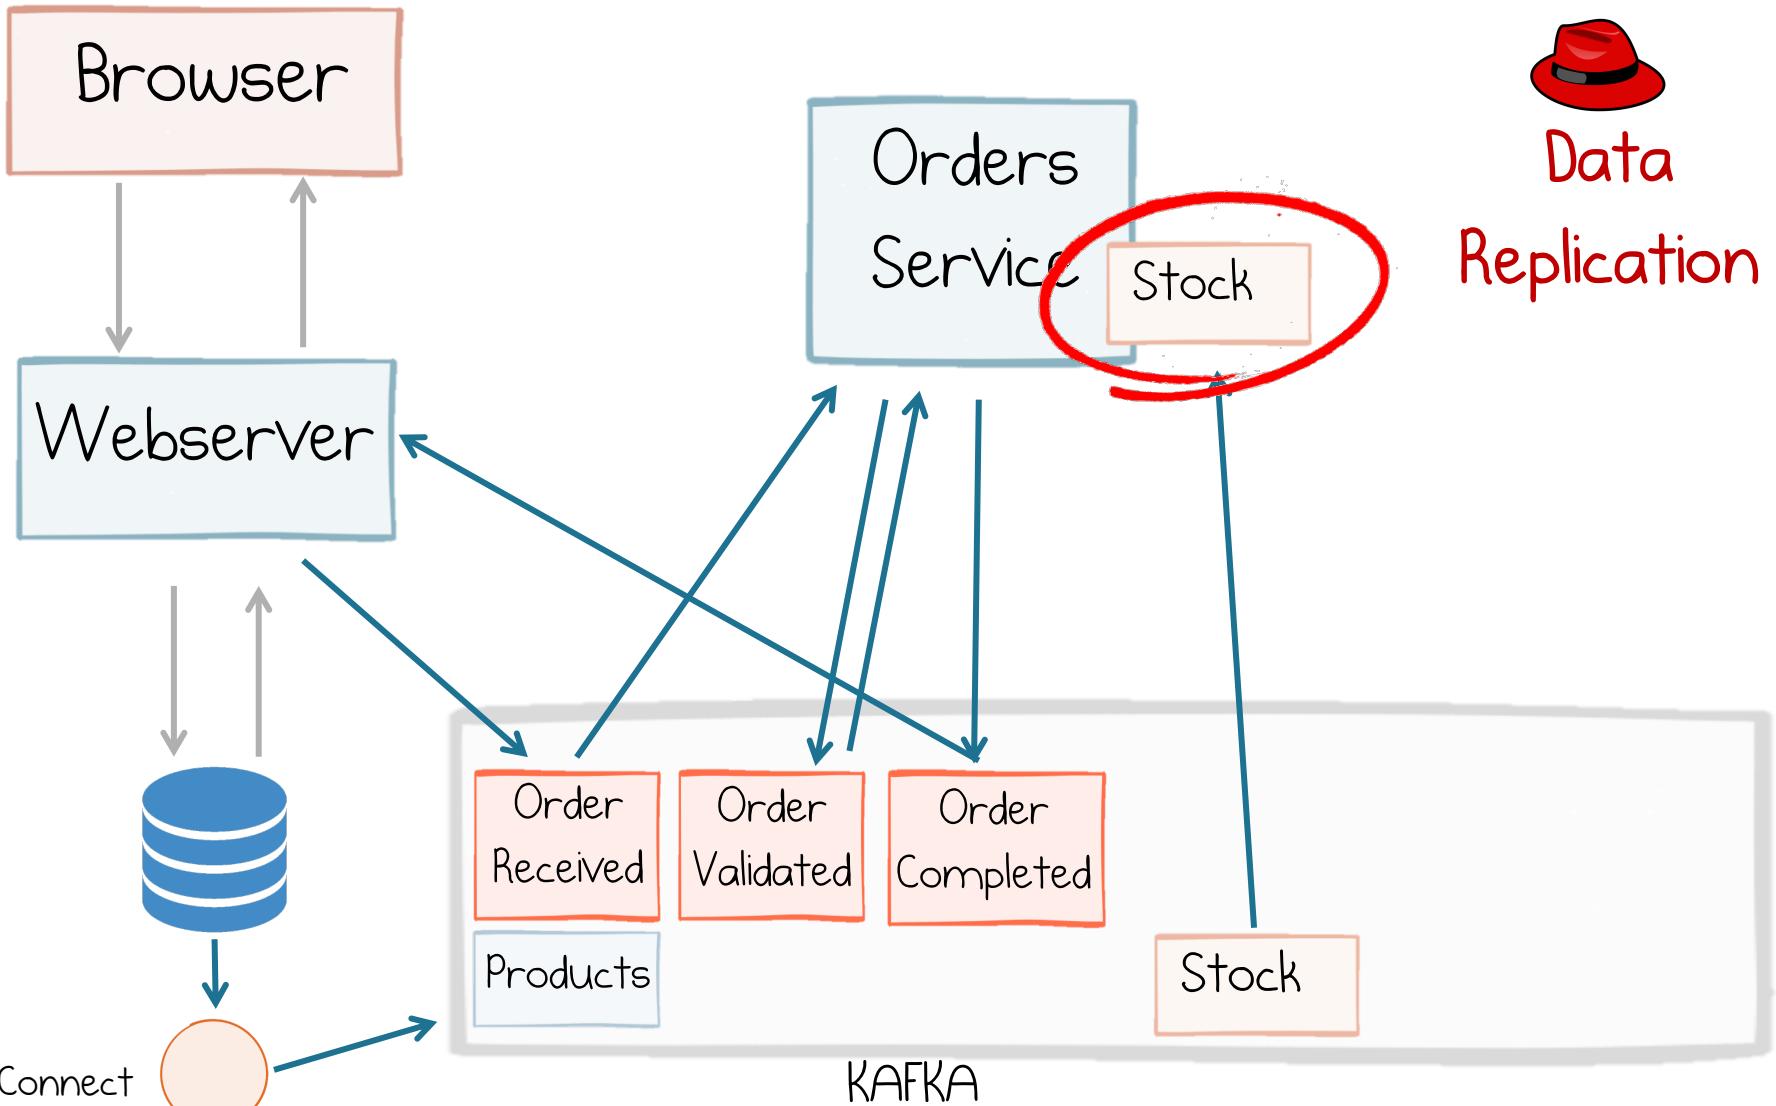


Product Catalogue stored in 3 places



5. Derive “Materialized Views” instead of caching

Materialize Stock 'View' Inside Service

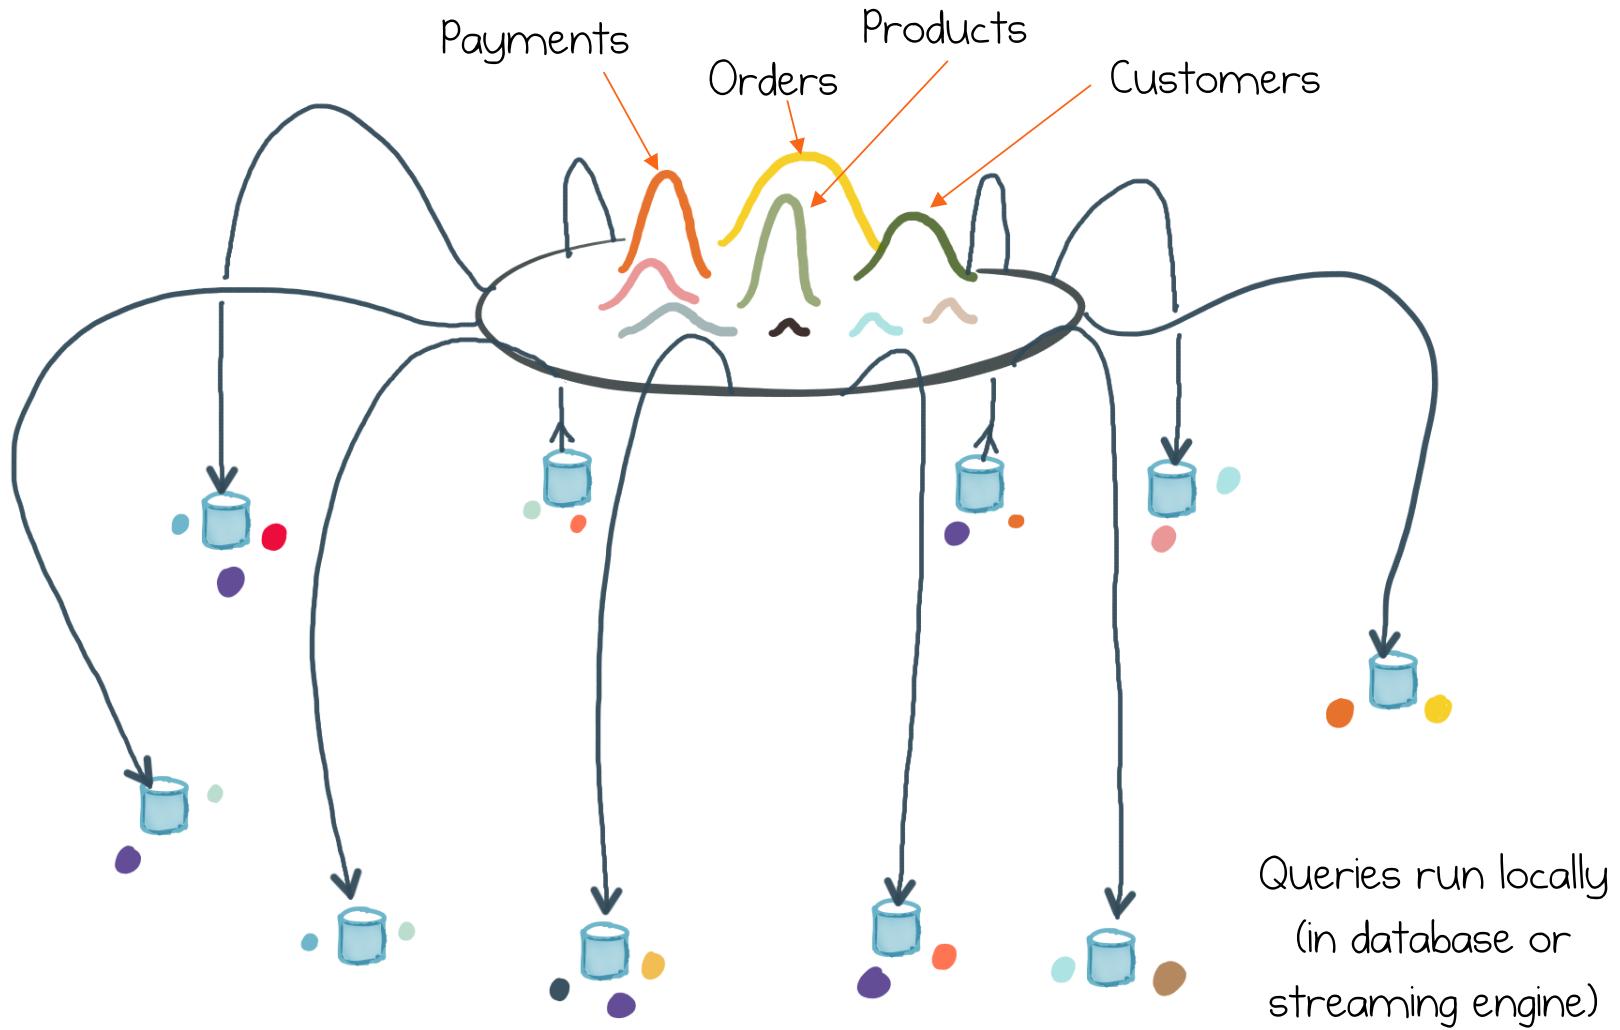


Is stateful a good idea?

- Standby Replicas
- Disk Checkpoints
- Compacted topics
- Or just use a database!

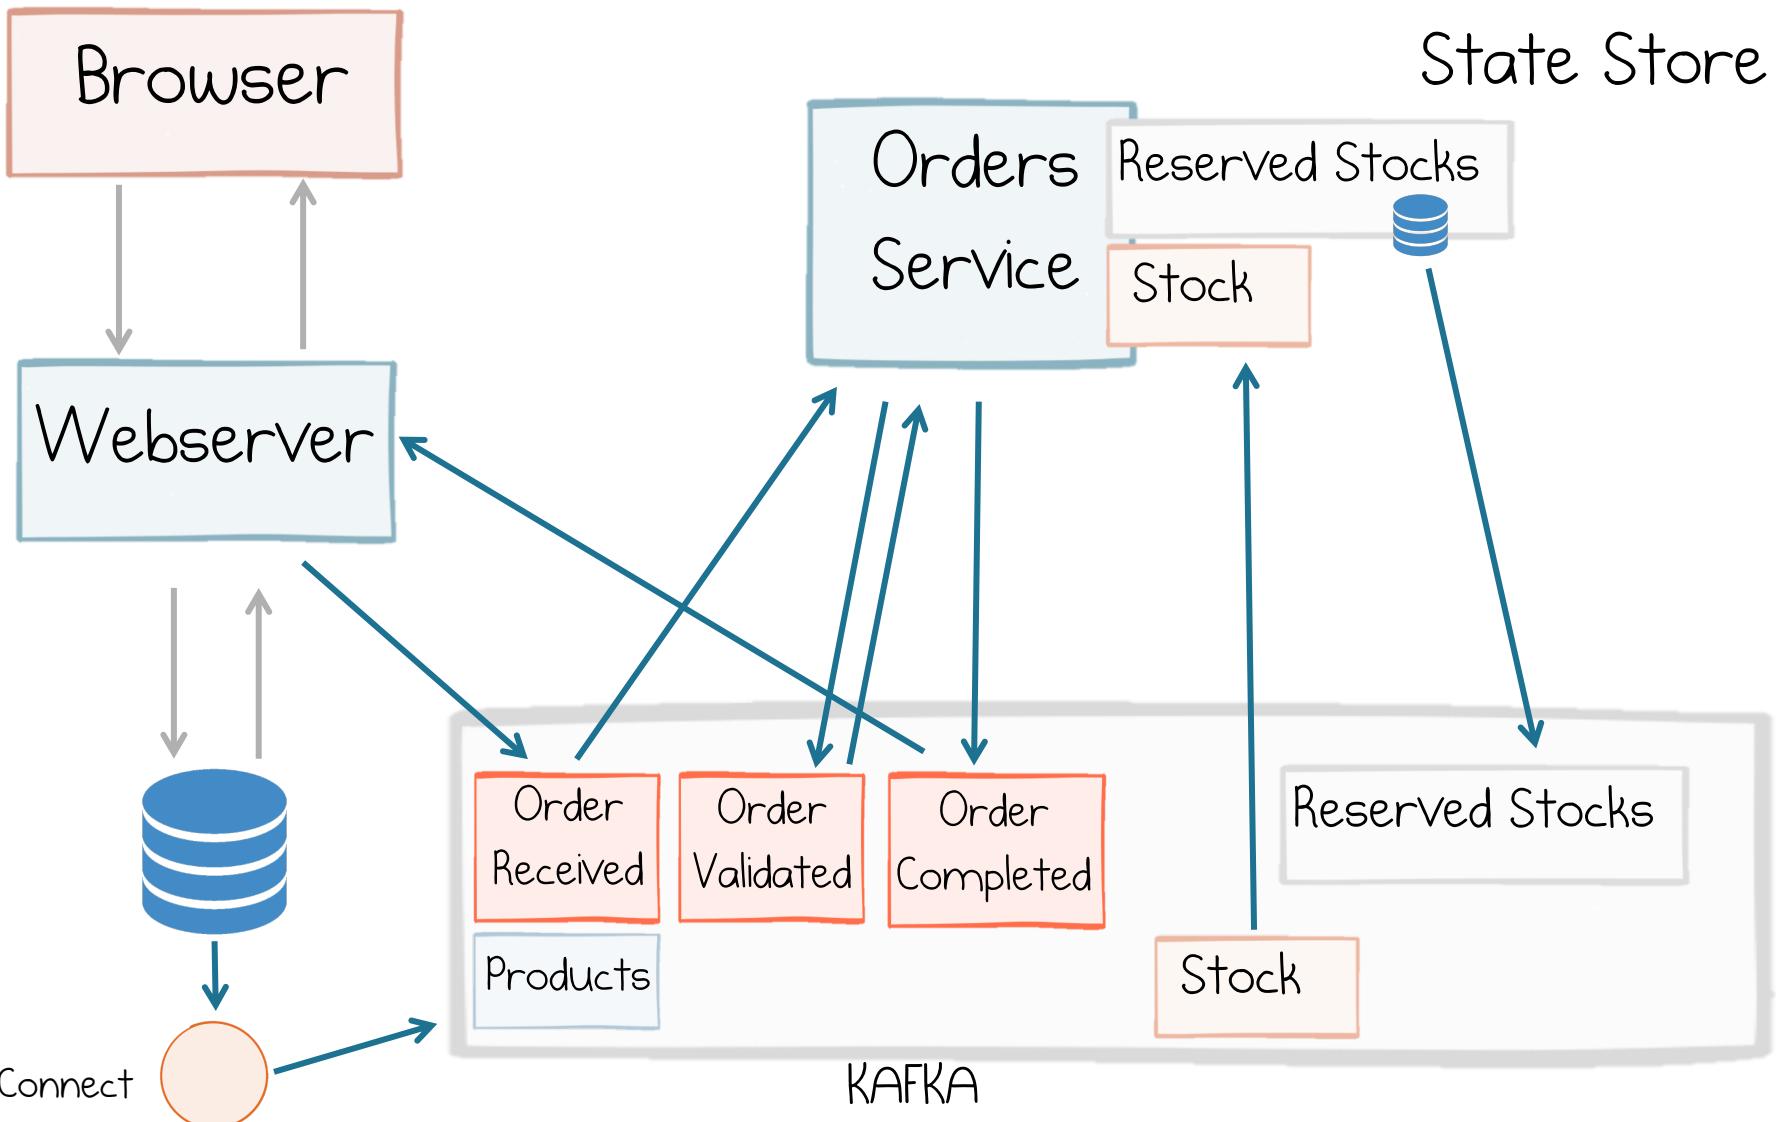
Database Inside Out Pattern

KAFKA: Retained Event Streams



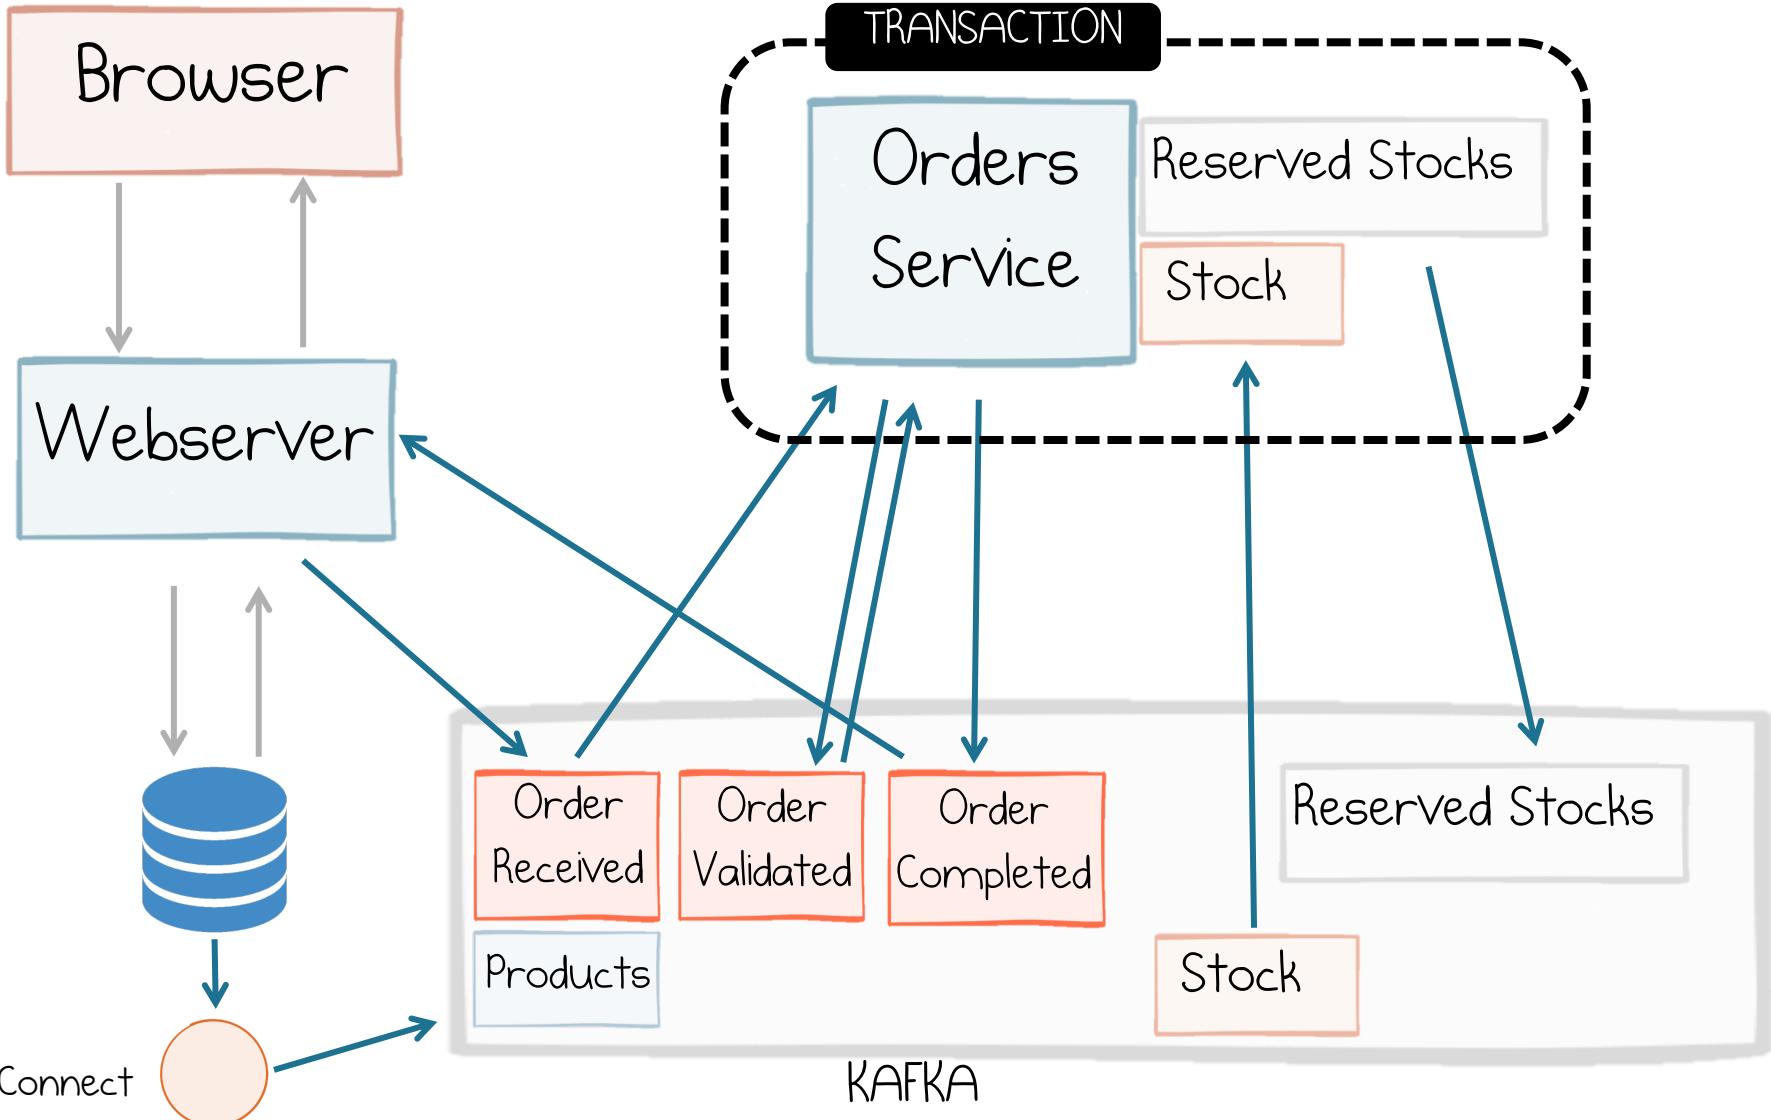
6. Write to State Stores, just like a local 'database', backed up in Kafka

State stores behave like local databases



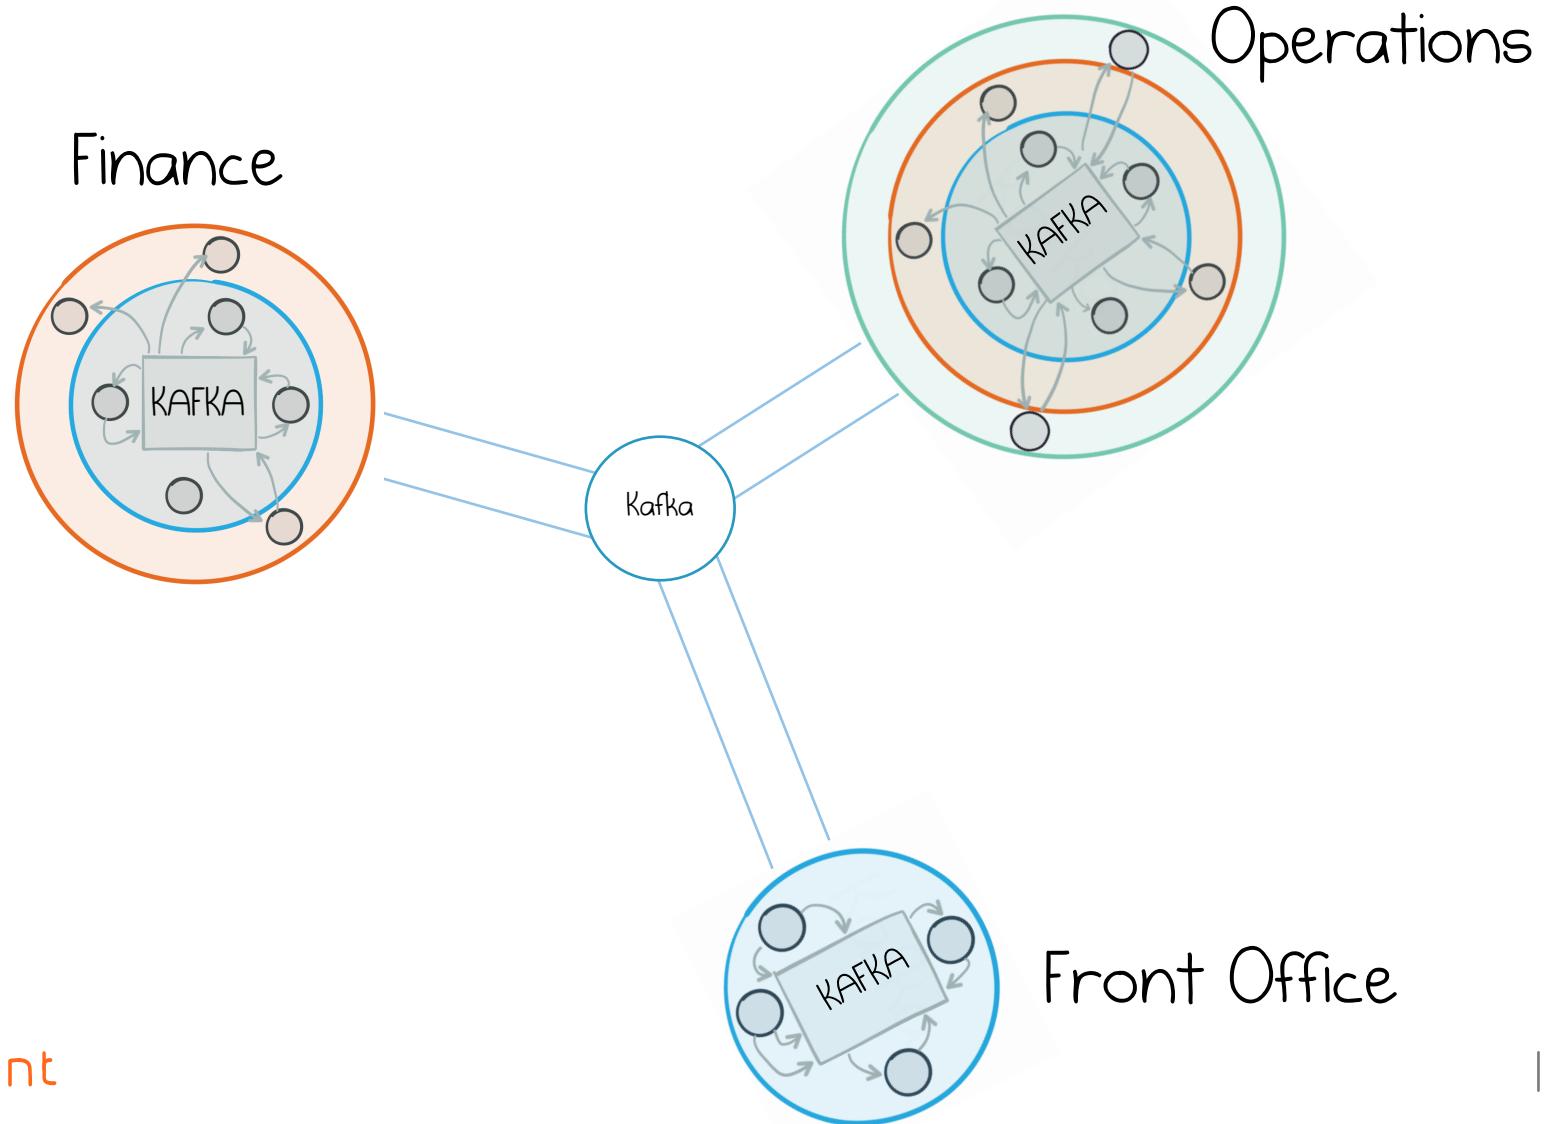
7. Use Transactions to tie All Interactions Together

Transactions

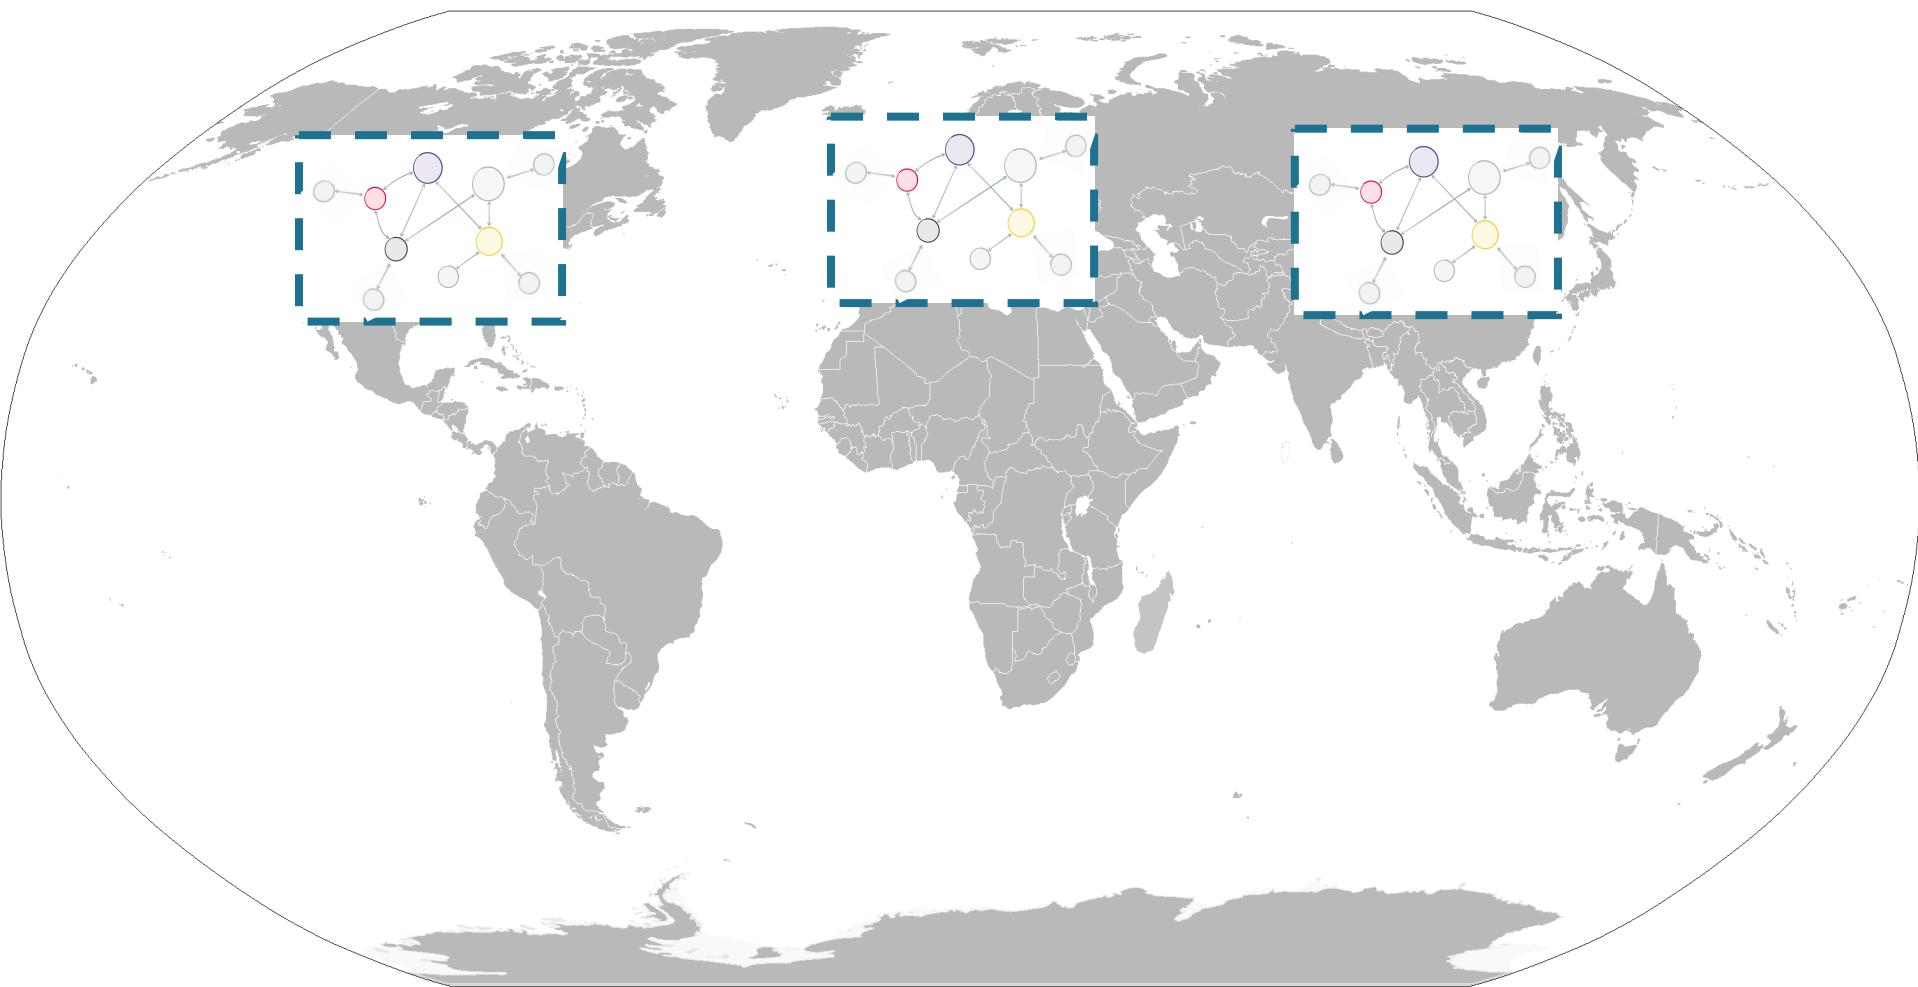


8. Evolve and Grow through Streaming Functions

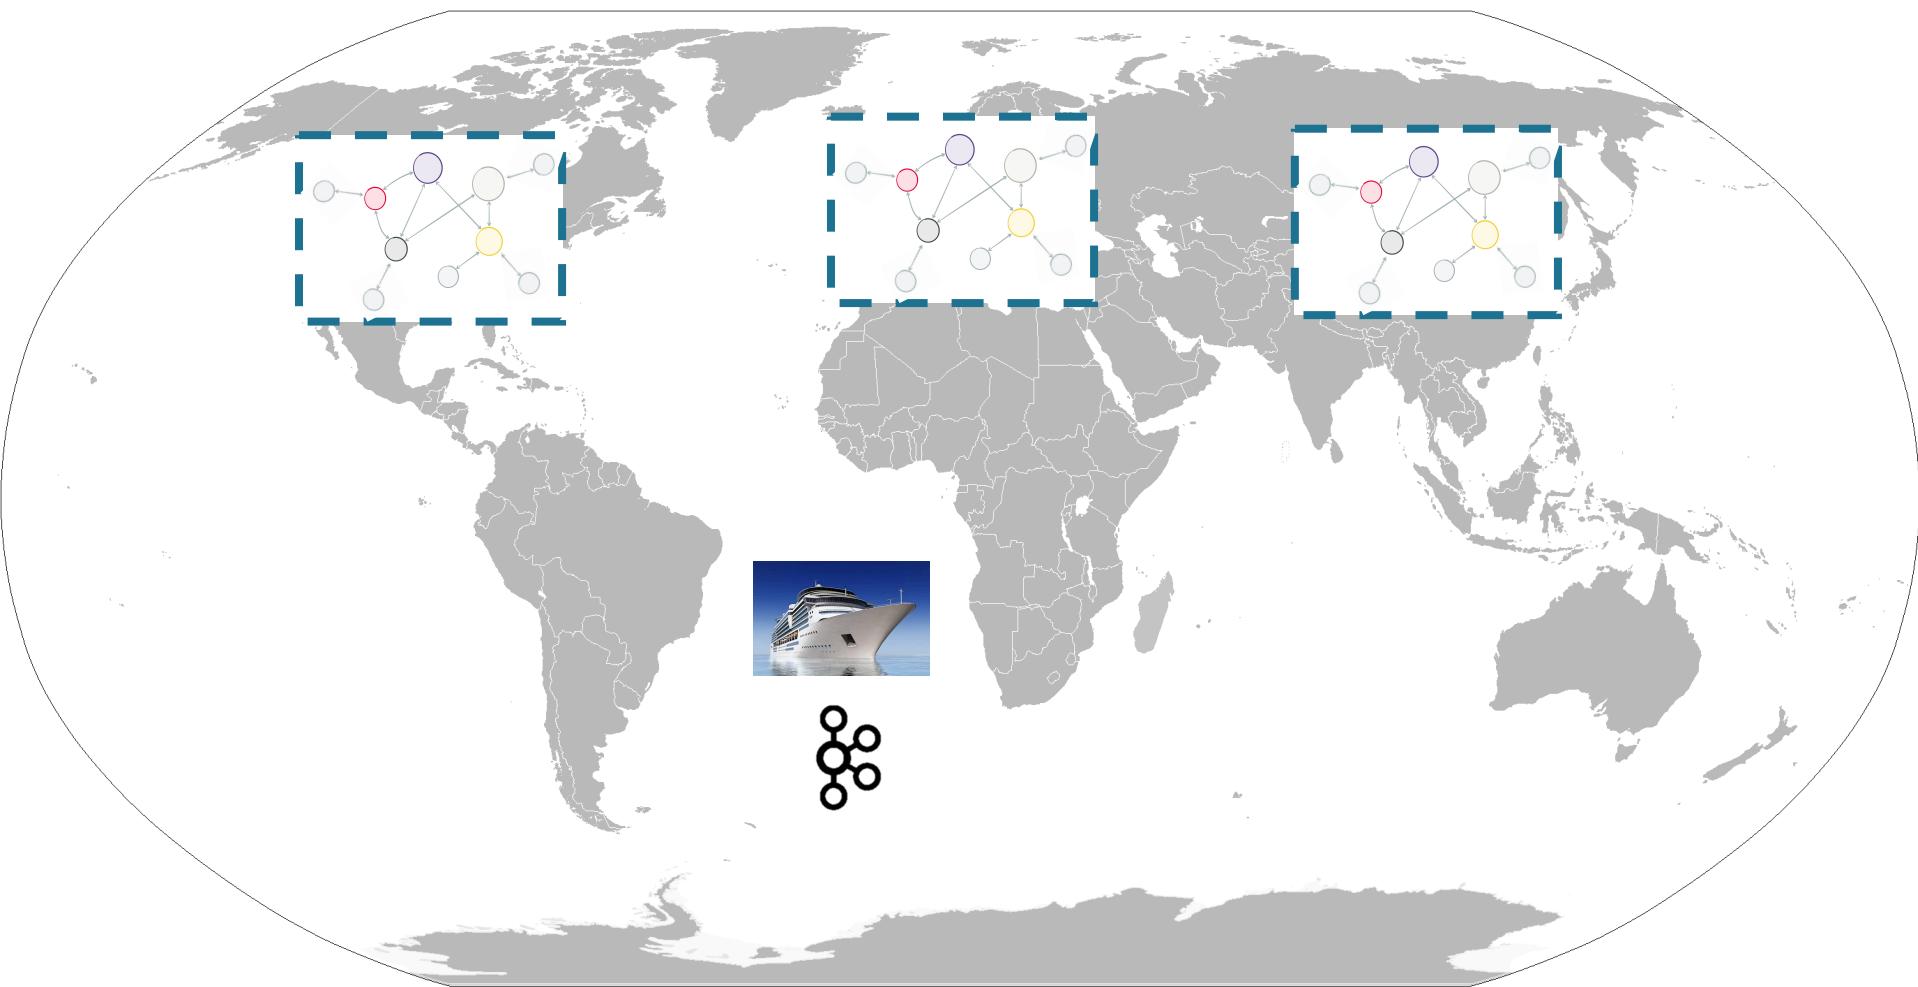
Tiered Contexts



Span regions or clouds



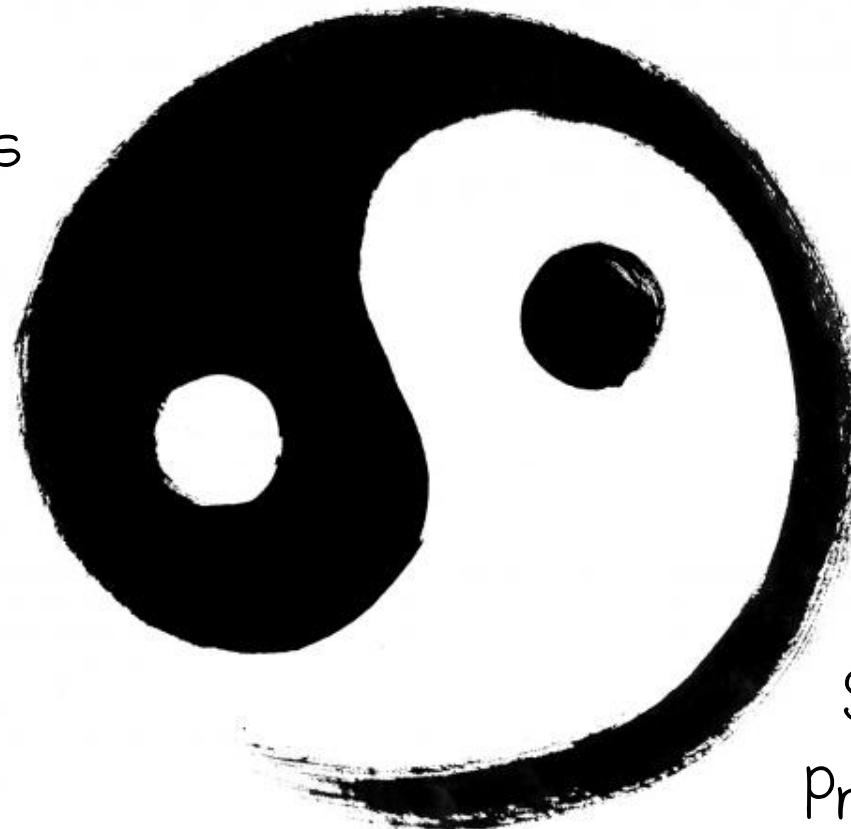
Handle Disconnectedness



So...

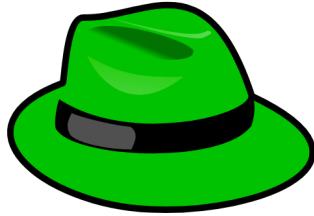
Optimize for complexity vs optimize for scale

Event Driven
Architectures

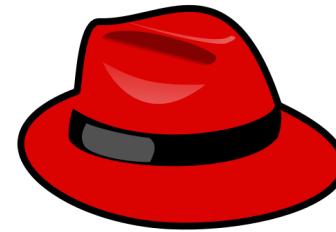


Stream
Processing

Events provide the key to evolutionary architectures

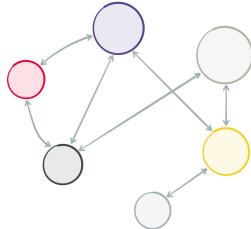


Notification

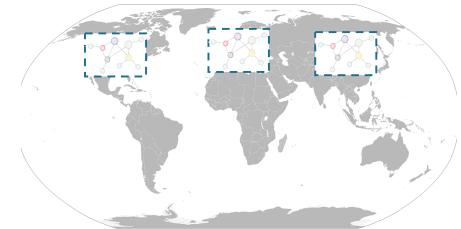


Data
replication

Spectrum of use cases



Finer Grained,
Collaborative,
Connected



Courser Grained,
Non-collaborative,
Disconnected

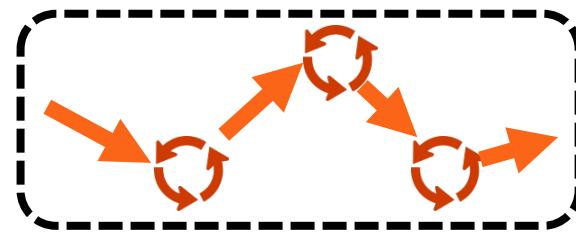


Notification

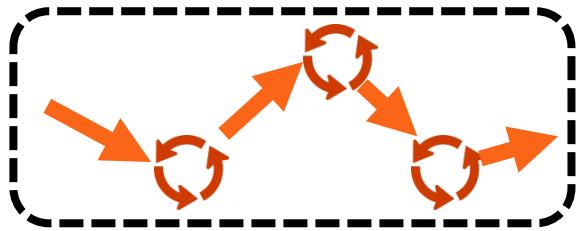


Data Replication

Events to transcend individual services



Notification



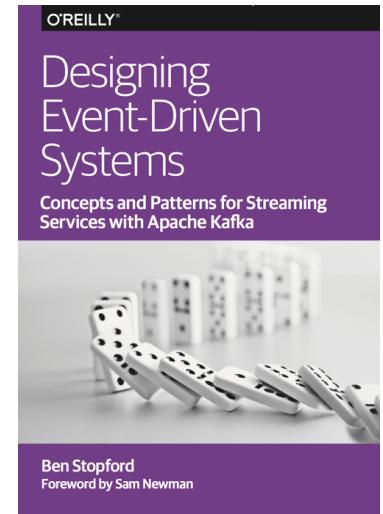
Data
replication

Start Simple and Evolve

1. Broadcast events
2. Retain them in the log
3. Evolve the event-stream with streaming functions
4. Recasting the event stream into views when you need to query.

Find out more

Book: <http://bit.ly/designing-event-driven-systems>



Software: <https://confluent.io/download/>

Cloud: <https://www.confluent.io/confluent-cloud/>

Kubernetes Operator: <https://www.confluent.io/confluent-operator/>

Example code: <http://bit.ly/kafka-microservice-examples>

Twitter: @benstopford

